

## Module 05 – Extra Class

# Multi-layer Perceptron

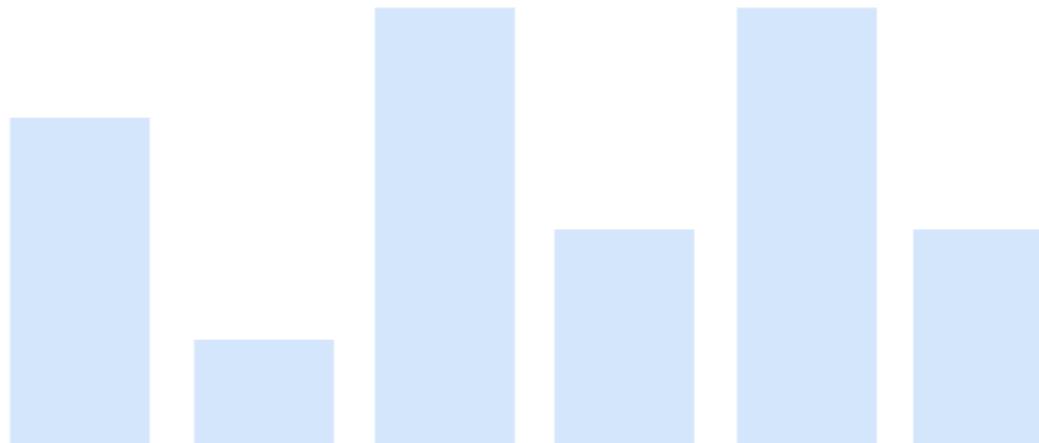
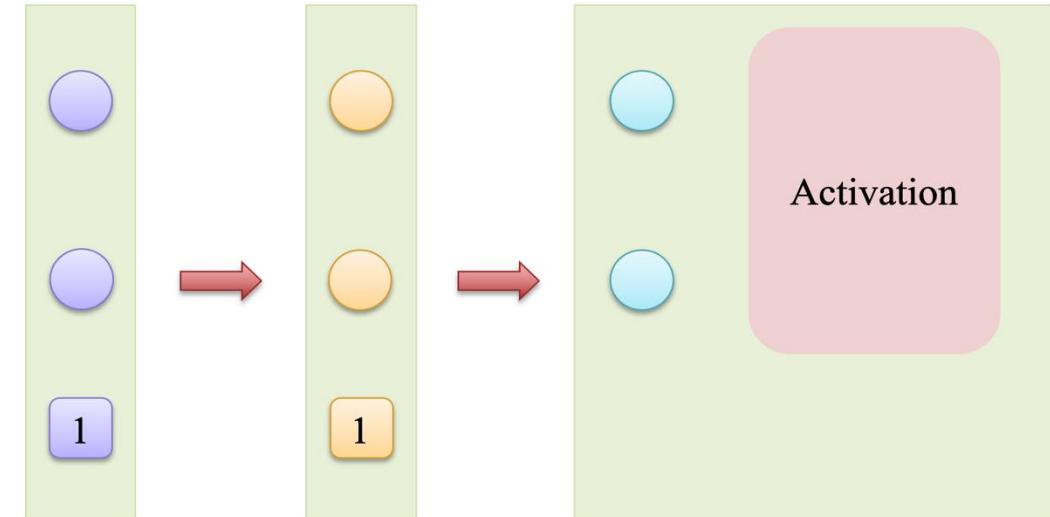
Nguyen Quoc Thai

Code - Data

# Objectives

## Softmax Regression (Review)

- ❖ Logistic Regression
- ❖ Softmax Function
- ❖ Softmax Regression using SGD



## Multi-layer Perceptron

- ❖ Hidden Layers
- ❖ Activation
- ❖ Loss Function
- ❖ Optimization

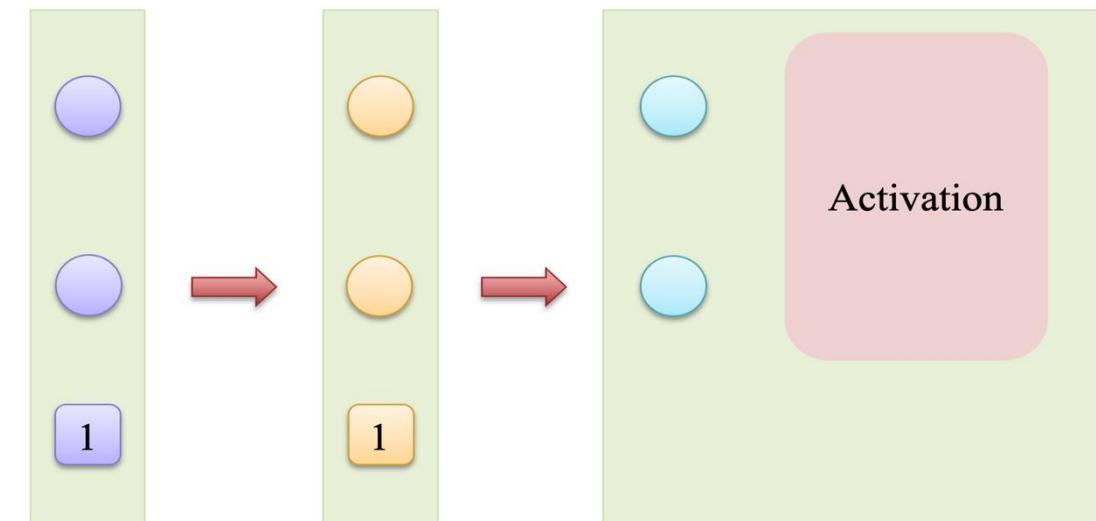
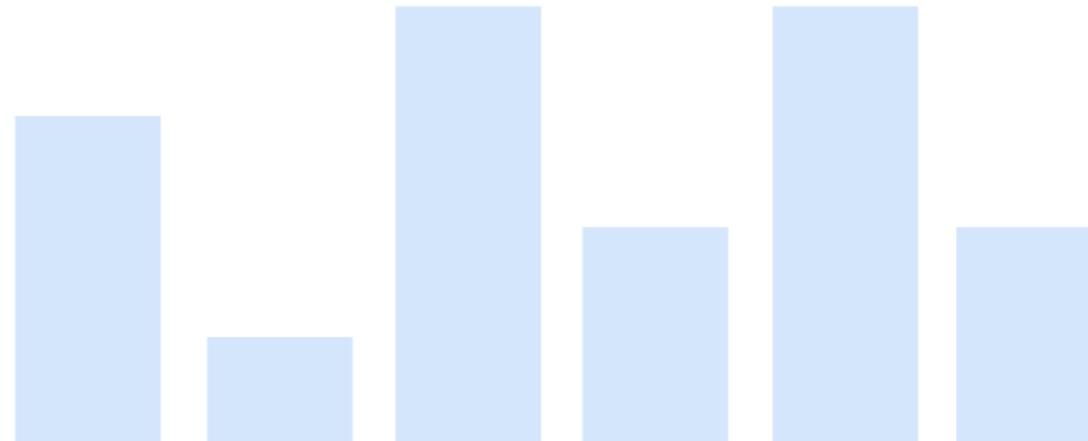
# Outline

SECTION 1

## Softmax Regression

SECTION 2

## Multi-layer Perceptron



# Softmax Regression

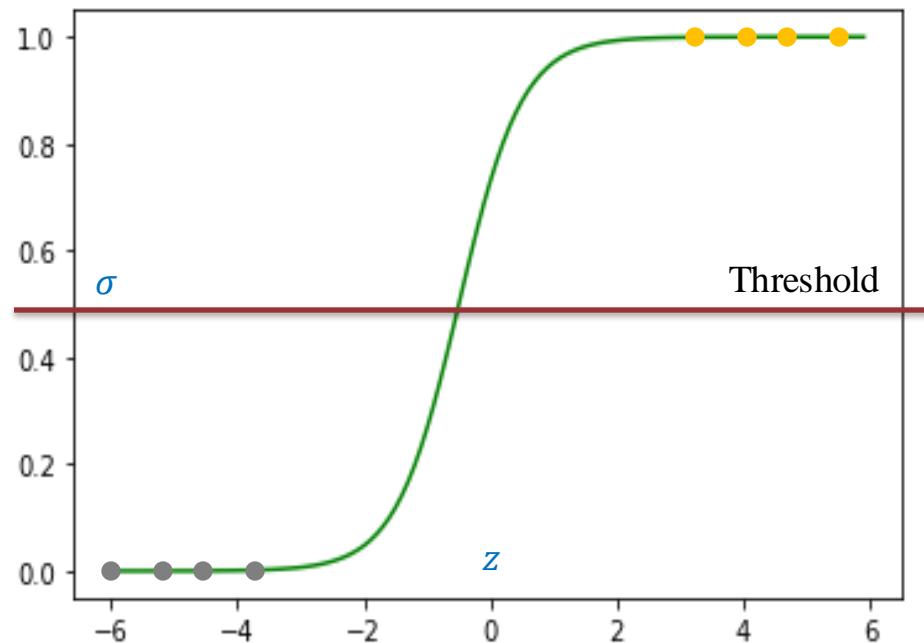


## Problem

Sigmoid function

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

$$z \in (-\infty, +\infty) \quad \sigma(z) \in (0, 1)$$



Hours	Pass
0.5	0
1.0	0
1.5	1
2.0	1

Classes: {0, 1}  
Binary Classification

Hours	Score
0.5	0
1.0	0
1.5	1
2.0	1
2.5	2
3.0	2
3.5	3
4.0	3

Classes: {0, 1, 2, 3}  
Multi-class Classification

# Softmax Regression

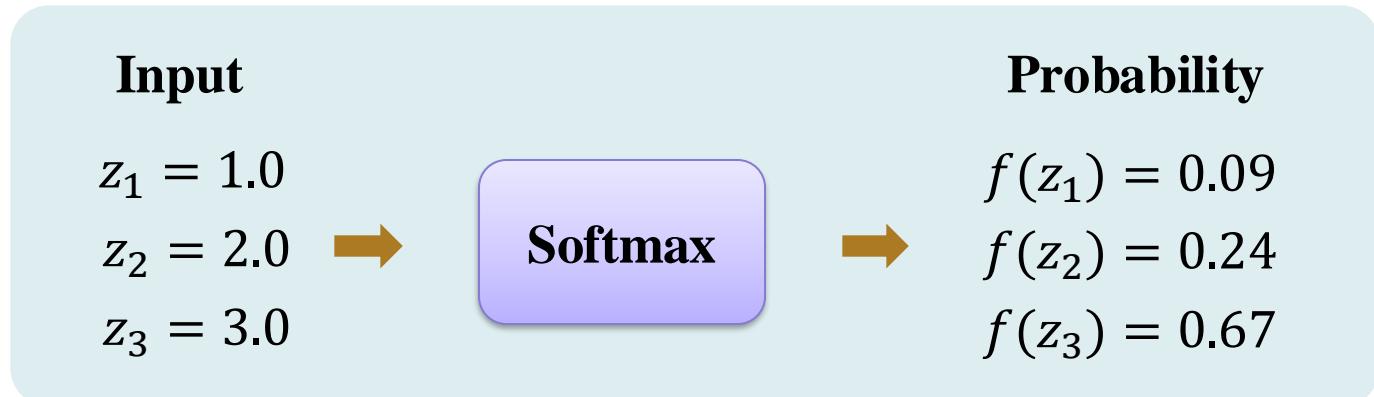
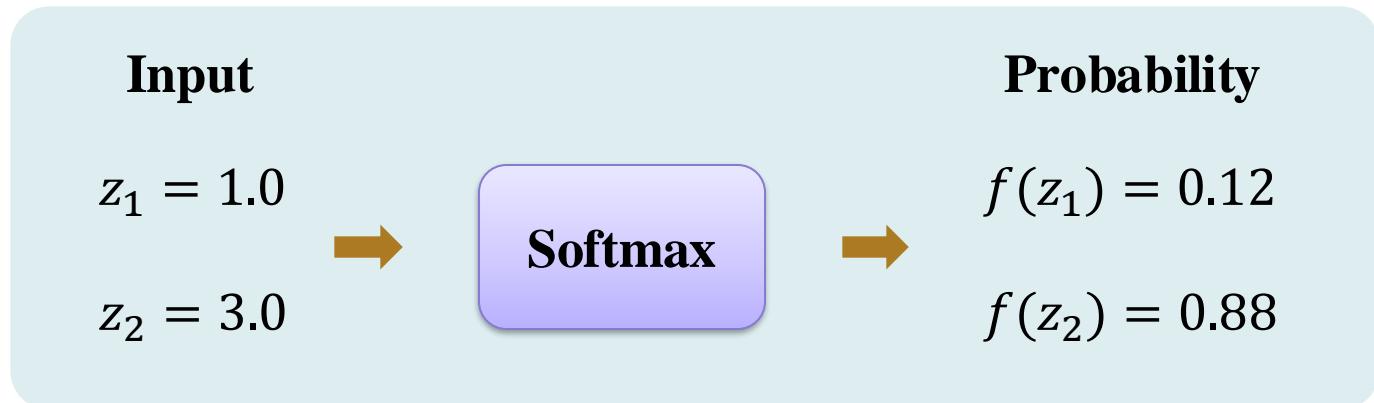


## Softmax Function

$$P_i = f(z_i) = \frac{e^{z_i}}{\sum_j e^{z_j}}$$

$$0 \leq f(z_i) \leq 1$$

$$\sum_i f(z_i) = 1$$



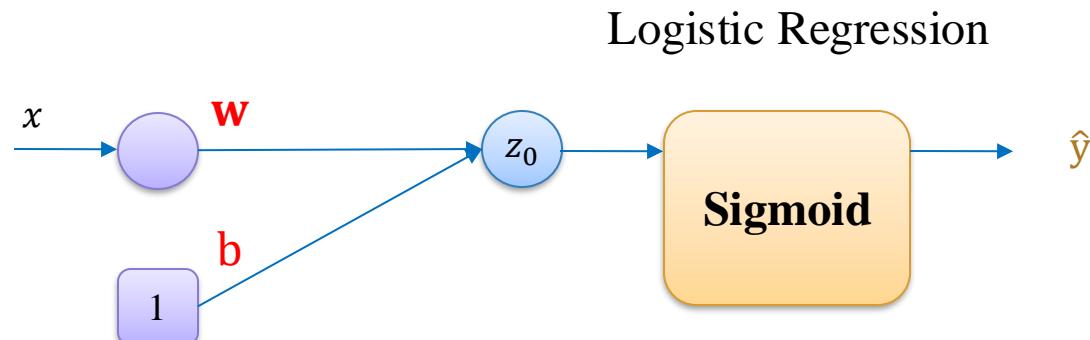
# Softmax Regression



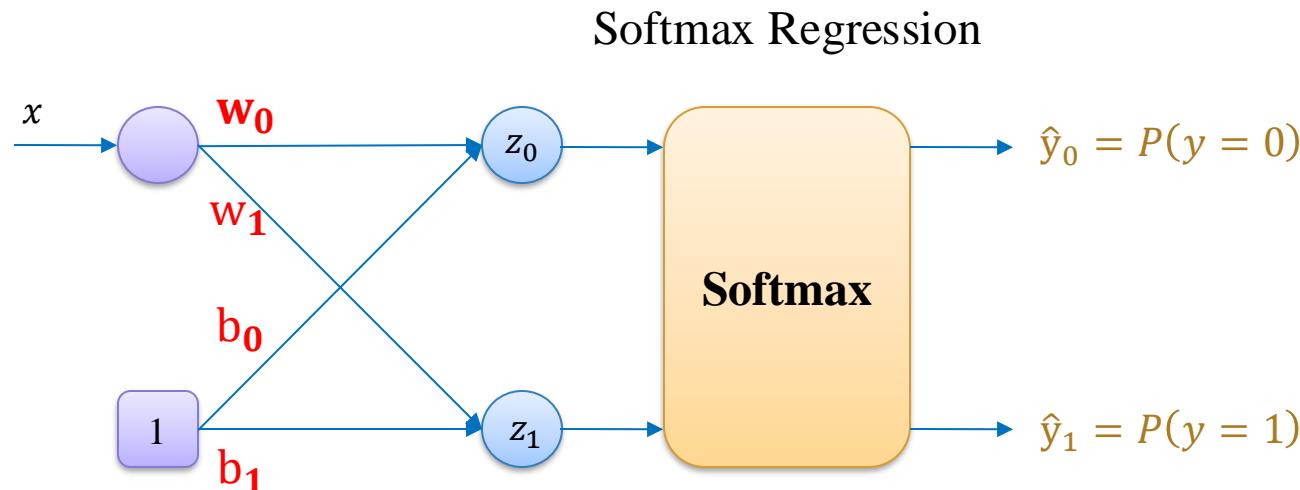
## Parameters

Classes: {0, 1}  
Binary Classification

Hours	Pass
0.5	0
2.0	1



#feature: 1  
#class: 2



# Softmax Regression



## One-hot Encoding

$$\mathbf{y} = \begin{bmatrix} y_0 \\ \dots \\ y_C \end{bmatrix} \quad y_i \in \{0,1\}$$

$$\sum_i y_i = 1 \quad C = \# \text{classes}$$

Classes: {0, 1}  
Binary Classification

Hours	Pass
0.5	0
2.0	1

#feature: 1  
#class: 2

$$y = 0 \rightarrow \mathbf{y} = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

$$y = 1 \rightarrow \mathbf{y} = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

Classes: {0, 1, 2}  
Multi-class Classification

Hours	Score
0.5	0
1.5	1
3.0	2

#feature: 1  
#class: 3

$$y = 0 \rightarrow \mathbf{y} = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}$$

$$y = 1 \rightarrow \mathbf{y} = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}$$

$$y = 2 \rightarrow \mathbf{y} = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$$

# Softmax Regression

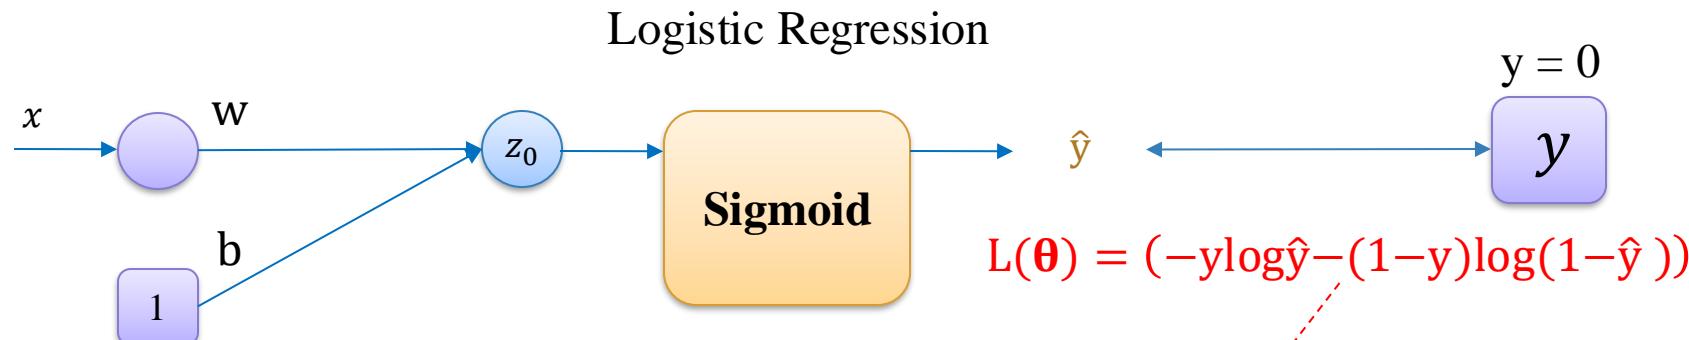


## Loss Function

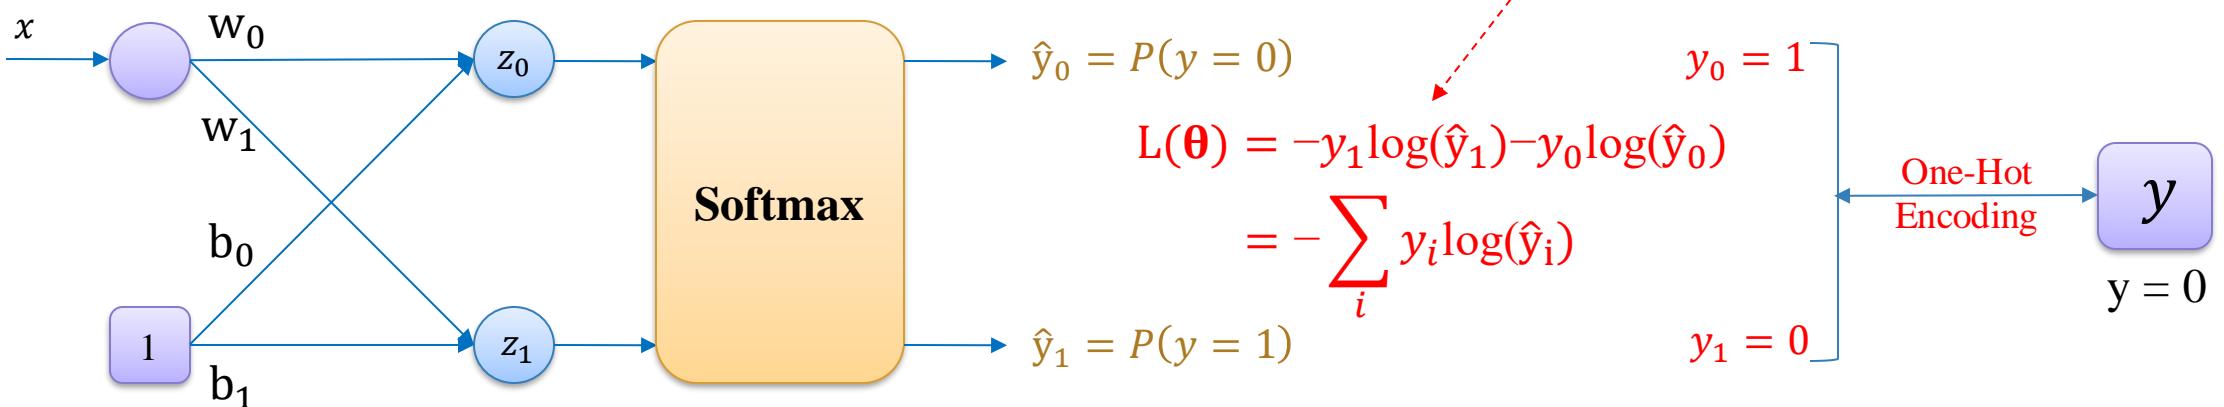
Classes: {0, 1}  
Binary Classification

Hours	Pass
0.5	0
2.0	1

#feature: 1  
#class: 2



Softmax Regression



# Softmax Regression



## Softmax Regression using Gradient Descent

- 1) Pick a sample from training data
- 2) Compute output  $\hat{y}$

$$\mathbf{z} = \boldsymbol{\theta}^T \mathbf{x}$$

$$\mathbf{d} = [1 \dots 1] \mathbf{e}^{\mathbf{z}}$$

$\emptyset$  is  
Hadamard  
division

$$\hat{y} = \mathbf{e}^{\mathbf{z}} \emptyset \mathbf{d}$$

- 3) Compute loss (cross-entropy)

$$L(\boldsymbol{\theta}) = -\mathbf{y}^T \log \hat{\mathbf{y}}$$

- 4) Compute derivative

$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{\mathbf{y}} - \mathbf{y})^T$$

- 5) Update parameters

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L$$

$\eta$  is learning rate

$$\mathbf{x} = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix} \quad \longleftarrow$$

$$\mathbf{y} = [0]$$

One-hot encoding for label

$$\mathbf{y} = 0 \rightarrow \mathbf{y} = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

$$\mathbf{y} = 1 \rightarrow \mathbf{y} = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

$$\boldsymbol{\theta} = \begin{bmatrix} b_0 & b_1 \\ w_0 & w_1 \end{bmatrix} = \begin{bmatrix} 0.1 & 0.2 \\ 0.3 & 0.4 \end{bmatrix}$$

$$\eta = 0.01$$

Hours	Pass
0.5	0
1.0	0
1.5	1
2.0	1

# Softmax Regression



## Softmax Regression using Gradient Descent

- 1) Pick a sample from training data
- 2) Compute output  $\hat{y}$

$$\mathbf{z} = \boldsymbol{\theta}^T \mathbf{x}$$

$$\mathbf{d} = [1 \dots 1]e^{\mathbf{z}}$$

$$\hat{y} = e^{\mathbf{z}} \emptyset \mathbf{d}$$

$\emptyset$  is  
Hadamard  
division

- 3) Compute loss (cross-entropy)

$$L(\boldsymbol{\theta}) = -\mathbf{y}^T \log \hat{\mathbf{y}}$$

- 4) Compute derivative

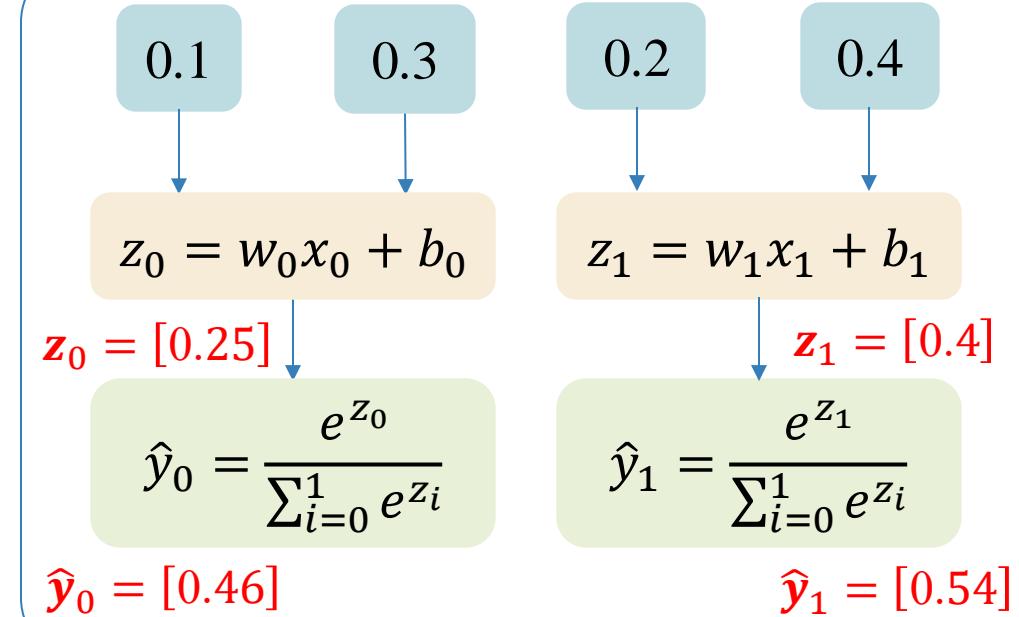
$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{\mathbf{y}} - \mathbf{y})^T$$

- 5) Update parameters

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L$$

$\eta$  is learning rate

$$\eta = 0.01 \quad \mathbf{x} = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix} \quad \mathbf{y} = 0 \rightarrow \mathbf{y} = \begin{bmatrix} 1 \\ 0 \end{bmatrix} \quad \boldsymbol{\theta} = \begin{bmatrix} 0.1 & 0.2 \\ 0.3 & 0.4 \end{bmatrix}$$



# Softmax Regression



## Softmax Regression using Gradient Descent

- 1) Pick a sample from training data
- 2) Compute output  $\hat{y}$

$$\mathbf{z} = \boldsymbol{\theta}^T \mathbf{x}$$

$$\mathbf{d} = [1 \dots 1]e^{\mathbf{z}}$$

$$\hat{y} = e^{\mathbf{z}} \emptyset \mathbf{d}$$

$\emptyset$  is  
Hadamard  
division

- 3) Compute loss (cross-entropy)

$$L(\boldsymbol{\theta}) = -\mathbf{y}^T \log \hat{\mathbf{y}}$$

- 4) Compute derivative

$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{\mathbf{y}} - \mathbf{y})^T$$

- 5) Update parameters

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L$$

$\eta$  is learning rate

$$\eta = 0.01$$

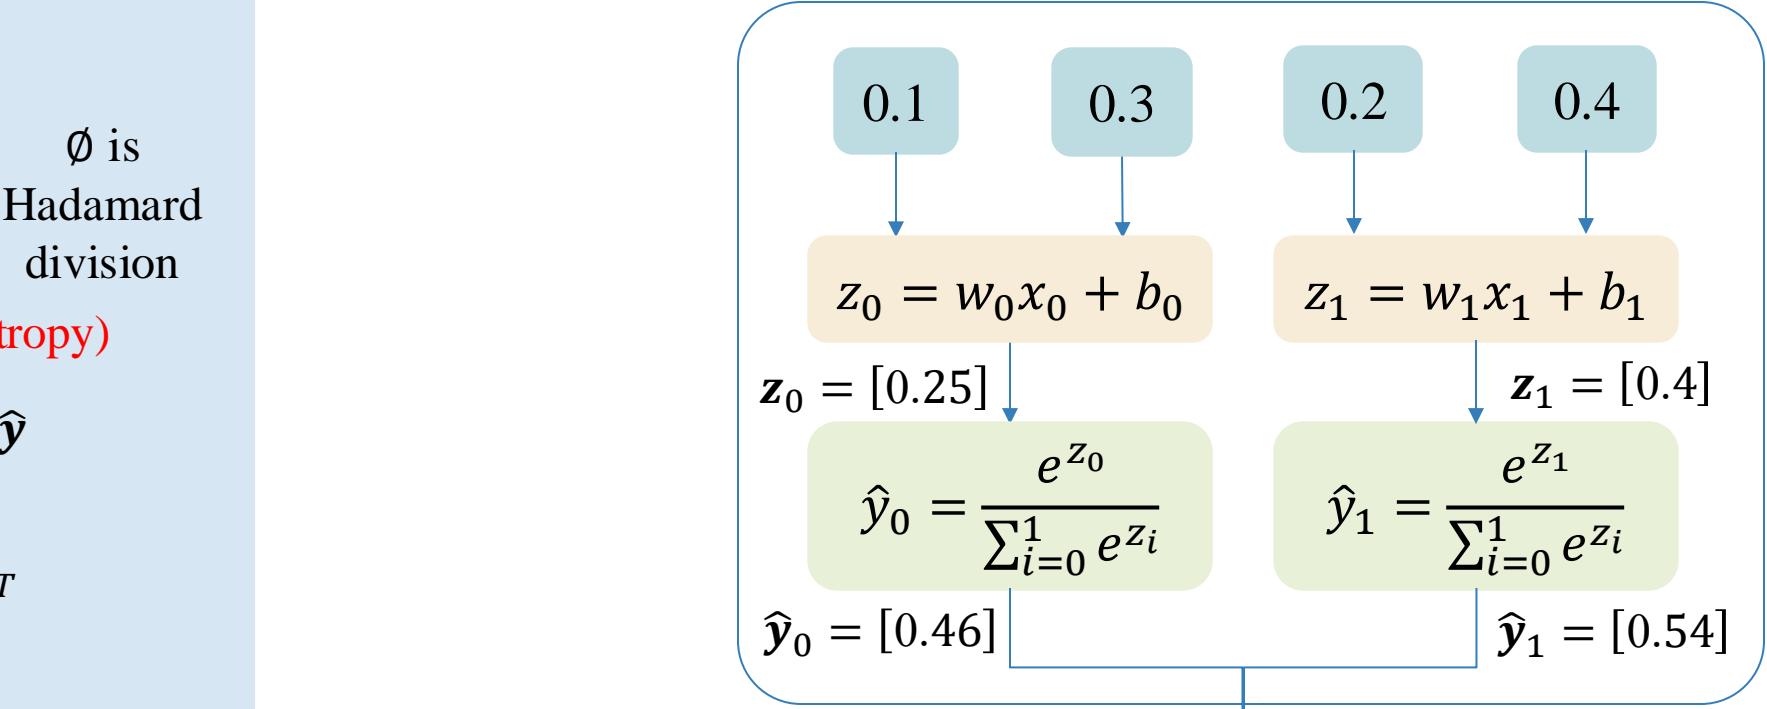
$$\mathbf{x} = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix}$$

$$\mathbf{y} = 0 \rightarrow \mathbf{y} = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

$$\boldsymbol{\theta} = \begin{bmatrix} 0.1 & 0.2 \\ 0.3 & 0.4 \end{bmatrix}$$

$$L = [0.77]$$

$$L = -y_0 \log \hat{y}_0 - y_1 \log \hat{y}_1$$



# Softmax Regression



## Softmax Regression using Gradient Descent

- 1) Pick a sample from training data
- 2) Compute output  $\hat{y}$

$$\mathbf{z} = \boldsymbol{\theta}^T \mathbf{x}$$

$$\mathbf{d} = [1 \dots 1]e^{\mathbf{z}}$$

$$\hat{y} = e^{\mathbf{z}} \emptyset \mathbf{d}$$

$\emptyset$  is  
Hadamard  
division

- 3) Compute loss (cross-entropy)

$$L(\boldsymbol{\theta}) = -\mathbf{y}^T \log \hat{\mathbf{y}}$$

- 4) Compute derivative

$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{\mathbf{y}} - \mathbf{y})^T$$

- 5) Update parameters

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L$$

$\eta$  is learning rate

$$\eta = 0.01$$

$$\mathbf{x} = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix}$$

$$\mathbf{y} = 0 \rightarrow \mathbf{y} = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

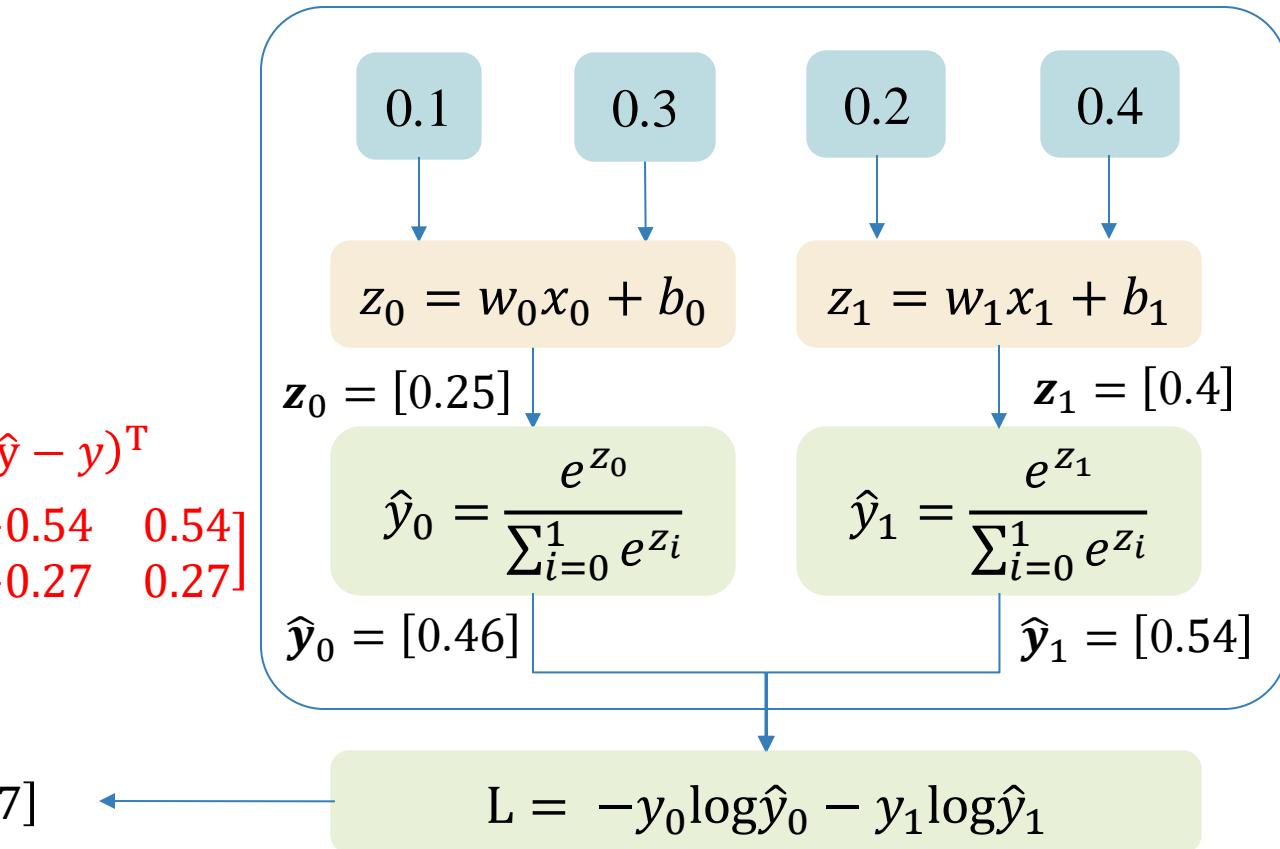
$$\boldsymbol{\theta} = \begin{bmatrix} 0.1 & 0.2 \\ 0.3 & 0.4 \end{bmatrix}$$

$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{\mathbf{y}} - \mathbf{y})^T$$

$$= \begin{bmatrix} -0.54 & 0.54 \\ -0.27 & 0.27 \end{bmatrix}$$

$$L = [0.77]$$

$$L = -y_0 \log \hat{y}_0 - y_1 \log \hat{y}_1$$



# Softmax Regression



## Softmax Regression using Gradient Descent

- 1) Pick a sample from training data
- 2) Compute output  $\hat{y}$

$$\mathbf{z} = \boldsymbol{\theta}^T \mathbf{x}$$

$$\mathbf{d} = [1 \dots 1]e^{\mathbf{z}}$$

$$\hat{y} = e^{\mathbf{z}} \emptyset \mathbf{d}$$

$\emptyset$  is Hadamard division

- 3) Compute loss (cross-entropy)

$$L(\boldsymbol{\theta}) = -\mathbf{y}^T \log \hat{\mathbf{y}}$$

- 4) Compute derivative

$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{\mathbf{y}} - \mathbf{y})^T$$

- 5) Update parameters

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L$$

$\eta$  is learning rate

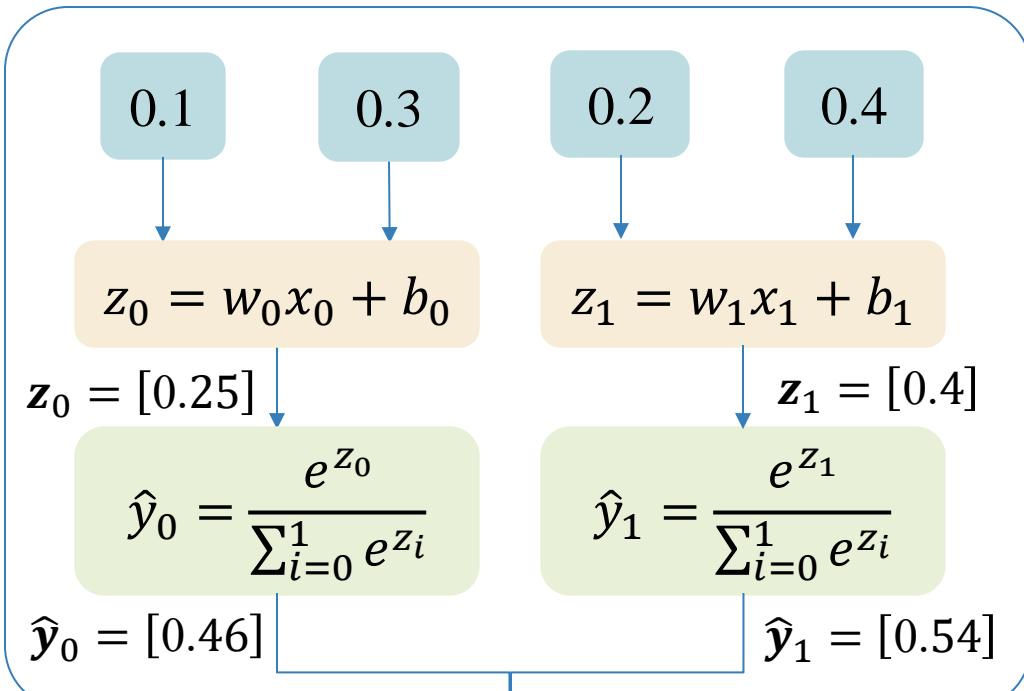
$$\eta = 0.01 \quad \mathbf{x} = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix}$$

$$\mathbf{y} = 0 \rightarrow \mathbf{y} = \begin{bmatrix} 1 \\ 0 \end{bmatrix} \quad \boldsymbol{\theta} = \begin{bmatrix} 0.1 & 0.2 \\ 0.3 & 0.4 \end{bmatrix}$$

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L \\ = \begin{bmatrix} 0.105 & 0.195 \\ 0.303 & 0.397 \end{bmatrix}$$

$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{\mathbf{y}} - \mathbf{y})^T \\ = \begin{bmatrix} -0.54 & 0.54 \\ -0.27 & 0.27 \end{bmatrix}$$

$$L = [0.77]$$



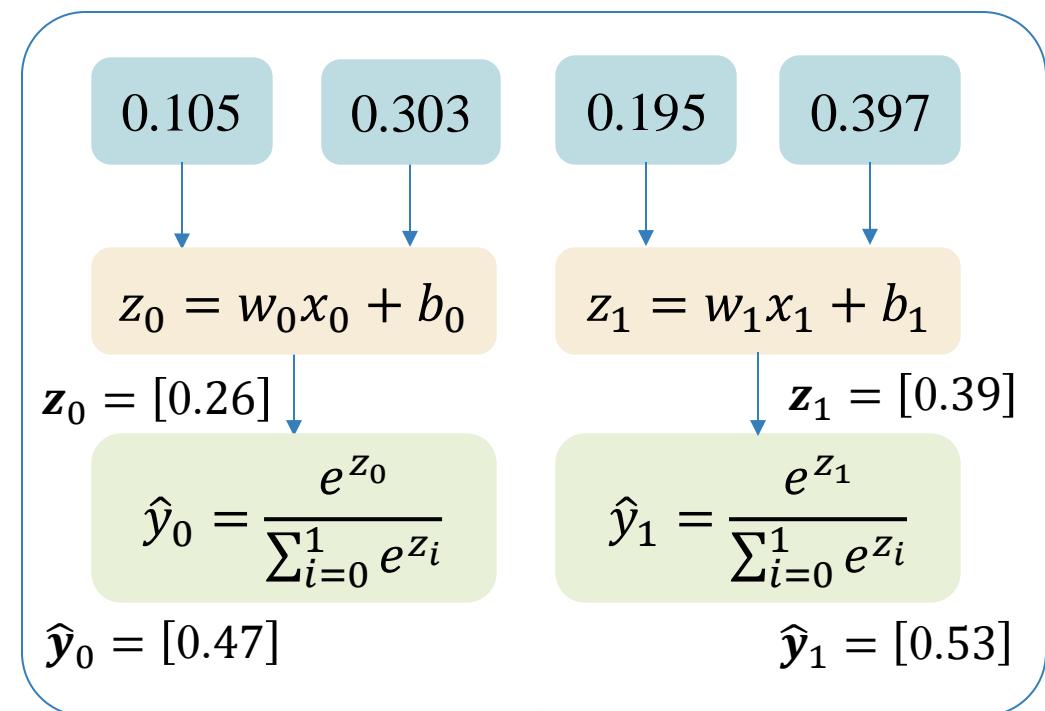
$$L = -y_0 \log \hat{y}_0 - y_1 \log \hat{y}_1$$

# Softmax Regression



## Softmax Regression using Gradient Descent

Hours	Pass
0.5	?

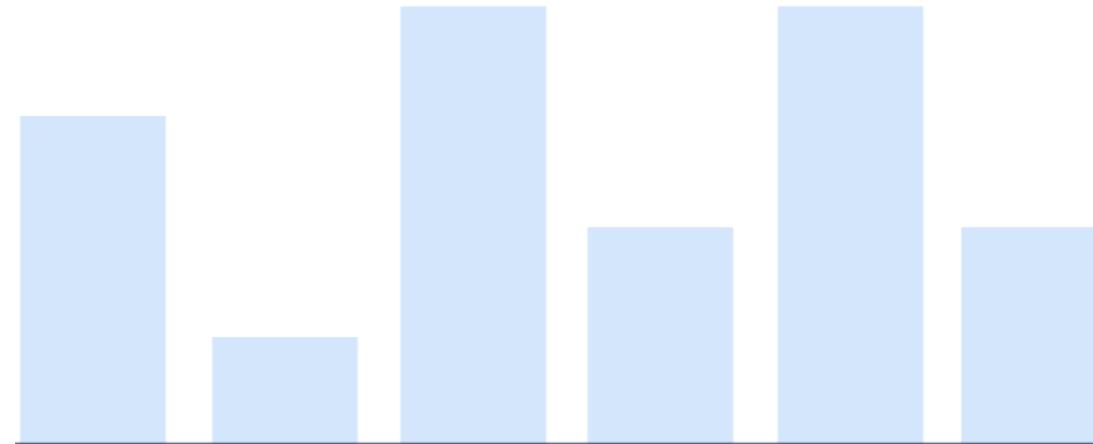


$$y_{pred} = 1$$

# Outline

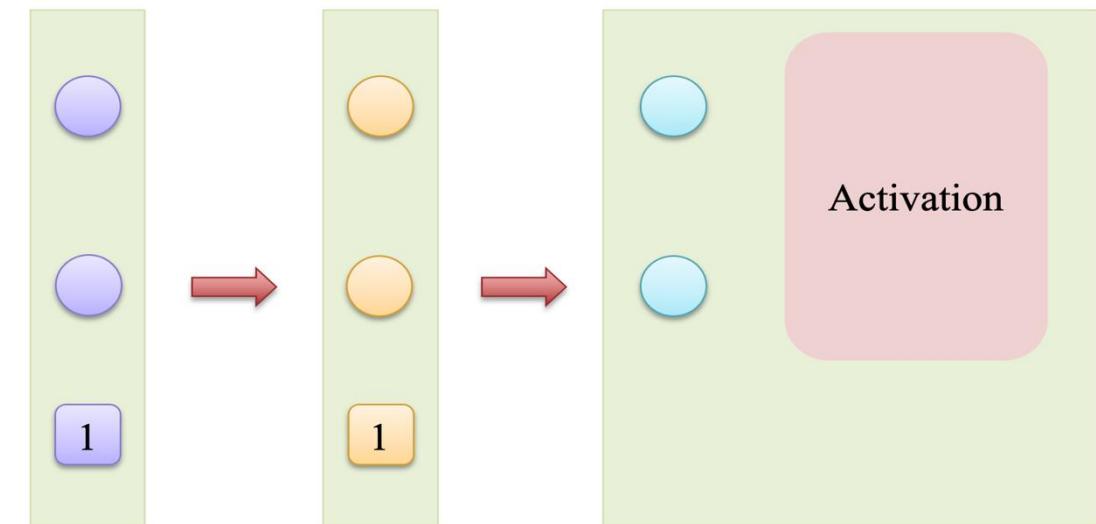
SECTION 1

## Softmax Regression



SECTION 2

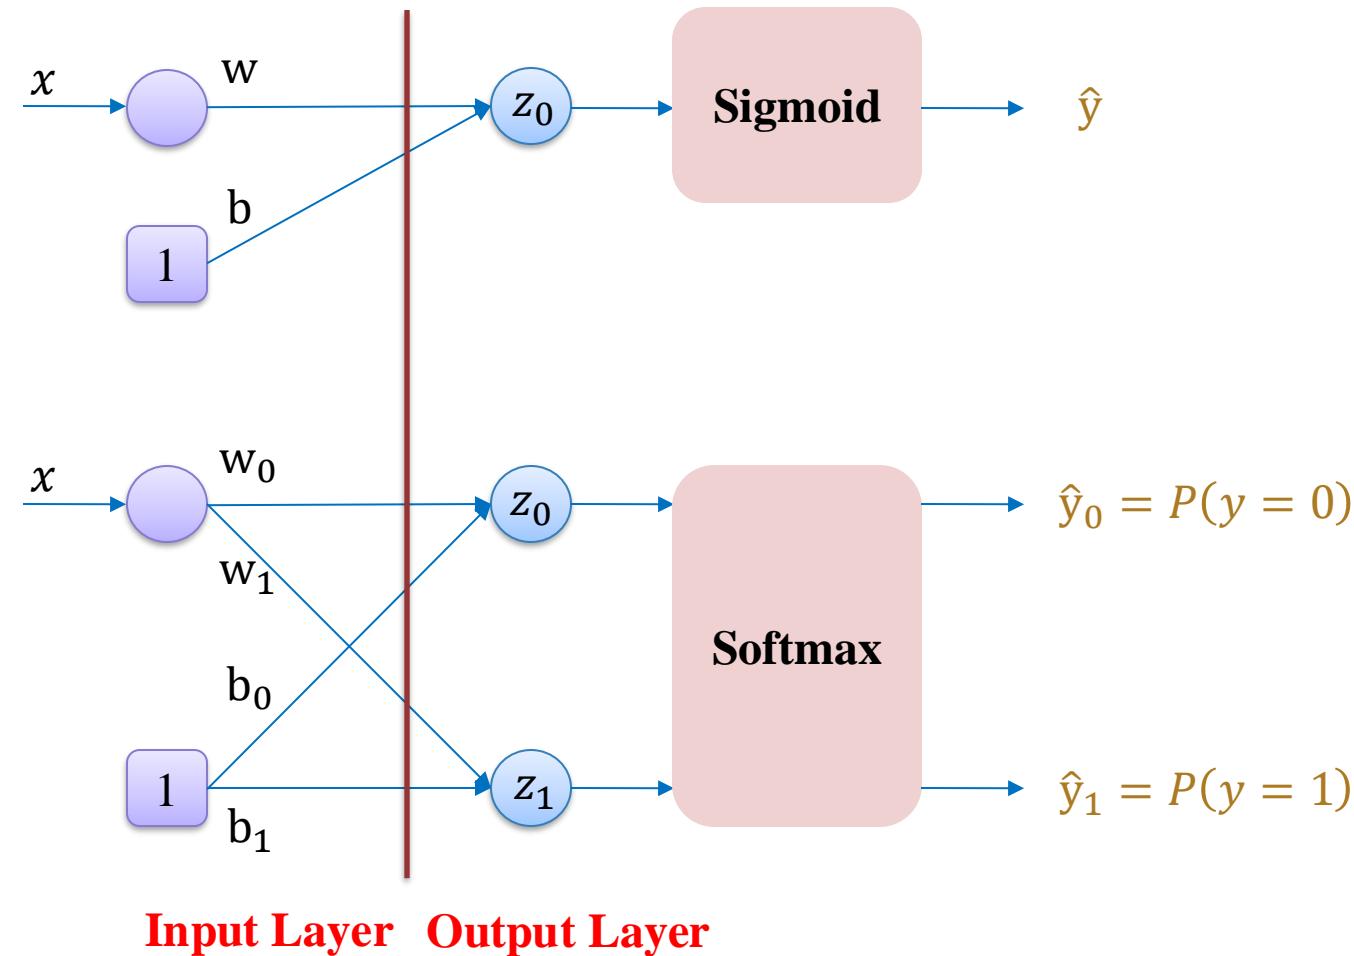
## Multi-layer Perceptron



# Multi-layer Perceptron



## Motivation

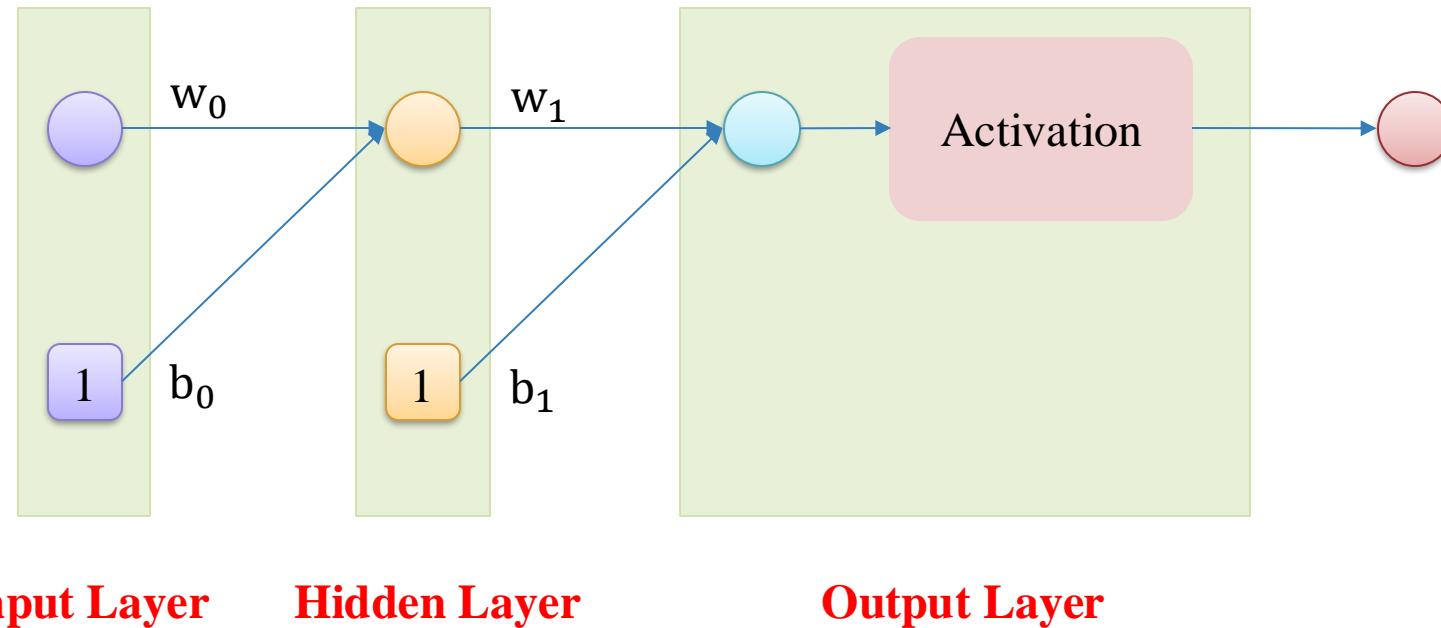


# Multi-layer Perceptron



## Hidden Layer

#parameters: 4

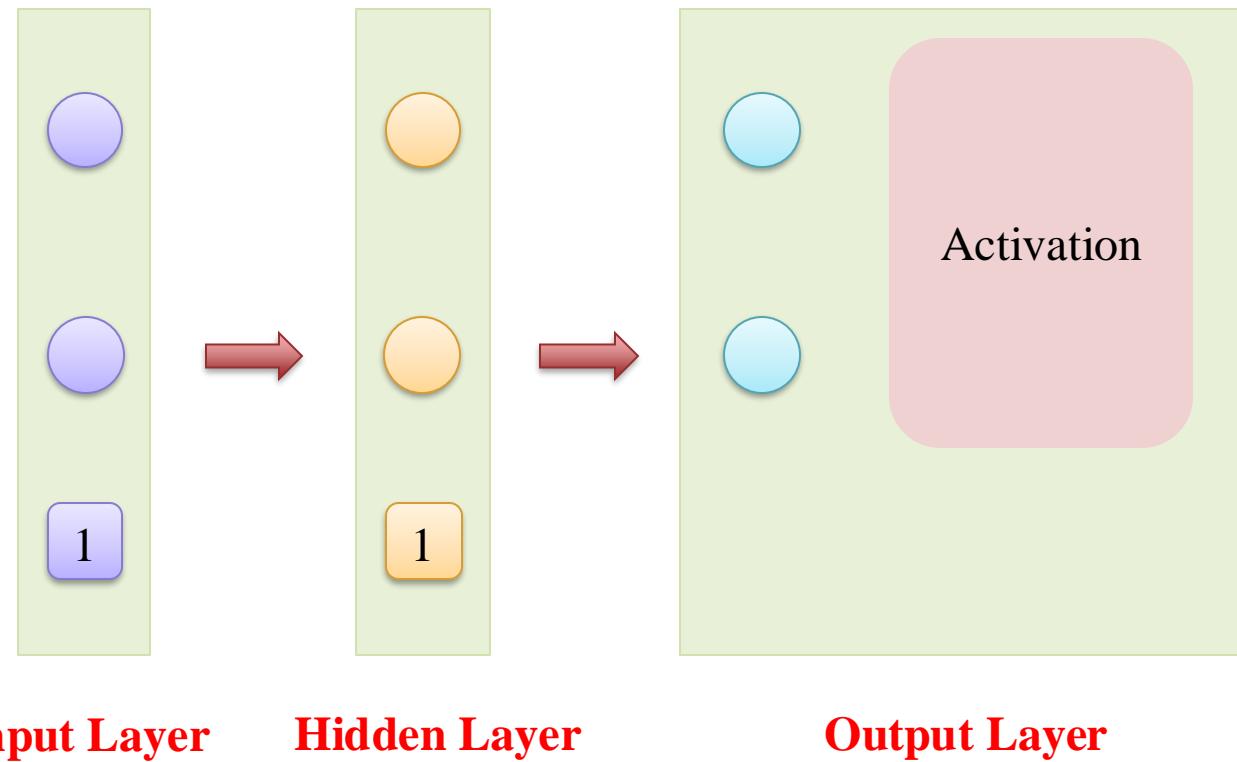


# Multi-layer Perceptron



## Multi-layer Perceptron

#parameters: 12

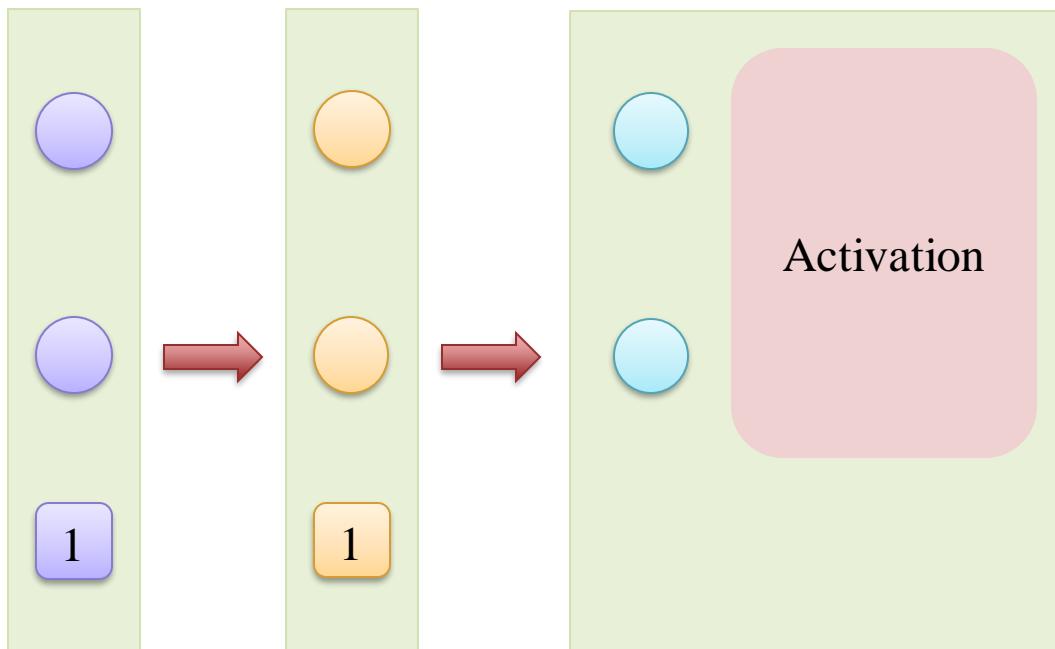


# Multi-layer Perceptron



## Multi-layer Perceptron using Pytorch

#parameters: 12



```
1 model = nn.Sequential(  
2     nn.Linear(2, 2),  
3     nn.Linear(2, 2),  
4     nn.Sigmoid()  
5 )
```

```
1 summary(model, (10000000, 2))
```

Layer (type)	Output Shape	Param #
Linear-1	[-1, 10000000, 2]	6
Linear-2	[-1, 10000000, 2]	6
Sigmoid-3	[-1, 10000000, 2]	0

Total params: 12

Trainable params: 12

Non-trainable params: 0

Input size (MB): 76.29

Forward/backward pass size (MB): 457.76

Params size (MB): 0.00

Estimated Total Size (MB): 534.06

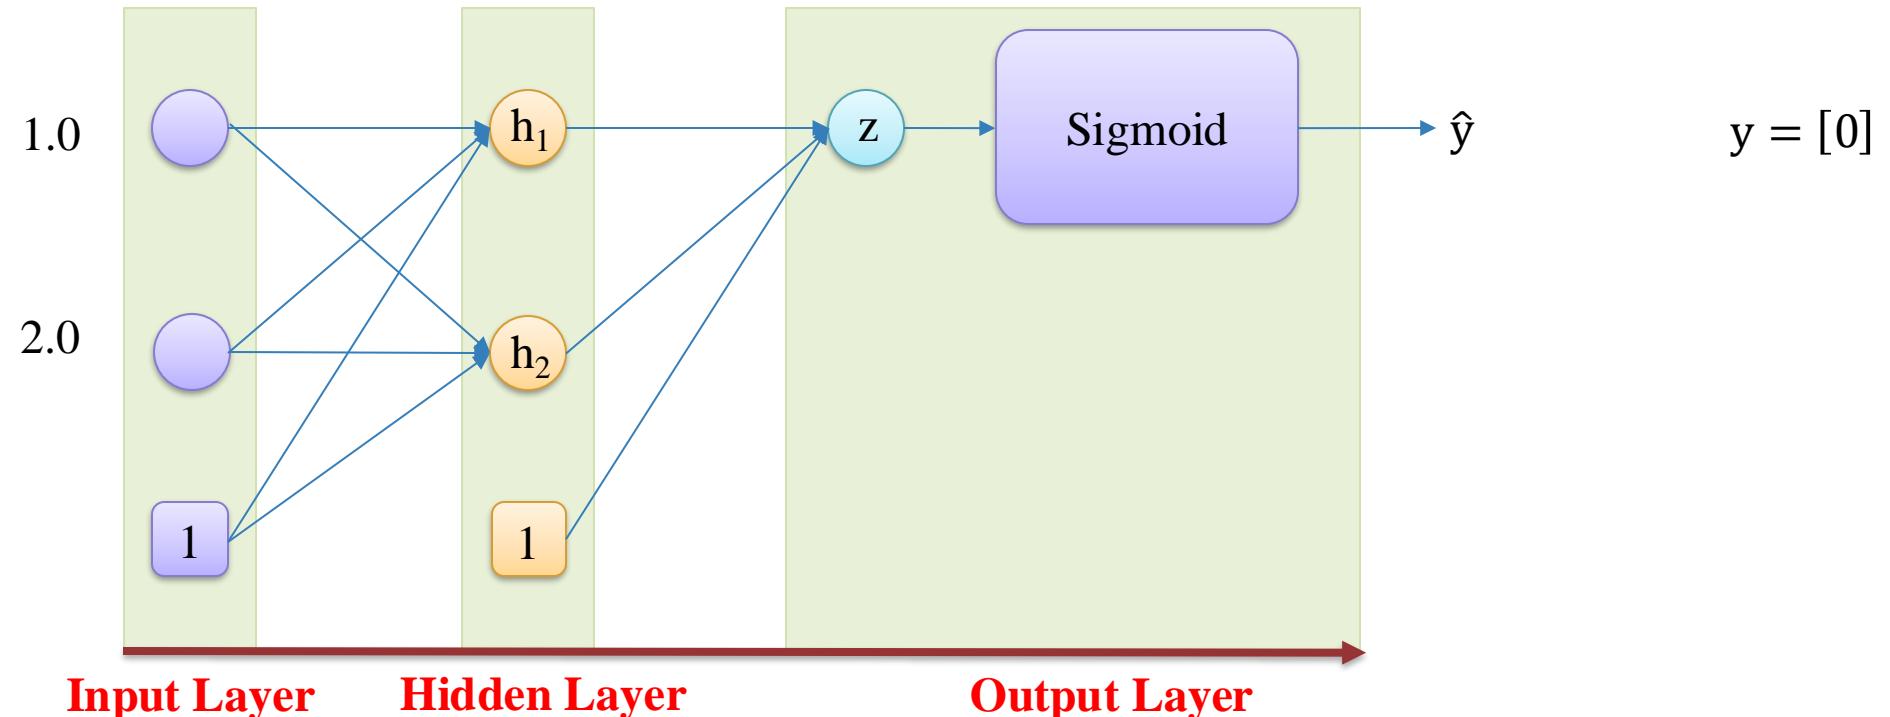
# Multi-layer Perceptron

! Forward

$$\mathbf{x} = [1.0 \quad 2.0]$$

$$\mathbf{W}_h = [W_{h1} \quad W_{h2}] \\ = \begin{bmatrix} 0.1 & 0.1 \\ 0.1 & 0.1 \\ 0.1 & 0.1 \end{bmatrix}$$

$$\mathbf{W}_z = [W_z] = \begin{bmatrix} 0.1 \\ 0.1 \\ 0.1 \end{bmatrix}$$



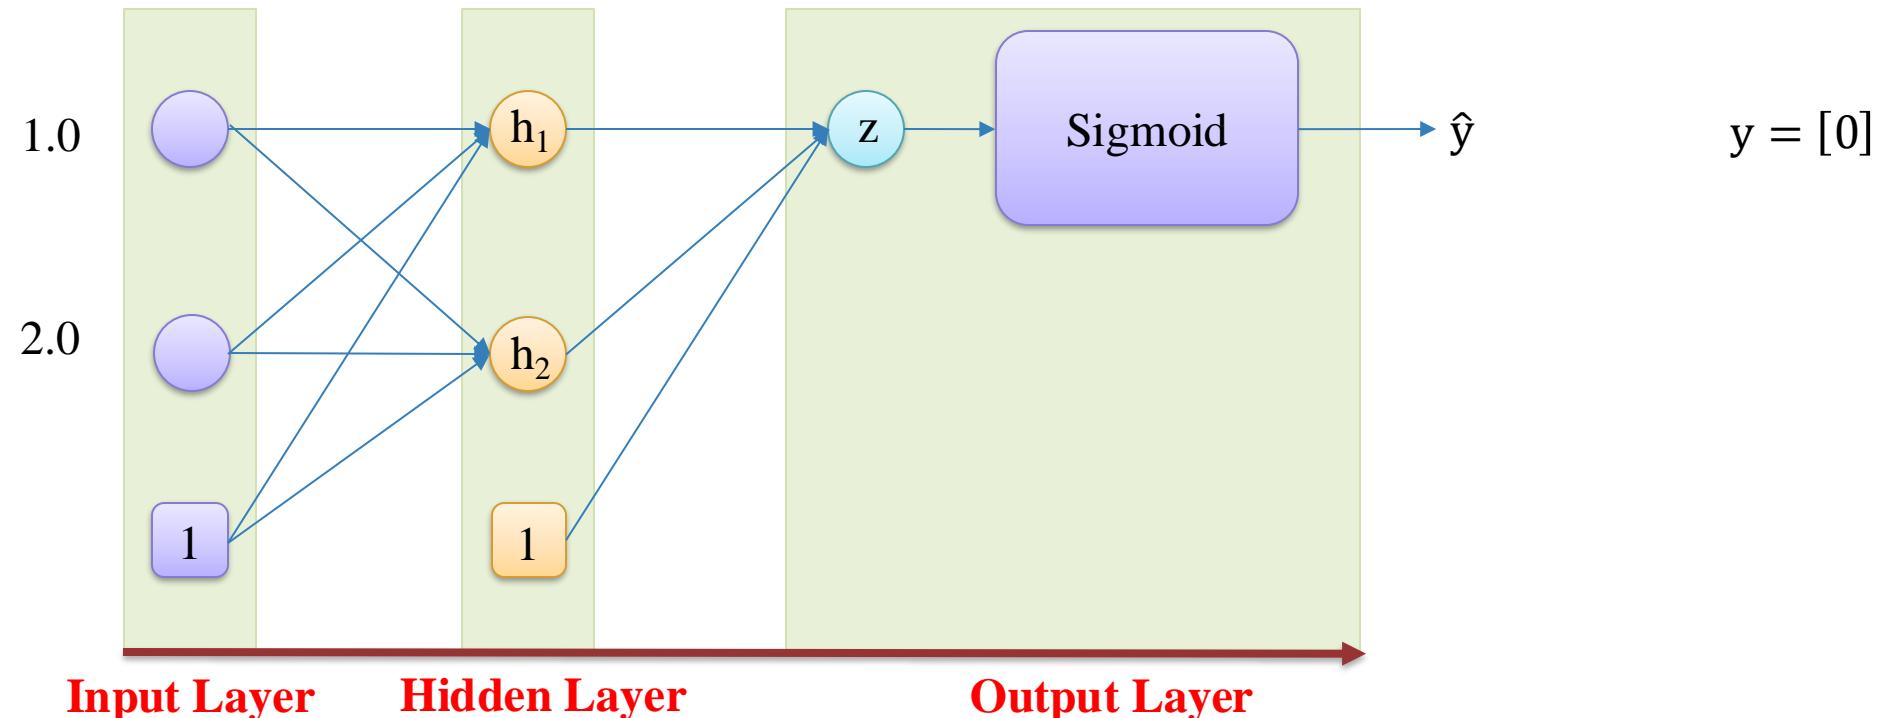
# Multi-layer Perceptron

! Forward

$$\mathbf{x} = [1.0 \quad 2.0]$$

$$\mathbf{W}_h = [W_{h1} \quad W_{h2}] \\ = \begin{bmatrix} 0.1 & 0.1 \\ 0.1 & 0.1 \\ 0.1 & 0.1 \end{bmatrix}$$

$$\mathbf{W}_z = [W_z] = \begin{bmatrix} 0.1 \\ 0.1 \\ 0.1 \end{bmatrix}$$



$$h = [1.0 \quad \mathbf{x}] \mathbf{W}_h = [1.0 \quad 1.0 \quad 2.0] \begin{bmatrix} 0.1 & 0.1 \\ 0.1 & 0.1 \\ 0.1 & 0.1 \end{bmatrix} = [0.4 \quad 0.4]$$

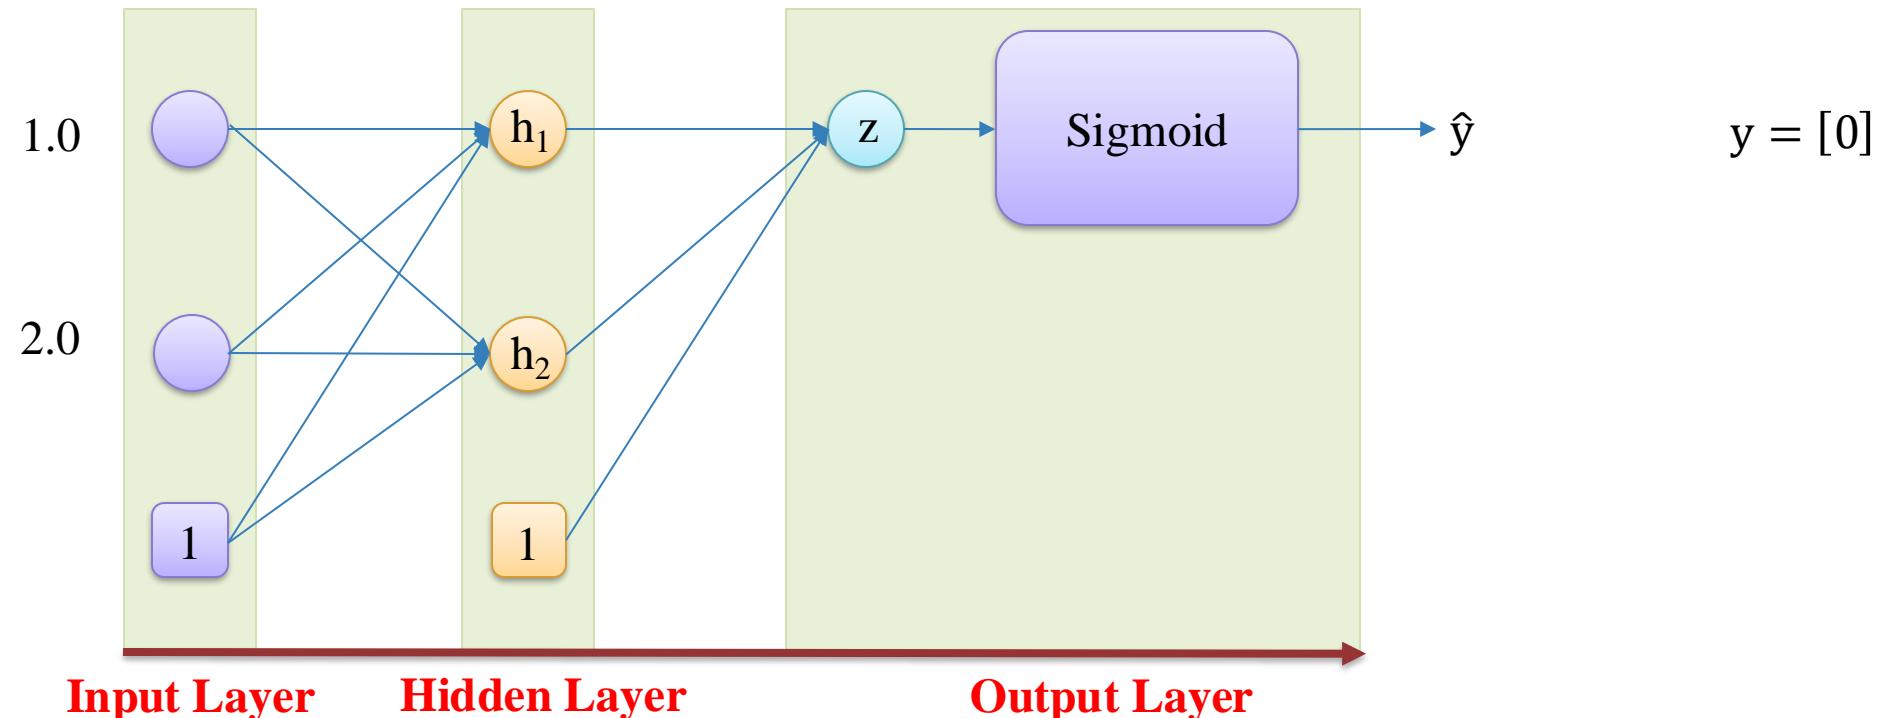
# Multi-layer Perceptron

! Forward

$$\mathbf{x} = [1.0 \quad 2.0]$$

$$\mathbf{W}_h = [W_{h1} \quad W_{h2}] \\ = \begin{bmatrix} 0.1 & 0.1 \\ 0.1 & 0.1 \\ 0.1 & 0.1 \end{bmatrix}$$

$$\mathbf{W}_z = [W_z] = \begin{bmatrix} 0.1 \\ 0.1 \\ 0.1 \end{bmatrix}$$



$$h = [1.0 \mathbf{x}] \mathbf{W}_h = [1.0 \quad 1.0 \quad 2.0] \begin{bmatrix} 0.1 & 0.1 \\ 0.1 & 0.1 \\ 0.1 & 0.1 \end{bmatrix} = [0.4 \quad 0.4]$$

$$z = [1.0 \mathbf{h}] \mathbf{W}_z = [1.0 \quad 0.4 \quad 0.4] \begin{bmatrix} 0.1 \\ 0.1 \\ 0.1 \end{bmatrix} = 0.18$$

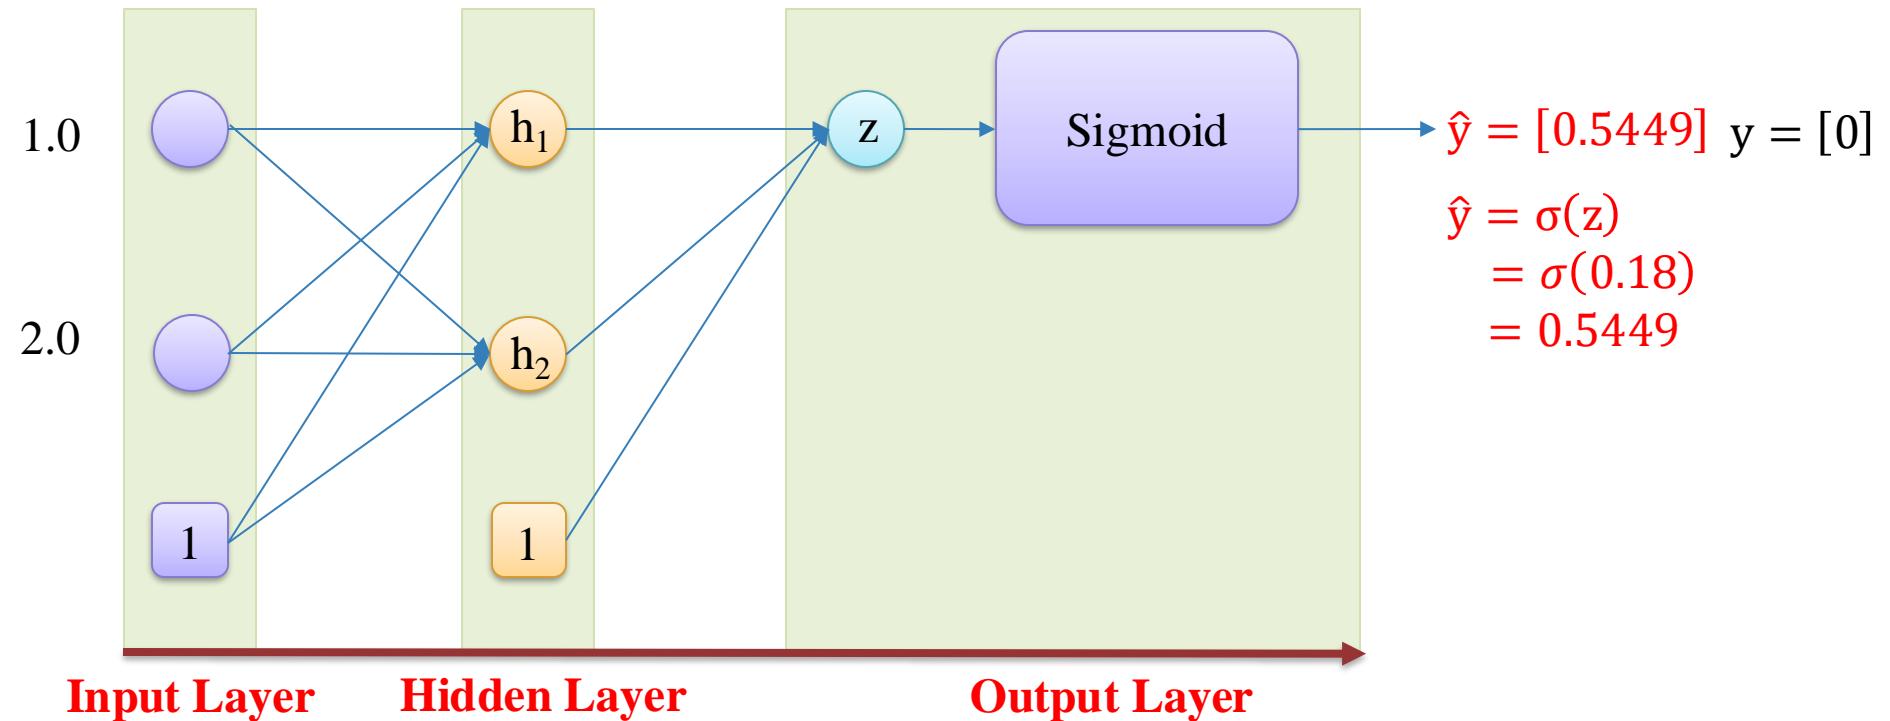
# Multi-layer Perceptron

## ! Forward

$$x = [1.0 \quad 2.0]$$

$$W_h = [W_{h1} \quad W_{h2}] \\ = \begin{bmatrix} 0.1 & 0.1 \\ 0.1 & 0.1 \\ 0.1 & 0.1 \end{bmatrix}$$

$$W_z = [W_z] = \begin{bmatrix} 0.1 \\ 0.1 \\ 0.1 \end{bmatrix}$$



$$h = [1.0 \quad x]W_h = [1.0 \quad 1.0 \quad 2.0] \begin{bmatrix} 0.1 & 0.1 \\ 0.1 & 0.1 \\ 0.1 & 0.1 \end{bmatrix} = [0.4 \quad 0.4] \quad z = [1.0 \quad h]W_z = [1.0 \quad 0.4 \quad 0.4] \begin{bmatrix} 0.1 \\ 0.1 \\ 0.1 \end{bmatrix} = 0.18$$

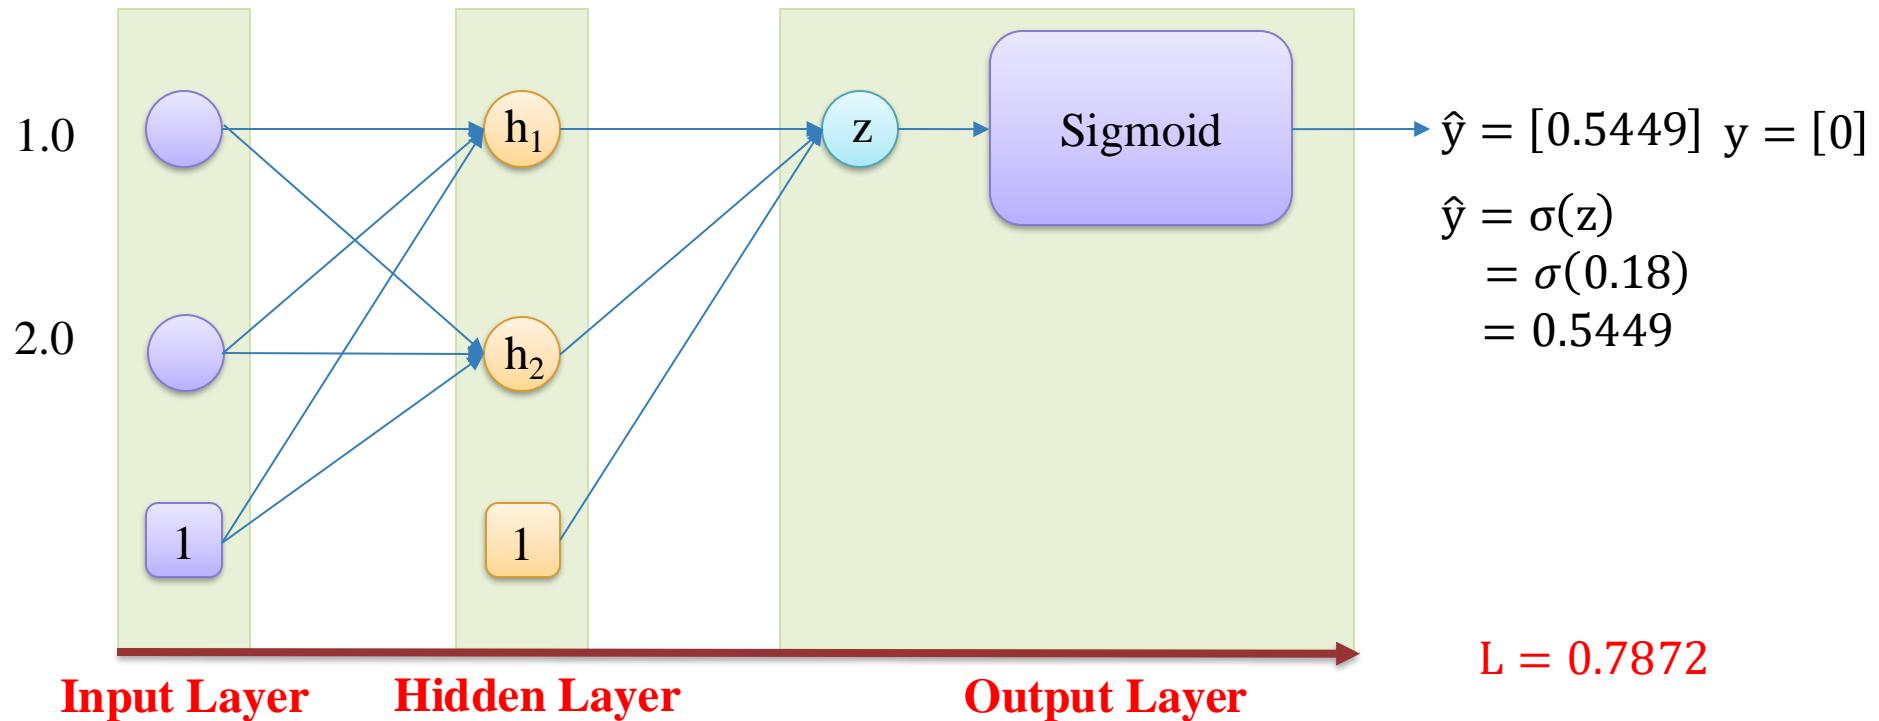
# Multi-layer Perceptron

## ! Forward

$$x = [1.0 \quad 2.0]$$

$$W_h = [W_{h1} \quad W_{h2}] \\ = \begin{bmatrix} 0.1 & 0.1 \\ 0.1 & 0.1 \\ 0.1 & 0.1 \end{bmatrix}$$

$$W_z = [W_z] = \begin{bmatrix} 0.1 \\ 0.1 \\ 0.1 \end{bmatrix}$$



$$h = [1.0 \quad x]W_h = [1.0 \quad 1.0 \quad 2.0] \begin{bmatrix} 0.1 & 0.1 \\ 0.1 & 0.1 \\ 0.1 & 0.1 \end{bmatrix} = [0.4 \quad 0.4]$$
$$z = [1.0 \quad h]W_z = [1.0 \quad 0.4 \quad 0.4] \begin{bmatrix} 0.1 \\ 0.1 \\ 0.1 \end{bmatrix} = 0.18$$

# Multi-layer Perceptron

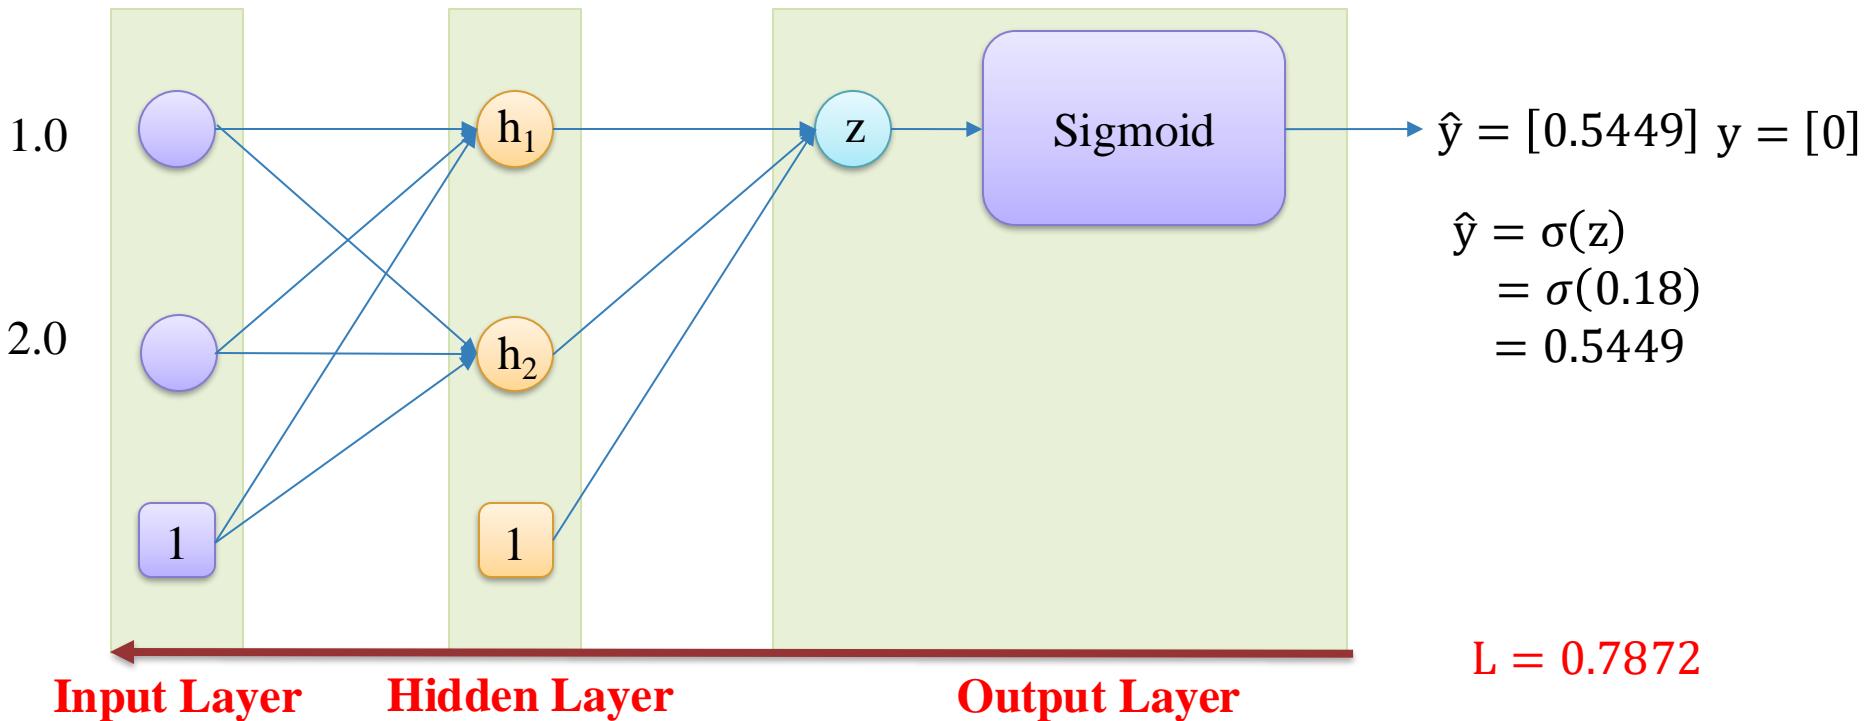


## Backward

$$x = [1.0 \quad 2.0]$$

$$\begin{aligned}W_h &= [W_{h1} \quad W_{h2}] \\&= \begin{bmatrix} 0.0946 & 0.0946 \\ 0.0946 & 0.0891 \\ 0.0946 & 0.0891 \end{bmatrix}\end{aligned}$$

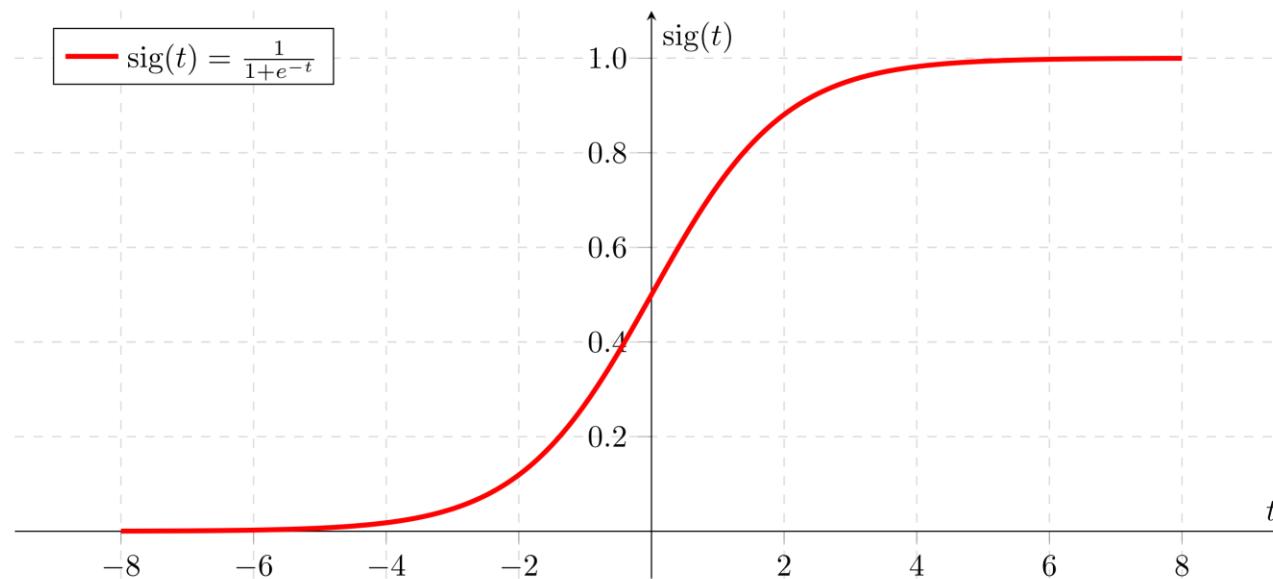
$$W_z = [W_z] = \begin{bmatrix} 0.0455 \\ 0.0782 \\ 0.0782 \end{bmatrix}$$



# Activation



## Sigmoid Function



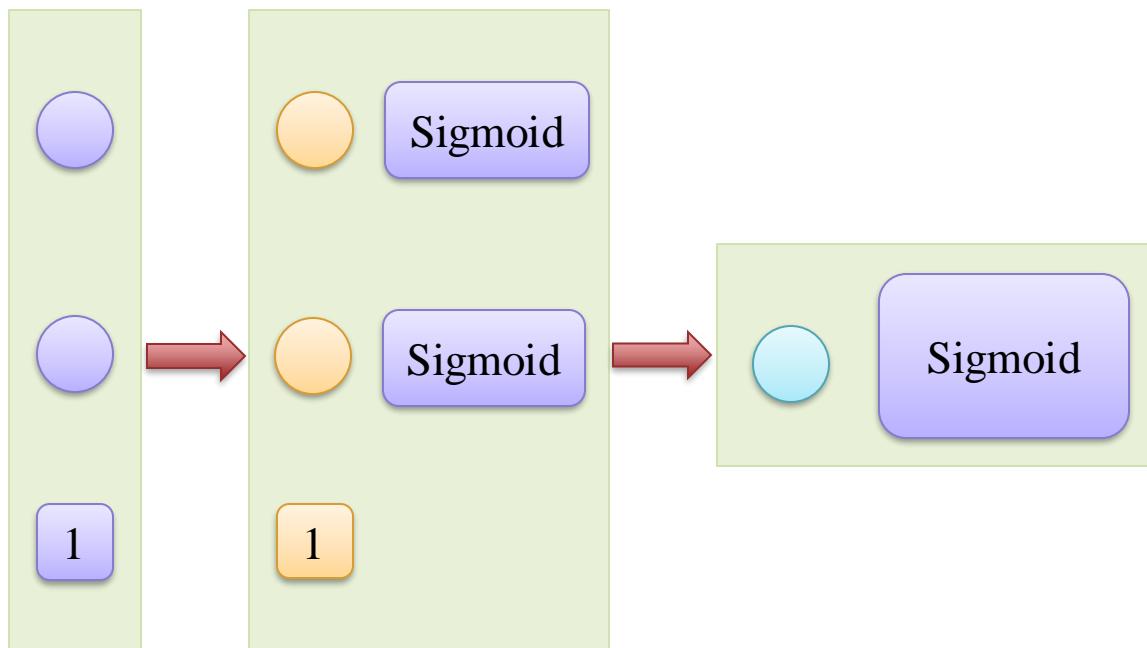
```
1 import torch.nn as nn  
2  
3 act = nn.Sigmoid()  
4 input = torch.tensor([0.18, -0.18])  
5 act(input)
```

tensor([0.5449, 0.4551])

# Activation



## Sigmoid Function



```
1 model = nn.Sequential(  
2     nn.Linear(2, 2),  
3     nn.Sigmoid(),  
4     nn.Linear(2, 1),  
5     nn.Sigmoid()  
6 )
```

```
1 summary(model, (1, 2))
```

Layer (type)	Output Shape	Param #
Linear-1	[-1, 1, 2]	6
Sigmoid-2	[-1, 1, 2]	0
Linear-3	[-1, 1, 1]	3
Sigmoid-4	[-1, 1, 1]	0

Total params: 9

Trainable params: 9

Non-trainable params: 0

Input size (MB): 0.00

Forward/backward pass size (MB): 0.00

Params size (MB): 0.00

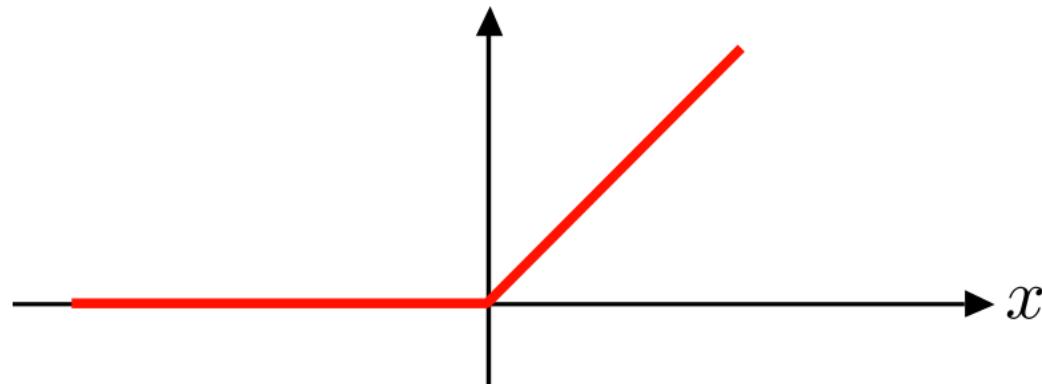
Estimated Total Size (MB): 0.00

# Activation



## ReLU Function

$$\text{ReLU}(x) \triangleq \max(0, x)$$



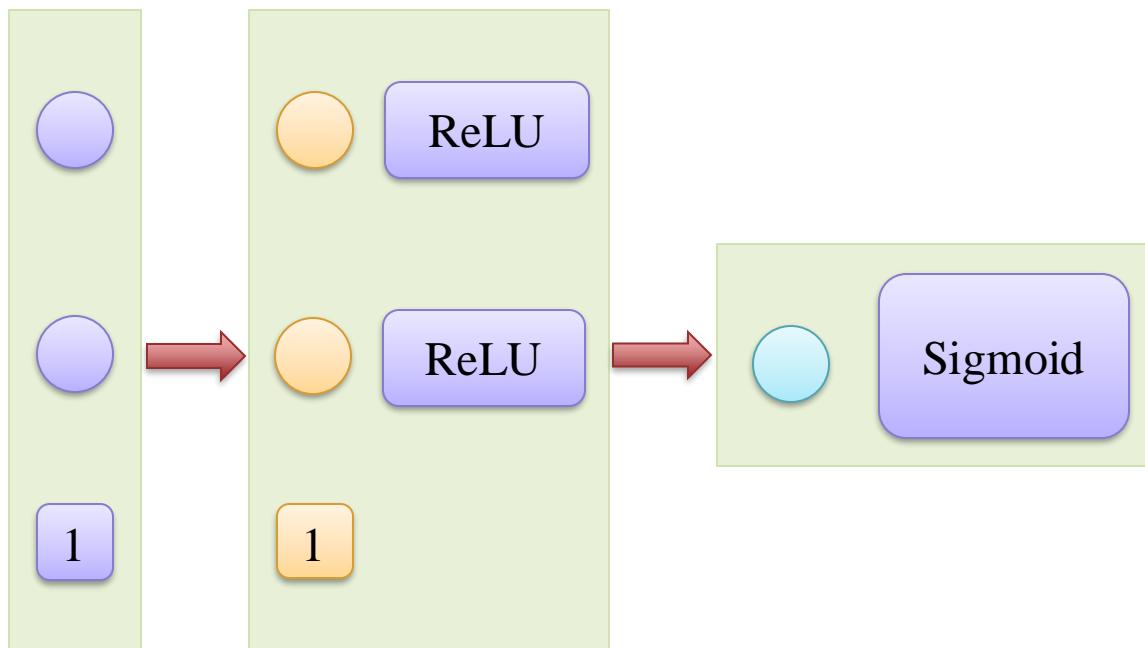
```
1 import torch.nn as nn  
2  
3 act = nn.ReLU()  
4 input = torch.tensor([0.18, -0.18])  
5 act(input)
```

```
tensor([0.1800, 0.0000])
```

# Activation



## ReLU Function



```
1 model = nn.Sequential(  
2     nn.Linear(2, 2),  
3     nn.ReLU(),  
4     nn.Linear(2, 1),  
5     nn.Sigmoid()  
6 )
```

```
1 summary(model, (1, 2))
```

Layer (type)	Output Shape	Param #
Linear-1	[-1, 1, 2]	6
ReLU-2	[-1, 1, 2]	0
Linear-3	[-1, 1, 1]	3
Sigmoid-4	[-1, 1, 1]	0

Total params: 9

Trainable params: 9

Non-trainable params: 0

Input size (MB): 0.00

Forward/backward pass size (MB): 0.00

Params size (MB): 0.00

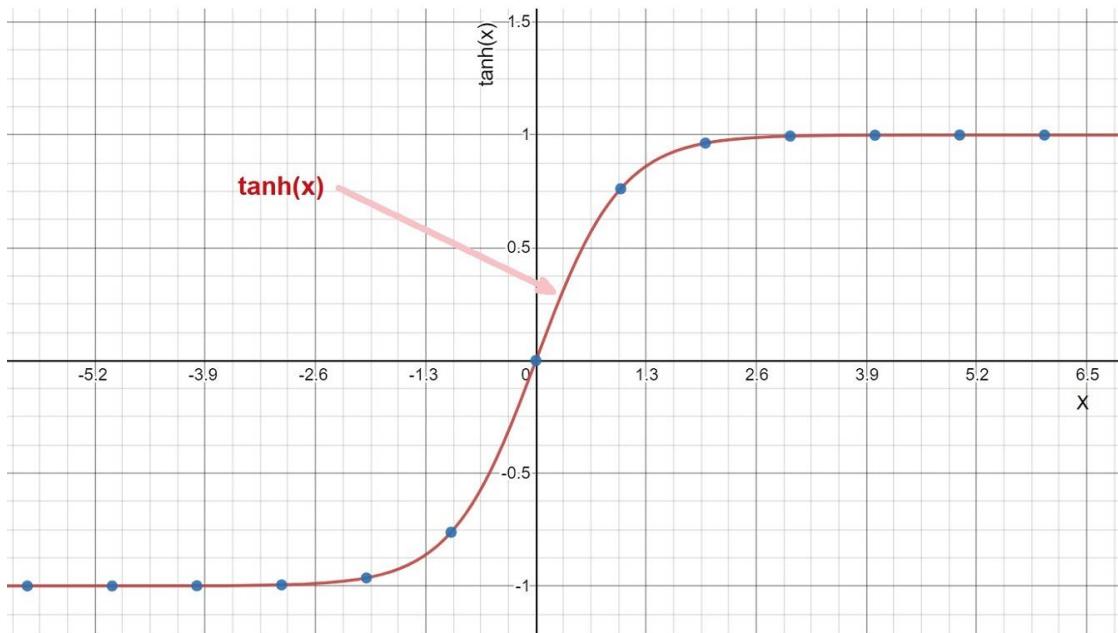
Estimated Total Size (MB): 0.00

# Activation



## Tanh Function

$$\tanh(x) = \frac{(e^x - e^{-x})}{(e^x + e^{-x})}$$



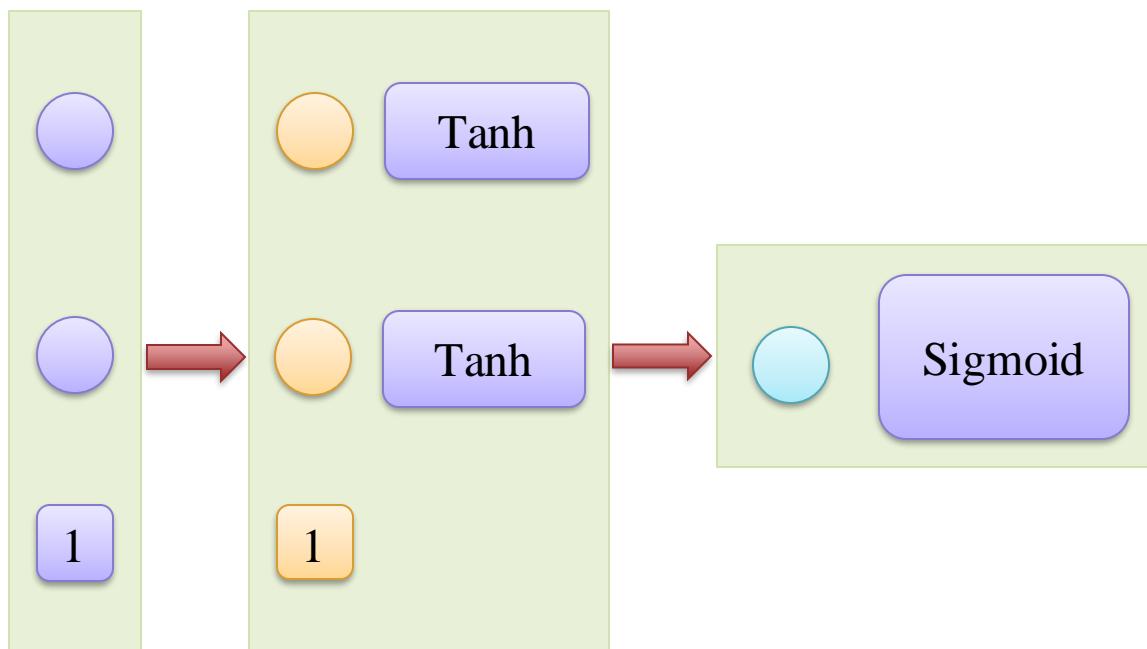
```
1 import torch.nn as nn  
2  
3 act = nn.Tanh()  
4 input = torch.tensor([0.18, -0.18])  
5 act(input)
```

tensor([ 0.1781, -0.1781])

# Activation



## ReLU Function



```
1 model = nn.Sequential(  
2     nn.Linear(2, 2),  
3     nn.Tanh(),  
4     nn.Linear(2, 1),  
5     nn.Sigmoid()  
6 )
```

```
1 summary(model, (1, 2))
```

Layer (type)	Output Shape	Param #
Linear-1	[ -1, 1, 2 ]	6
Tanh-2	[ -1, 1, 2 ]	0
Linear-3	[ -1, 1, 1 ]	3
Sigmoid-4	[ -1, 1, 1 ]	0

Total params: 9

Trainable params: 9

Non-trainable params: 0

Input size (MB): 0.00

Forward/backward pass size (MB): 0.00

Params size (MB): 0.00

Estimated Total Size (MB): 0.00

# Loss Function



## BCELoss

```
1 import torch.nn as nn  
2 loss_fn = nn.BCELoss()
```

```
1 y_pred = model(x)  
2 y_pred
```

```
tensor([0.5449], grad_fn=<SigmoidBackward0>)
```

```
1 y
```

```
tensor([0.])
```

```
1 loss = loss_fn(y_pred, y)  
2 loss
```

```
tensor(0.7872, grad_fn=<BinaryCrossEntropyBackward0>)
```

```
1 model = nn.Sequential(  
2     nn.Linear(2, 2),  
3     nn.Linear(2, 2),  
4     nn.Sigmoid()  
5 )
```

```
1 summary(model, (10000000, 2))
```

Layer (type)	Output Shape	Param #
<hr/>		
Linear-1	[-1, 10000000, 2]	6
Linear-2	[-1, 10000000, 2]	6
Sigmoid-3	[-1, 10000000, 2]	0
<hr/>		

```
Total params: 12
```

```
Trainable params: 12
```

```
Non-trainable params: 0
```

```
Input size (MB): 76.29
```

```
Forward/backward pass size (MB): 457.76
```

```
Params size (MB): 0.00
```

```
Estimated Total Size (MB): 534.06
```

# Loss Function



## CrossEntropyLoss

```
1 import torch.nn as nn
2 loss_fn = nn.CrossEntropyLoss()
```

```
1 y = torch.tensor(0)
2 y
```

```
tensor(0)
```

```
1 y_pred
```

```
tensor([ 0.9005, -0.1670], grad_fn=<ViewBackward0>)
```

```
1 loss_fn(y_pred, y)
```

```
tensor(0.2955, grad_fn=<NllLossBackward0>)
```

```
1 model = nn.Sequential(
2     nn.Linear(2, 2),
3     nn.ReLU(),
4     nn.Linear(2, 2),
5 )
```

```
1 summary(model, (1, 2))
```

Layer (type)	Output Shape	Param #
Linear-1	[ -1, 1, 2 ]	6
ReLU-2	[ -1, 1, 2 ]	0
Linear-3	[ -1, 1, 2 ]	6

```
Total params: 12
```

```
Trainable params: 12
```

```
Non-trainable params: 0
```

```
Input size (MB): 0.00
```

```
Forward/backward pass size (MB): 0.00
```

```
Params size (MB): 0.00
```

```
Estimated Total Size (MB): 0.00
```

# Optimizer



## SGD

```
1 learning_rate = 0.1
2 optimizer = optim.SGD(model.parameters(), learning_rate)
```

```
1 for layer in model.children():
2     print(layer.state_dict())
```

```
OrderedDict([('weight', tensor([[0.1000, 0.1000],
                               [0.1000, 0.1000]])), ('bias', tensor([0.1000, 0.1000]))])
OrderedDict()
OrderedDict([('weight', tensor([[0.1000, 0.1000]])), ('bias', tensor([0.1000]))])
OrderedDict()
```

```
1 loss.backward()
```

```
1 optimizer.step()
```

```
1 for layer in model.children():
2     print(layer.state_dict())
```

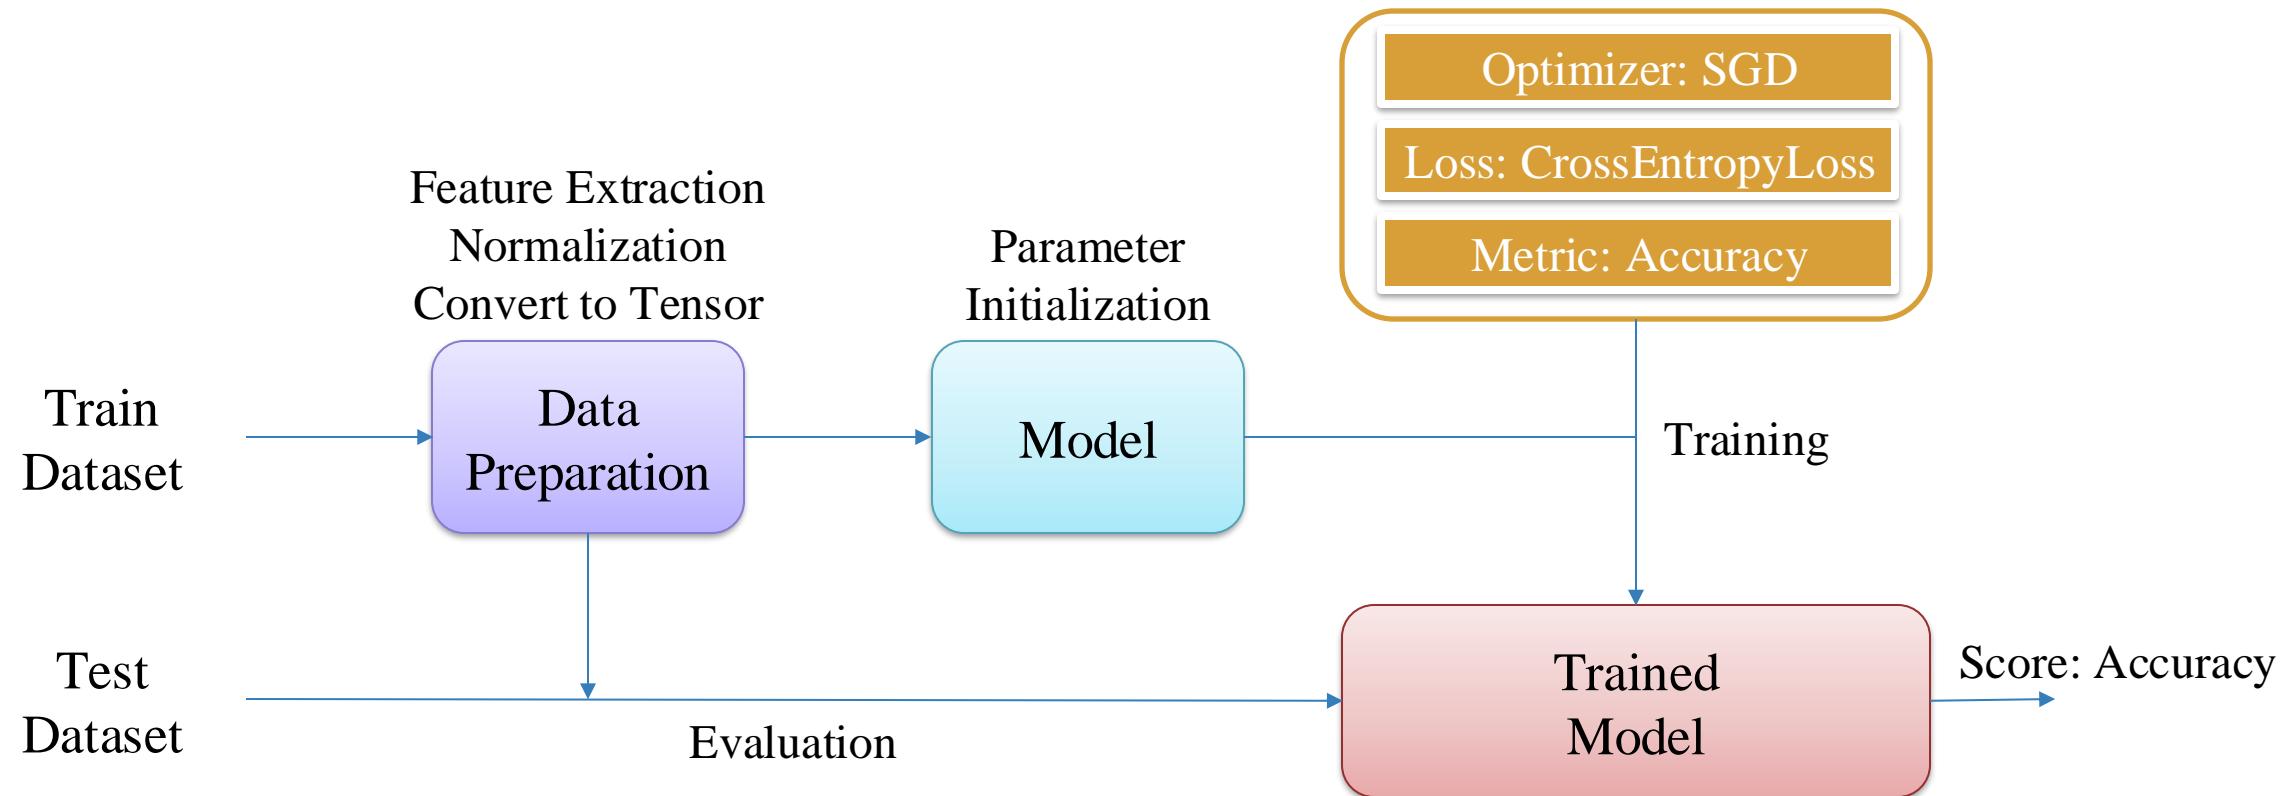
```
OrderedDict([('weight', tensor([[0.0946, 0.0891],
                               [0.0946, 0.0891]])), ('bias', tensor([0.0946, 0.0946]))])
OrderedDict()
OrderedDict([('weight', tensor([[0.0782, 0.0782]])), ('bias', tensor([0.0455]))])
OrderedDict()
```

**QUIZ TIME**

# Classification using MLP



## IRIS Dataset



# Classification using MLP



## Load the IRIS dataset

```
1 from sklearn.datasets import load_iris  
2 from sklearn.model_selection import train_test_split
```

```
1 data = load_iris()
```

```
1 data.data.shape
```

(150, 4)

```
1 data.target.shape
```

(150,)

# Classification using MLP



## Train Test Split

```
1 X_train, X_test, Y_train, Y_test = train_test_split(  
2     data.data,  
3     data.target,  
4     test_size=0.4,  
5     random_state=7  
6 )
```

```
1 X_valid, X_test, Y_valid, Y_test = train_test_split(  
2     X_test,  
3     Y_test,  
4     test_size=0.5,  
5     random_state=7  
6 )
```

```
1 X_train.shape, X_valid.shape, X_test.shape
```

```
((90, 4), (30, 4), (30, 4))
```

# Classification using MLP



## Preprocessing

```
1 from sklearn.preprocessing import StandardScaler
```

```
1 scaler = StandardScaler()
```

```
1 scaler.fit(X_train)
```

▼ StandardScaler ⓘ ?

StandardScaler()

```
1 X_train = scaler.transform(X_train)
2 X_valid = scaler.transform(X_valid)
3 X_test = scaler.transform(X_test)
```

# Classification using MLP



## Model

```
1 model_classifier = nn.Sequential(  
2     nn.Linear(4, 8),  
3     nn.ReLU(),  
4     nn.Linear(8, 3)  
5 )
```

```
1 summary(model_classifier, (1, 4))
```

Layer (type)	Output Shape	Param #
Linear-1	[-1, 1, 8]	40
ReLU-2	[-1, 1, 8]	0
Linear-3	[-1, 1, 3]	27

Total params: 67

Trainable params: 67

Non-trainable params: 0

Input size (MB): 0.00

Forward/backward pass size (MB): 0.00

Params size (MB): 0.00

Estimated Total Size (MB): 0.00

# Classification using MLP



## Loss Function

```
1 X_train[0]
```

```
tensor([-1.0341,  0.9957, -1.2020, -0.7777])
```

```
1 y_pred = model_classifier(X_train[0])
2 y_pred
```

```
tensor([ 0.1499,  0.0141, -0.4355], grad_fn=<ViewBackward0>)
```

```
1 loss_fn = nn.CrossEntropyLoss()
```

```
1 loss = loss_fn(y_pred, Y_train[0])
2 loss
```

```
tensor(0.8878, grad_fn=<NllLossBackward0>)
```

```
1 loss.item()
```

```
0.8878384828567505
```

# Classification using MLP



## Optimizer

```
1 learning_rate = 0.01
2
3 optimizer = optim.SGD(
4     model_classifier.parameters(),
5     learning_rate
6 )
```

```
1 def evaluate(model_classifier, X_valid, Y_valid):
2     with torch.no_grad():
3         Y_pred = model_classifier(X_valid)
4
5     Y_pred = torch.argmax(Y_pred, dim=1)
6     return sum(Y_pred == Y_valid)/len(Y_valid)
```

```
1 evaluate(model_classifier, X_valid, Y_valid)
```

```
tensor(0.5000)
```

# Classification using MLP



## Training

```
1 num_epochs = 20
2 losses = []
3 for epoch in range(num_epochs):
4     epoch_loss = []
5     for x_train, y_train in zip(X_train, Y_train):
6         y_pred = model_classifier(x_train)
7         loss = loss_fn(y_pred, y_train)
8         epoch_loss.append(loss.item())
9
10    optimizer.zero_grad()
11    loss.backward()
12    optimizer.step()
13    avg_loss = sum(epoch_loss)/len(epoch_loss)
14    losses.append(avg_loss)
15    acc = evaluate(model_classifier, X_valid, Y_valid)
16    print(f'{avg_loss} -- {acc}')
```

# Classification using MLP



## Evaluation

```
1 with torch.no_grad():
2     Y_pred = model_classifier(X_test)
```

```
1 Y_pred = torch.argmax(Y_pred, dim=1)
```

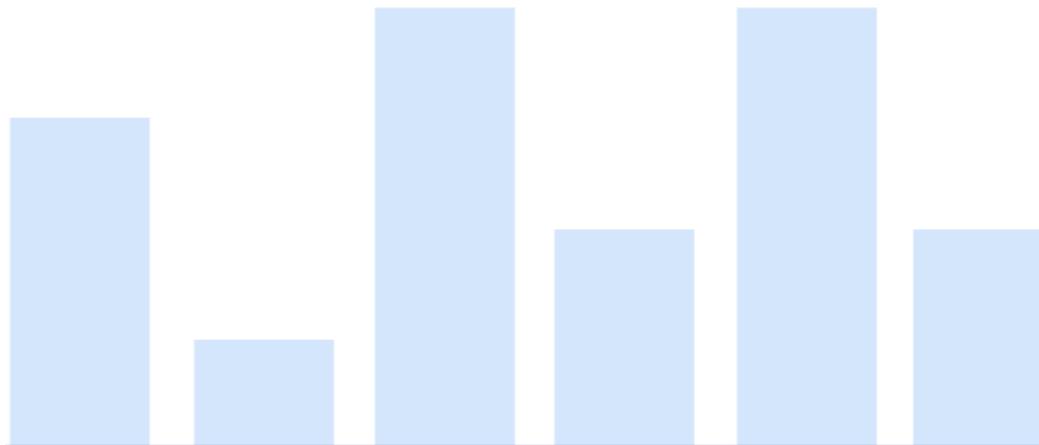
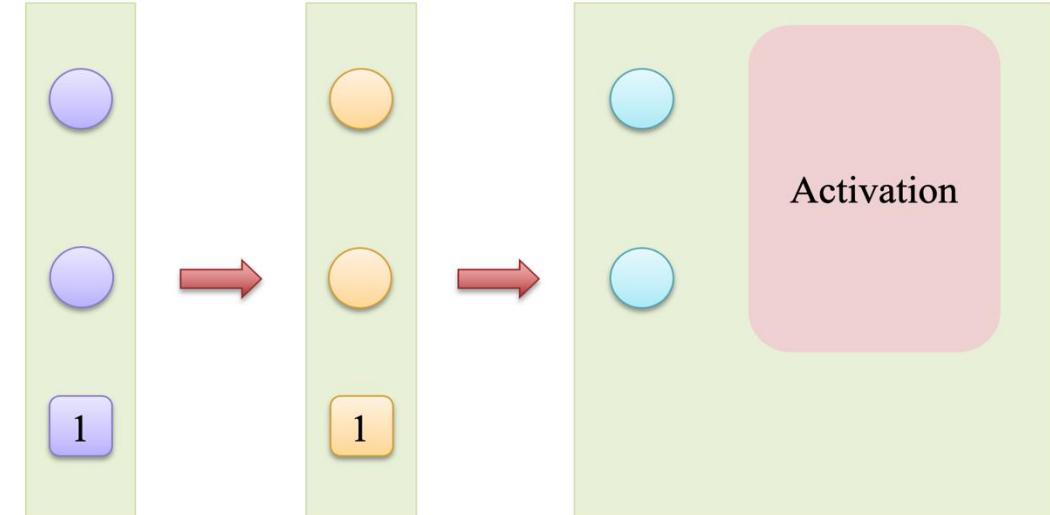
```
1 sum(Y_pred == Y_test)/len(Y_test)
```

tensor(0.9000)

# Objectives

## Softmax Regression (Review)

- ❖ Logistic Regression
- ❖ Softmax Function
- ❖ Softmax Regression using SGD



## Multi-layer Perceptron

- ❖ Hidden Layers
- ❖ Activation
- ❖ Loss Function
- ❖ Optimization



# Thanks!

Any questions?