

Module 05 – Extra Class

MOMENTUM

Gradient Descent with Momentum



Nguyen Quoc Thai

Code - Data

Objectives

Multi-layer Perceptron (Review)

- ❖ Hidden Layers
- ❖ Activation, Loss Function, Optimization
- ❖ Classification using MLP




Momentum

- ❖ Gradient Descent (Review)
- ❖ Gradient Descent with Momentum
- ❖ Nesterov Momentum

Outline


SECTION 1

Multi-layer Perceptron



SECTION 2

Momentum




Multi-layer Perceptron



Hidden Layer

#parameters: 4




Multi-layer Perceptron



Multi-layer Perceptron

#parameters: 12




Multi-layer Perceptron



Multi-layer Perceptron using Pytorch

#parameters: 12



```
1 model = nn.Sequential(  
2     nn.Linear(2, 2),  
3     nn.Linear(2, 2),  
4     nn.Sigmoid()  
5 )
```

```
1 summary(model, (10000000, 2))
```

Layer (type)	Output Shape	Param #
Linear-1	[-1, 10000000, 2]	6
Linear-2	[-1, 10000000, 2]	6
Sigmoid-3	[-1, 10000000, 2]	0

Total params: 12

Trainable params: 12

Non-trainable params: 0

Input size (MB): 76.29

Forward/backward pass size (MB): 457.76

Params size (MB): 0.00

Estimated Total Size (MB): 534.06


Multi-layer Perceptron

! Forward

$$\mathbf{x} = [1.0 \quad 2.0]$$

$$\mathbf{W}_h = [W_{h1} \quad W_{h2}] \\ = \begin{bmatrix} 0.1 & 0.1 \\ 0.1 & 0.1 \\ 0.1 & 0.1 \end{bmatrix}$$

$$\mathbf{W}_z = [W_z] = \begin{bmatrix} 0.1 \\ 0.1 \\ 0.1 \end{bmatrix}$$




Multi-layer Perceptron

! Forward

$$x = [1.0 \quad 2.0]$$

$$W_h = [W_{h1} \quad W_{h2}] \\ = \begin{bmatrix} 0.1 & 0.1 \\ 0.1 & 0.1 \\ 0.1 & 0.1 \end{bmatrix}$$

$$W_z = [W_z] = \begin{bmatrix} 0.1 \\ 0.1 \\ 0.1 \end{bmatrix}$$



$$h = [1.0 \ x]W_h = [1.0 \quad 1.0 \quad 2.0] \begin{bmatrix} 0.1 & 0.1 \\ 0.1 & 0.1 \\ 0.1 & 0.1 \end{bmatrix} = [0.4 \quad 0.4]$$

Multi-layer Perceptron




Forward

$$\mathbf{x} = [1.0 \quad 2.0]$$

$$\mathbf{W}_h = [W_{h1} \quad W_{h2}] \\ = \begin{bmatrix} 0.1 & 0.1 \\ 0.1 & 0.1 \\ 0.1 & 0.1 \end{bmatrix}$$

$$\mathbf{W}_z = [W_z] = \begin{bmatrix} 0.1 \\ 0.1 \\ 0.1 \end{bmatrix}$$



$$h = [1.0 \mathbf{x}] \mathbf{W}_h = [1.0 \quad 1.0 \quad 2.0] \begin{bmatrix} 0.1 & 0.1 \\ 0.1 & 0.1 \\ 0.1 & 0.1 \end{bmatrix} = [0.4 \quad 0.4]$$

$$z = [1.0 \mathbf{h}] \mathbf{W}_z = [1.0 \quad 0.4 \quad 0.4] \begin{bmatrix} 0.1 \\ 0.1 \\ 0.1 \end{bmatrix} = 0.18$$


Multi-layer Perceptron

! Forward

$$x = [1.0 \quad 2.0]$$

$$W_h = [W_{h1} \quad W_{h2}] \\ = \begin{bmatrix} 0.1 & 0.1 \\ 0.1 & 0.1 \\ 0.1 & 0.1 \end{bmatrix}$$

$$W_z = [W_z] = \begin{bmatrix} 0.1 \\ 0.1 \\ 0.1 \end{bmatrix}$$



$$h = [1.0 \quad x]W_h = [1.0 \quad 1.0 \quad 2.0] \begin{bmatrix} 0.1 & 0.1 \\ 0.1 & 0.1 \\ 0.1 & 0.1 \end{bmatrix} = [0.4 \quad 0.4]$$
$$z = [1.0 \quad h]W_z = [1.0 \quad 0.4 \quad 0.4] \begin{bmatrix} 0.1 \\ 0.1 \\ 0.1 \end{bmatrix} = 0.18$$


Multi-layer Perceptron

! Forward

$$x = [1.0 \quad 2.0]$$

$$W_h = [W_{h1} \quad W_{h2}] \\ = \begin{bmatrix} 0.1 & 0.1 \\ 0.1 & 0.1 \\ 0.1 & 0.1 \end{bmatrix}$$

$$W_z = [W_z] = \begin{bmatrix} 0.1 \\ 0.1 \\ 0.1 \end{bmatrix}$$



$$h = [1.0 \quad x]W_h = [1.0 \quad 1.0 \quad 2.0] \begin{bmatrix} 0.1 & 0.1 \\ 0.1 & 0.1 \\ 0.1 & 0.1 \end{bmatrix} = [0.4 \quad 0.4]$$
$$z = [1.0 \quad h]W_z = [1.0 \quad 0.4 \quad 0.4] \begin{bmatrix} 0.1 \\ 0.1 \\ 0.1 \end{bmatrix} = 0.18$$


Multi-layer Perceptron

! Backward

$$x = [1.0 \quad 2.0]$$

$$W_h = [W_{h1} \quad W_{h2}] \\ = \begin{bmatrix} 0.0946 & 0.0946 \\ 0.0946 & 0.0891 \\ 0.0946 & 0.0891 \end{bmatrix}$$

$$W_z = [W_z] = \begin{bmatrix} 0.0455 \\ 0.0782 \\ 0.0782 \end{bmatrix}$$



Multi-layer Perceptron




Activation

```
import torch.nn as nn

act = nn.Sigmoid()
input = torch.tensor([0.18, -0.18])
act(input)


tensor([0.5449, 0.4551])
```



```
import torch.nn as nn

act = nn.ReLU()
input = torch.tensor([0.18, -0.18])
act(input)


tensor([0.1800, 0.0000])
```



```
import torch.nn as nn

act = nn.Tanh()
input = torch.tensor([0.18, -0.18])
act(input)

tensor([ 0.1781, -0.1781])
```



Multi-layer Perceptron



Loss Function

```
import torch.nn as nn
loss_fn = nn.BCELoss()

y_pred
tensor([0.5449], grad_fn=<SigmoidBackward0>

y
tensor([0.])

loss = loss_fn(y_pred, y)
loss
tensor(0.7872, grad_fn=<BinaryCrossEntropyBackward0>)
```

```
import torch.nn as nn
loss_fn = nn.CrossEntropyLoss()

y = torch.tensor(0)
y
tensor(0)

y_pred
tensor([0.0580, 0.4275], grad_fn=<AddBackward0>

loss_fn(y_pred, y)
tensor(0.8949, grad_fn=<NllLossBackward0>)
```

Multi-layer Perceptron



Optimizer (SGD)

```
learning_rate = 0.1
optimizer = optim.SGD(model.parameters(), learning_rate)
```

```
loss.backward()
```

```
optimizer.step()
```


```
for layer in model.children():
    print(layer.state_dict())
```

```
OrderedDict([('weight', tensor([[0.0946, 0.0891],
                               [0.0946, 0.0891]])), ('bias', tensor([0.0946, 0.0946]))])
OrderedDict([('weight', tensor([[0.0782, 0.0782]])), ('bias', tensor([0.0455]))])
OrderedDict()
```

Classification using MLP




IRIS Dataset



Outline


SECTION 1

Multi-layer Perceptron



SECTION 2

Momentum



Momentum



SGD

1) Pick a sample (x, y) from training data

2) Compute output \hat{y}

$$\hat{y} = \boldsymbol{\theta}^T \mathbf{x} = \mathbf{x}^T \boldsymbol{\theta}$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\nabla_{\boldsymbol{\theta}} L = 2\mathbf{x}(\hat{y} - y)$$

5) Update parameters

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L$$

η is learning rate

1) Pick a sample (x, y) from training

2) Compute output \hat{y}

$$z = \boldsymbol{\theta}^T \mathbf{x}$$

$$\hat{y} = \sigma(z) = \frac{1}{1 + e^{-z}}$$

3) Compute loss

$$L(\boldsymbol{\theta}) = (-y \log \hat{y} - (1-y) \log(1-\hat{y}))$$

4) Compute derivative

$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{y} - y)$$

5) Update parameters

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L$$


η is learning rate

Momentum



- Objective function: x^2
- Derivative: $2x$
- Gradient:

$$x = x - \eta * f'(x)$$



Momentum

! SGD


- Objective function: x^2
- Derivative: $2x$
- Gradient:

$$x = x - \eta * f'(x)$$

$$x = -2.0$$

$$\eta = 0.1$$

$$\begin{aligned}x &= -2.0 - 0.1 * (2 * -2.0) \\&= -1.6\end{aligned}$$




Momentum



SGD

- Objective function: x^2
- Derivative: $2x$
- Gradient:

$$x = x - \eta * f'(x)$$



Momentum



SGD

```
1 def init_inputs(r_min=-2.0, r_max=2.0):
2     inputs = torch.arange(r_min, r_max+0.1, 0.1)
3     return inputs
```

```
1 init_inputs(r_min=-2.0, r_max=2.0)
```

```
tensor([-2.0000e+00, -1.9000e+00, -1.8000e+00, -1.7000e+00, -1.6000e+00,
       -1.5000e+00, -1.4000e+00, -1.3000e+00, -1.2000e+00, -1.1000e+00,
       -1.0000e+00, -9.0000e-01, -8.0000e-01, -7.0000e-01, -6.0000e-01,
       -5.0000e-01, -4.0000e-01, -3.0000e-01, -2.0000e-01, -1.0000e-01,
       -5.9605e-09, 1.0000e-01, 2.0000e-01, 3.0000e-01, 4.0000e-01,
       5.0000e-01, 6.0000e-01, 7.0000e-01, 8.0000e-01, 9.0000e-01,
       1.0000e+00, 1.1000e+00, 1.2000e+00, 1.3000e+00, 1.4000e+00,
       1.5000e+00, 1.6000e+00, 1.7000e+00, 1.8000e+00, 1.9000e+00,
       2.0000e+00])
```

```
1 def objective(x):
2     return x**2.0
3
4 def derivative(x):
5     return 2.0*x
6
7 def sgd(input, num_epochs, learning_rate):
8     solutions, scores = [], []
9     solution = input
10    solution_eval = objective(solution)
11    solutions.append(solution)
12    scores.append(solution_eval)
13
14    for epoch in range(num_epochs):
15        # calculate gradient
16        gradient = derivative(solution)
17        # take a step
18        solution = solution - learning_rate * gradient
19        # evaluate candidate point
20        solution_eval = objective(solution)
21        # store solution
22        solutions.append(solution)
23        scores.append(solution_eval)
24        # report progress
25        print('Epoch: %0.2d -- f(%0.3f) = %.5f' % (
26            epoch, solution, solution_eval))
27
return solutions, scores
```


Momentum



SGD

```
1 num_epochs = 20
2 learning_rate = 0.1
3 inputs = init_inputs()
4 input = inputs[0]
5 solutions, scores = sgd(input, num_epochs, learning_rate)
```

```
Epoch: 00 -- f(-1.600) = 2.56000
Epoch: 01 -- f(-1.280) = 1.63840
Epoch: 02 -- f(-1.024) = 1.04858
Epoch: 03 -- f(-0.819) = 0.67109
Epoch: 04 -- f(-0.655) = 0.42950
Epoch: 05 -- f(-0.524) = 0.27488
Epoch: 06 -- f(-0.419) = 0.17592
Epoch: 07 -- f(-0.336) = 0.11259
Epoch: 08 -- f(-0.268) = 0.07206
Epoch: 09 -- f(-0.215) = 0.04612
Epoch: 10 -- f(-0.172) = 0.02951
Epoch: 11 -- f(-0.137) = 0.01889
Epoch: 12 -- f(-0.110) = 0.01209
Epoch: 13 -- f(-0.088) = 0.00774
Epoch: 14 -- f(-0.070) = 0.00495
Epoch: 15 -- f(-0.056) = 0.00317
Epoch: 16 -- f(-0.045) = 0.00203
Epoch: 17 -- f(-0.036) = 0.00130
Epoch: 18 -- f(-0.029) = 0.00083
Epoch: 19 -- f(-0.023) = 0.00053
```




Momentum



SGD

- Objective function:
 $x^4 + x^3 - 2x^2 - 2x + 2$
- Derivative:
 $4x^3 + 3x^2 - 4x - 2$
- Gradient:
 $x = x - \eta * f'(x)$



Momentum

! SGD

➤ Objective function:

$$x^4 + x^3 - 2x^2 - 2x + 2$$

➤ Derivative:

$$4x^3 + 3x^2 - 4x - 2$$


➤ Gradient:

$$x = x - \eta * f'(x)$$

$$x = -2.0$$

$$\eta = 0.1$$

$$\begin{aligned}x &= -2.0 - 0.1 * (-14) \\&= -0.6\end{aligned}$$



Momentum

! SGD

➤ Objective function:

$$x^4 + x^3 - 2x^2 - 2x + 2$$

➤ Derivative:

$$4x^3 + 3x^2 - 4x - 2$$

➤ Gradient:


$$x = x - \eta * f'(x)$$

$$x = -2.0$$

$$\eta = 0.1$$

$$\text{num_epochs} = 5$$

$$\text{Cost} = 2.05$$



Momentum




SGD

```
# objective function
def objective(x):
    return x**4+x**3-2*x**2-2*x + 2

# derivative of objective function
def derivative(x):
    return 4*x**3 + 3*x**2 - 4*x - 2

num_epochs = 20
learning_rate = 0.1
inputs = init_inputs()
input = inputs[0]
solutions, scores = sgd(input, num_epochs, learning_rate)
```



```
Epoch: 00 -- f(-0.600) = 2.39360
Epoch: 01 -- f(-0.662) = 2.34977
Epoch: 02 -- f(-0.742) = 2.27775
Epoch: 03 -- f(-0.840) = 2.17372
Epoch: 04 -- f(-0.951) = 2.05123
Epoch: 05 -- f(-1.059) = 1.94550
Epoch: 06 -- f(-1.144) = 1.88629
Epoch: 07 -- f(-1.195) = 1.86667
Epoch: 08 -- f(-1.219) = 1.86278
Epoch: 09 -- f(-1.228) = 1.86225
Epoch: 10 -- f(-1.231) = 1.86219
Epoch: 11 -- f(-1.232) = 1.86218
Epoch: 12 -- f(-1.232) = 1.86218
Epoch: 13 -- f(-1.232) = 1.86218
Epoch: 14 -- f(-1.232) = 1.86218
Epoch: 15 -- f(-1.232) = 1.86218
Epoch: 16 -- f(-1.232) = 1.86218
Epoch: 17 -- f(-1.232) = 1.86218
Epoch: 18 -- f(-1.232) = 1.86218
Epoch: 19 -- f(-1.232) = 1.86218
```



Momentum



SGD



Momentum



SGD with Momentum

- Objective function:

$$x^4 + x^3 - 2x^2 - 2x + 2$$

- Derivative:


$$4x^3 + 3x^2 - 4x - 2$$

- Gradient with Momentum:

$$v_t = \gamma v_{t-1} + \eta f'(x)$$

$$x = x - v_t$$

γ : momentum



Momentum



SGD with Momentum

➤ Objective function:

$$x^4 + x^3 - 2x^2 - 2x + 2$$

➤ Derivative:

$$4x^3 + 3x^2 - 4x - 2$$

➤ Gradient with Momentum:


$$v_t = \gamma v_{t-1} + \eta f'(x)$$

$$x = x - v_t$$

$$x = -2.0 \quad \eta = 0.1 \quad \gamma = 0.8 \quad v_0 = 0.0$$

$$v_1 = 0.8 * 0.0 + 0.1 * (-14) = -1.4$$

$$x = -2.0 - (-1.4) = -0.6$$



Momentum



SGD with Momentum

➤ Objective function:

$$x^4 + x^3 - 2x^2 - 2x + 2$$

➤ Derivative:

$$4x^3 + 3x^2 - 4x - 2$$

➤ Gradient with Momentum:


$$v_t = \gamma v_{t-1} + \eta f'(x)$$

$$x = x - v_t$$

$$x = -2.0 \quad \eta = 0.1 \quad \gamma = 0.8 \quad v_1 = -1.2$$

$$v_1 = 0.8 * (-1.4) + 0.1 * (0.62) = -1.05$$

$$x = -0.6 - (-1.05) = 0.45$$



Momentum



SGD with Momentum

➤ Objective function:

$$x^4 + x^3 - 2x^2 - 2x + 2$$

➤ Derivative:

$$4x^3 + 3x^2 - 4x - 2$$


➤ Gradient with Momentum:

$$v_t = \gamma v_{t-1} + \eta f'(x)$$

$$x = x - v_t$$

$$x = -2.0 \quad \eta = 0.1 \quad \gamma = 0.8 \quad v_0 = 0.0$$

$$\text{num_epochs} = 5$$




Momentum



SGD with Momentum

$$v_t = \gamma v_{t-1} + \eta f'(x)$$
$$x = x - v_t$$



Movement = Negative of Gradient + Momentum

Momentum



SGD with Momentum

```
1 # gradient descent algorithm
2 def gradient_descent_with_momentum(input, num_epochs, learning_rate, momentum):
3     solutions, scores = [], []
4     solution = input
5     # keep track of the change
6     change = 0.0
7     solution_eval = objective(solution)
8     solutions.append(solution)
9     scores.append(solution_eval)
10
11    for epoch in range(num_epochs):
12        # calculate gradient
13        gradient = derivative(solution)
14        # calculate update
15        new_change = learning_rate * gradient + momentum * change
16        # take a step
17        solution = solution - new_change
18        # save the change
19        change = new_change
20        # evaluate candidate point
21        solution_eval = objective(solution)
22        # store solution
23        solutions.append(solution)
24        scores.append(solution_eval)
25        # report progress
26        print('Epoch: %0.2d -- f(%0.3f) = %.5f' %
27              epoch, solution, solution_eval))
28    return solutions, scores
```


Momentum



SGD with Momentum

```
1 num_epochs = 20
2 learning_rate = 0.1
3 inputs = init_inputs()
4 input = inputs[0]
5 # define momentum
6 momentum = 0.8
7 solutions, scores = gradient_descent_with_momentum(
8     input, num_epochs, learning_rate, momentum
9 )
```

```
Epoch: 00 -- f(-0.800) = 2.21760
Epoch: 01 -- f(-0.147) = 2.24834
Epoch: 02 -- f(0.311) = 1.22418
Epoch: 03 -- f(0.761) = 0.09618
Epoch: 04 -- f(1.075) = 0.11696
Epoch: 05 -- f(0.913) = -0.03723
Epoch: 06 -- f(0.594) = 0.44088
Epoch: 07 -- f(0.387) = 1.00798
Epoch: 08 -- f(0.308) = 1.23381
Epoch: 09 -- f(0.327) = 1.17774
Epoch: 10 -- f(0.428) = 0.89005
Epoch: 11 -- f(0.593) = 0.44244
Epoch: 12 -- f(0.774) = 0.07689
Epoch: 13 -- f(0.863) = -0.01795
Epoch: 14 -- f(0.799) = 0.04293
Epoch: 15 -- f(0.672) = 0.26025
Epoch: 16 -- f(0.582) = 0.46945
Epoch: 17 -- f(0.563) = 0.51926
Epoch: 18 -- f(0.606) = 0.41078
Epoch: 19 -- f(0.684) = 0.23563
```



QUIZ TIME

Momentum



Problem of SGD with Momentum

➤ Objective function:

$$x^4 + x^3 - 2x^2 - 2x + 2$$

➤ Derivative:

$$4x^3 + 3x^2 - 4x - 2$$


➤ Gradient with Momentum:

$$v_t = \gamma v_{t-1} + \eta f'(x)$$

$$x = x - v_t$$

$$x = -2.0 \quad \eta = 0.1 \quad \gamma = 0.9 \quad v_0 = 0.0$$

$$\text{num_epochs} = 10$$



Momentum



Nesterov Momentum

- Objective function:

$$x^4 + x^3 - 2x^2 + 2$$

- Derivative:


$$4x^3 + 3x^2 - 4x - 2$$

- Gradient with Momentum:

$$v_t = \gamma v_{t-1} - \eta f'(x + \gamma * v_{t-1})$$

$$x = x + v_t$$


$$x = -2.0 \quad \eta = 0.1 \quad \gamma = 0.9 \quad v_0 = 0.0$$



Momentum




Nesterov Momentum



(a) Momentum-Based Gradient Descent

$$\text{Green Circle} \Rightarrow \frac{\partial L}{\partial w_0} = \frac{\text{Negative}(-)}{\text{Positive}(+)}$$



(b) Nesterov Accelerated Gradient Descent

$$\text{Red Circle} \Rightarrow \frac{\partial L}{\partial w_0} = \frac{\text{Negative}(-)}{\text{Negative}(-)}$$

Momentum



Nesterov Momentum

```
1 # gradient descent algorithm
2 def nesterov_momentum(input, num_epochs, learning_rate, momentum):
3     solutions, scores = [], []
4     solution = input
5     # keep track of the change
6     change = 0.0
7     solution_eval = objective(solution)
8     solutions.append(solution)
9     scores.append(solution_eval)
10
11    for epoch in range(num_epochs):
12        # calculate the projected solution
13        projected = solution + momentum * change
14        # calculate gradient
15        gradient = derivative(projected)
16        # calculate update
17        new_change = momentum * change - learning_rate * gradient
18        # take a step
19        solution = solution + new_change
20        # save the change
21        change = new_change
22        # evaluate candidate point
23        solution_eval = objective(solution)
24        # store solution
25        solutions.append(solution)
26        scores.append(solution_eval)
27        # report progress
28        print('Epoch: %0.2d -- f(%0.3f) = %.5f' %
29              epoch, solution, solution_eval)
30    return solutions, scores
```


Momentum



Nesterov Momentum

```
1 num_epochs = 20
2 learning_rate = 0.1
3 inputs = init_inputs()
4 input = inputs[0]
5 # define momentum
6 momentum = 0.9
7 solutions, scores = nesterov_momentum(
8     input, num_epochs, learning_rate, momentum
9 )
```


```
Epoch: 00 -- f(-0.800) = 2.21760
Epoch: 01 -- f(0.360) = 1.08511
Epoch: 02 -- f(0.268) = 1.34411
Epoch: 03 -- f(0.247) = 1.40184
Epoch: 04 -- f(0.300) = 1.25650
Epoch: 05 -- f(0.432) = 0.87709
Epoch: 06 -- f(0.614) = 0.39136
Epoch: 07 -- f(0.719) = 0.16670
Epoch: 08 -- f(0.725) = 0.15602
Epoch: 09 -- f(0.707) = 0.19025
Epoch: 10 -- f(0.692) = 0.21938
Epoch: 11 -- f(0.687) = 0.22943
Epoch: 12 -- f(0.689) = 0.22596
Epoch: 13 -- f(0.692) = 0.21937
Epoch: 14 -- f(0.694) = 0.21555
Epoch: 15 -- f(0.694) = 0.21499
Epoch: 16 -- f(0.694) = 0.21595
Epoch: 17 -- f(0.693) = 0.21690
Epoch: 18 -- f(0.693) = 0.21728
Epoch: 19 -- f(0.693) = 0.21723
```



Classification





IRIS Dataset



Objectives

Multi-layer Perceptron (Review)

- ❖ Hidden Layers
- ❖ Activation, Loss Function, Optimization
- ❖ Classification using MLP



Momentum

- ❖ Gradient Descent
- ❖ Gradient Descent with Momentum
- ❖ Nesterov Momentum



Thanks!

Any questions?