

## Module 05 – Extra Class

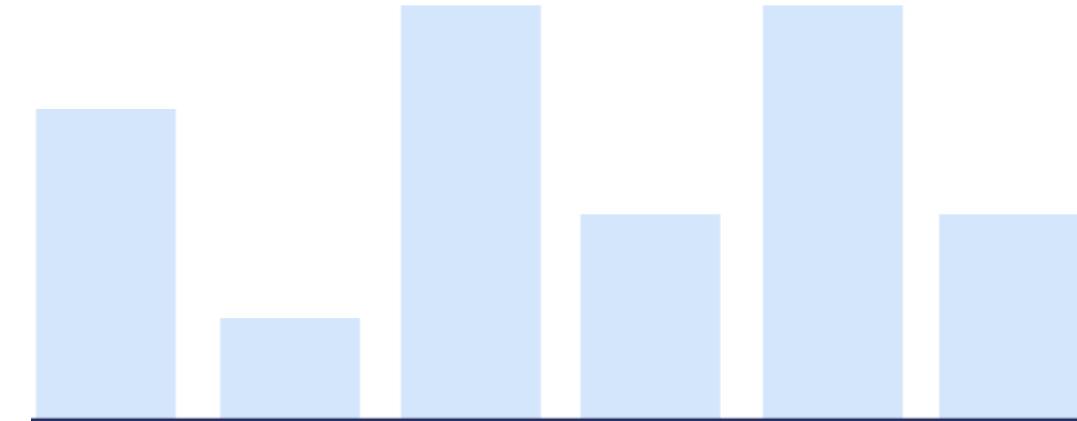
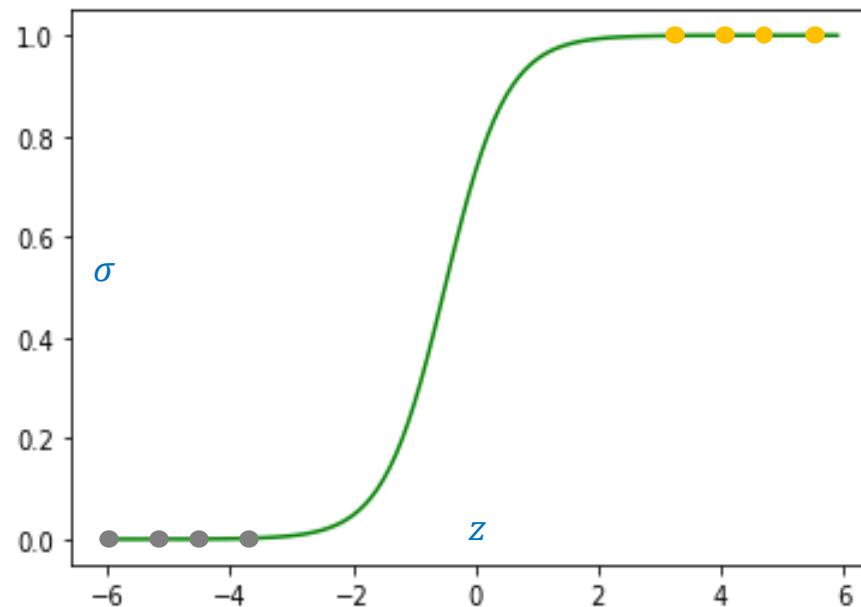
# SOFTMAX REGRESSION

Nguyen Quoc Thai

# Objectives

## Logistic Regression (Review)

- ❖ Logistic Regression
- ❖ Sigmoid Function
- ❖ Gradient Descent



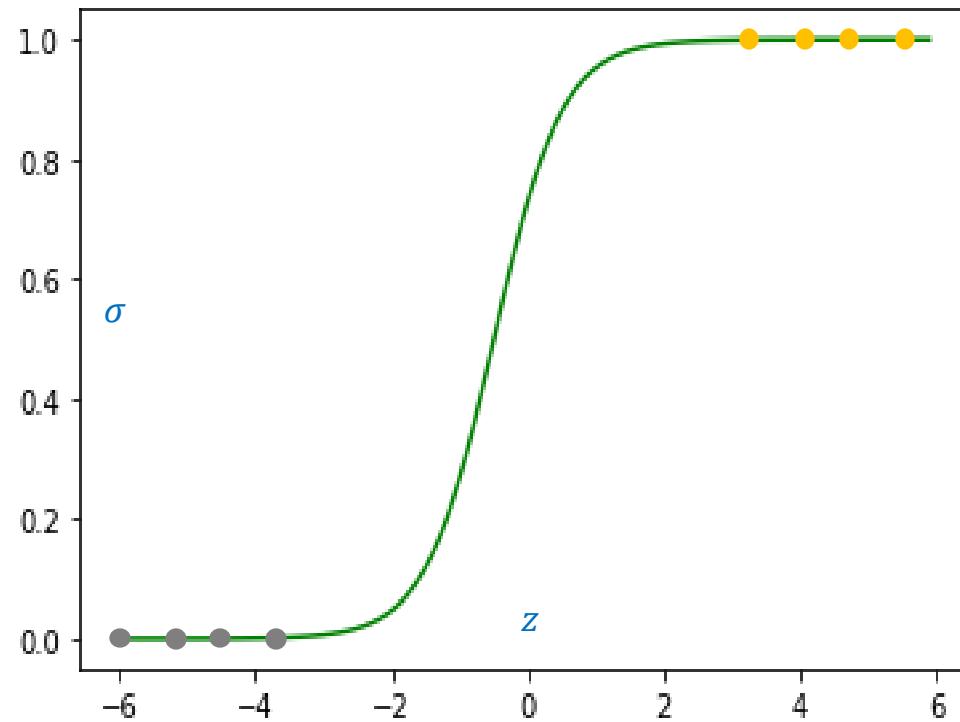
## Softmax Regression

- ❖ Softmax Regression
- ❖ Softmax Function
- ❖ One Sample
- ❖ N Sample

# Outline

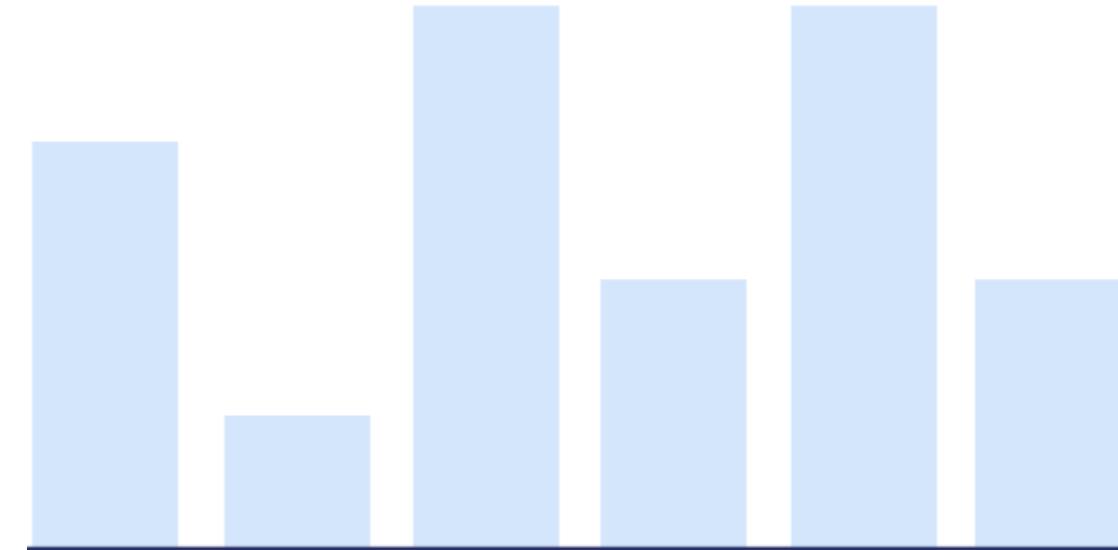
SECTION 1

## Logistic Regression



SECTION 2

## Softmax Regression



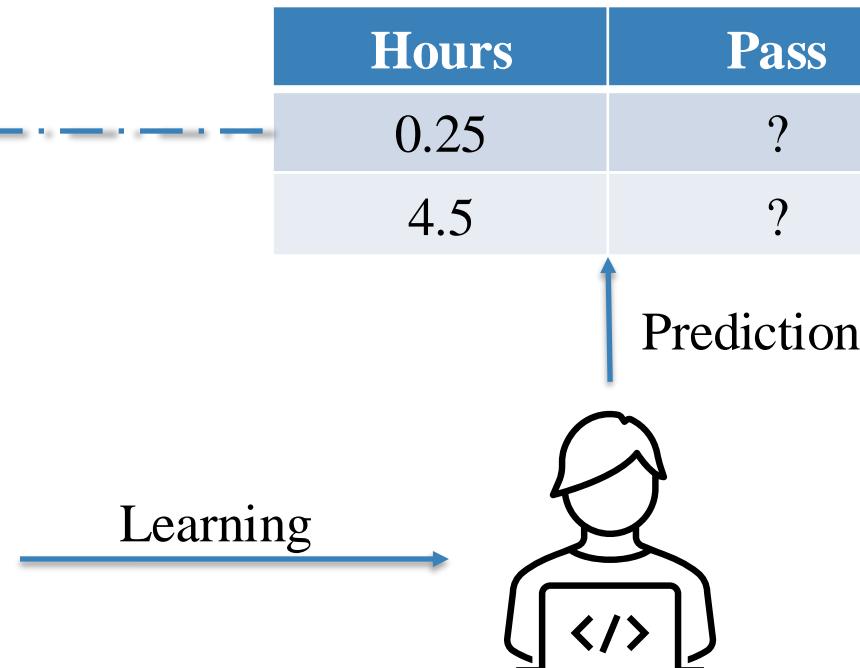
# Logistic Regression



## Classification Task

| Hours | Pass |
|-------|------|
| 0.5   | 0    |
| 1.0   | 0    |
| 1.5   | 0    |
| 2.0   | 0    |
| 2.5   | 1    |
| 3.0   | 1    |
| 3.5   | 1    |
| 4.0   | 1    |

| Hours | Pass |
|-------|------|
| 0.25  | ?    |
| 4.5   | ?    |



# Logistic Regression



## Classification Task

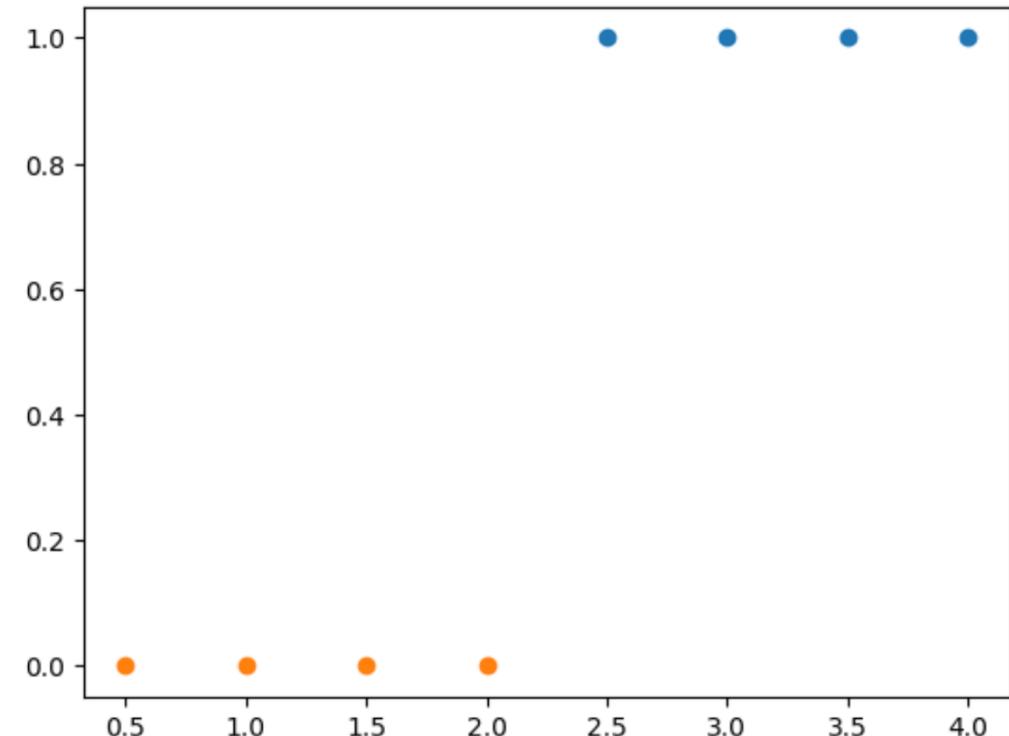
| Hours | Pass |
|-------|------|
| 0.5   | 0    |
| 1.0   | 0    |
| 1.5   | 0    |
| 2.0   | 0    |
| 2.5   | 1    |
| 3.0   | 1    |
| 3.5   | 1    |
| 4.0   | 1    |

$$y = f(x)$$

Find a function to fit the data

Sigmoid function

## Visualization



# Logistic Regression

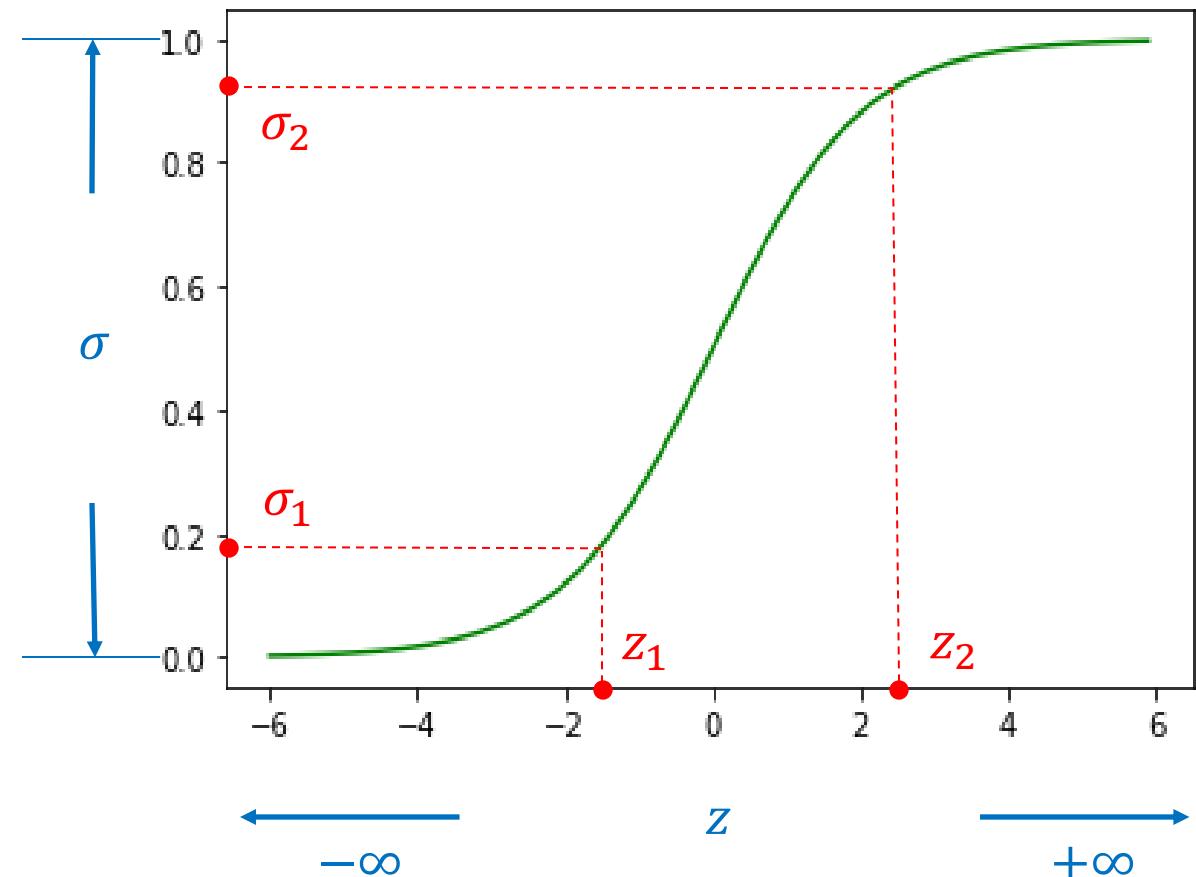


## Sigmoid Function

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

$$z \in (-\infty, +\infty)$$

$$\sigma(z) \in (0, 1)$$



# Logistic Regression



## Logistic Regression using Gradient Descent

1) Pick a sample  $(x, y)$  from training

2) Compute output  $\hat{y}$

$$z = \theta^T x$$

$$\hat{y} = \sigma(z) = \frac{1}{1 + e^{-z}}$$

3) Compute loss

$$L(\theta) = (-y \log \hat{y} - (1-y) \log(1-\hat{y}))$$

4) Compute derivative

$$\nabla_{\theta} L = x(\hat{y} - y)$$

5) Update parameters

$$\theta = \theta - \eta \nabla_{\theta} L$$

$\eta$  is learning rate

$$\theta = \begin{bmatrix} b \\ w \end{bmatrix} = \begin{bmatrix} 0.1 \\ 0.1 \end{bmatrix}$$

$$x = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix}$$

$$y = [0]$$

$$\eta = 0.1$$

| Hours | Pass |
|-------|------|
| 0.5   | 0    |
| 1.0   | 0    |
| 1.5   | 0    |
| 2.0   | 0    |
| 2.5   | 1    |
| 3.0   | 1    |
| 3.5   | 1    |
| 4.0   | 1    |

# Logistic Regression



## Logistic Regression using Gradient Descent

1) Pick a sample  $(x, y)$  from training

2) Compute output  $\hat{y}$

$$z = \theta^T x$$

$$\hat{y} = \sigma(z) = \frac{1}{1 + e^{-z}}$$

3) Compute loss

$$L(\theta) = (-y \log \hat{y} - (1-y) \log(1-\hat{y}))$$

4) Compute derivative

$$\nabla_{\theta} L = x(\hat{y} - y)$$

5) Update parameters

$$\theta = \theta - \eta \nabla_{\theta} L$$

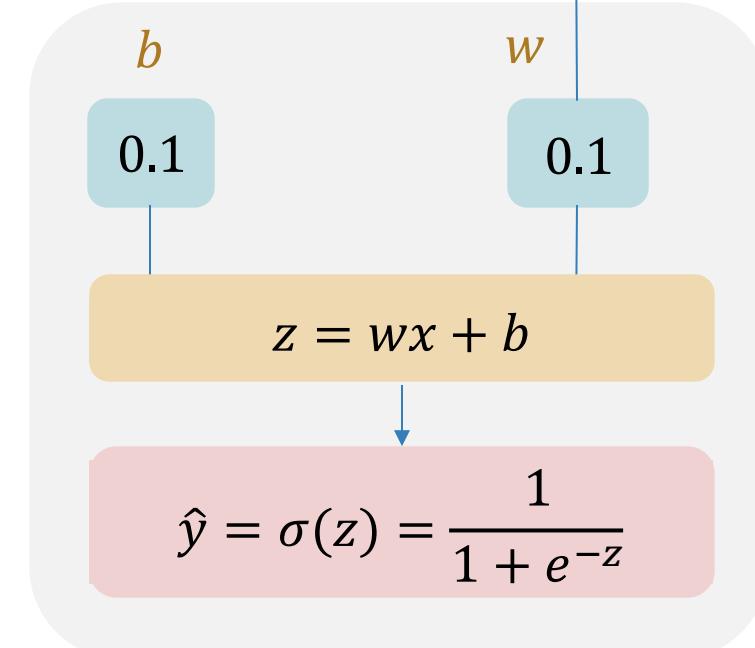
$\eta$  is learning rate

$$\eta = 0.1$$

$$\theta = \begin{bmatrix} 0.1 \\ 0.1 \end{bmatrix}$$

$$x = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix}$$

$$y = [0]$$



$$z = 0.15$$

$$\hat{y} = 0.54$$

# Logistic Regression



## Logistic Regression using Gradient Descent

1) Pick a sample  $(x, y)$  from training

2) Compute output  $\hat{y}$

$$z = \theta^T x$$

$$\hat{y} = \sigma(z) = \frac{1}{1 + e^{-z}}$$

3) Compute loss

$$L(\theta) = (-y \log \hat{y} - (1-y) \log(1-\hat{y}))$$

4) Compute derivative

$$\nabla_{\theta} L = x(\hat{y} - y)$$

5) Update parameters

$$\theta = \theta - \eta \nabla_{\theta} L$$

$\eta$  is learning rate

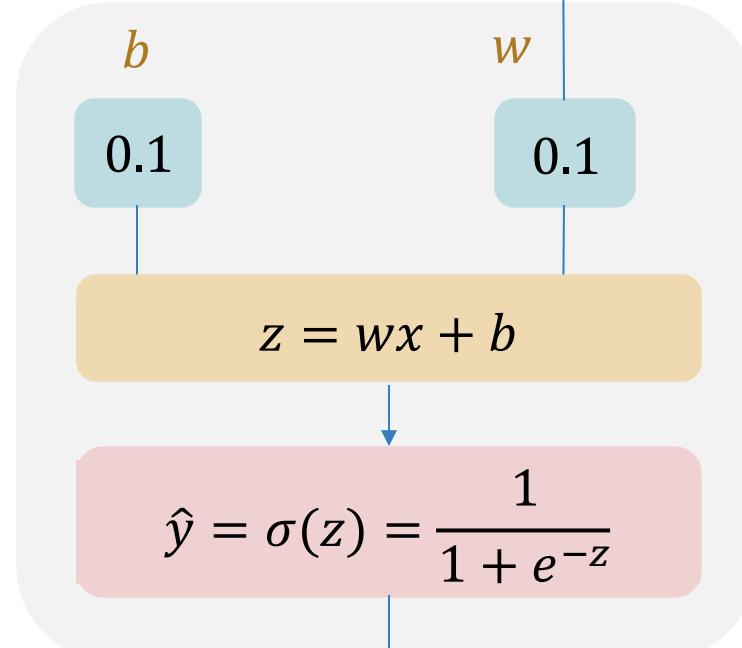
$$\eta = 0.1$$

$$\theta = \begin{bmatrix} 0.1 \\ 0.1 \end{bmatrix}$$

$$x = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix}$$

$$y = [0]$$

$$L = 0.771$$



$$y = 0$$

# Logistic Regression



## Logistic Regression using Gradient Descent

1) Pick a sample  $(x, y)$  from training

2) Compute output  $\hat{y}$

$$z = \theta^T x$$

$$\hat{y} = \sigma(z) = \frac{1}{1 + e^{-z}}$$

3) Compute loss

$$L(\theta) = (-y \log \hat{y} - (1-y) \log(1-\hat{y}))$$

4) Compute derivative

$$\nabla_{\theta} L = x(\hat{y} - y)$$

5) Update parameters

$$\theta = \theta - \eta \nabla_{\theta} L$$

$\eta$  is learning rate

$$\eta = 0.1$$

$$\theta = \begin{bmatrix} 0.1 \\ 0.1 \end{bmatrix}$$

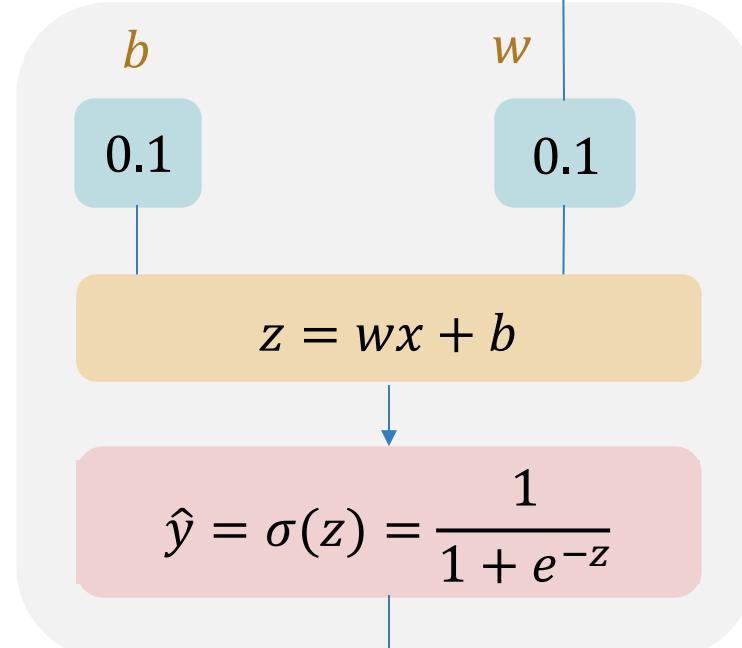
$$x = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix}$$

$$y = [0]$$

$$\nabla_{\theta} L = x(\hat{y} - y)$$

$$\begin{aligned} &= \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix} * 0.54 \\ &= \begin{bmatrix} 0.54 \\ 0.27 \end{bmatrix} \end{aligned}$$

$$L = 0.771$$



# Logistic Regression



## Logistic Regression using Gradient Descent

1) Pick a sample  $(x, y)$  from training

2) Compute output  $\hat{y}$

$$z = \theta^T x$$

$$\hat{y} = \sigma(z) = \frac{1}{1 + e^{-z}}$$

3) Compute loss

$$L(\theta) = (-y \log \hat{y} - (1-y) \log(1-\hat{y}))$$

4) Compute derivative

$$\nabla_{\theta} L = x(\hat{y} - y)$$

5) Update parameters

$$\theta = \theta - \eta \nabla_{\theta} L$$

$\eta$  is learning rate

$$\eta = 0.1$$

$$\theta = \begin{bmatrix} 0.1 \\ 0.1 \end{bmatrix}$$

$$x = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix}$$

$$y = [0]$$

$$\theta = \theta - \eta \nabla_{\theta} L$$

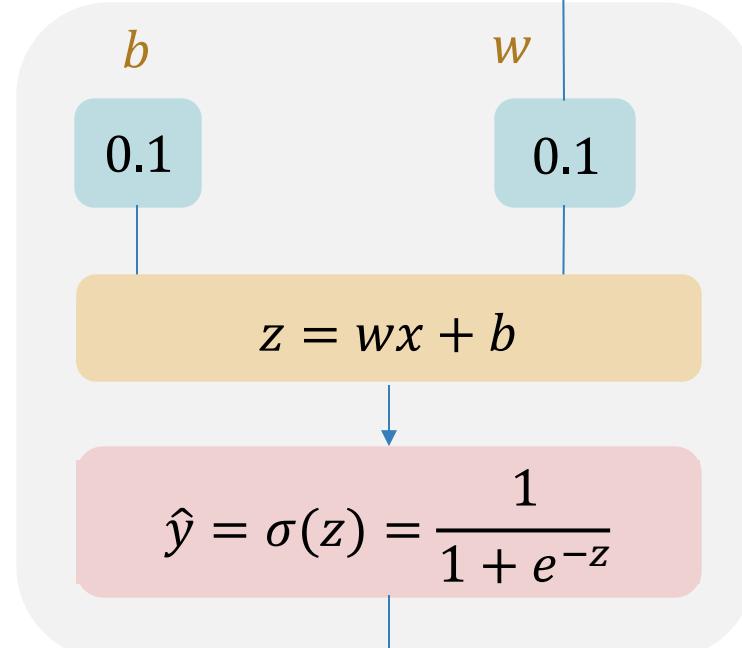
$$= \begin{bmatrix} 0.046 \\ 0.073 \end{bmatrix}$$

$$\nabla_{\theta} L = x(\hat{y} - y)$$

$$= \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix} * 0.54$$

$$= \begin{bmatrix} 0.54 \\ 0.27 \end{bmatrix}$$

$$L = 0.771$$



$$z = 0.15$$

$$\hat{y} = 0.54$$

$$y = 0$$

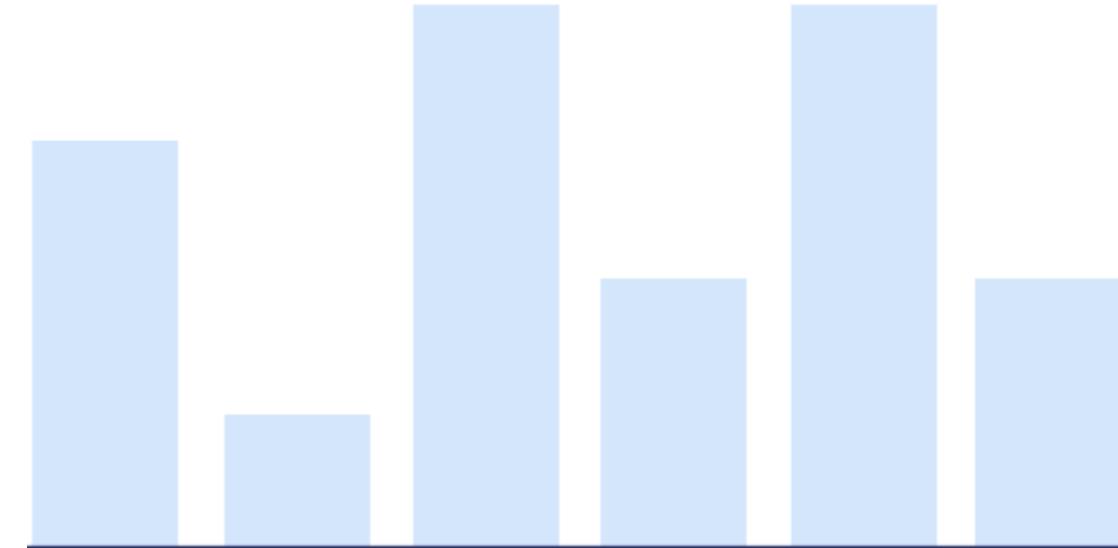
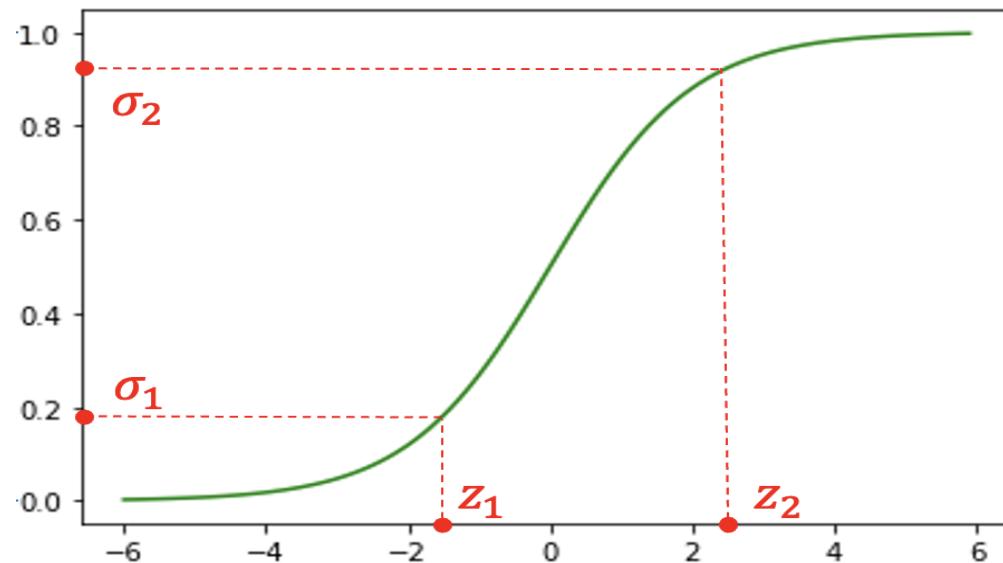
# Outline

SECTION 1

## Logistic Regression

SECTION 2

## Softmax Regression



# Softmax Regression



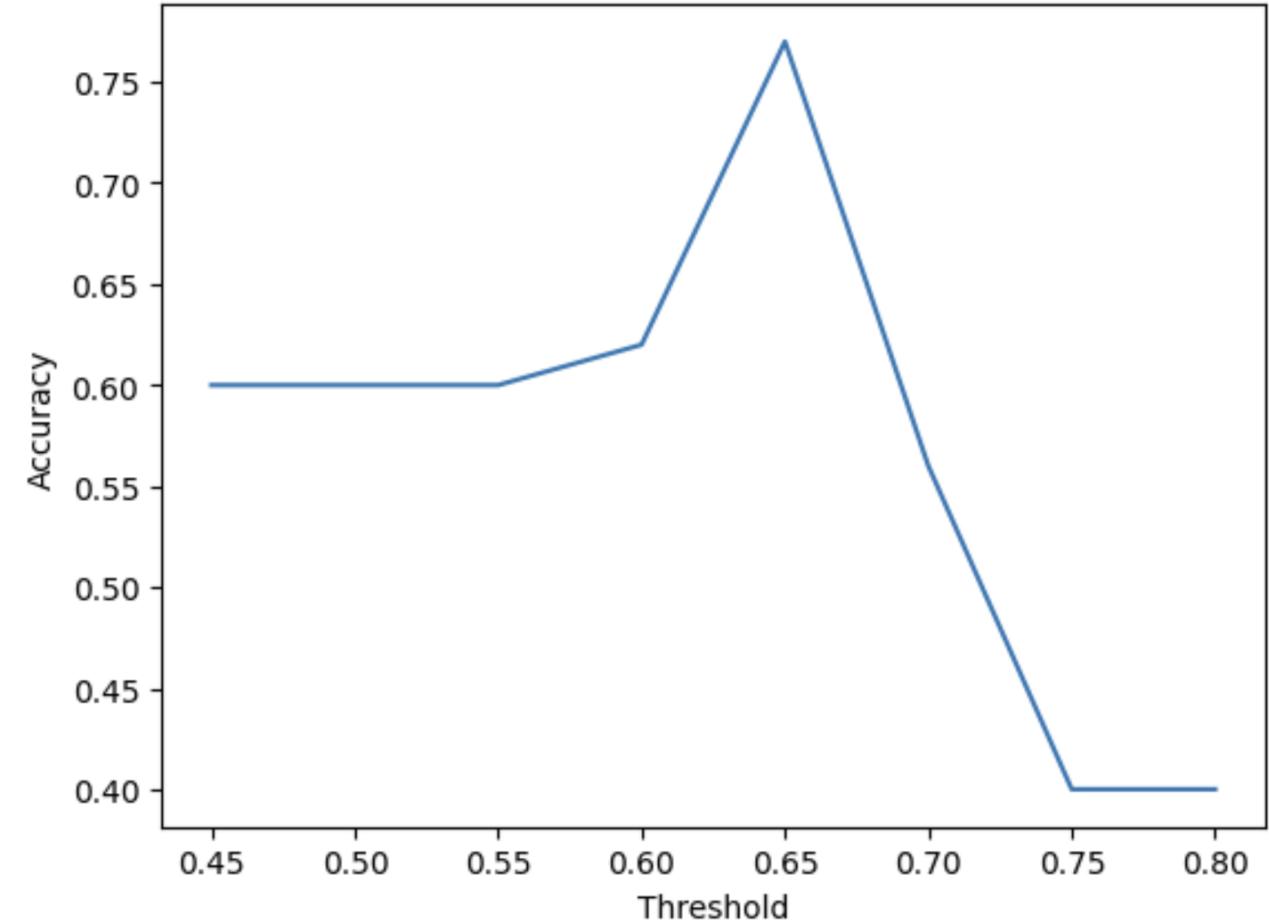
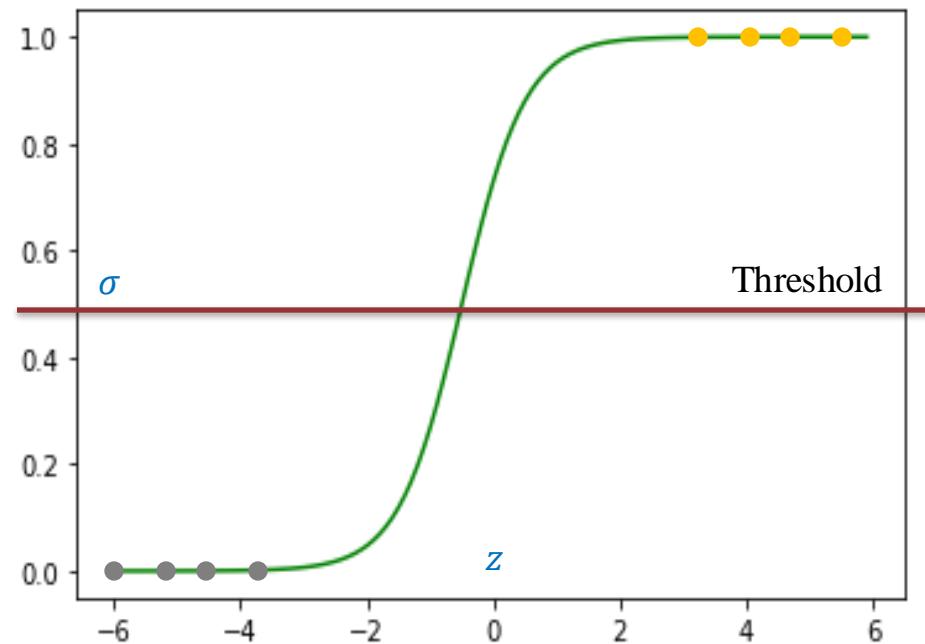
## Problem

Sigmoid function

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

$z \in (-\infty, +\infty)$

$\sigma(z) \in (0, 1)$



# Softmax Regression

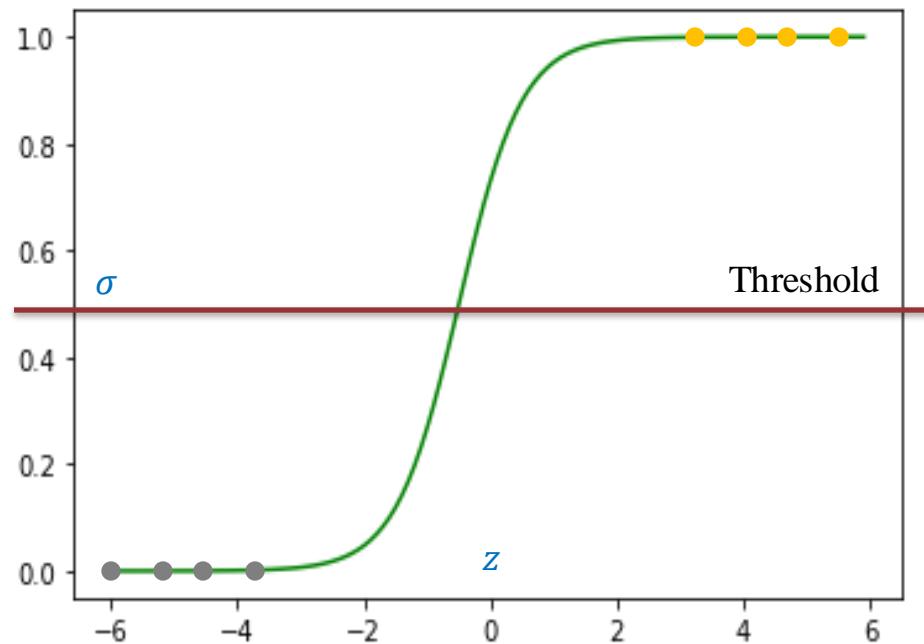


## Problem

Sigmoid function

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

$$z \in (-\infty, +\infty) \quad \sigma(z) \in (0, 1)$$



| Hours | Pass |
|-------|------|
| 0.5   | 0    |
| 1.0   | 0    |
| 1.5   | 1    |
| 2.0   | 1    |

Classes: {0, 1}  
Binary Classification

| Hours | Score |
|-------|-------|
| 0.5   | 0     |
| 1.0   | 0     |
| 1.5   | 1     |
| 2.0   | 1     |
| 2.5   | 2     |
| 3.0   | 2     |
| 3.5   | 3     |
| 4.0   | 3     |

Classes: {0, 1, 2, 3}  
Multi-class Classification

# Softmax Regression

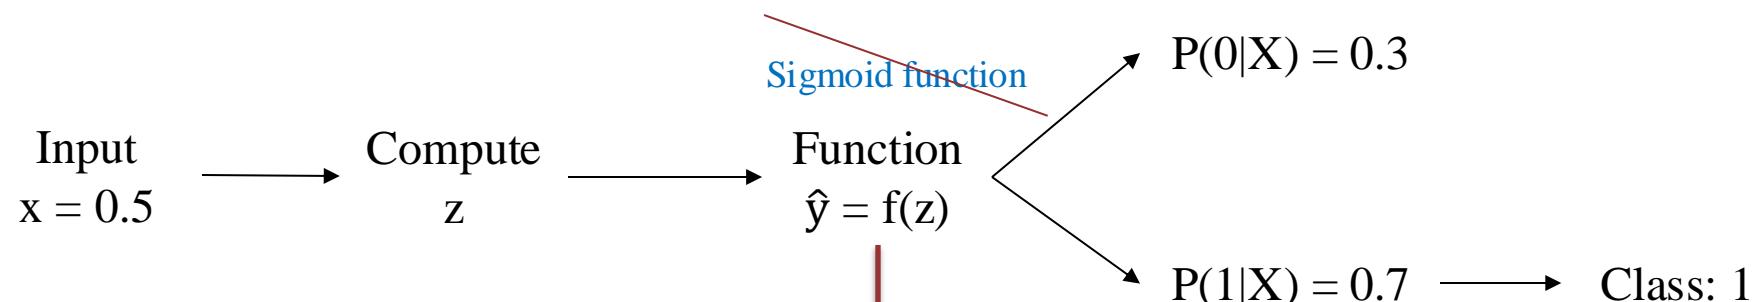


## Problem

Classes: {0, 1}

Binary Classification

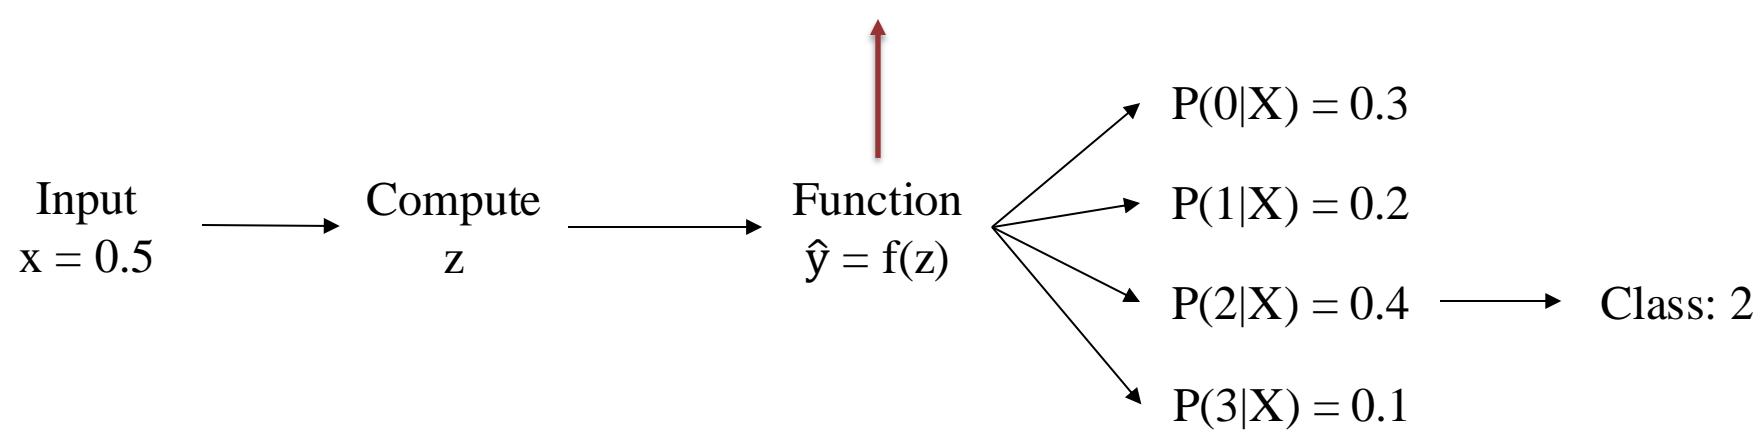
| Hours | Pass |
|-------|------|
| 0.5   | 0    |
| 2.0   | 1    |



Classes: {0, 1, 2, 3}

Multi-class Classification

| Hours | Score |
|-------|-------|
| 0.5   | 0     |
| 1.5   | 1     |
| 3.0   | 2     |
| 4.0   | 3     |



# Softmax Regression

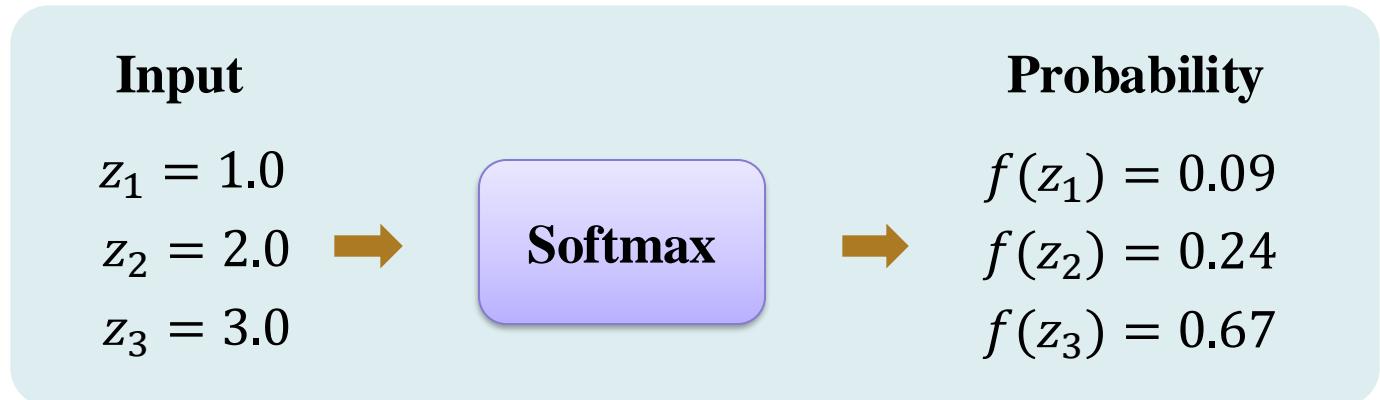
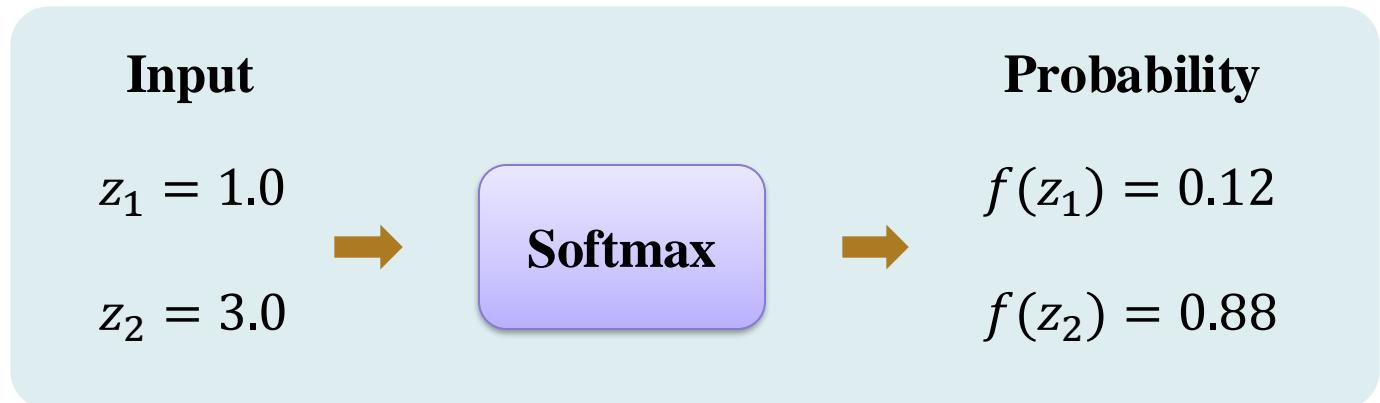


## Softmax Function

$$P_i = f(z_i) = \frac{e^{z_i}}{\sum_j e^{z_j}}$$

$$0 \leq f(z_i) \leq 1$$

$$\sum_i f(z_i) = 1$$



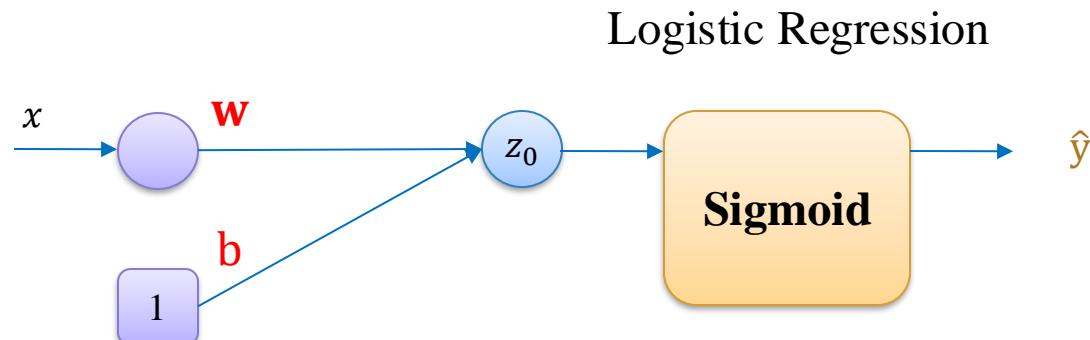
# Softmax Regression



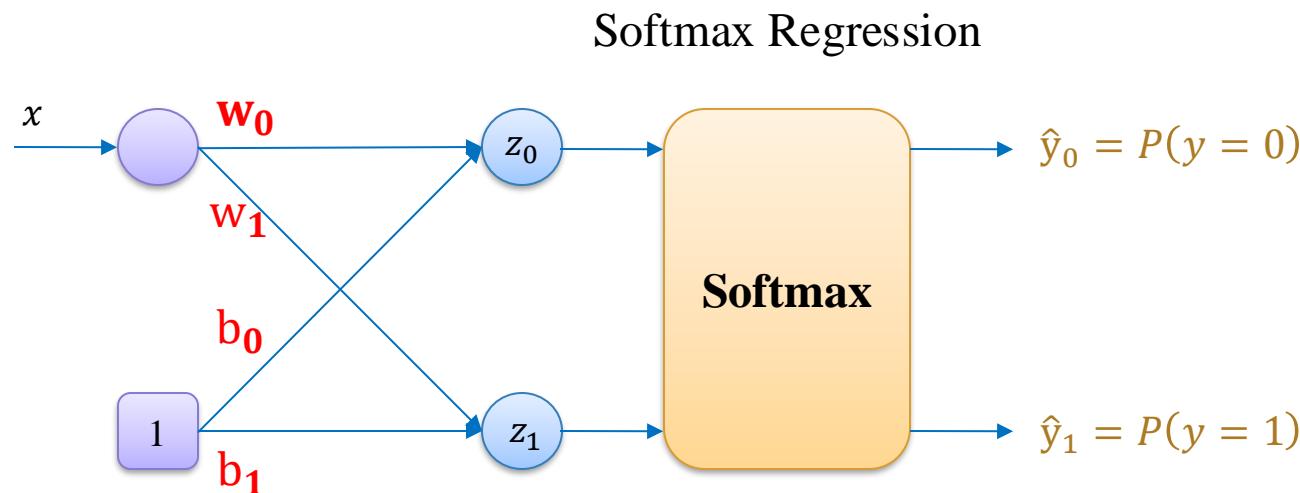
## Parameters

Classes: {0, 1}  
Binary Classification

| Hours | Pass |
|-------|------|
| 0.5   | 0    |
| 2.0   | 1    |



#feature: 1  
#class: 2



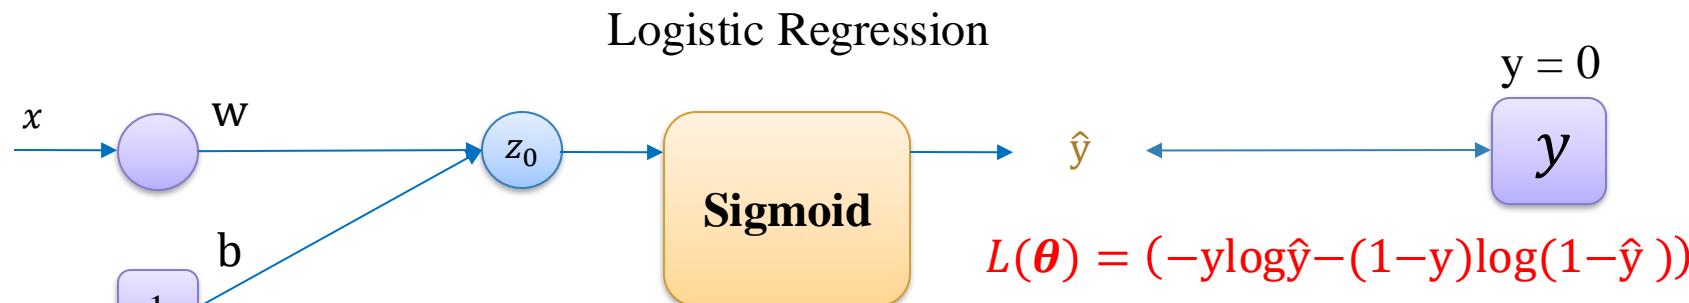
# Softmax Regression



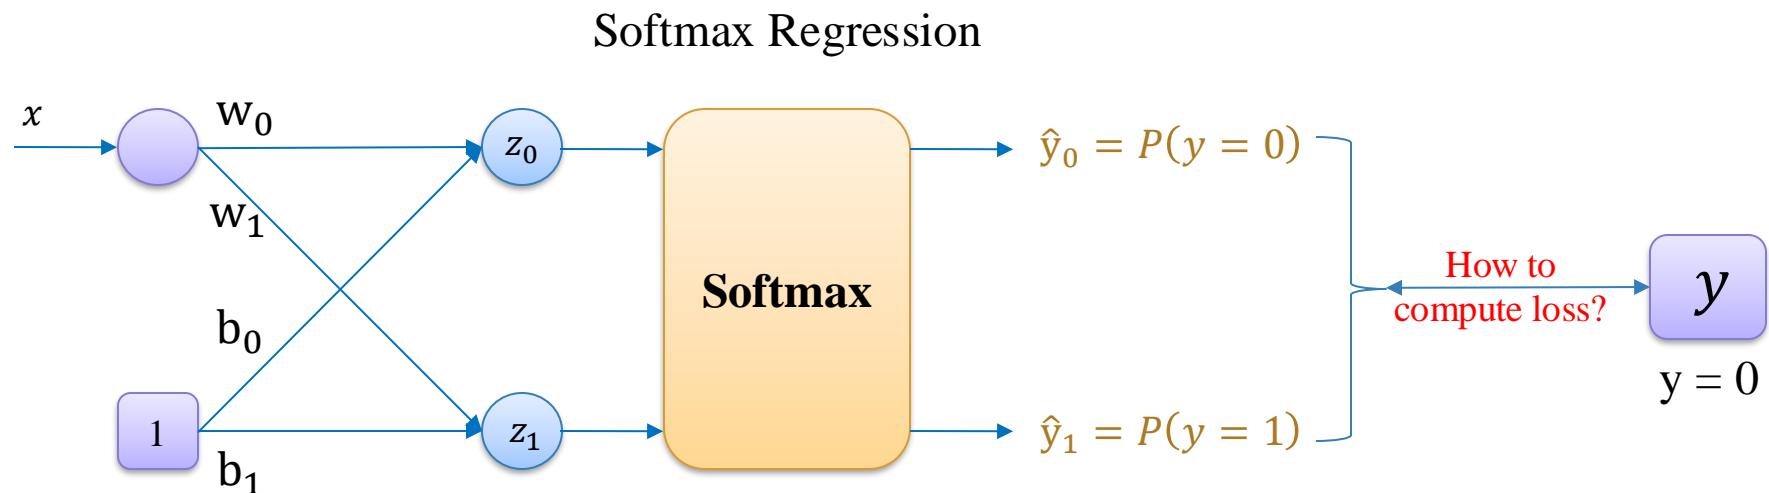
## Loss Function

Classes: {0, 1}  
Binary Classification

| Hours | Pass |
|-------|------|
| 0.5   | 0    |
| 2.0   | 1    |



#feature: 1  
#class: 2



# Softmax Regression



## One-hot Encoding

$$\mathbf{y} = \begin{bmatrix} y_0 \\ \dots \\ y_C \end{bmatrix} \quad y_i \in \{0,1\}$$

$$\sum_i y_i = 1 \quad C = \# \text{classes}$$

Classes: {0, 1}  
Binary Classification

| Hours | Pass |
|-------|------|
| 0.5   | 0    |
| 2.0   | 1    |

#feature: 1  
#class: 2

$$y = 0 \rightarrow \mathbf{y} = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

$$y = 1 \rightarrow \mathbf{y} = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

Classes: {0, 1, 2}  
Multi-class Classification

| Hours | Score |
|-------|-------|
| 0.5   | 0     |
| 1.5   | 1     |
| 3.0   | 2     |

#feature: 1  
#class: 3

$$y = 0 \rightarrow \mathbf{y} = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}$$

$$y = 1 \rightarrow \mathbf{y} = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}$$

$$y = 2 \rightarrow \mathbf{y} = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$$

# Softmax Regression

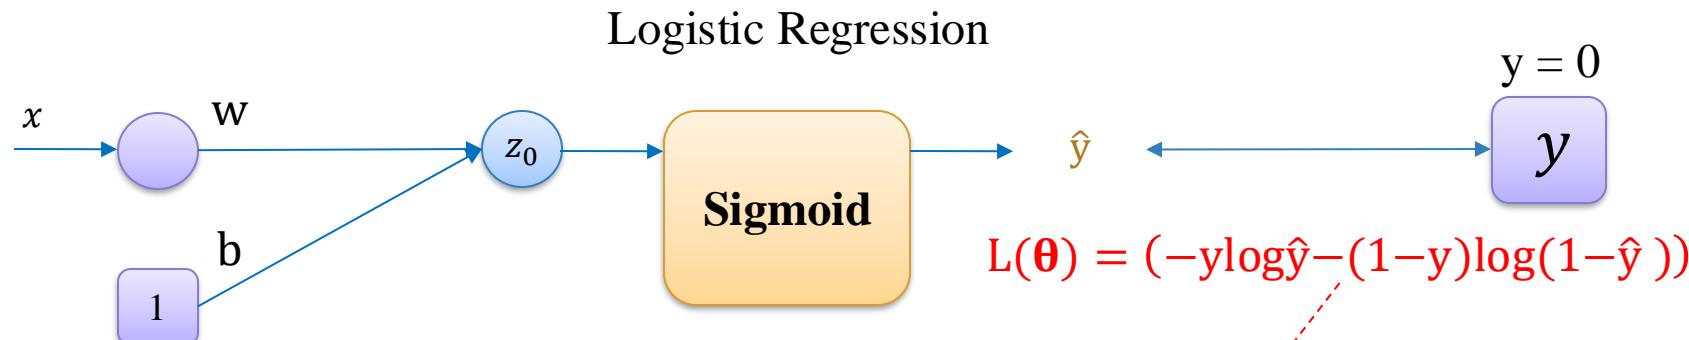


## Loss Function

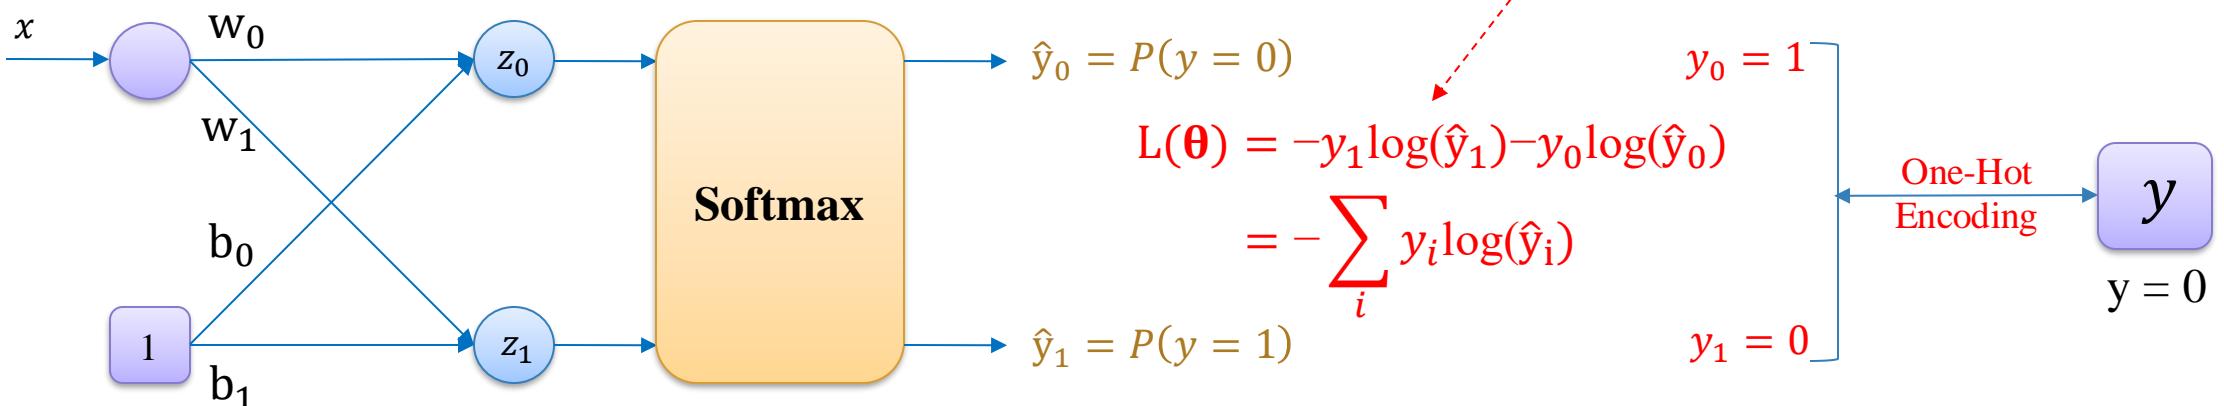
Classes: {0, 1}  
Binary Classification

| Hours | Pass |
|-------|------|
| 0.5   | 0    |
| 2.0   | 1    |

#feature: 1  
#class: 2



Softmax Regression



# Softmax Regression



## Softmax Regression using Gradient Descent

- 1) Pick a sample from training data
- 2) Compute output  $\hat{y}$

$$\mathbf{z} = \boldsymbol{\theta}^T \mathbf{x}$$

$$\mathbf{d} = [1 \dots 1] \mathbf{e}^{\mathbf{z}}$$

$\emptyset$  is  
Hadamard  
division

$$\hat{y} = \mathbf{e}^{\mathbf{z}} \emptyset \mathbf{d}$$

- 3) Compute loss (cross-entropy)

$$L(\boldsymbol{\theta}) = -\mathbf{y}^T \log \hat{\mathbf{y}}$$

- 4) Compute derivative

$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{\mathbf{y}} - \mathbf{y})^T$$

- 5) Update parameters

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L$$

$\eta$  is learning rate

$$\mathbf{x} = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix} \quad \longleftarrow$$

$$\mathbf{y} = [0]$$

One-hot encoding for label

$$\mathbf{y} = 0 \rightarrow \mathbf{y} = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

$$\mathbf{y} = 1 \rightarrow \mathbf{y} = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

$$\boldsymbol{\theta} = \begin{bmatrix} b_0 & b_1 \\ w_0 & w_1 \end{bmatrix} = \begin{bmatrix} 0.1 & 0.2 \\ 0.3 & 0.4 \end{bmatrix}$$

$$\eta = 0.1$$

| Hours | Pass |
|-------|------|
| 0.5   | 0    |
| 1.0   | 0    |
| 1.5   | 1    |
| 2.0   | 1    |

# Softmax Regression



## Softmax Regression using Gradient Descent

- 1) Pick a sample from training data
- 2) Compute output  $\hat{y}$

$$\mathbf{z} = \boldsymbol{\theta}^T \mathbf{x}$$

$$\mathbf{d} = [1 \dots 1]e^{\mathbf{z}}$$

$\emptyset$  is  
Hadamard  
division

$$\hat{\mathbf{y}} = e^{\mathbf{z}} \emptyset \mathbf{d}$$

- 3) Compute loss (cross-entropy)

$$L(\boldsymbol{\theta}) = -\mathbf{y}^T \log \hat{\mathbf{y}}$$

- 4) Compute derivative

$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{\mathbf{y}} - \mathbf{y})^T$$

- 5) Update parameters

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L$$

$\eta$  is learning rate

$$\eta = 0.1 \quad \mathbf{x} = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix} \quad \mathbf{y} = 0 \rightarrow \mathbf{y} = \begin{bmatrix} 1 \\ 0 \end{bmatrix} \quad \boldsymbol{\theta} = \begin{bmatrix} 0.1 & 0.2 \\ 0.3 & 0.4 \end{bmatrix}$$

# Softmax Regression



## Softmax Regression using Gradient Descent

- 1) Pick a sample from training data
- 2) Compute output  $\hat{y}$

$$\mathbf{z} = \boldsymbol{\theta}^T \mathbf{x}$$

$$\mathbf{d} = [1 \dots 1]e^{\mathbf{z}}$$

$$\hat{y} = e^{\mathbf{z}} \emptyset \mathbf{d}$$

$\emptyset$  is  
Hadamard  
division

- 3) Compute loss (cross-entropy)

$$L(\boldsymbol{\theta}) = -\mathbf{y}^T \log \hat{\mathbf{y}}$$

- 4) Compute derivative

$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{\mathbf{y}} - \mathbf{y})^T$$

- 5) Update parameters

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L$$

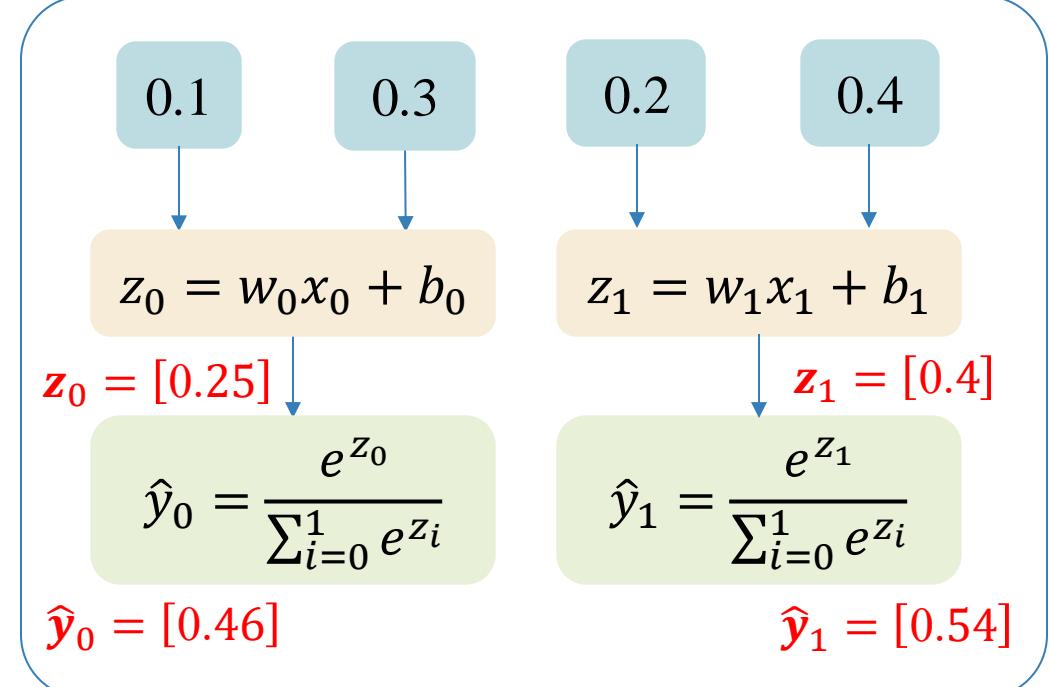
$\eta$  is learning rate

$$\eta = 0.1$$

$$\mathbf{x} = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix}$$

$$\mathbf{y} = 0 \rightarrow \mathbf{y} = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

$$\boldsymbol{\theta} = \begin{bmatrix} 0.1 & 0.2 \\ 0.3 & 0.4 \end{bmatrix}$$



# Softmax Regression



## Softmax Regression using Gradient Descent

- 1) Pick a sample from training data
- 2) Compute output  $\hat{y}$

$$\mathbf{z} = \boldsymbol{\theta}^T \mathbf{x}$$

$$\mathbf{d} = [1 \dots 1]e^{\mathbf{z}}$$

$$\hat{y} = e^{\mathbf{z}} \emptyset \mathbf{d}$$

$\emptyset$  is  
Hadamard  
division

- 3) Compute loss (cross-entropy)

$$L(\boldsymbol{\theta}) = -\mathbf{y}^T \log \hat{\mathbf{y}}$$

- 4) Compute derivative

$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{\mathbf{y}} - \mathbf{y})^T$$

- 5) Update parameters

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L$$

$\eta$  is learning rate

$$\eta = 0.1$$

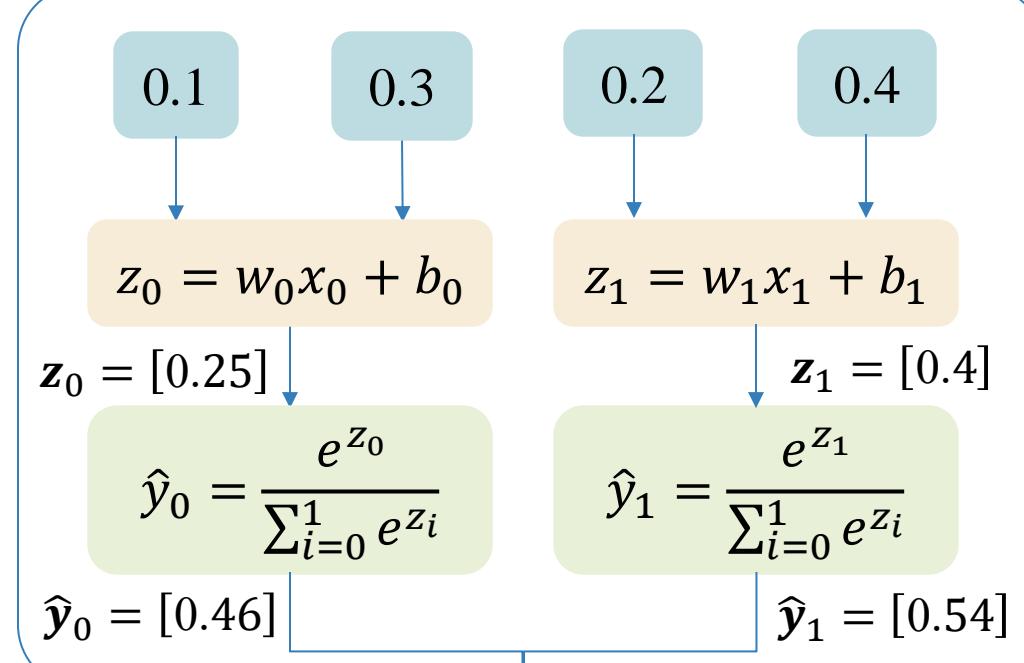
$$\mathbf{x} = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix}$$

$$\mathbf{y} = 0 \rightarrow \mathbf{y} = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

$$\boldsymbol{\theta} = \begin{bmatrix} 0.1 & 0.2 \\ 0.3 & 0.4 \end{bmatrix}$$

$$L = [0.77]$$

$$L = -y_0 \log \hat{y}_0 - y_1 \log \hat{y}_1$$



# Softmax Regression



## Softmax Regression using Gradient Descent

- 1) Pick a sample from training data
- 2) Compute output  $\hat{y}$

$$\mathbf{z} = \boldsymbol{\theta}^T \mathbf{x}$$

$$\mathbf{d} = [1 \dots 1]e^{\mathbf{z}}$$

$$\hat{y} = e^{\mathbf{z}} \emptyset \mathbf{d}$$

$\emptyset$  is  
Hadamard  
division

- 3) Compute loss (cross-entropy)

$$L(\boldsymbol{\theta}) = -\mathbf{y}^T \log \hat{\mathbf{y}}$$

- 4) Compute derivative

$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{\mathbf{y}} - \mathbf{y})^T$$

- 5) Update parameters

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L$$

$\eta$  is learning rate

$$\eta = 0.1$$

$$\mathbf{x} = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix}$$

$$\mathbf{y} = 0 \rightarrow \mathbf{y} = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

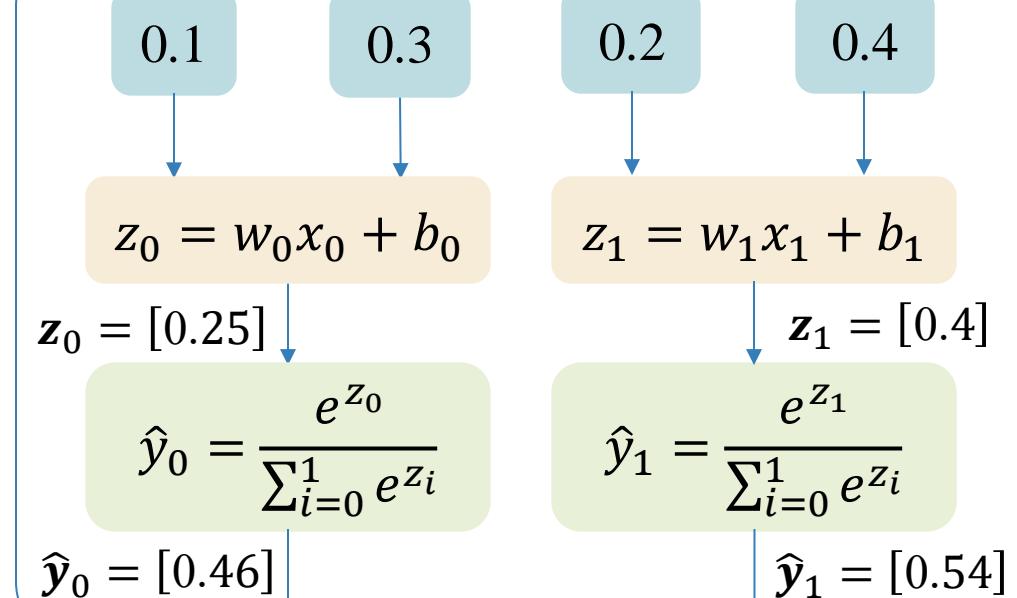
$$\boldsymbol{\theta} = \begin{bmatrix} 0.1 & 0.2 \\ 0.3 & 0.4 \end{bmatrix}$$

$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{\mathbf{y}} - \mathbf{y})^T$$

$$= \begin{bmatrix} -0.54 & 0.54 \\ -0.27 & 0.27 \end{bmatrix}$$

$$L = [0.77]$$

$$L = -y_0 \log \hat{y}_0 - y_1 \log \hat{y}_1$$



# Softmax Regression



## Softmax Regression using Gradient Descent

- 1) Pick a sample from training data
- 2) Compute output  $\hat{y}$

$$\mathbf{z} = \boldsymbol{\theta}^T \mathbf{x}$$

$$\mathbf{d} = [1 \dots 1]e^{\mathbf{z}}$$

$$\hat{y} = e^{\mathbf{z}} \emptyset \mathbf{d}$$

$\emptyset$  is Hadamard division

- 3) Compute loss (cross-entropy)

$$L(\boldsymbol{\theta}) = -\mathbf{y}^T \log \hat{\mathbf{y}}$$

- 4) Compute derivative

$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{\mathbf{y}} - \mathbf{y})^T$$

- 5) Update parameters

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L$$

$\eta$  is learning rate

$$\eta = 0.1$$

$$\mathbf{x} = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix}$$

$$\mathbf{y} = 0 \rightarrow \mathbf{y} = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

$$\boldsymbol{\theta} = \begin{bmatrix} 0.1 & 0.2 \\ 0.3 & 0.4 \end{bmatrix}$$

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L$$

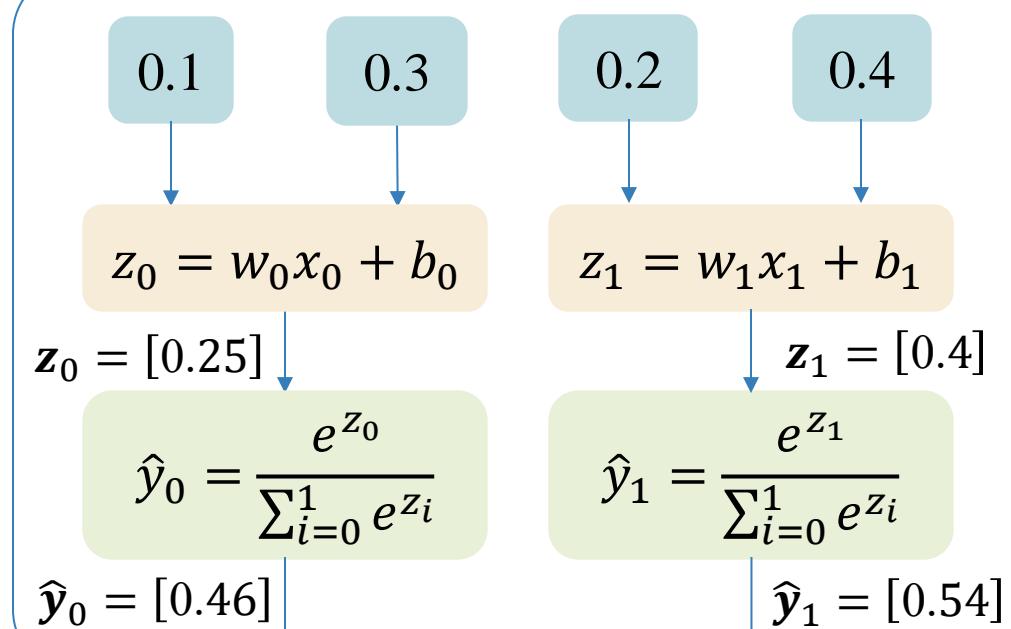
$$= \begin{bmatrix} 0.105 & 0.195 \\ 0.303 & 0.397 \end{bmatrix}$$

$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{\mathbf{y}} - \mathbf{y})^T$$

$$= \begin{bmatrix} -0.54 & 0.54 \\ -0.27 & 0.27 \end{bmatrix}$$

$$L = [0.77]$$

$$L = -y_0 \log \hat{y}_0 - y_1 \log \hat{y}_1$$



# Softmax Regression



## Softmax Regression using Gradient Descent

- 1) Pick a sample from training data
- 2) Compute output  $\hat{y}$

$$\mathbf{z} = \boldsymbol{\theta}^T \mathbf{x}$$

$$\mathbf{d} = [1 \dots 1] \mathbf{e}^{\mathbf{z}}$$

$\emptyset$  is  
Hadamard  
division

$$\hat{y} = \mathbf{e}^{\mathbf{z}} \emptyset \mathbf{d}$$

- 3) Compute loss (cross-entropy)

$$L(\boldsymbol{\theta}) = -\mathbf{y}^T \log \hat{\mathbf{y}}$$

- 4) Compute derivative

$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{\mathbf{y}} - \mathbf{y})^T$$

- 5) Update parameters

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L$$

$\eta$  is learning rate

$$\mathbf{x} = \begin{bmatrix} 1.0 \\ 1.5 \end{bmatrix} \quad \longleftarrow$$

$$\mathbf{y} = [1]$$

One-hot encoding for label

$$\mathbf{y} = 0 \rightarrow \mathbf{y} = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

$$\mathbf{y} = 1 \rightarrow \mathbf{y} = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

$$\boldsymbol{\theta} = \begin{bmatrix} b_0 & b_1 \\ w_0 & w_1 \end{bmatrix} = \begin{bmatrix} 0.1 & 0.2 \\ 0.3 & 0.4 \end{bmatrix}$$

$$\eta = 0.1$$

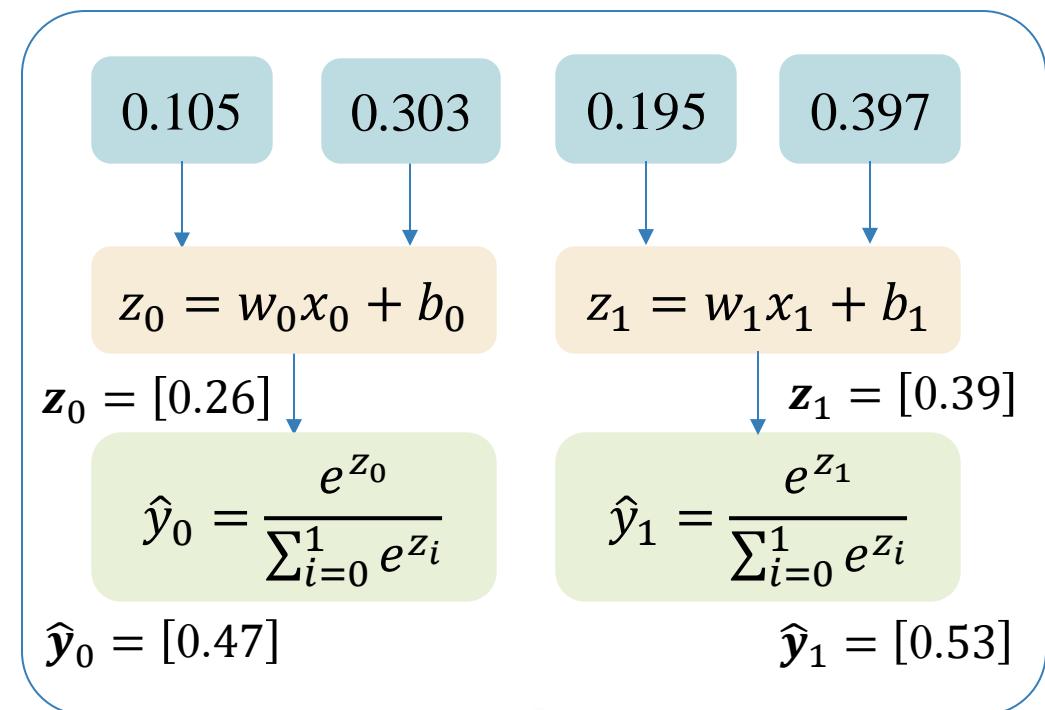
| Hours | Pass |
|-------|------|
| 0.5   | 0    |
| 1.0   | 0    |
| 1.5   | 1    |
| 2.0   | 1    |

# Softmax Regression



## Softmax Regression using Gradient Descent

| Hours | Pass |
|-------|------|
| 0.5   | ?    |



**QUIZ TIME**

# Softmax Regression



## Softmax Regression using Gradient Descent

- 1) Pick a sample from training data
- 2) Compute output  $\hat{y}$

$$\mathbf{z} = \boldsymbol{\theta}^T \mathbf{x}$$

$$\mathbf{d} = [1 \dots 1]e^{\mathbf{z}}$$

$\emptyset$  is  
Hadamard  
division

$$\hat{y} = e^{\mathbf{z}} \emptyset \mathbf{d}$$

- 3) Compute loss (cross-entropy)

$$L(\boldsymbol{\theta}) = -\mathbf{y}^T \log \hat{\mathbf{y}}$$

- 4) Compute derivative

$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{\mathbf{y}} - \mathbf{y})^T$$

- 5) Update parameters

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L$$

$\eta$  is learning rate

```
1 X = np.array([
2     [0.5],
3     [1.0],
4     [1.5],
5     [2.0],
6     [2.5],
7     [3.0],
8     [3.5],
9     [4.0],
10 ])
11 Y = np.array([0, 0, 1, 0, 0, 1, 1, 1])
```

```
1 def convert_one_hot(y, k):
2     one_hot = np.zeros(len(y), k)
3     one_hot[np.arange(len(y)), y] = 1
4     return one_hot
5
6 n_classes = 2
7 Y_onehot = convert_one_hot(Y, n_classes)
```

```
1 X = np.hstack([np.ones(X.shape[0], 1), X])
```

# Softmax Regression



## Softmax Regression using Gradient Descent

- 1) Pick a sample from training data
- 2) Compute output  $\hat{y}$

$$\mathbf{z} = \boldsymbol{\theta}^T \mathbf{x}$$

$$\mathbf{d} = [1 \dots 1]e^{\mathbf{z}}$$

$\emptyset$  is  
Hadamard  
division

$$\hat{y} = e^{\mathbf{z}} \emptyset \mathbf{d}$$

- 3) Compute loss (cross-entropy)

$$L(\boldsymbol{\theta}) = -\mathbf{y}^T \log \hat{\mathbf{y}}$$

- 4) Compute derivative

$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{\mathbf{y}} - \mathbf{y})^T$$

- 5) Update parameters

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L$$

$\eta$  is learning rate

$$\eta = 0.1 \quad \mathbf{x} = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix} \quad \mathbf{y} = 0 \rightarrow \mathbf{y} = \begin{bmatrix} 1 \\ 0 \end{bmatrix} \quad \boldsymbol{\theta} = \begin{bmatrix} 0.1 & 0.2 \\ 0.3 & 0.4 \end{bmatrix}$$

```
1 x = X[0].reshape(2,1)
2 y = Y_onehot[0].reshape(2,1)
3 x, y
```

```
(array([[1. ,  
       0.5]]),  
 array([[1.],  
       [0.]]))
```

```
1 theta = np.array([[0.1, 0.2], [0.3, 0.4]])
2 theta
```

```
array([[0.1, 0.2],  
       [0.3, 0.4]])
```

# Softmax Regression



## Softmax Regression using Gradient Descent

- 1) Pick a sample from training data
- 2) Compute output  $\hat{y}$

$$\mathbf{z} = \boldsymbol{\theta}^T \mathbf{x}$$

$$\mathbf{d} = [1 \dots 1] \mathbf{e}^{\mathbf{z}} \quad \emptyset \text{ is Hadamard division}$$

$$\hat{\mathbf{y}} = \mathbf{e}^{\mathbf{z}} \emptyset \mathbf{d}$$

- 3) Compute loss (cross-entropy)

$$L(\boldsymbol{\theta}) = -\mathbf{y}^T \log \hat{\mathbf{y}}$$

- 4) Compute derivative

$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{\mathbf{y}} - \mathbf{y})^T$$

- 5) Update parameters

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L$$

$\eta$  is learning rate

$$\eta = 0.1 \quad \mathbf{x} = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix} \quad \mathbf{y} = 0 \rightarrow \mathbf{y} = \begin{bmatrix} 1 \\ 0 \end{bmatrix} \quad \boldsymbol{\theta} = \begin{bmatrix} 0.1 & 0.2 \\ 0.3 & 0.4 \end{bmatrix}$$

```
1 # define softmax function
2 def softmax_function(z):
3     return np.exp(z) / np.sum(np.exp(z))
```

```
1 # compute y_hat
2 def predict(x, theta):
3     z = theta.T.dot(x)
4     y_hat = softmax_function(z)
5     return z, y_hat
6
7 z, y_hat = predict(x, theta)
8 z, y_hat
```

```
(array([[0.25],
       [0.4]]),
array([[0.46257015],
       [0.53742985]]))
```

# Softmax Regression



## Softmax Regression using Gradient Descent

- 1) Pick a sample from training data
- 2) Compute output  $\hat{y}$

$$\mathbf{z} = \boldsymbol{\theta}^T \mathbf{x}$$

$$\mathbf{d} = [1 \dots 1]e^{\mathbf{z}}$$

$\emptyset$  is  
Hadamard  
division

$$\hat{\mathbf{y}} = e^{\mathbf{z}} \emptyset \mathbf{d}$$

- 3) Compute loss (cross-entropy)

$$L(\boldsymbol{\theta}) = -\mathbf{y}^T \log \hat{\mathbf{y}}$$

- 4) Compute derivative

$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{\mathbf{y}} - \mathbf{y})^T$$

- 5) Update parameters

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L$$

$\eta$  is learning rate

$$\eta = 0.1 \quad \mathbf{x} = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix} \quad \mathbf{y} = 0 \rightarrow \mathbf{y} = \begin{bmatrix} 1 \\ 0 \end{bmatrix} \quad \boldsymbol{\theta} = \begin{bmatrix} 0.1 & 0.2 \\ 0.3 & 0.4 \end{bmatrix}$$

```
1 # compute loss
2 def compute_loss(y_hat, y):
3     loss = -np.log(y.T.dot(y_hat))
4     return loss
5
6 loss = compute_loss(y_hat, y)
7 loss
```

array([[0.77095705]])

# Softmax Regression



## Softmax Regression using Gradient Descent

- 1) Pick a sample from training data
- 2) Compute output  $\hat{y}$

$$\mathbf{z} = \boldsymbol{\theta}^T \mathbf{x}$$

$$\begin{aligned}\mathbf{d} &= [1 \dots 1]\mathbf{e}^{\mathbf{z}} & \emptyset \text{ is} \\ \hat{\mathbf{y}} &= \mathbf{e}^{\mathbf{z}} \emptyset \mathbf{d} & \text{Hadamard} \\ && \text{division}\end{aligned}$$

- 3) Compute loss (cross-entropy)

$$L(\boldsymbol{\theta}) = -\mathbf{y}^T \log \hat{\mathbf{y}}$$

- 4) Compute derivative

$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{\mathbf{y}} - \mathbf{y})^T$$

- 5) Update parameters

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L$$

$\eta$  is learning rate

$$\eta = 0.1 \quad \mathbf{x} = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix} \quad \mathbf{y} = 0 \rightarrow \mathbf{y} = \begin{bmatrix} 1 \\ 0 \end{bmatrix} \quad \boldsymbol{\theta} = \begin{bmatrix} 0.1 & 0.2 \\ 0.3 & 0.4 \end{bmatrix}$$

```
1 # compute gradient
2 def compute_gradient(y_hat, y, x):
3     gradient = x.dot((y_hat - y).T)
4     return gradient
5
6 gradient = compute_gradient(y_hat, y, x)
7 gradient
```

```
array([[-0.53742985,  0.53742985],
       [-0.26871492,  0.26871492]])
```

# Softmax Regression



## Softmax Regression using Gradient Descent

- 1) Pick a sample from training data
- 2) Compute output  $\hat{y}$

$$\mathbf{z} = \boldsymbol{\theta}^T \mathbf{x}$$

$$\mathbf{d} = [1 \dots 1]e^{\mathbf{z}}$$

$\emptyset$  is  
Hadamard  
division

$$\hat{\mathbf{y}} = e^{\mathbf{z}} \emptyset \mathbf{d}$$

- 3) Compute loss (cross-entropy)

$$L(\boldsymbol{\theta}) = -\mathbf{y}^T \log \hat{\mathbf{y}}$$

- 4) Compute derivative

$$\nabla_{\boldsymbol{\theta}} L = \mathbf{x}(\hat{\mathbf{y}} - \mathbf{y})^T$$

- 5) Update parameters

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L$$

$\eta$  is learning rate

$$\eta = 0.1 \quad \mathbf{x} = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix} \quad \mathbf{y} = 0 \rightarrow \mathbf{y} = \begin{bmatrix} 1 \\ 0 \end{bmatrix} \quad \boldsymbol{\theta} = \begin{bmatrix} 0.1 & 0.2 \\ 0.3 & 0.4 \end{bmatrix}$$

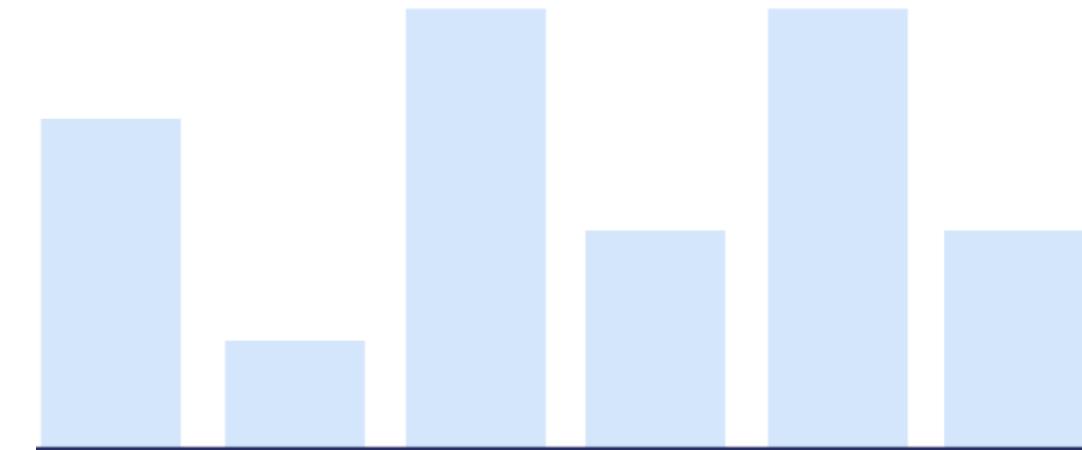
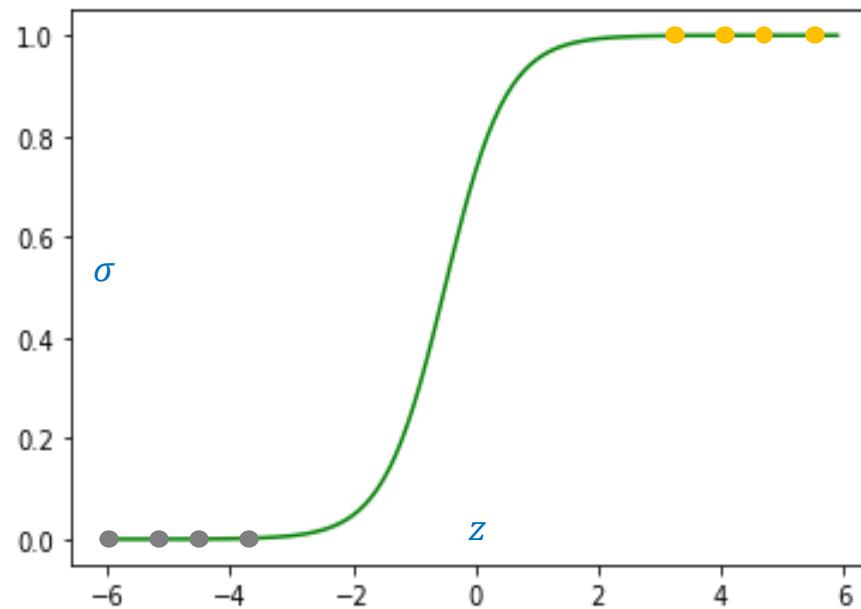
```
1 # update weights
2 learning_rate = 0.01
3
4 def update_weight(theta, gradient, learning_rate):
5     theta -= (learning_rate * gradient)
6     return theta
7
8 theta = update_weight(theta, gradient, learning_rate)
9 theta
```

```
array([[0.1053743 , 0.1946257 ],
       [0.30268715, 0.39731285]])
```

# Objectives

## Logistic Regression (Review)

- ❖ Logistic Regression
- ❖ Sigmoid Function
- ❖ Gradient Descent



## Softmax Regression

- ❖ Softmax Regression
- ❖ Softmax Function
- ❖ One Sample
- ❖ N Sample



# Thanks!

Any questions?