

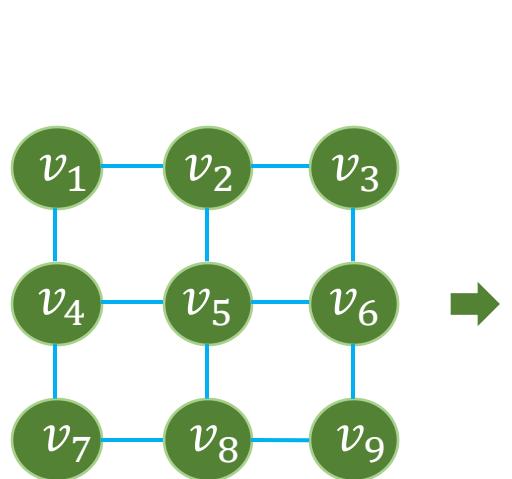
From MLP to CNN

Introduction to CNN

Quang-Vinh Dinh
Ph.D. in Computer Science

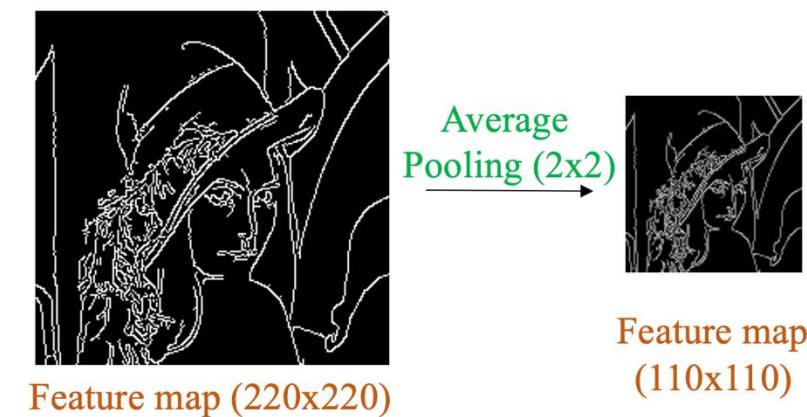
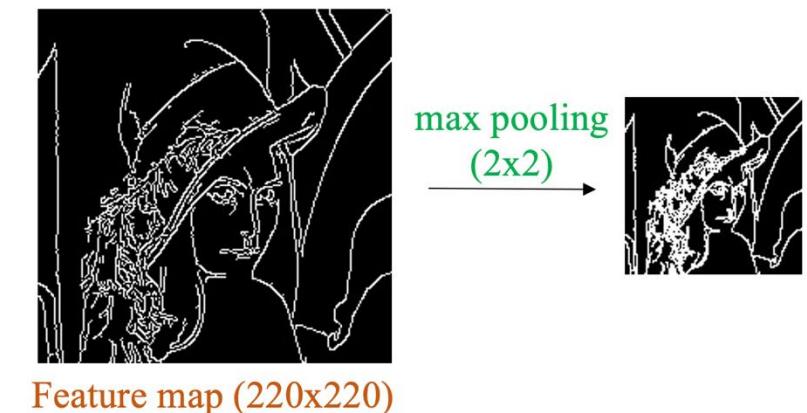
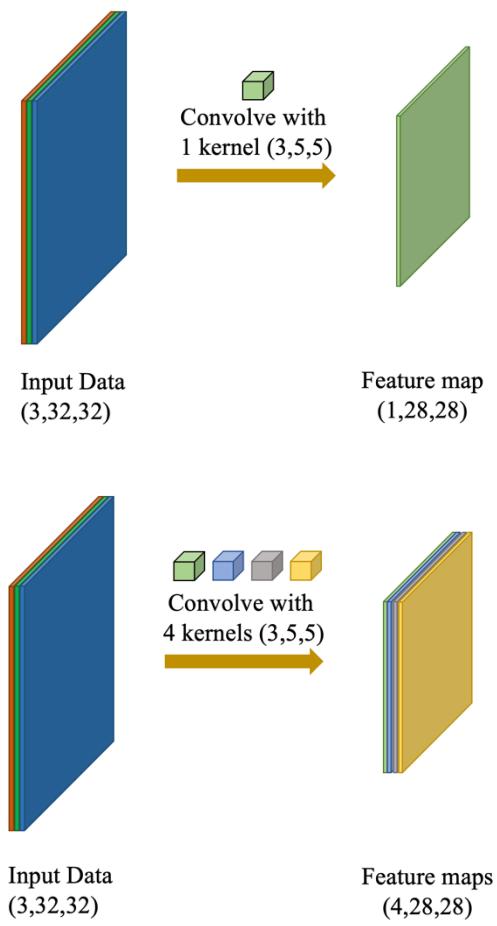
Objectives

MLP Limitations



Flattened Data

Convolutional Layer



Standard CNNs

Outline

SECTION 1

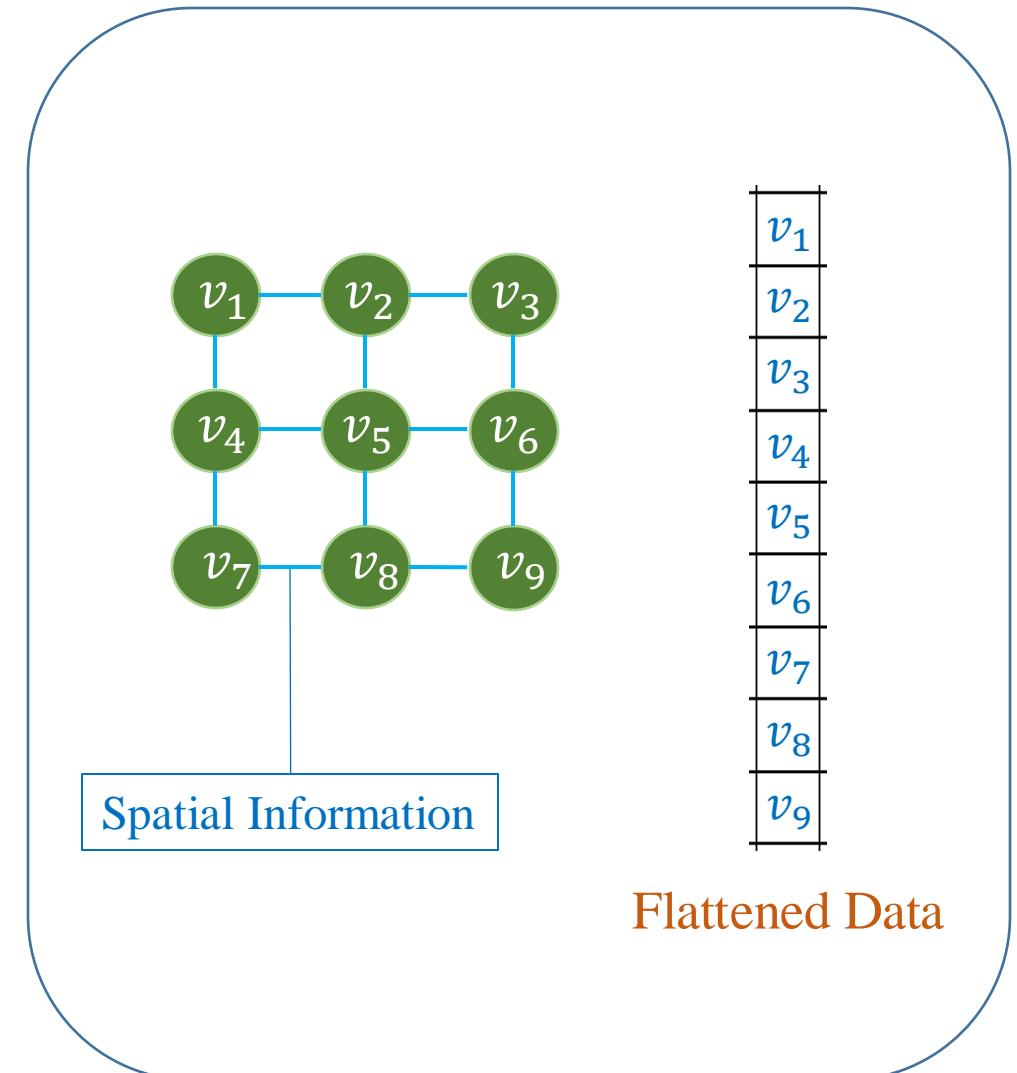
MLP Limitations

SECTION 2

Convolutional Layer

SECTION 3

Standard CNNs



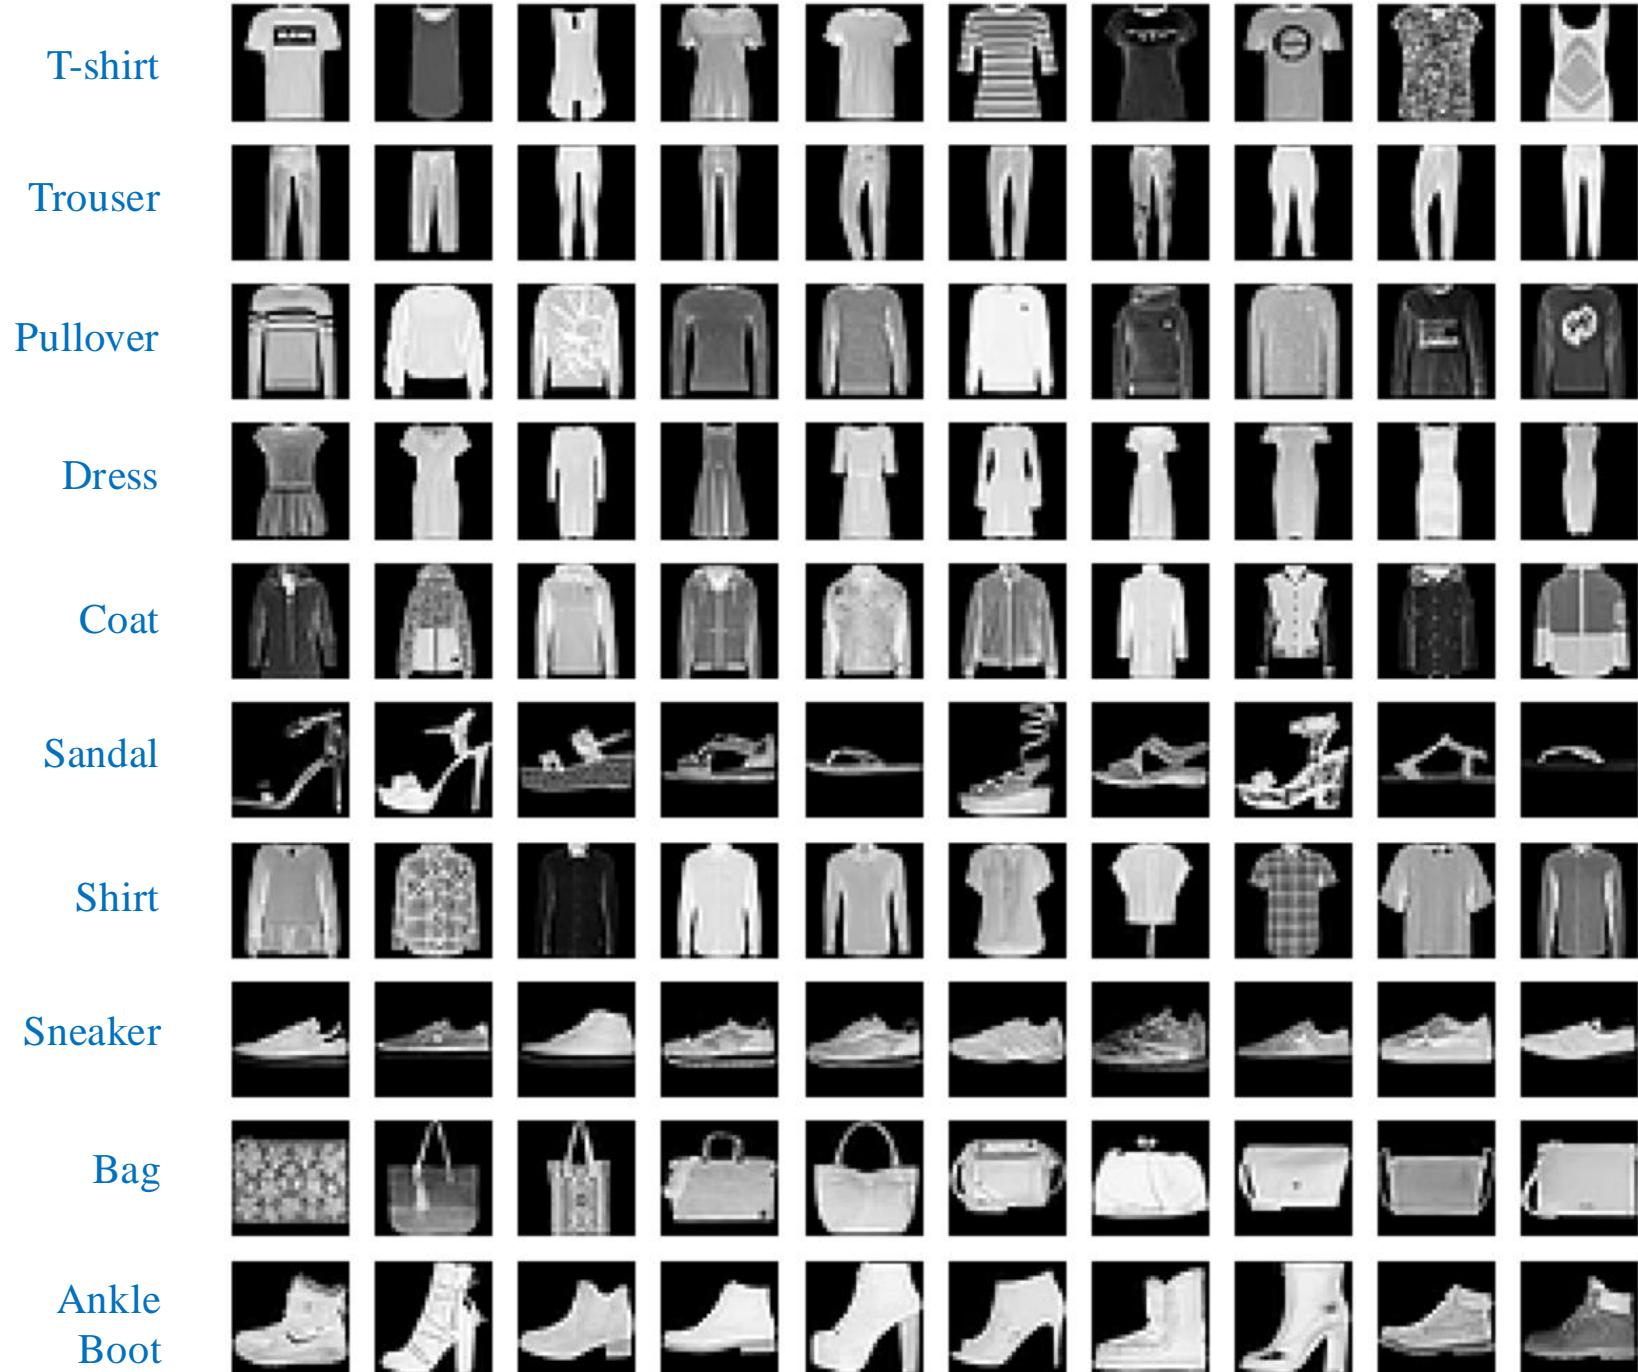
Fashion-MNIST dataset

Grayscale images

Resolution=28x28

Training set: 60000 samples

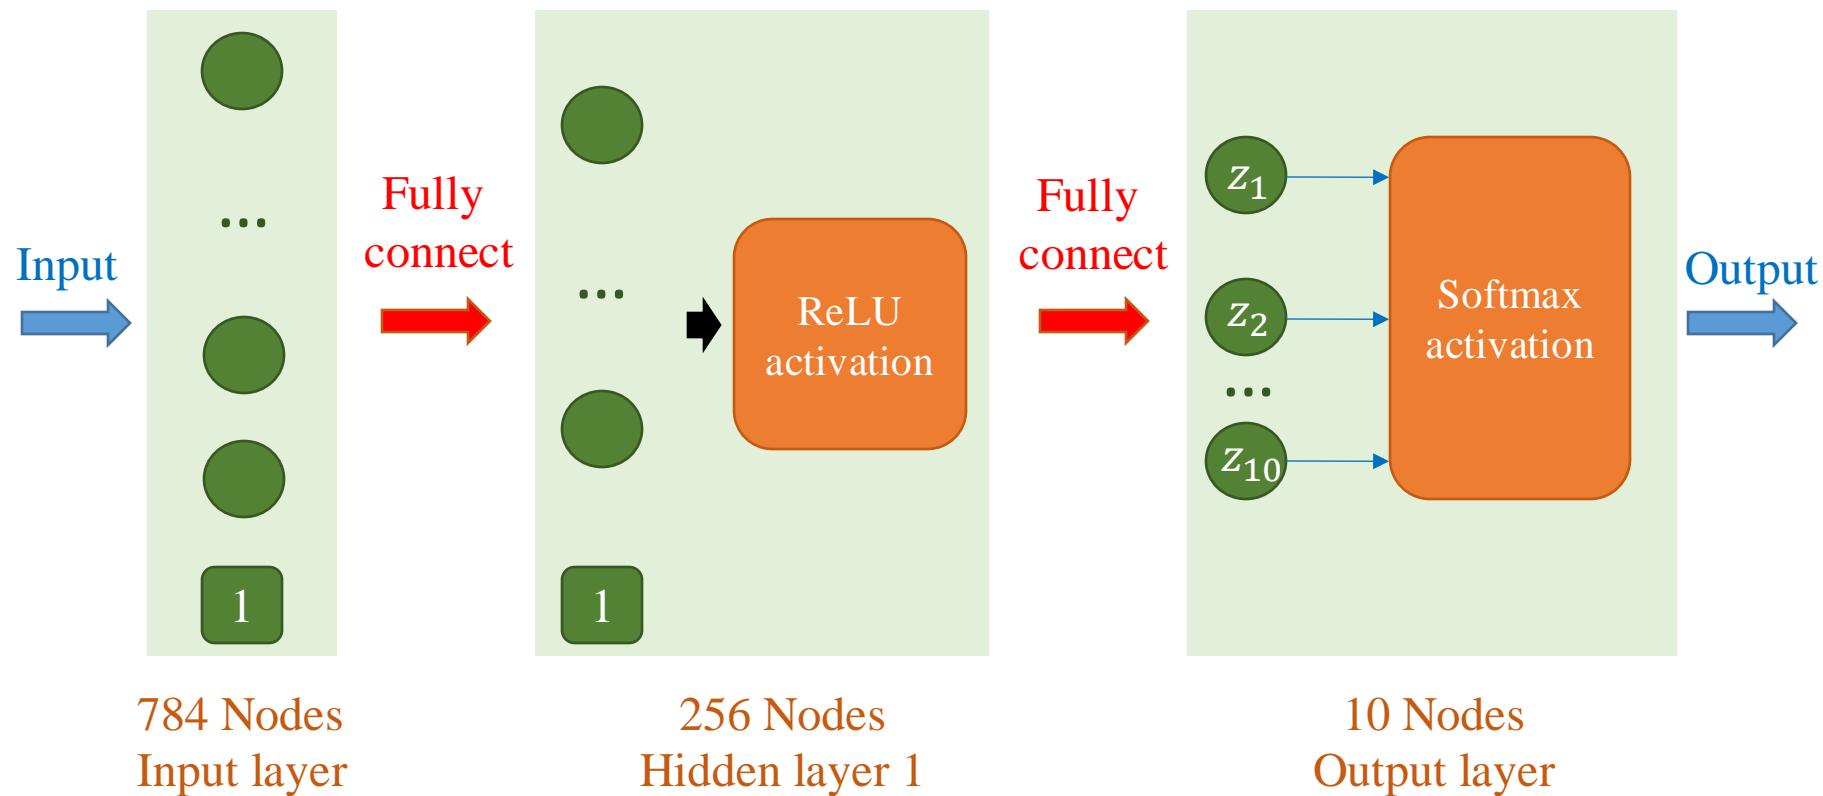
Testing set: 10000 samples



MLP for Fashion-MNIST

❖ ReLU, He and Adam

Case 1



❖ ReLU, He and Adam

```
# model
model = nn.Sequential(
    nn.Flatten(),
    nn.Linear(784, 256),
    nn.ReLU(),
    nn.Linear(256, 10)
)
```

```
# Initialize the weights
for layer in model:
    if isinstance(layer, nn.Linear):
        init.kaiming_uniform_(layer.weight,
                              nonlinearity='relu')
    if layer.bias is not None:
        layer.bias.data.fill_(0)
```

```
# Loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(),
                      lr=0.001)
```

```
# Load CFashionMNIST dataset
transform = Compose([transforms.ToTensor(),
                     transforms.Normalize((0.5, ),
                                         (0.5,))])

trainset = FashionMNIST(root='data',
                        train=True,
                        download=True,
                        transform=transform)

trainloader = DataLoader(trainset,
                        batch_size=1024,
                        num_workers=10,
                        shuffle=True,
                        drop_last=True)

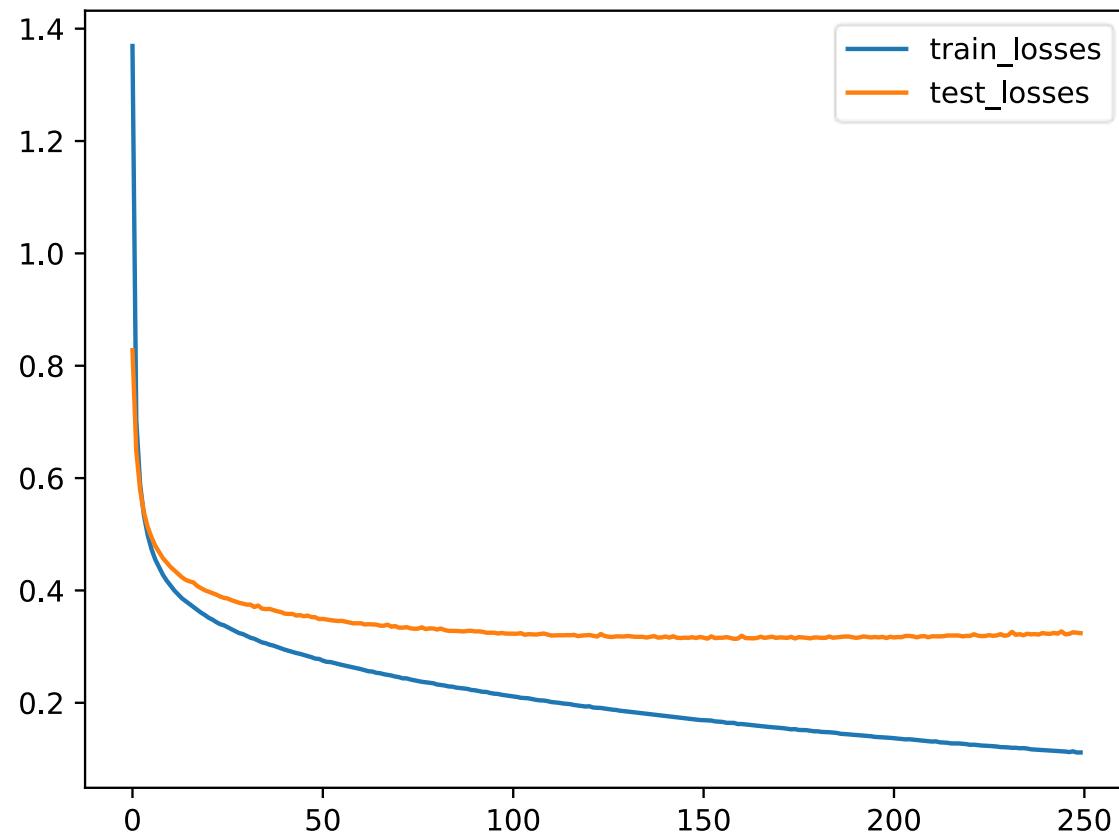
testset = FashionMNIST(root='data',
                       train=False,
                       download=True,
                       transform=transform)

testloader = DataLoader(testset,
                        batch_size=1024,
                        num_workers=10,
                        shuffle=False)
```

MLP for Fashion-MNIST

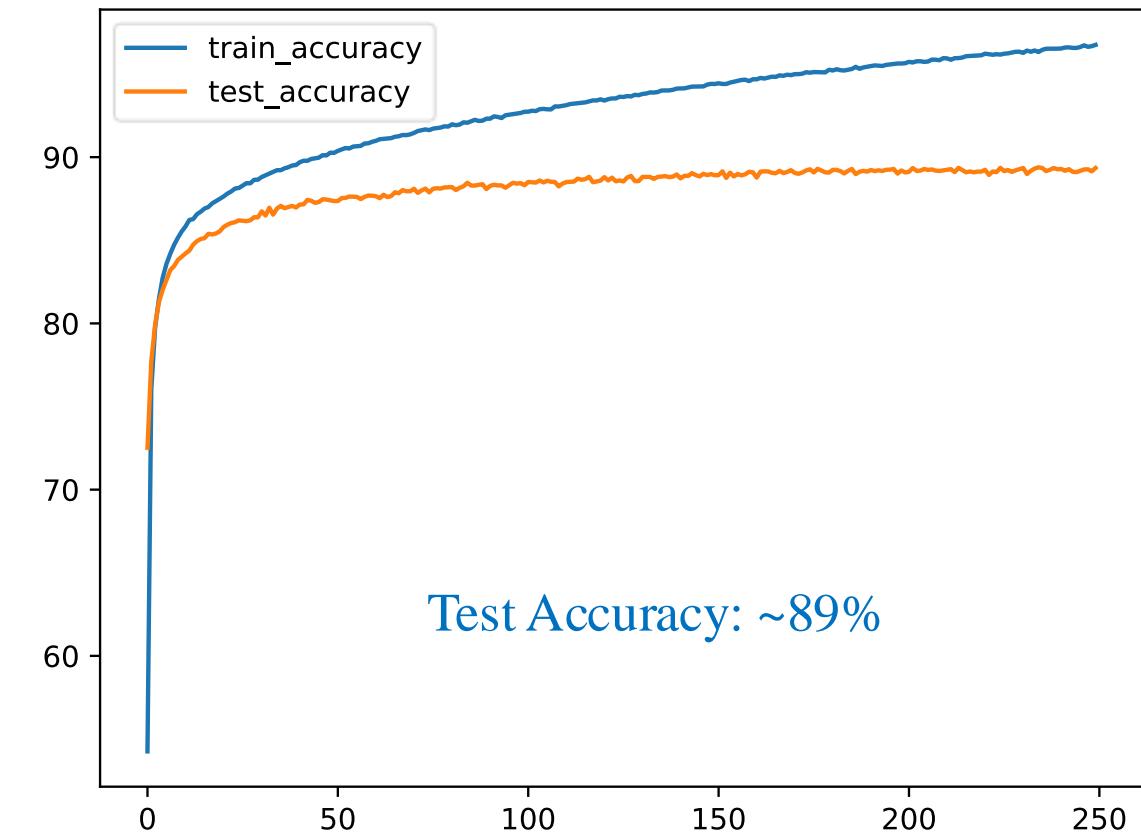
❖ ReLU, He and Adam

Case 1



Adam with learning rate of 1e-4

Perform reasonably



Test Accuracy: ~89%

Cifar-10 dataset

Color images

Resolution=32x32

Training set: 50000 samples

Testing set: 10000 samples

airplane



automobile



bird



cat



deer



dog



frog



horse



ship



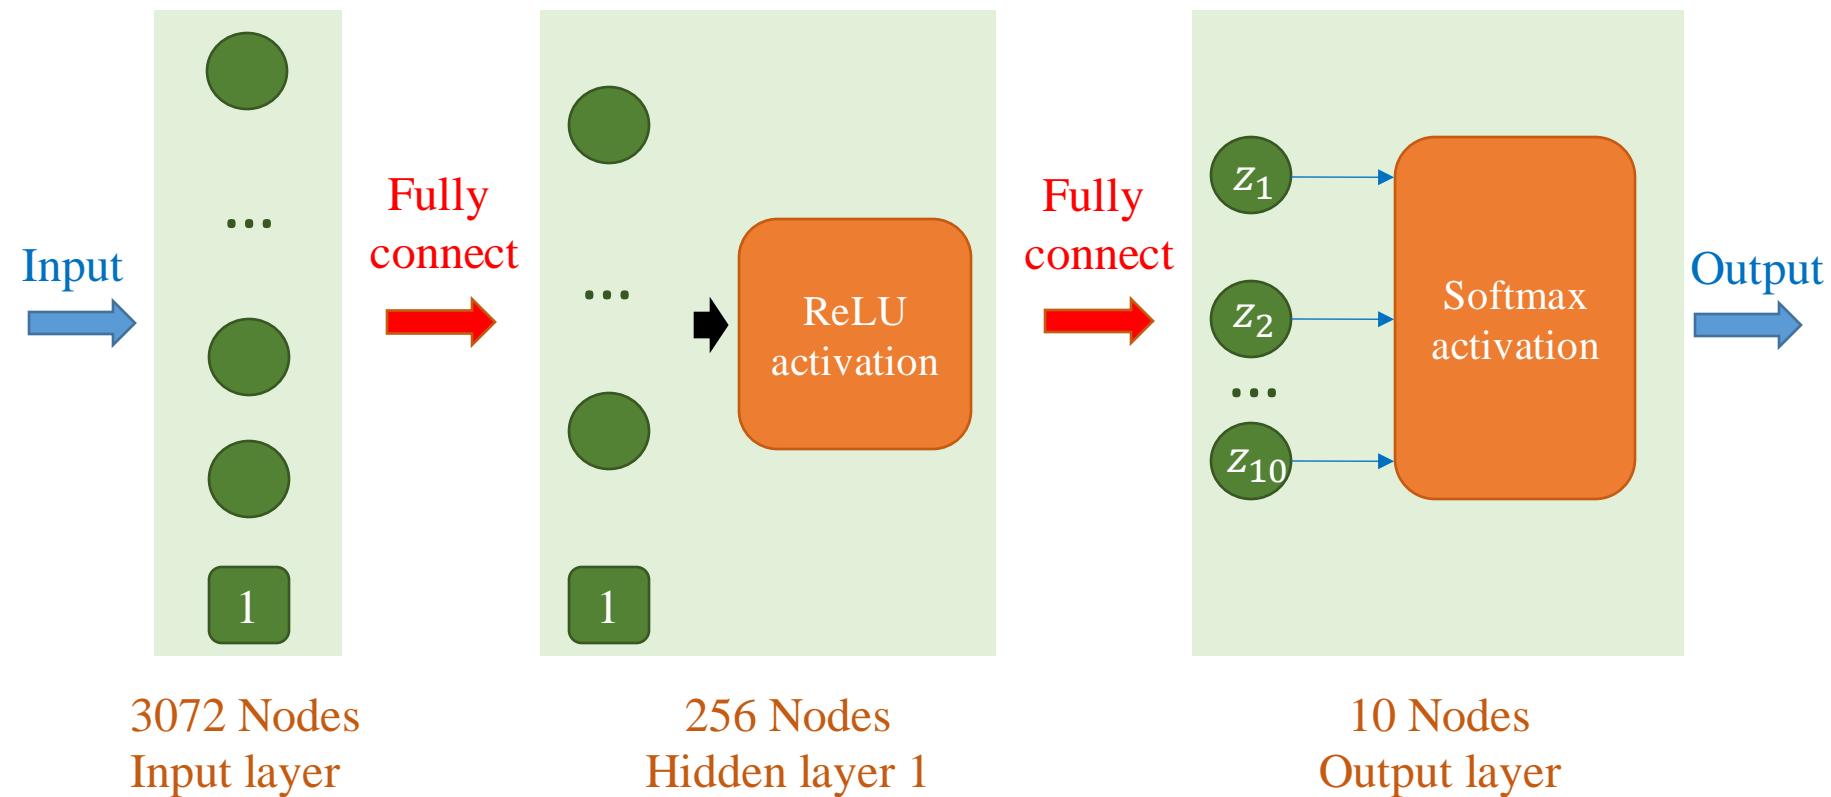
truck



MLP for Cifar-10

❖ ReLU, He and Adam

Case 2



MLP for Cifar-10

❖ ReLU, He and Adam

```
# model
model = nn.Sequential(
    nn.Flatten(),
    nn.Linear(32*32*3, 256),
    nn.ReLU(),
    nn.Linear(256, 10)
)

# Initialize the weights
for layer in model:
    if isinstance(layer, nn.Linear):
        init.kaiming_uniform_(layer.weight,
                              nonlinearity='relu')
        if layer.bias is not None:
            layer.bias.data.fill_(0)

# loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(),
                      lr=0.001)
```

```
# Load CIFAR10 dataset
transform = Compose([ToTensor(),
                     Normalize((0.5,0.5, 0.5),
                               (0.5,0.5, 0.5))])

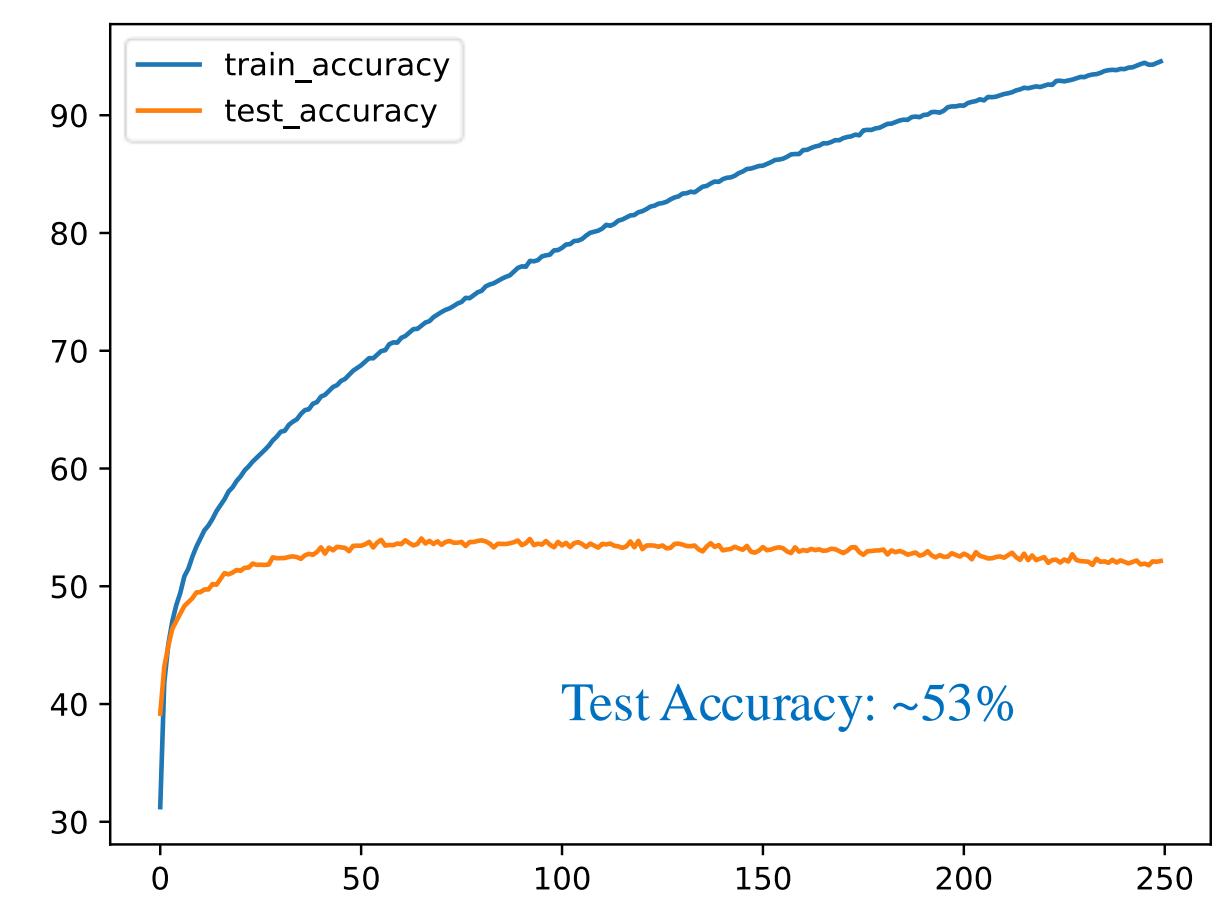
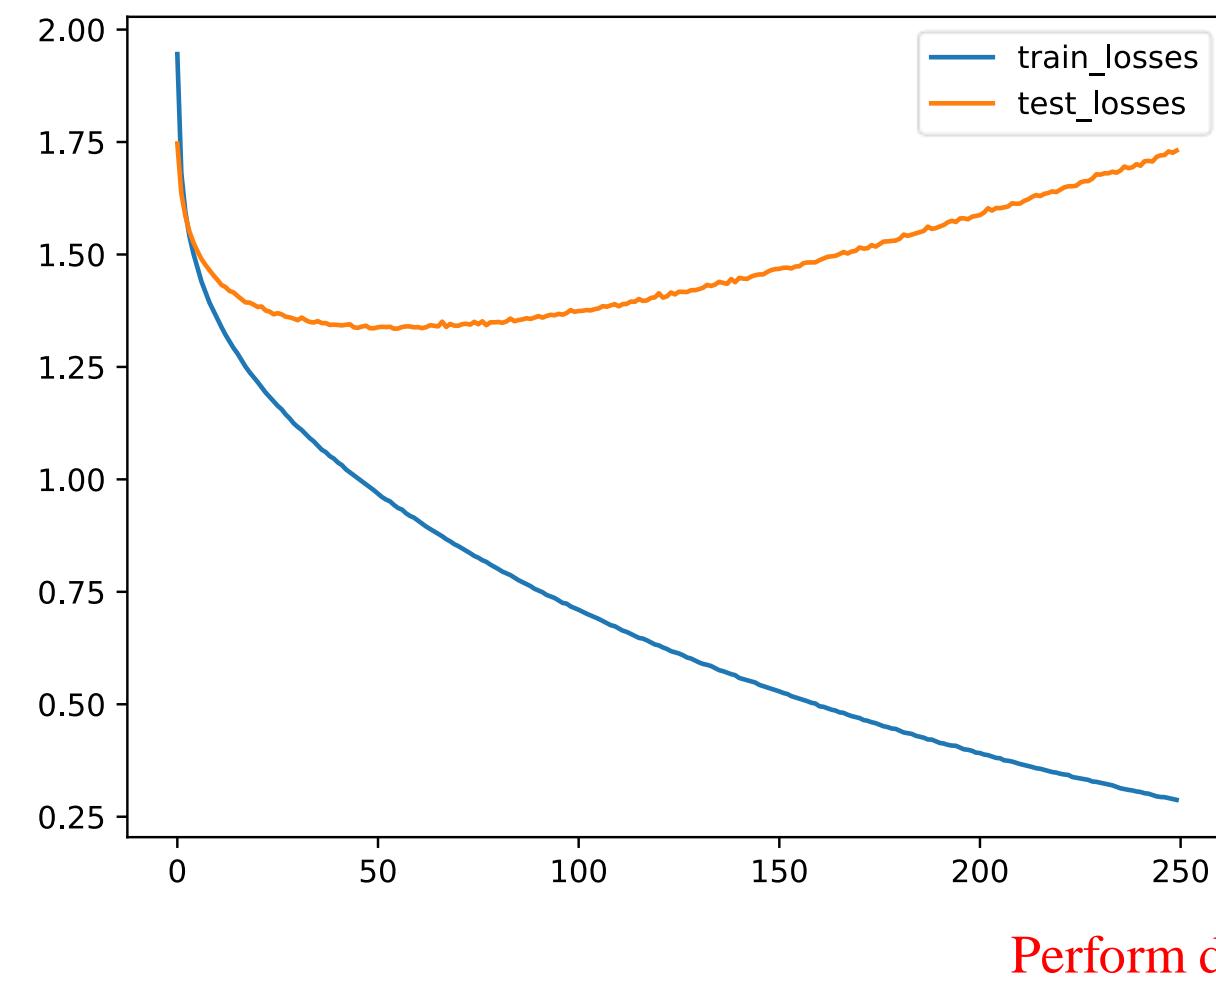
trainset = CIFAR10(root='data',
                    train=True,
                    download=True,
                    transform=transform)
trainloader = DataLoader(trainset,
                        batch_size=1024,
                        num_workers=10,
                        shuffle=True,
                        drop_last=True)

testset = CIFAR10(root='data',
                  train=False,
                  download=True,
                  transform=transform)
testloader = DataLoader(testset,
                        batch_size=1024,
                        num_workers=10,
                        shuffle=False)
```

MLP for Cifar-10

❖ ReLU, He and Adam

Case 2

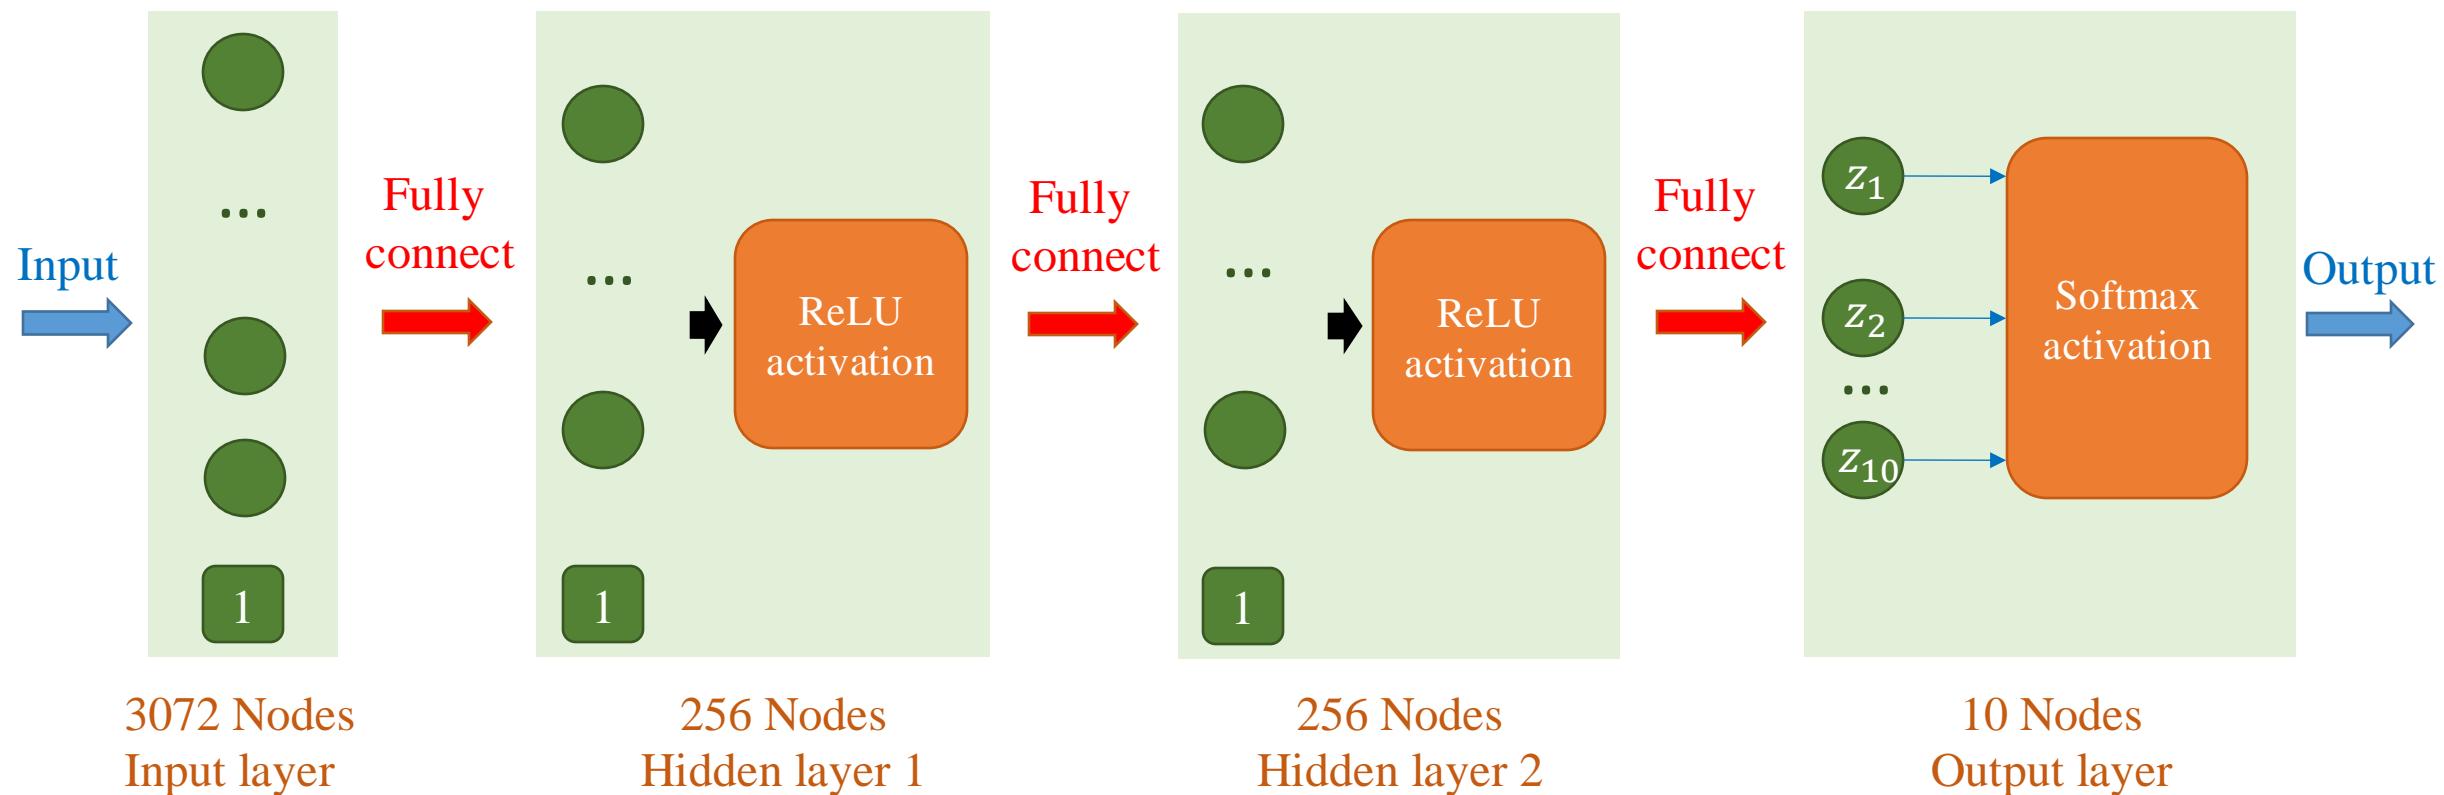


Perform disappointedly

MLP for Cifar-10

❖ ReLU, He and Adam: add more layers

Case 3



❖ ReLU, He and Adam

```
# model
model = nn.Sequential(
    nn.Flatten(),
    nn.Linear(32*32*3, 256),
    nn.ReLU(),
    nn.Linear(256, 256),
    nn.ReLU(),
    nn.Linear(256, 10)
)

# Initialize the weights
for layer in model:
    if isinstance(layer, nn.Linear):
        init.kaiming_uniform_(layer.weight,
                              nonlinearity='relu')
        if layer.bias is not None:
            layer.bias.data.fill_(0)

# loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(),
                      lr=0.001)
```

```
# Load CIFAR10 dataset
transform = Compose([ToTensor(),
                     Normalize((0.5,0.5, 0.5),
                               (0.5,0.5, 0.5))])

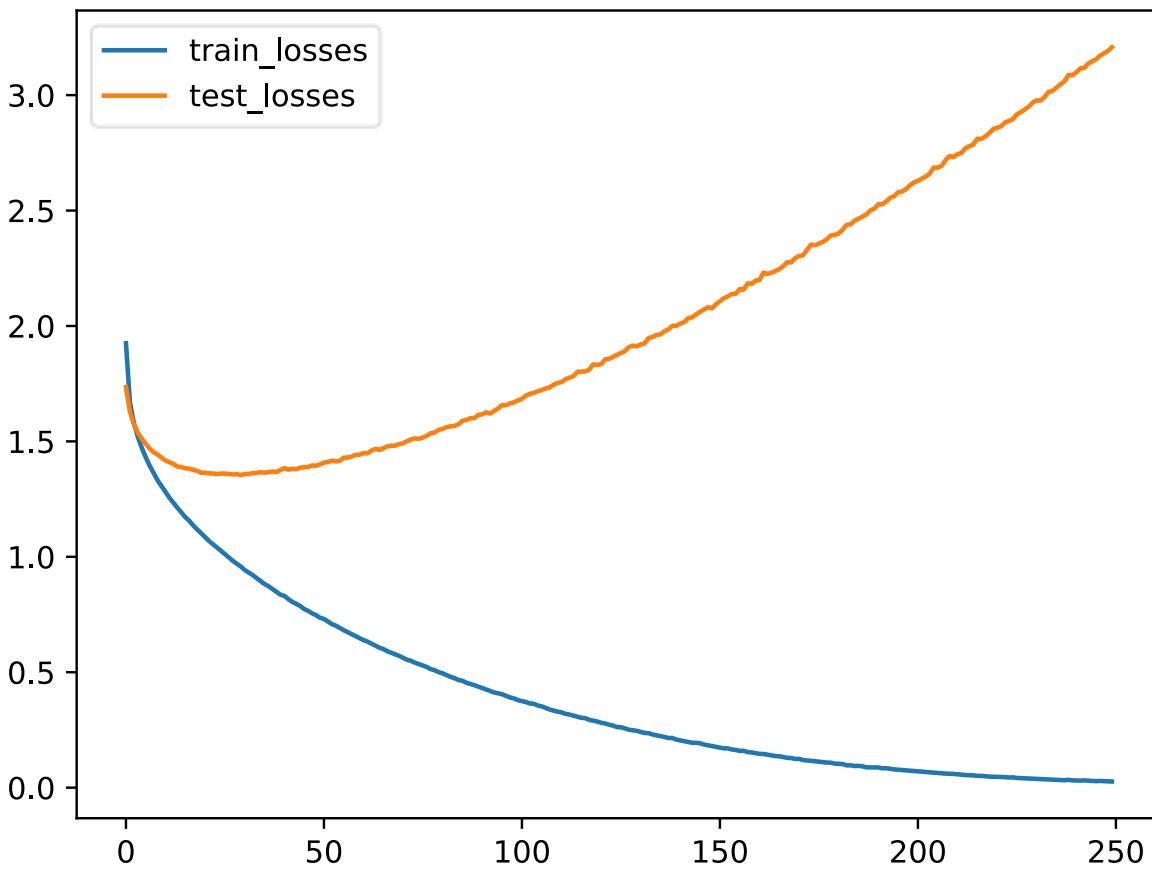
trainset = CIFAR10(root='data',
                    train=True,
                    download=True,
                    transform=transform)
trainloader = DataLoader(trainset,
                        batch_size=1024,
                        num_workers=10,
                        shuffle=True,
                        drop_last=True)

testset = CIFAR10(root='data',
                  train=False,
                  download=True,
                  transform=transform)
testloader = DataLoader(testset,
                        batch_size=1024,
                        num_workers=10,
                        shuffle=False)
```

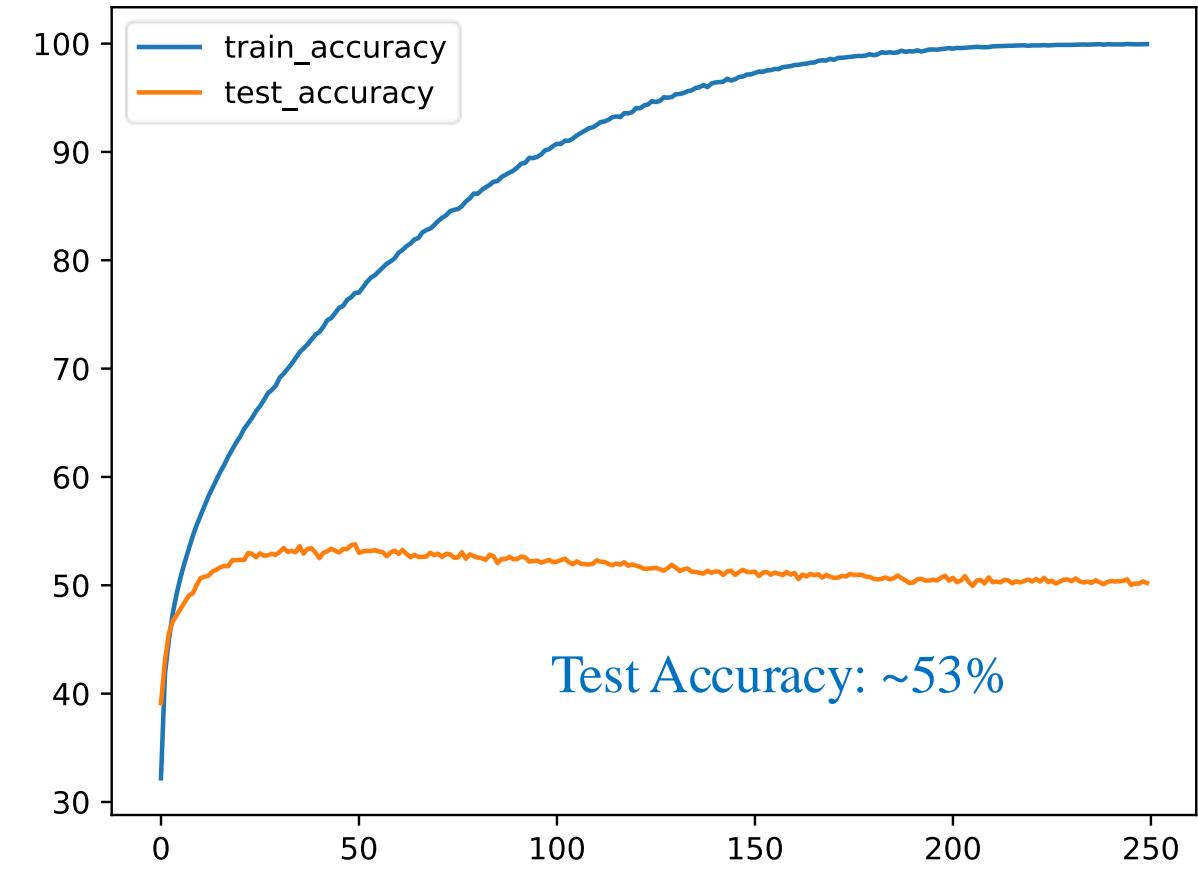
MLP for Cifar-10

❖ ReLU, He and Adam

Case 3



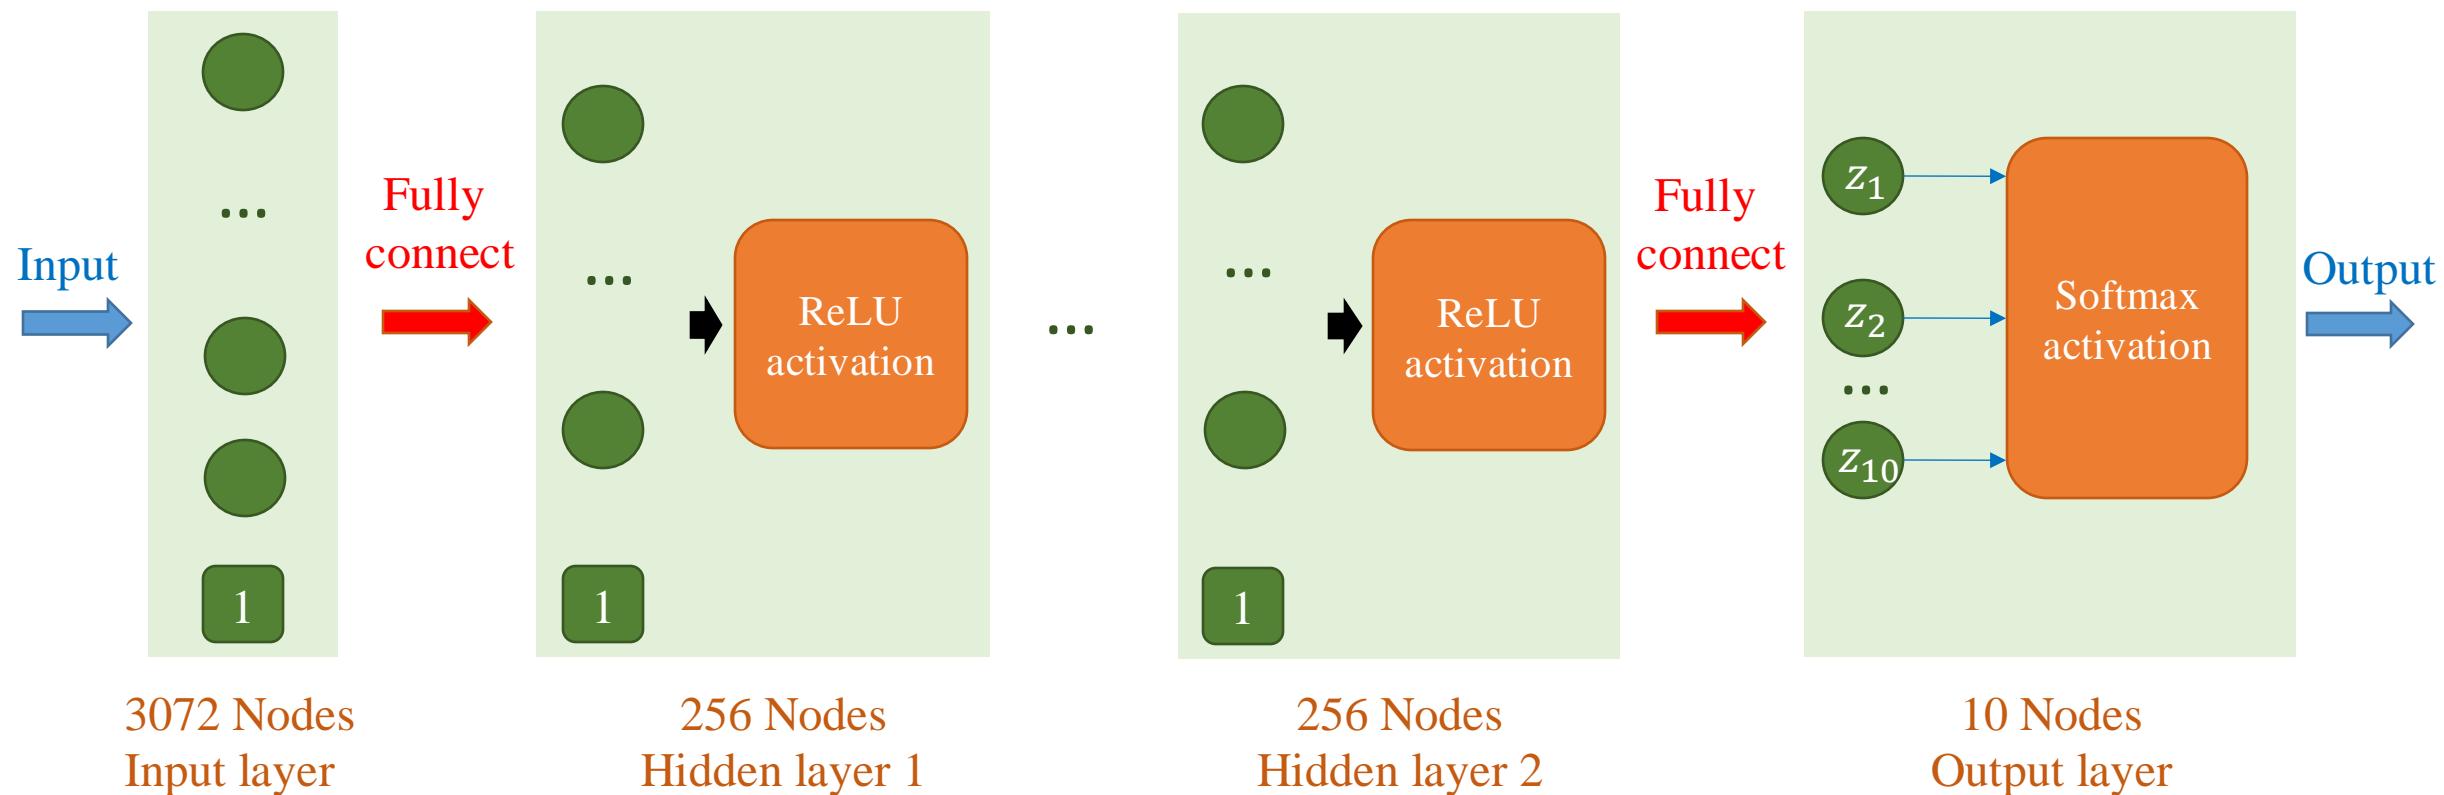
Still Perform poorly



MLP for Cifar-10

❖ ReLU, He and Adam

Case 4



3072 Nodes
Input layer

256 Nodes
Hidden layer 1

256 Nodes
Hidden layer 2

10 Nodes
Output layer

❖ ReLU, He and Adam

```
# model
model = nn.Sequential(
    nn.Flatten(), nn.Linear(32*32*3, 256),
    nn.ReLU(), nn.Linear(256, 256),
    nn.ReLU(), nn.Linear(256, 256),
    nn.ReLU(), nn.Linear(256, 10)
)

# Initialize the weights
for layer in model:
    if isinstance(layer, nn.Linear):
        init.kaiming_uniform_(layer.weight,
                              nonlinearity='relu')
        if layer.bias is not None:
            layer.bias.data.fill_(0)

# loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(),
                      lr=0.001)
```

```
# Load CIFAR10 dataset
transform = Compose([ToTensor(),
                     Normalize((0.5,0.5, 0.5),
                               (0.5,0.5, 0.5))])

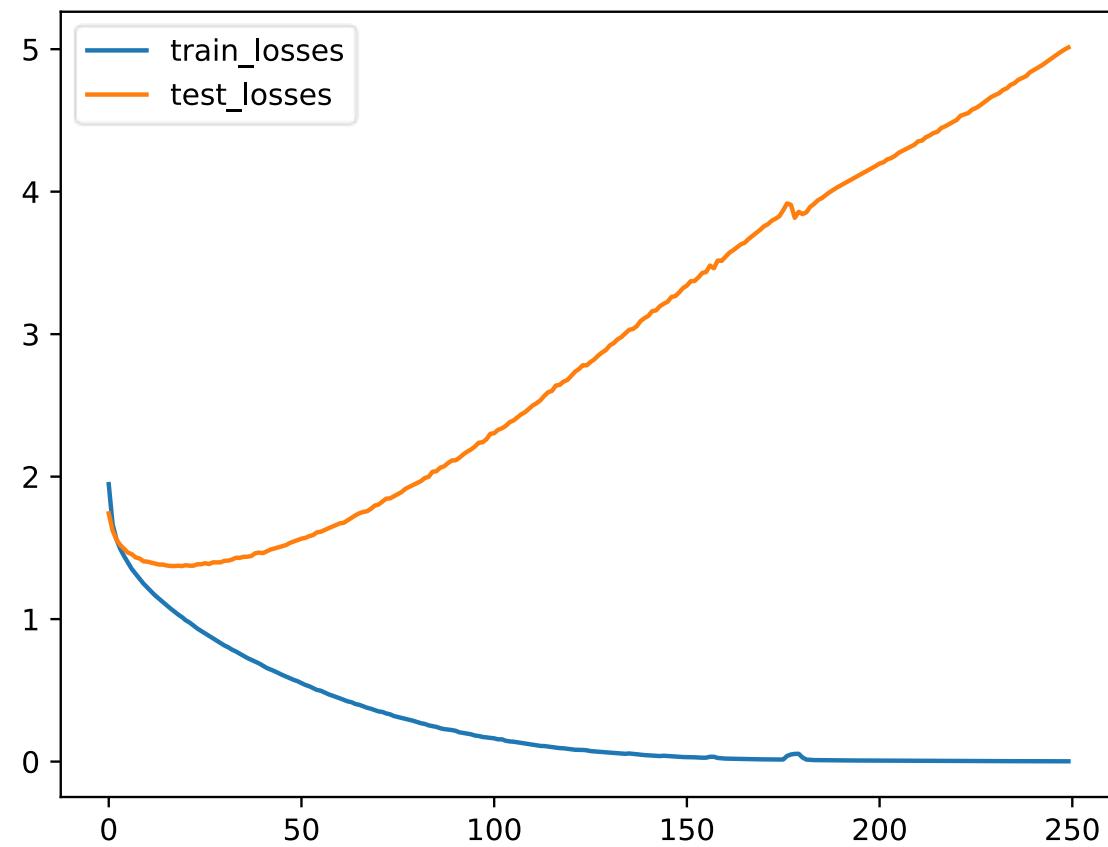
trainset = CIFAR10(root='data',
                    train=True,
                    download=True,
                    transform=transform)
trainloader = DataLoader(trainset,
                        batch_size=1024,
                        num_workers=10,
                        shuffle=True,
                        drop_last=True)

testset = CIFAR10(root='data',
                  train=False,
                  download=True,
                  transform=transform)
testloader = DataLoader(testset,
                        batch_size=1024,
                        num_workers=10,
                        shuffle=False)
```

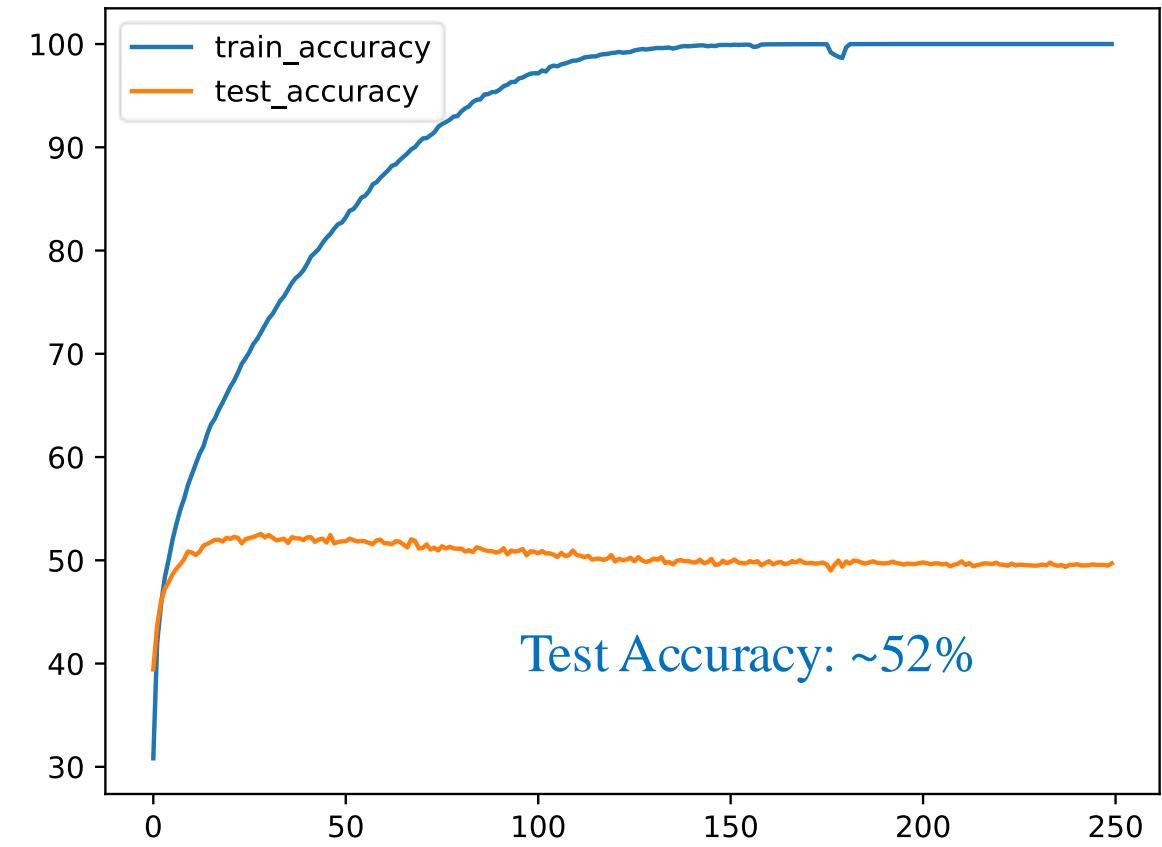
MLP for Cifar-10

❖ ReLU, He and Adam: Using 3 hidden layers

Case 4

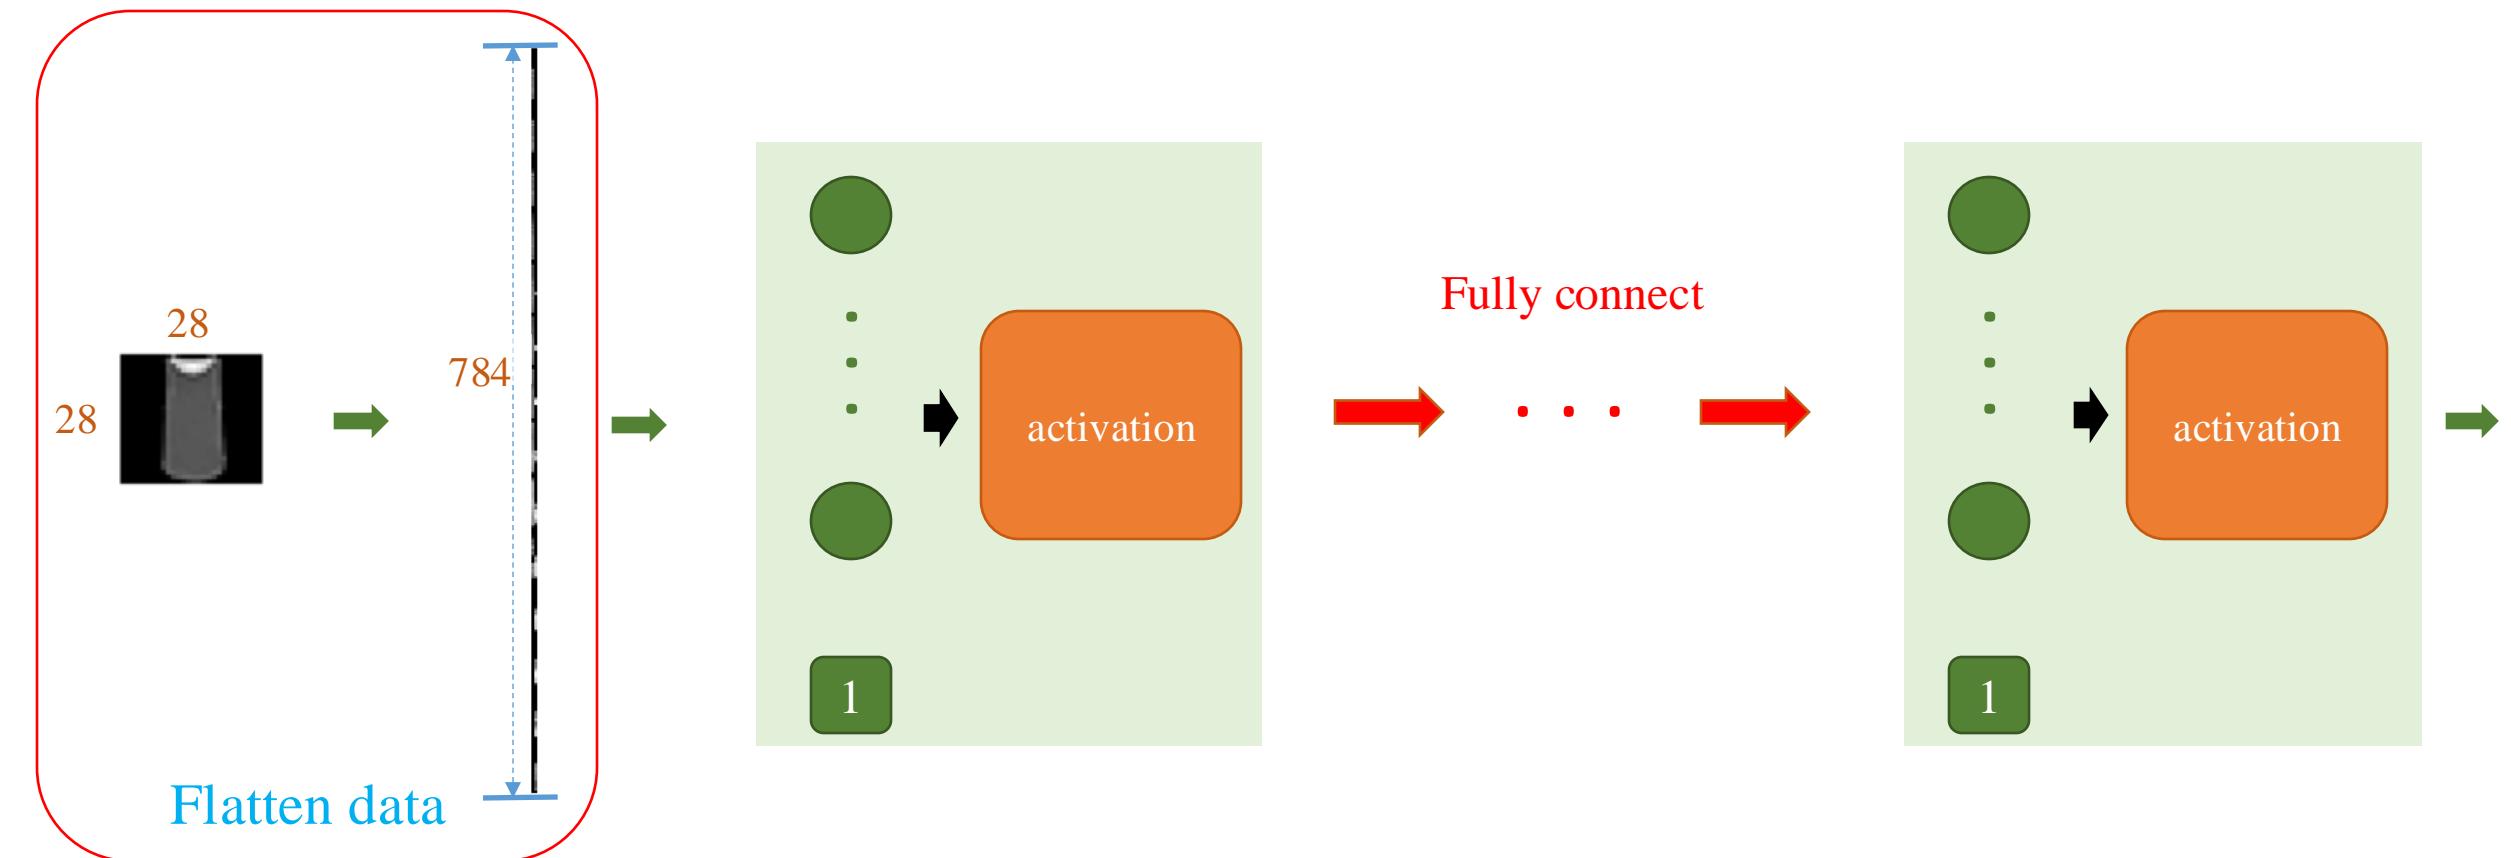


Perform even worse



MLP Limitations

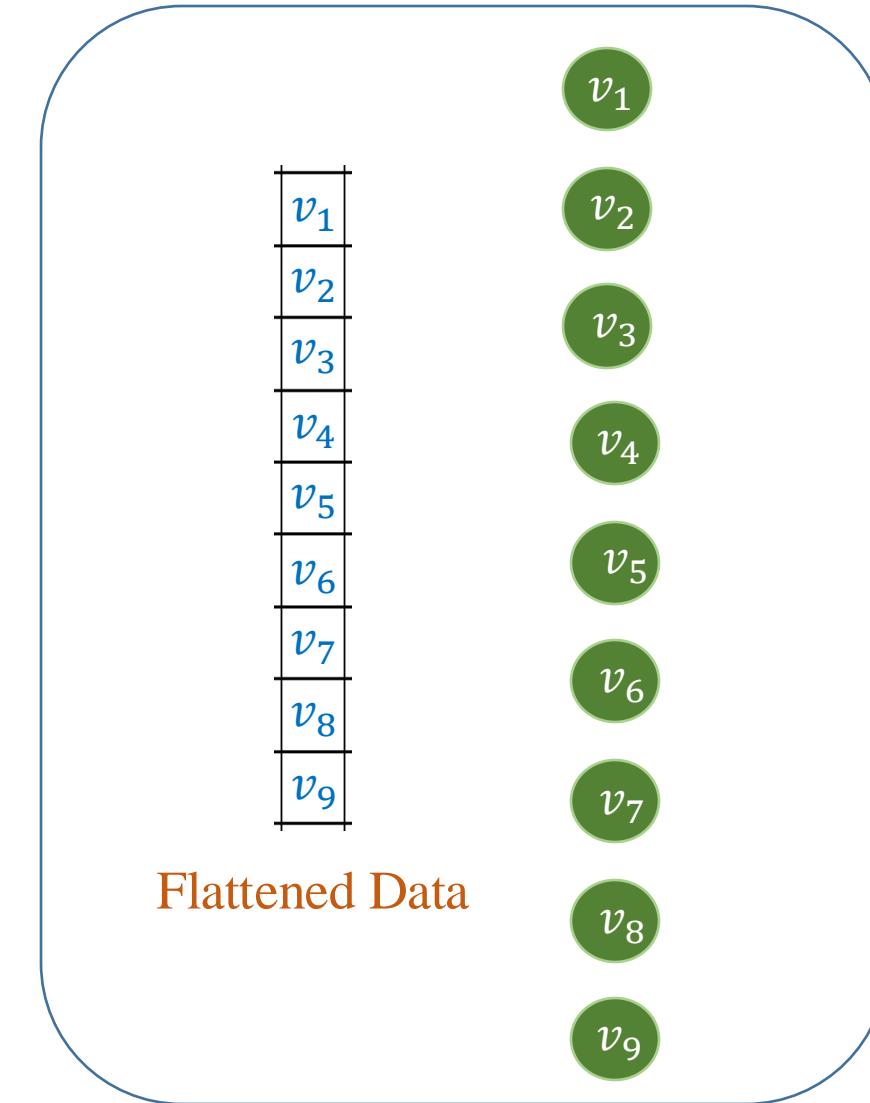
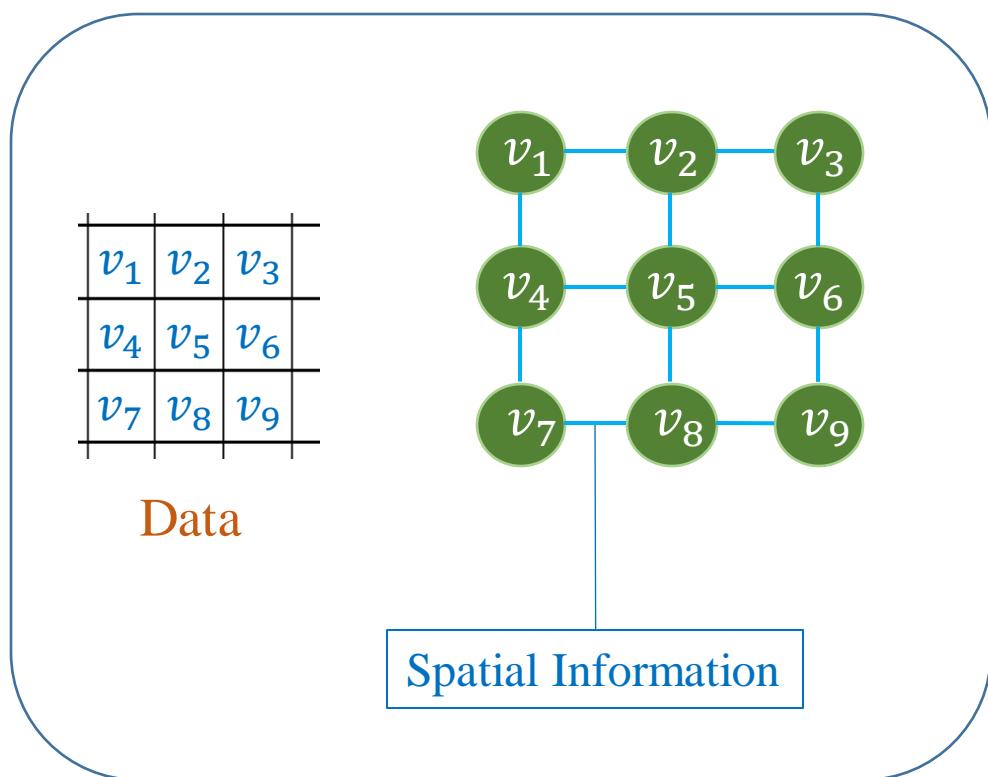
❖ Multi-layer Perceptron



Problem: Remove spatial information of the data
Inefficiency have a large amount of parameters

MLP Limitations

❖ Problem of flattening data



Outline

SECTION 1

MLP Limitations

SECTION 2

Convolutional Layer

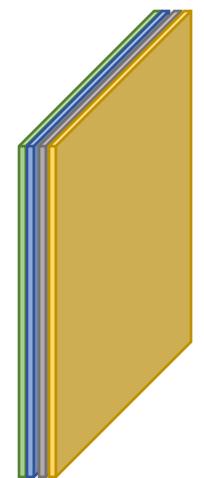
SECTION 3

Standard CNNs



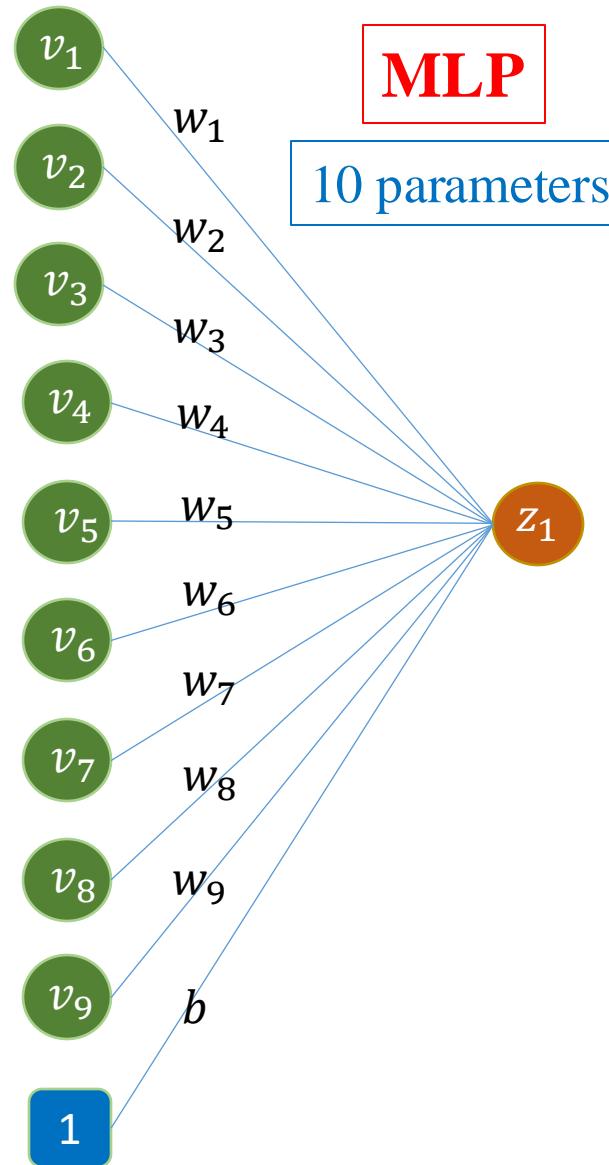
Input Data
(3,32,32)

Convolve with
4 kernels (3,5,5)
→

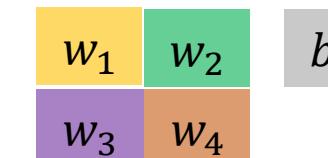


Feature maps
(4,28,28)

From MLP to CNN

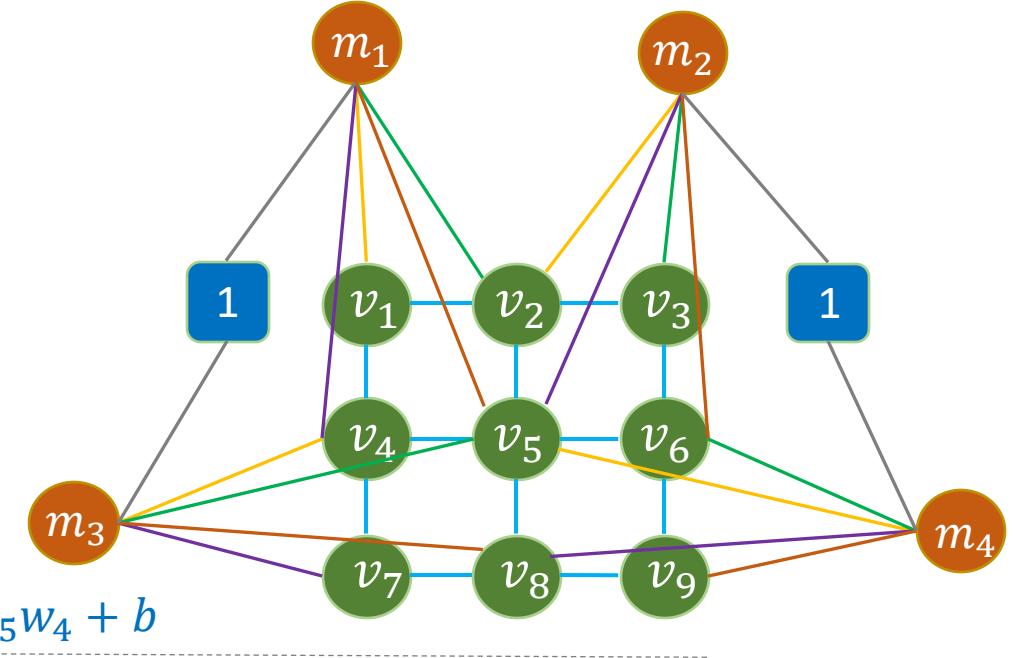
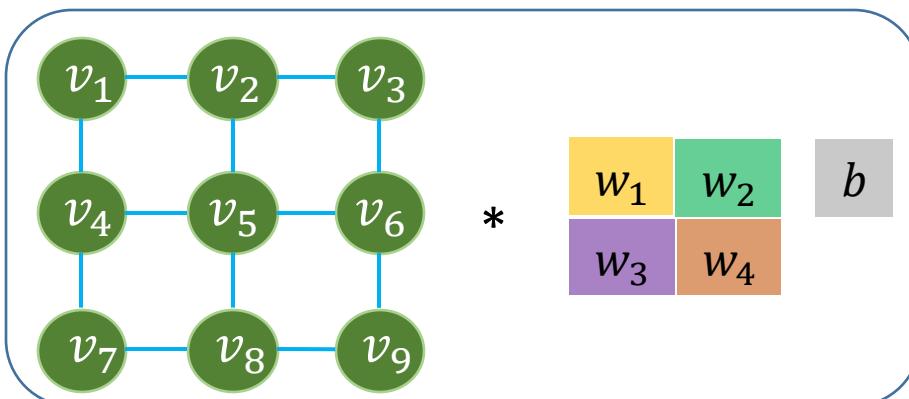


CNN 5 parameters



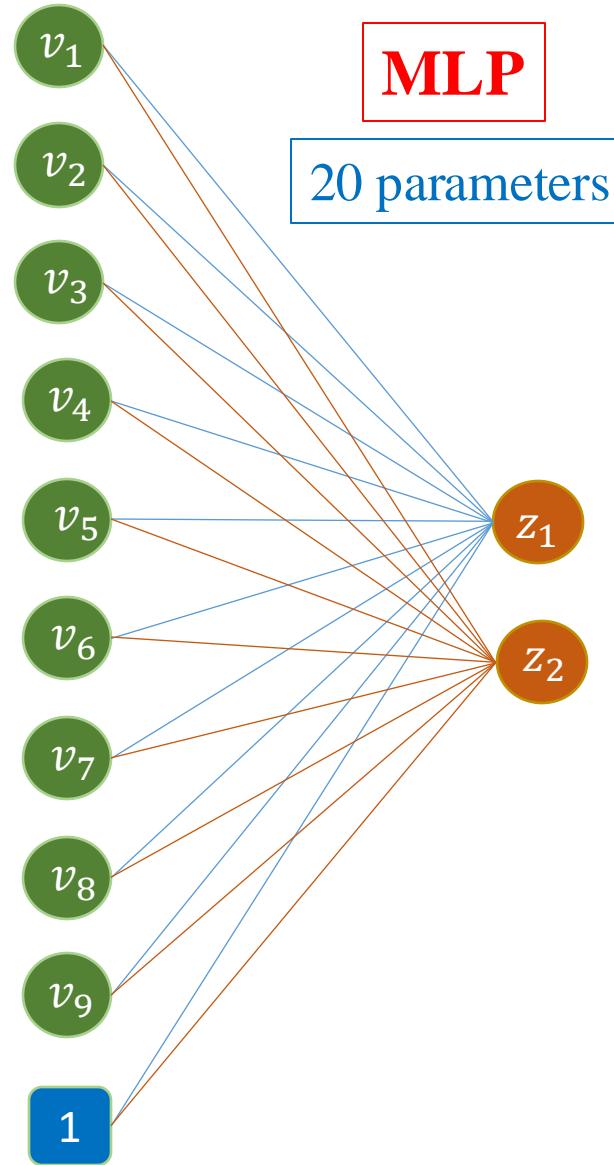
Kernel of parameters

$$m_1 = v_1w_1 + v_2w_2 + v_4w_3 + v_5w_4 + b$$



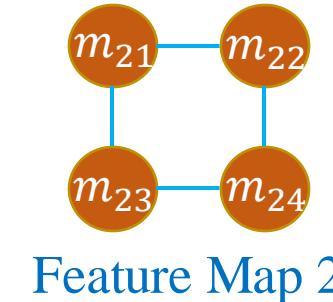
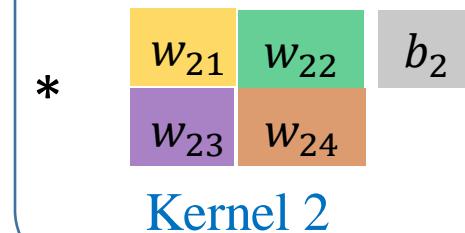
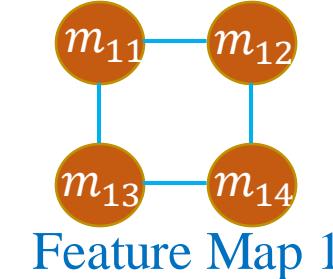
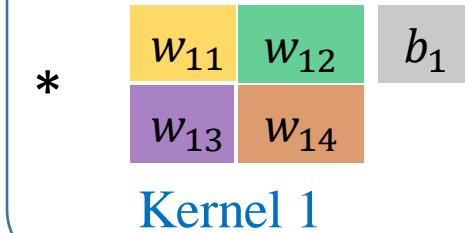
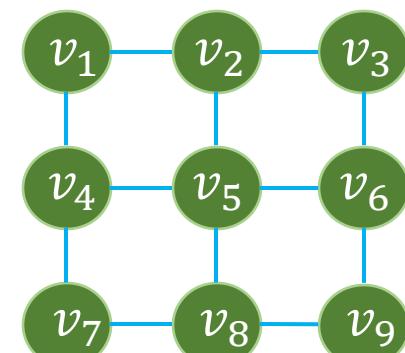
Feature Map

From MLP to CNN



CNN 10 parameters

Kernel 1 \neq Kernel 2

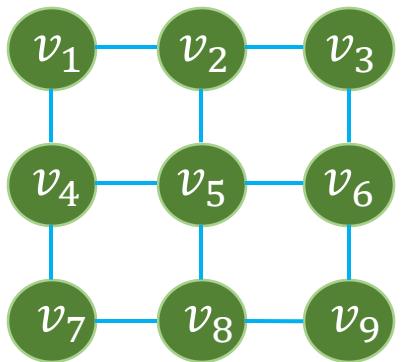


Global vs. Local?

Parameter size?

From MLP to CNN

❖ Understand convolution



Shape=(1,3,3)

(Channel=1, Height=3, Width=3)

#channels of data must
 =

#channels of kernel

$$\begin{matrix} w_1 & w_2 \\ w_3 & w_4 \end{matrix} \quad b$$

Shape=(1,2,2)

#parameters (+bias) = 5



Convolutional Layer

❖ How many cases?

0	0	1	2	2
1	2	2	1	2
0	2	0	2	1
0	1	1	1	0
1	0	0	0	1

Data D

0.0	0.1	-0.1
-0.2	0.0	0.1
0.0	0.0	0.1

Bias b = 0.0

0	0	1	2	2
1	2	2	1	2
0	2	0	2	1
0	1	1	1	0
1	0	0	0	1

0	0	1	2	2
1	2	2	1	2
0	2	0	2	1
0	1	1	1	0
1	0	0	0	1

0	0	1	2	2
1	2	2	1	2
0	2	0	2	1
0	1	1	1	0
1	0	0	0	1

0	0	1	2	2
1	2	2	1	2
0	2	0	2	1
0	1	1	1	0
1	0	0	0	1

0	0	1	2	2
1	2	2	1	2
0	2	0	2	1
0	1	1	1	0
1	0	0	0	1

0	0	1	2	2
1	2	2	1	2
0	2	0	2	1
0	1	1	1	0
1	0	0	0	1

0	0	1	2	2
1	2	2	1	2
0	2	0	2	1
0	1	1	1	0
1	0	0	0	1

0	0	1	2	2
1	2	2	1	2
0	2	0	2	1
0	1	1	1	0
1	0	0	0	1

0	0	1	2	2
1	2	2	1	2
0	2	0	2	1
0	1	1	1	0
1	0	0	0	1

Convolutional Layer

❖ Example

0	0	1	2	2
1	2	2	1	2
0	2	0	2	1
0	1	1	1	0
1	0	0	0	1

Data D

Bias $b = 0.0$

0.0	0.1	-0.1
-0.2	0.0	0.1
0.0	0.0	0.1

Kernel K

m_1		

Output

Data size = 5×5

Kernel size = 3×3

Stride = 1

$$\begin{aligned}m_1 &= 0 \times 0.0 + 0 \times 0.1 + 1 \times -0.1 + \\&\quad 1 \times -0.2 + 2 \times 0.0 + 2 \times 0.1 + \\&\quad 0 \times 0.0 + 2 \times 0.0 + 0 \times 0.1\end{aligned}$$



$$m_1 = -0.1$$

Convolutional Layer

❖ Example

0	0	1	2	2
1	2	2	1	2
0	2	0	2	1
0	1	1	1	0
1	0	0	0	1

Data D

Bias $b = 0.0$

0.0	0.1	-0.1
-0.2	0.0	0.1
0.0	0.0	0.1

Kernel K

-0.1	-0.1	-0.2
0.3	-0.2	0.1
0.3	-0.3	0.1

Output

Data size = 5×5

Kernel size = 3×3

Stride = 1

$$S_o = \left\lfloor \frac{S_D - K}{S} \right\rfloor + 1$$

Convolutional Layer

❖ Example

0	0	1	2	2
1	2	2	1	2
0	2	0	2	1
0	1	1	1	0
1	0	0	0	1

Data D

Bias $b = 0.0$

0.0	0.1	-0.1
-0.2	0.0	0.1
0.0	0.0	0.1

Kernel K

m_1	

Output

Data size = 5×5

Kernel size = 3×3

Stride = 2

$$m_1 = 0 \times 0.0 + 0 \times 0.1 + 1 \times -0.1 +$$

$$1 \times -0.2 + 2 \times 0.0 + 2 \times 0.1 +$$

$$0 \times 0.0 + 2 \times 0.0 + 0 \times 0.1$$



$$m_1 = -0.1$$

Convolutional Layer

❖ Example

0	0	1	2	2
1	2	2	1	2
0	2	0	2	1
0	1	1	1	0
1	0	0	0	1

Data D

Bias $b = 0.0$

0.0	0.1	-0.1
-0.2	0.0	0.1
0.0	0.0	0.1

Kernel K

-0.1	-0.2
0.3	0.1

Output

Data size = 5×5

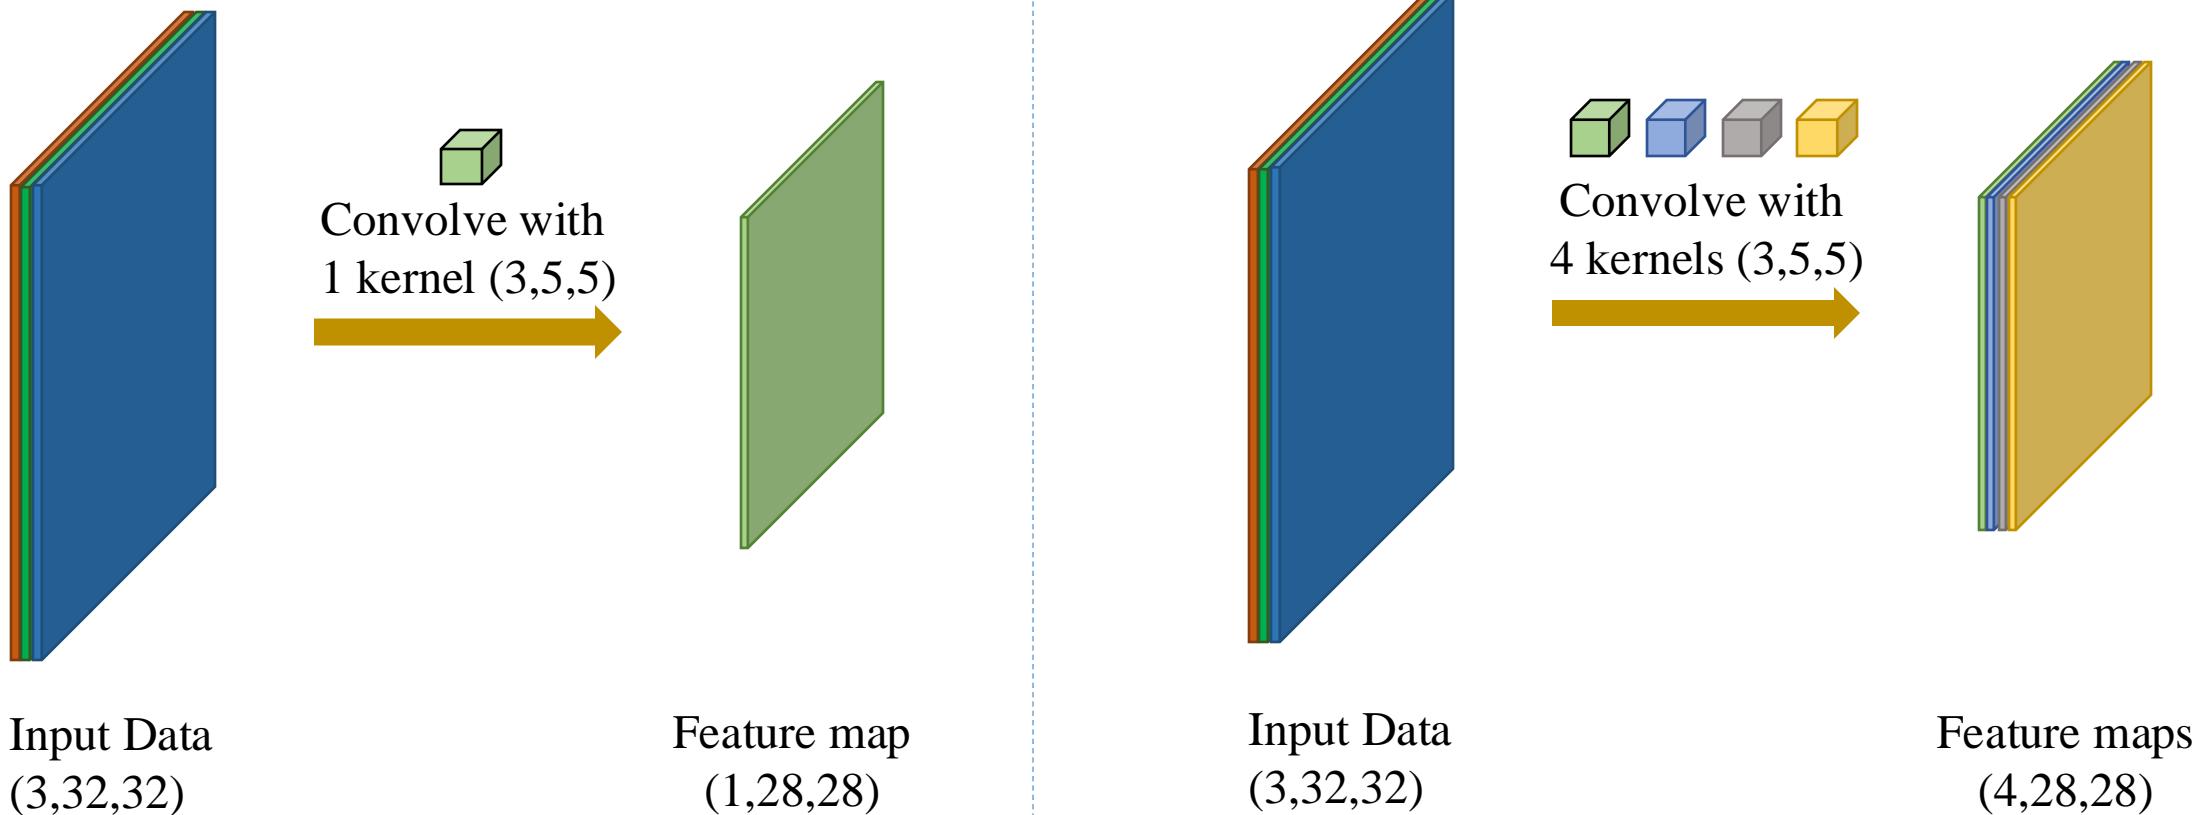
Kernel size = 3×3

Stride = 2

$$S_o = \left\lfloor \frac{S_D - K}{S} \right\rfloor + 1$$

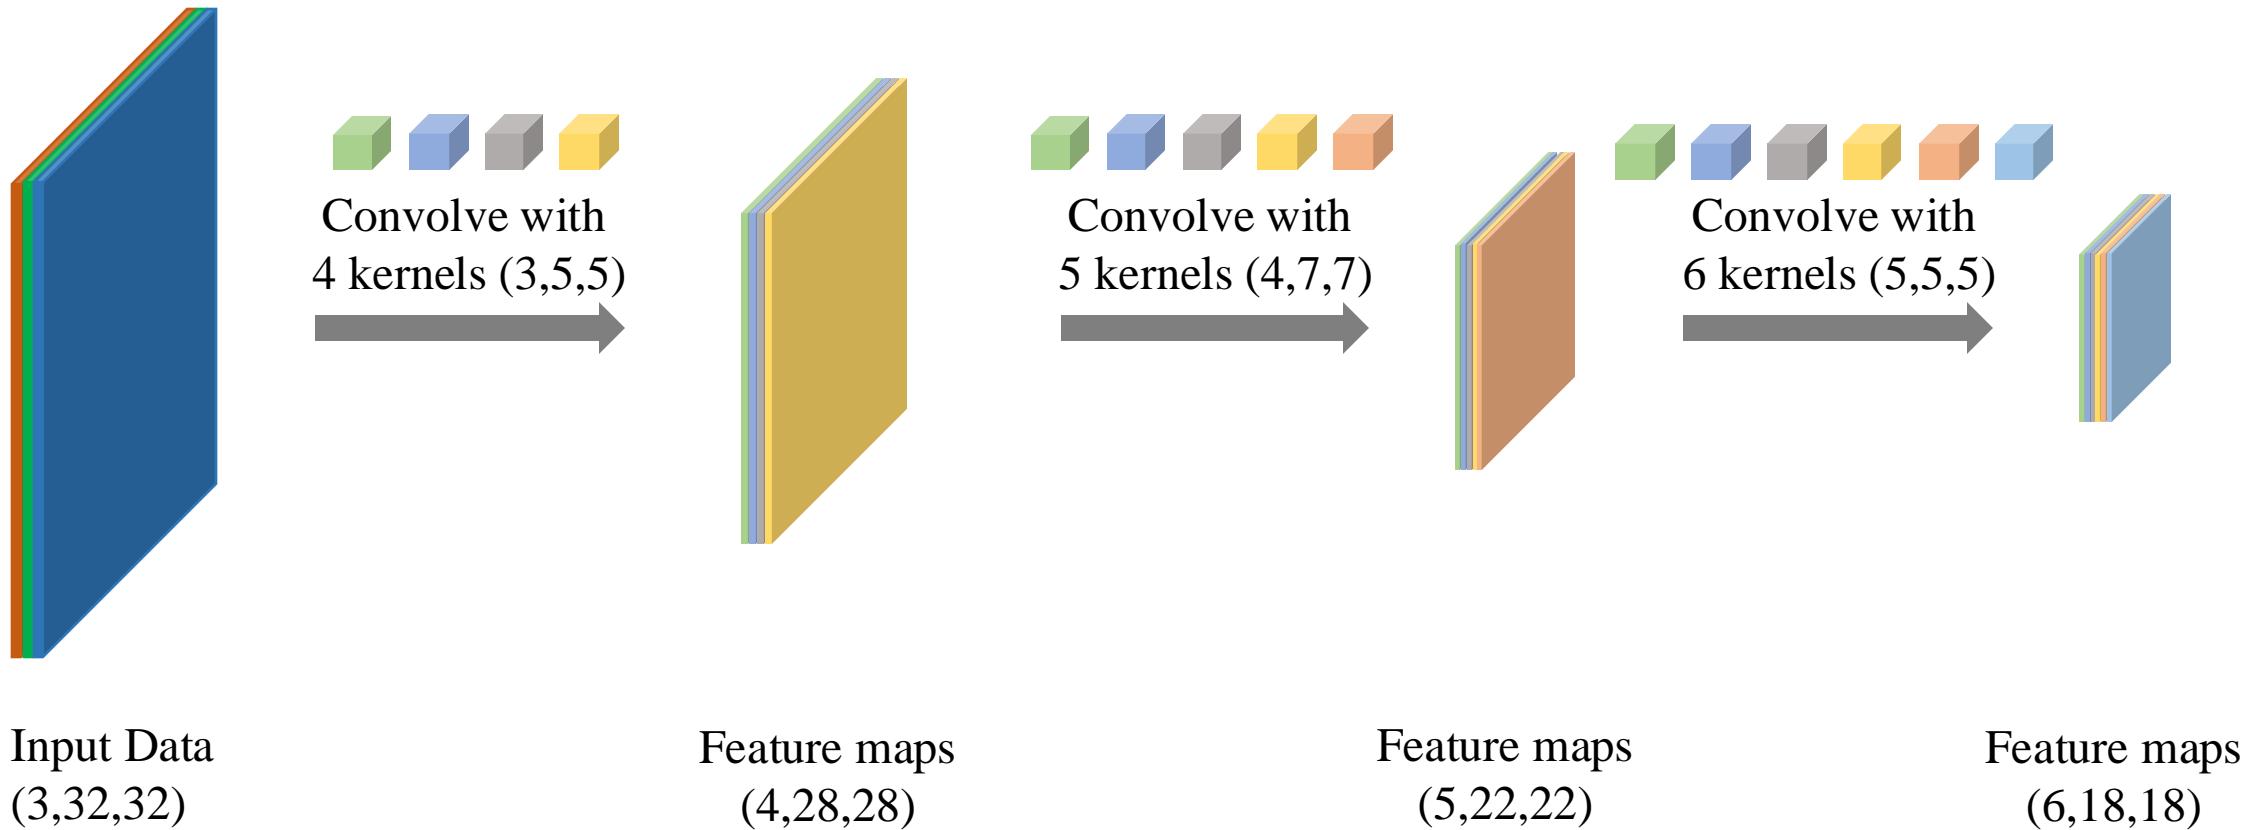
Convolutional Neural Network

❖ Understand convolutional layers



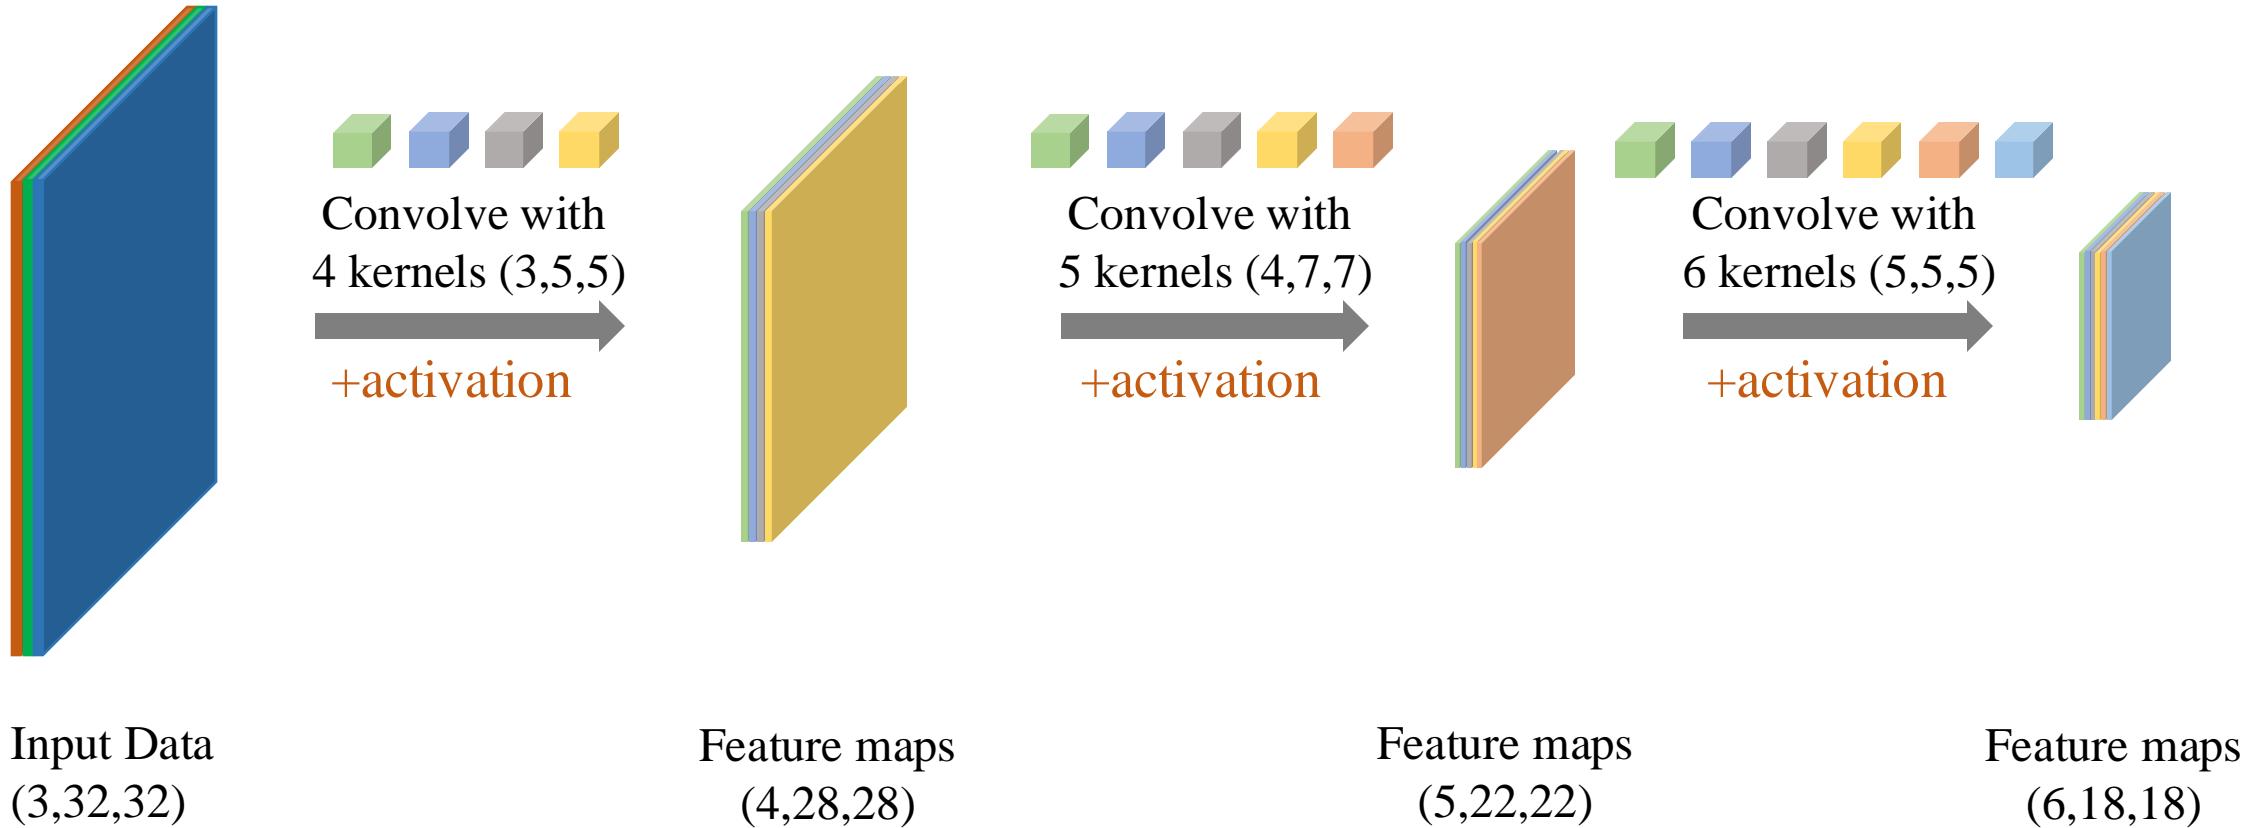
Convolutional Neural Network

❖ A stack of convolutional layers



Convolutional Neural Network

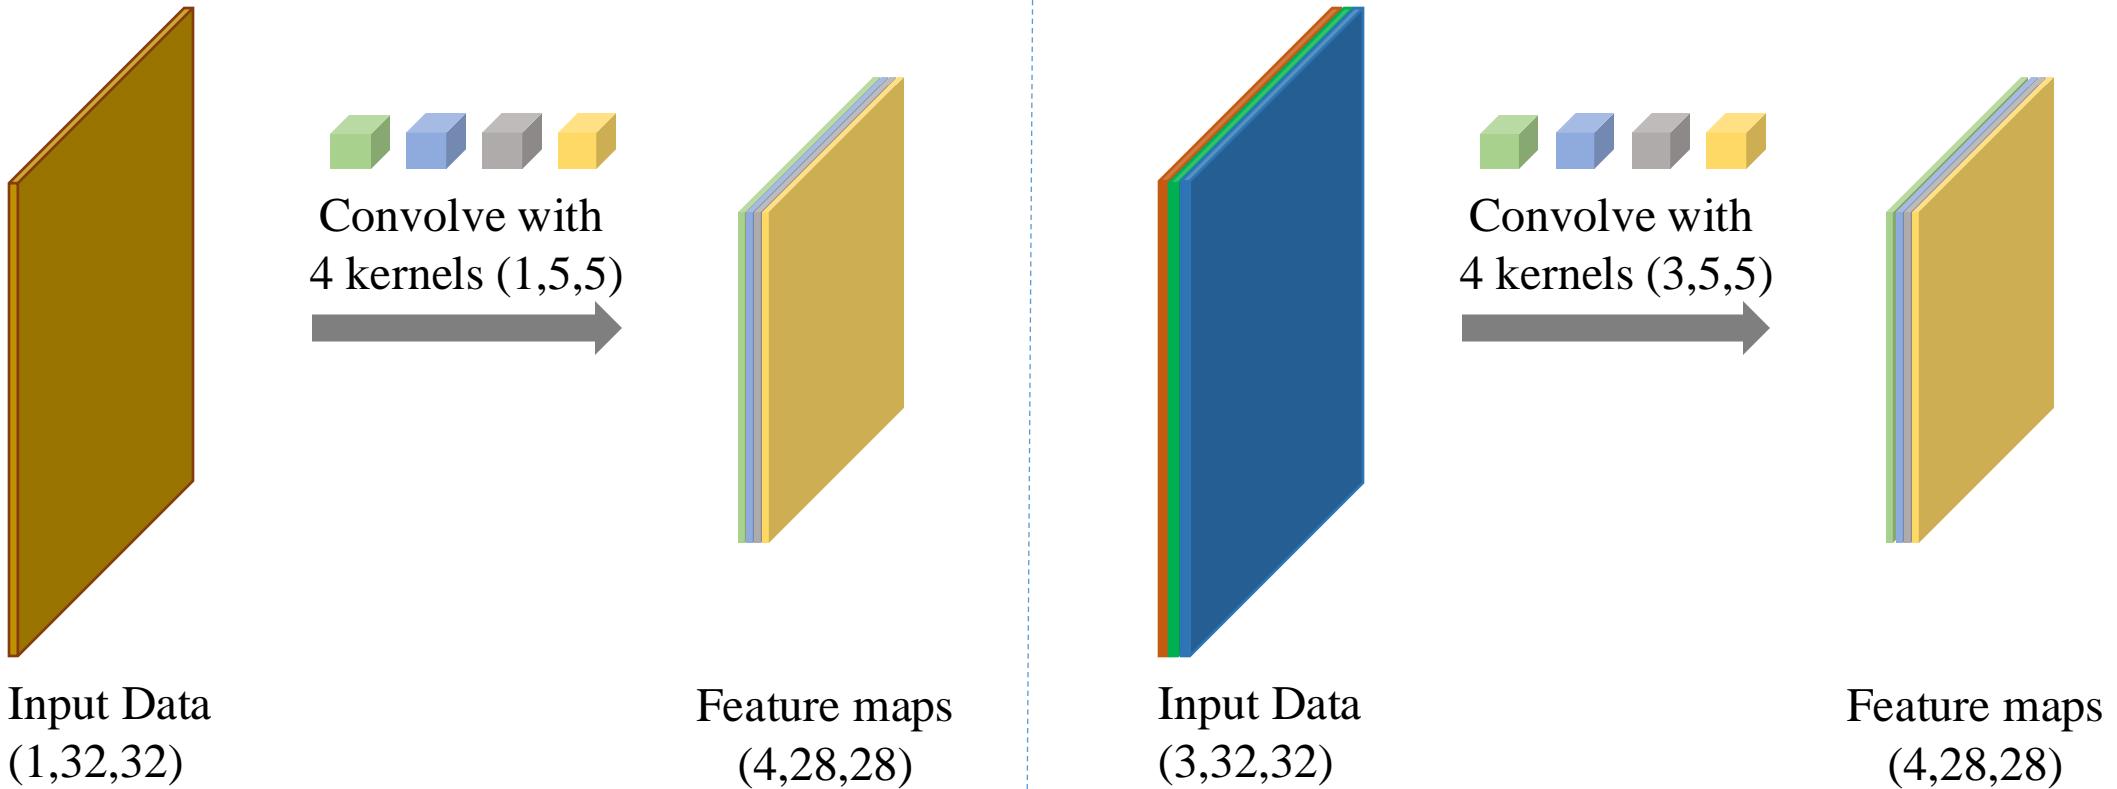
- ❖ A stack of pairs of convolutional layer + activation



Convolutional Neural Network

❖ Convolutional layer in PyTorch

```
nn.Conv2d(in_channels, out_channels, kernel_size)
```

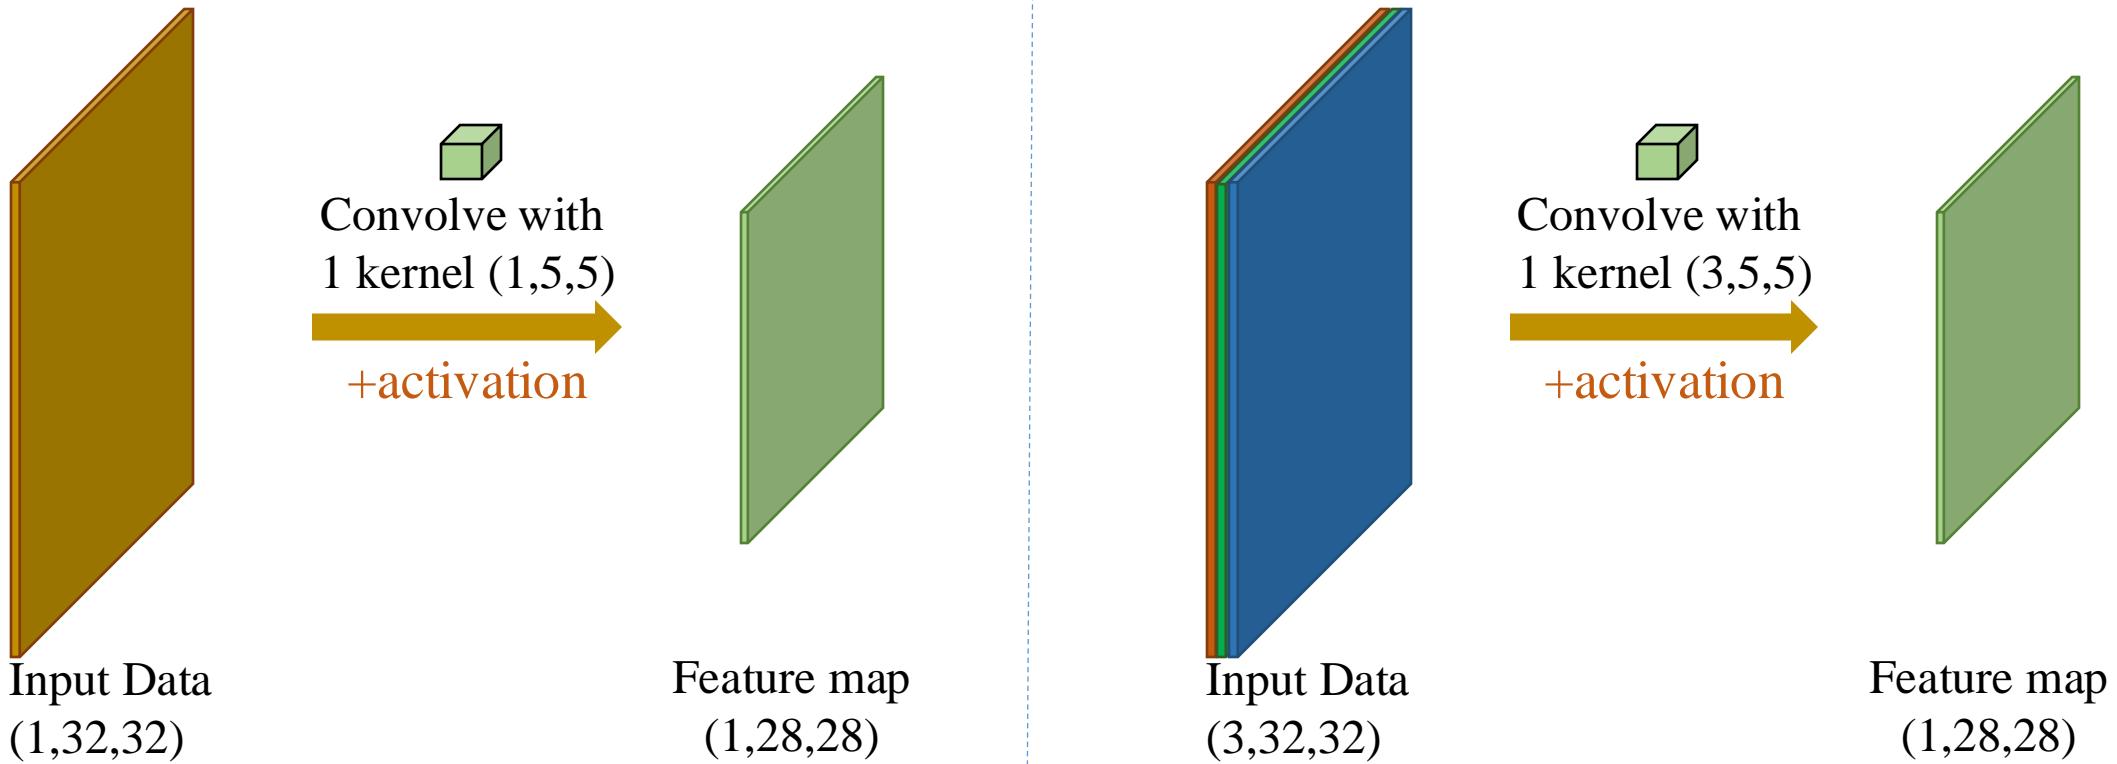


Convolutional Neural Network

❖ Convolutional layer in PyTorch

demo

```
nn.Conv2d(in_channels, out_channels, kernel_size)  
nn.ReLU()
```



Convolutional Neural Network

Fashion-MNIST dataset

Grayscale images

Resolution=28x28

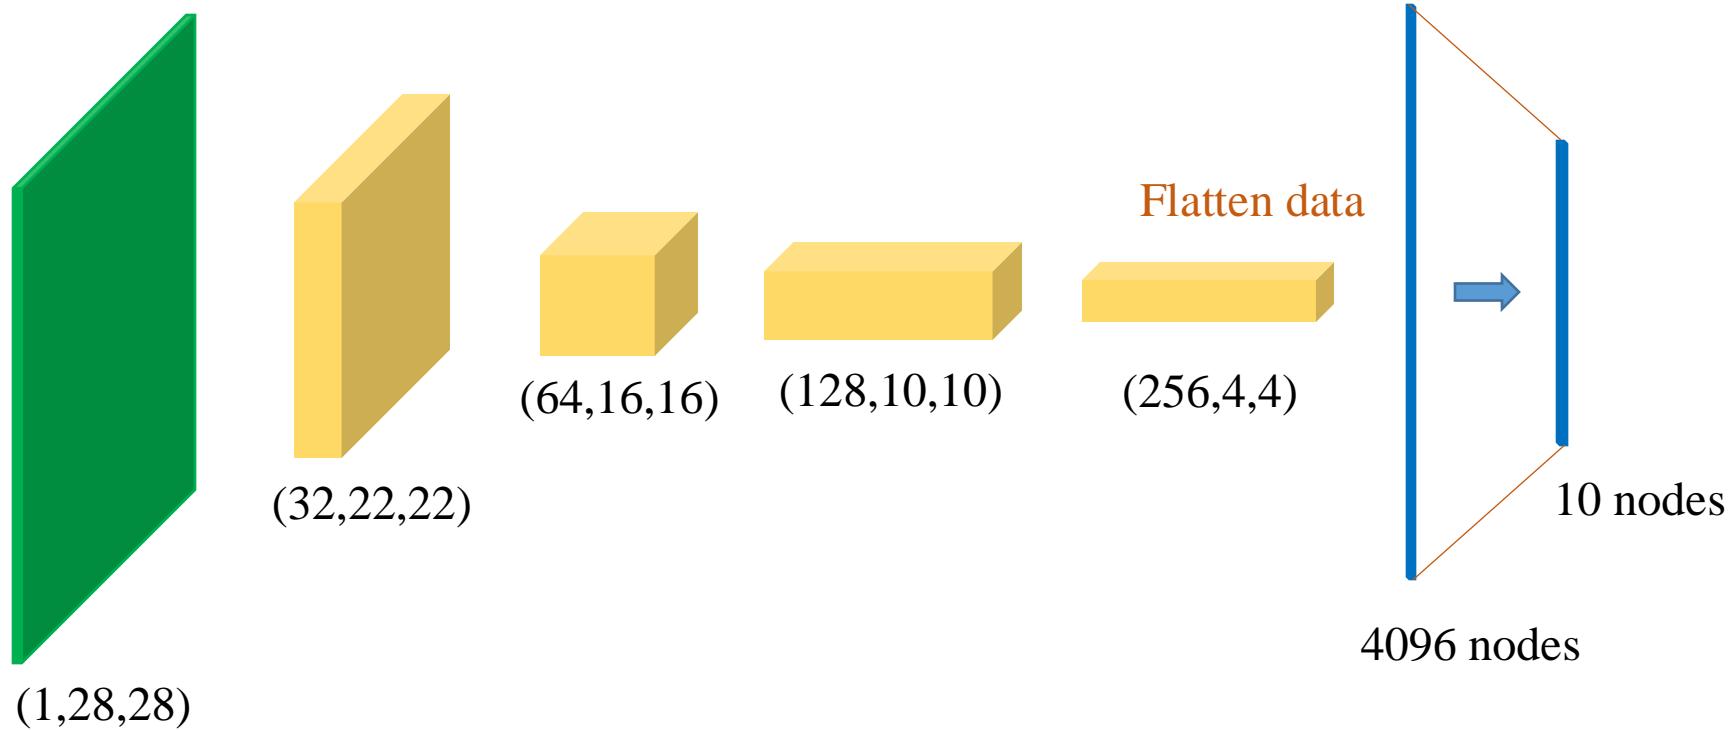
Training set: 60000 samples

Testing set: 10000 samples



Convolutional Neural Network

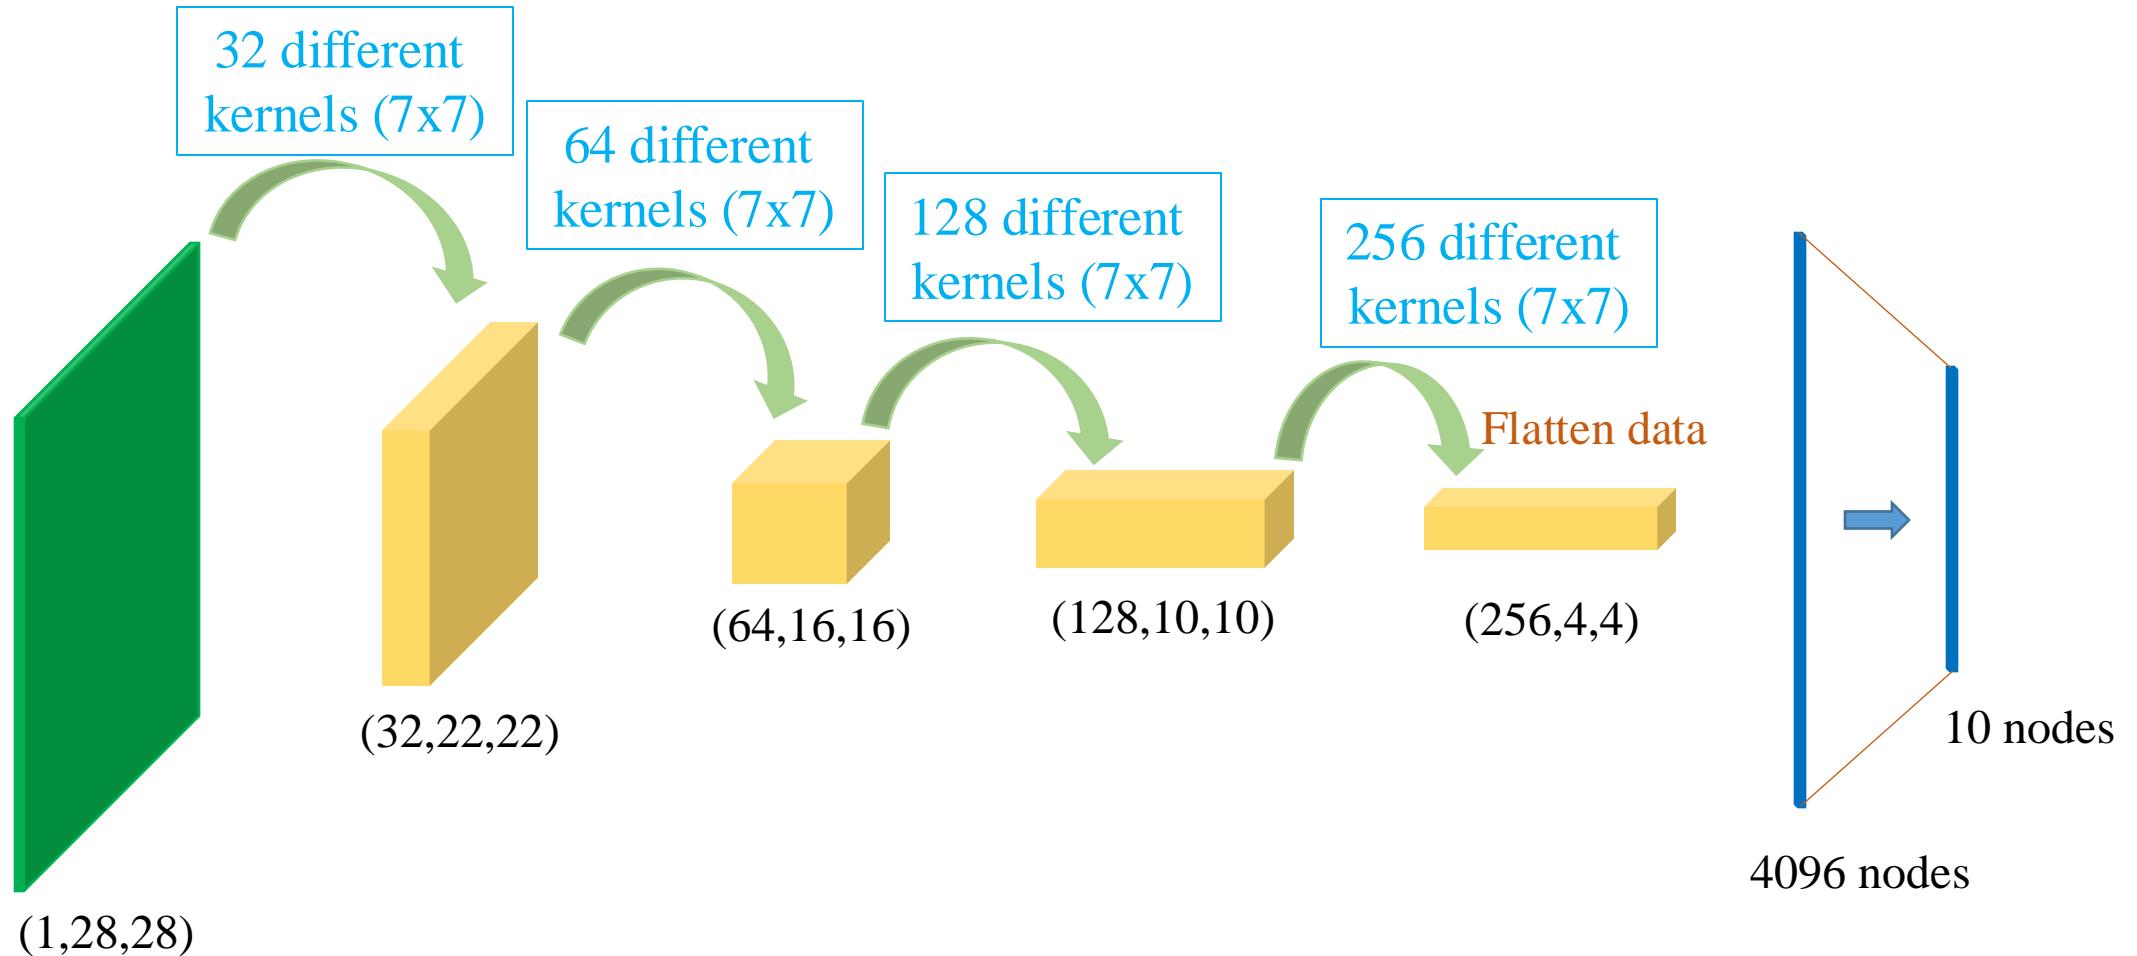
❖ Apply for Fashion-MNIST dataset



Convolutional Neural Network

❖ Apply for Fashion-MNIST dataset

demo



Simple Convolutional Neural Network

```
class CustomModel(nn.Module):
    def __init__(self):
        super(CustomModel, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, kernel_size=7)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=7)
        self.conv3 = nn.Conv2d(64, 128, kernel_size=7)
        self.conv4 = nn.Conv2d(128, 256, kernel_size=7)
        self.flatten = nn.Flatten()
        self.dense = nn.Linear(4*4*256, 10)
        self.relu = nn.ReLU()

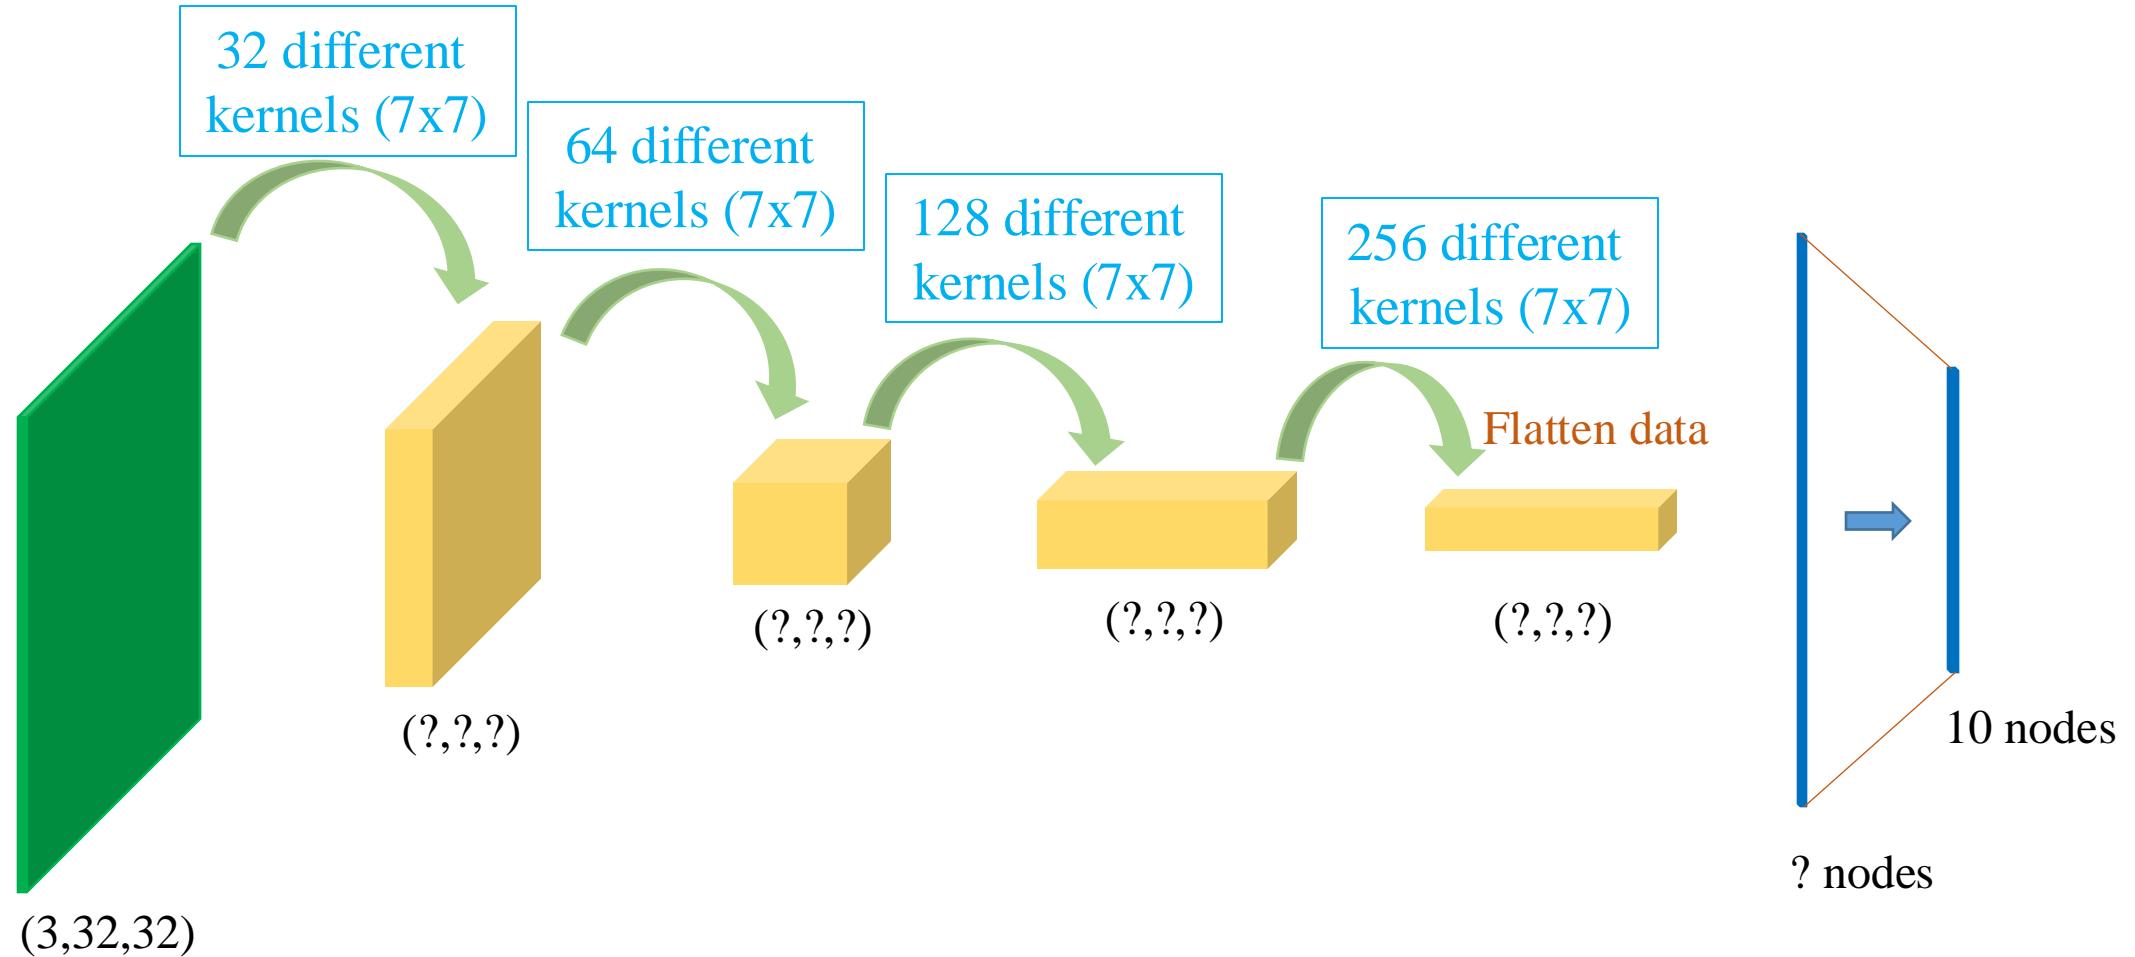
    def forward(self, x):
        x = self.relu(self.conv1(x))
        x = self.relu(self.conv2(x))
        x = self.relu(self.conv3(x))
        x = self.relu(self.conv4(x))
        x = self.flatten(x)
        x = self.dense(x)
        return x

model = CustomModel()
```

Convolutional Neural Network

❖ Apply for Cifar-10 dataset

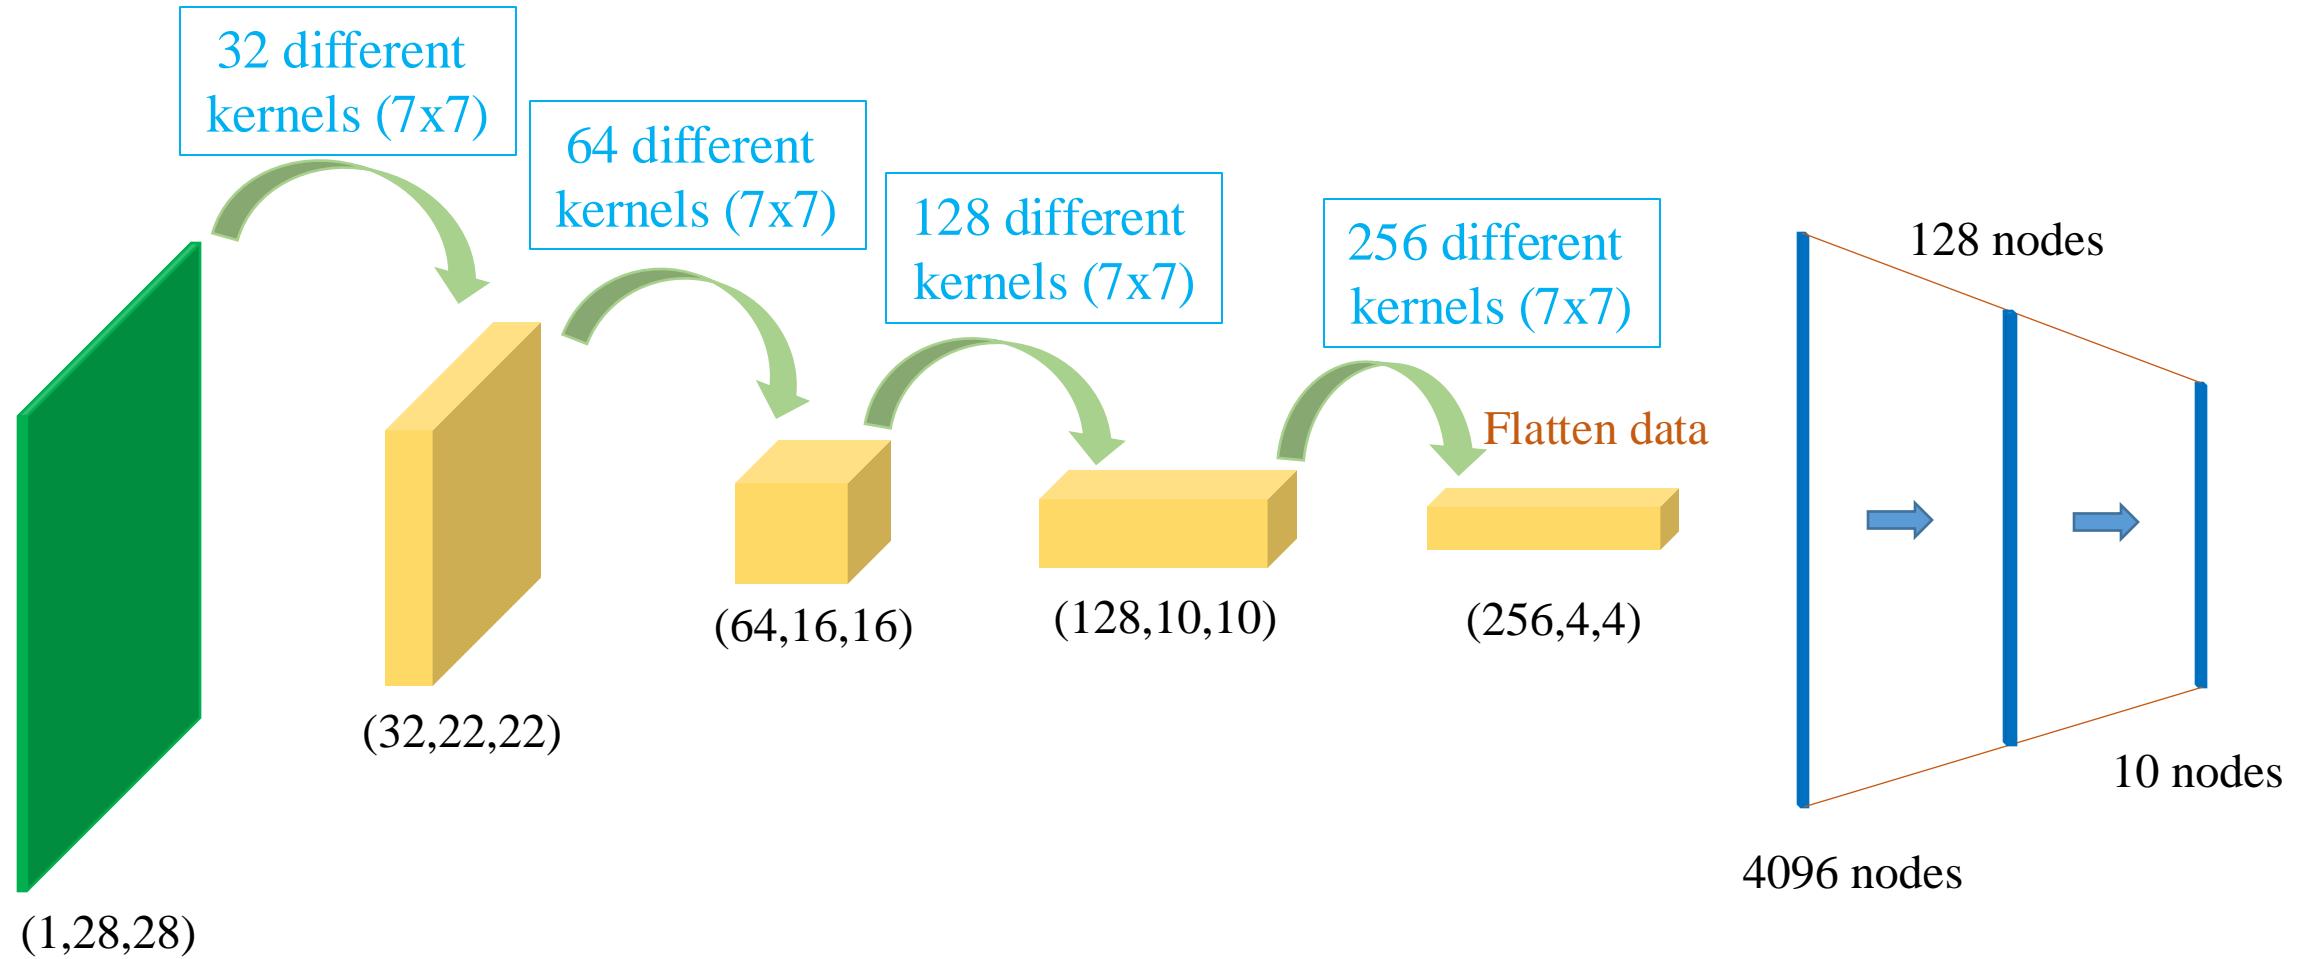
demo



More Examples

Convolutional Neural Network

❖ Apply for Fashion-MNIST dataset: case 1



```
class CustomModel(nn.Module):
    def __init__(self):
        super(CustomModel, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, kernel_size=7)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=7)
        self.conv3 = nn.Conv2d(64, 128, kernel_size=7)
        self.conv4 = nn.Conv2d(128, 256, kernel_size=7)
        self.flatten = nn.Flatten()
        self.dense1 = nn.Linear(4*4*256, 128)
        self.dense2 = nn.Linear(128, 10)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.relu(self.conv1(x))
        x = self.relu(self.conv2(x))
        x = self.relu(self.conv3(x))
        x = self.relu(self.conv4(x))
        x = self.flatten(x)
        x = self.relu(self.dense1(x))
        x = self.dense2(x)
        return x

model = CustomModel()
model = model.to(device)
```

```
# Load FashionMNIST dataset
transform = Compose([ToTensor(),
                     Normalize((0.5, ),
                               (0.5,))])

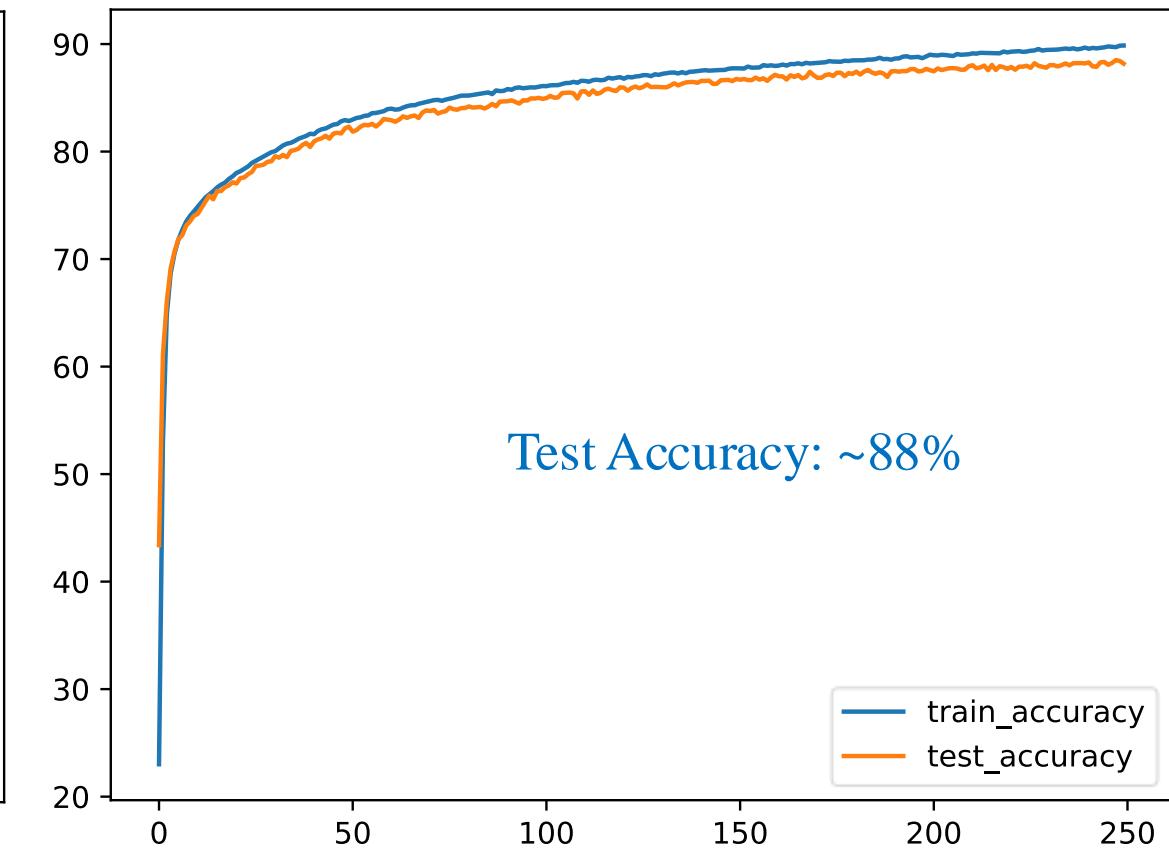
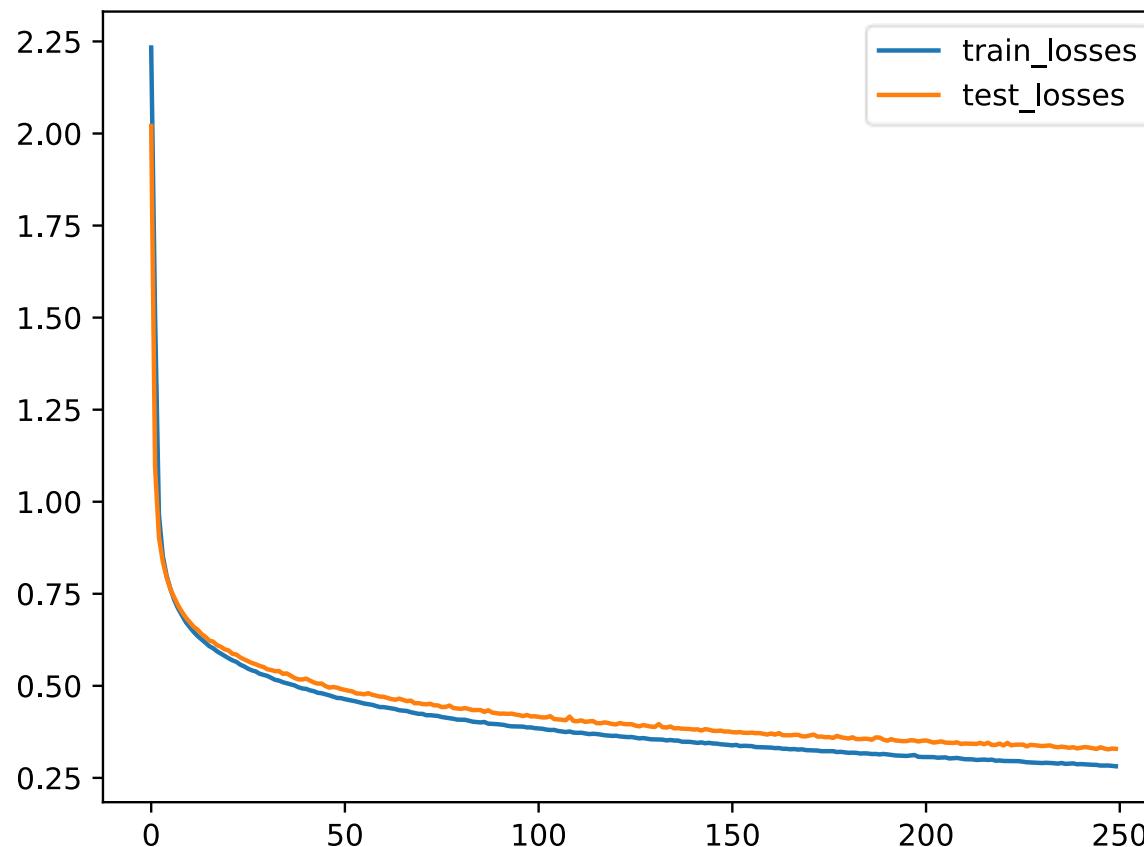
trainset = FashionMNIST(root='data',
                        train=True,
                        download=True,
                        transform=transform)
trainloader = DataLoader(trainset,
                        batch_size=1024,
                        num_workers=10,
                        shuffle=True,
                        drop_last=True)

testset = FashionMNIST(root='data',
                       train=False,
                       download=True,
                       transform=transform)
testloader = DataLoader(testset,
                        batch_size=1024,
                        num_workers=10,
                        shuffle=False)

# Loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=1e-5)
```

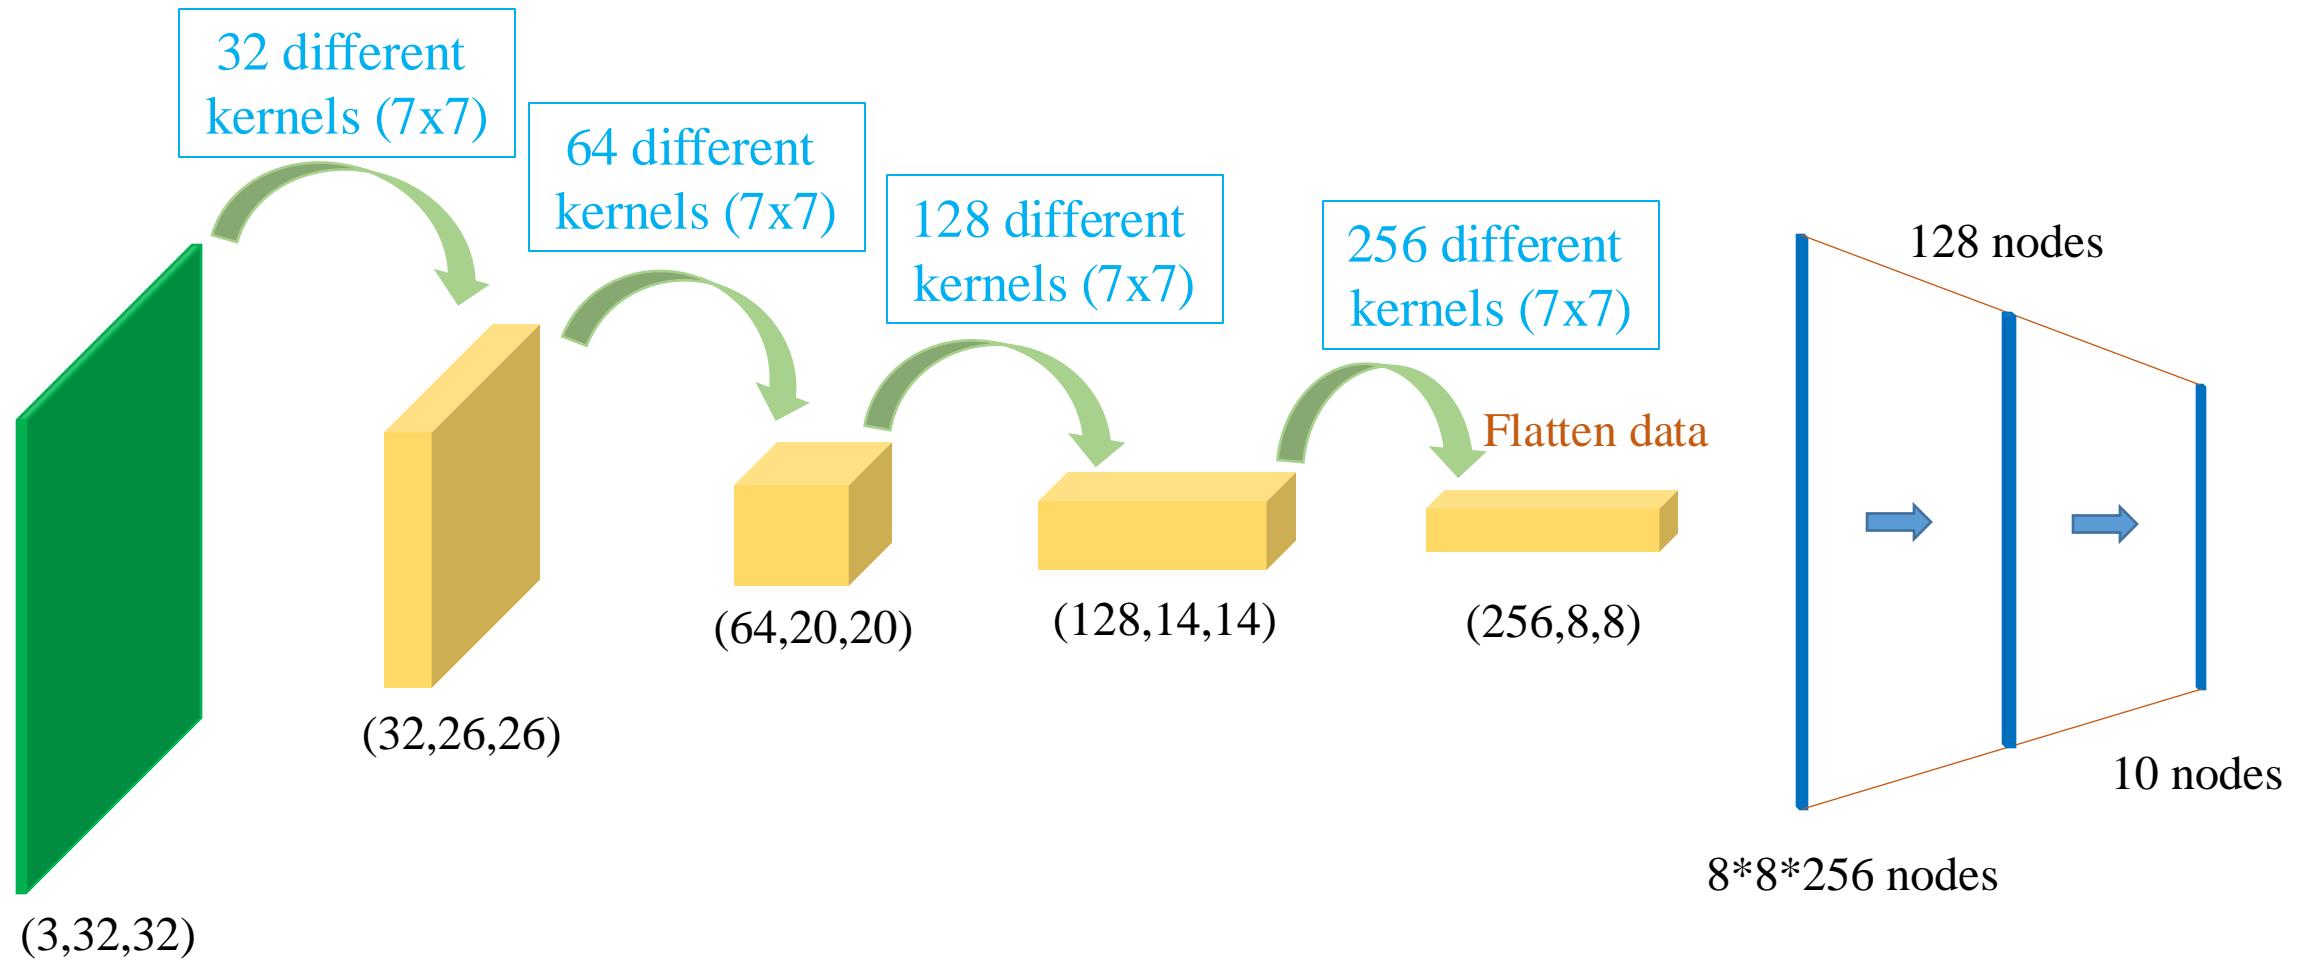
Convolutional Neural Network

❖ Apply for Fashion-MNIST dataset: case 1



Convolutional Neural Network

❖ Apply for Cifar-10 dataset: case 2



```
class CustomModel(nn.Module):
    def __init__(self):
        super(CustomModel, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, kernel_size=7)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=7)
        self.conv3 = nn.Conv2d(64, 128, kernel_size=7)
        self.conv4 = nn.Conv2d(128, 256, kernel_size=7)
        self.flatten = nn.Flatten()
        self.dense1 = nn.Linear(8*8*256, 128)
        self.dense2 = nn.Linear(128, 10)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.relu(self.conv1(x))
        x = self.relu(self.conv2(x))
        x = self.relu(self.conv3(x))
        x = self.relu(self.conv4(x))
        x = self.flatten(x)
        x = self.relu(self.dense1(x))
        x = self.dense2(x)
        return x

model = CustomModel()
model = model.to(device)
```

```
# Load CIFAR10 dataset
transform = Compose([ToTensor(),
                     Normalize((0.5, 0.5, 0.5),
                               (0.5, 0.5, 0.5))])

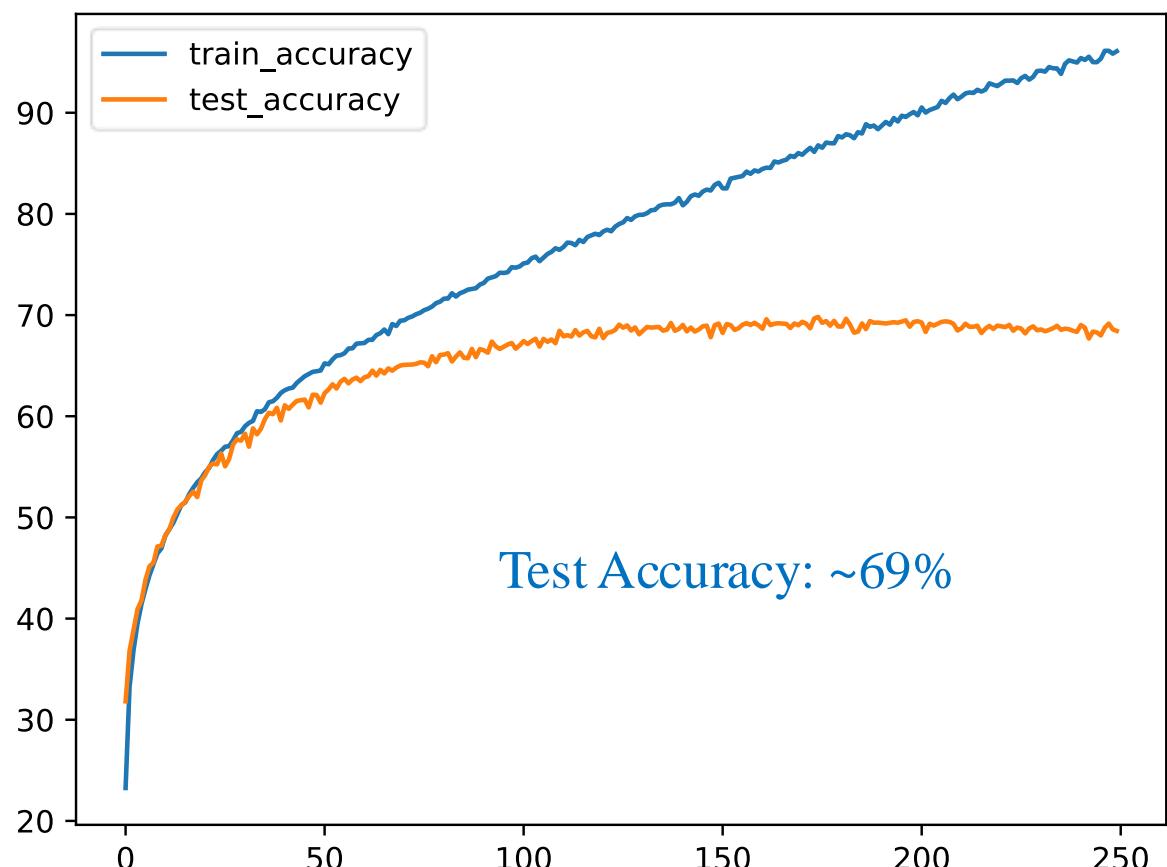
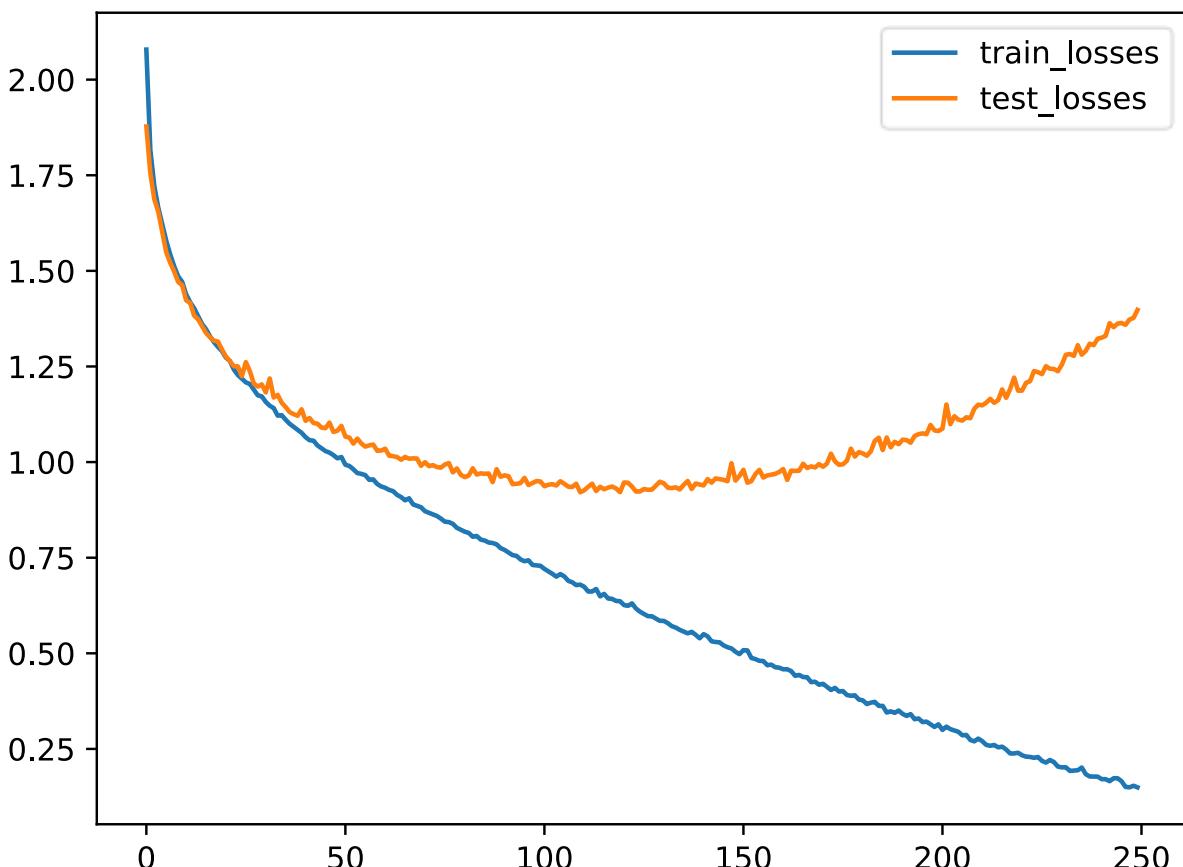
trainset = CIFAR10(root='data',
                   train=True,
                   download=True,
                   transform=transform)
trainloader = DataLoader(trainset,
                        batch_size=1024,
                        num_workers=10,
                        shuffle=True,
                        drop_last=True)

testset = CIFAR10(root='data',
                  train=False,
                  download=True,
                  transform=transform)
testloader = DataLoader(testset,
                        batch_size=1024,
                        num_workers=10,
                        shuffle=False)

# Loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=1e-5)
```

Convolutional Neural Network

❖ Apply for Cifar-10 dataset: case 2



Test Accuracy from MLP: ~53%

Outline

SECTION 1

MLP Limitations

SECTION 2

Convolutional Layer

SECTION 3

Standard CNNs



Feature map (220x220)

max pooling
(2x2)



Feature map (220x220)

Average
Pooling (2x2)



Feature map
(110x110)

Down-sample Feature Map

❖ Max pooling: Features are preserved

v_1	v_2	v_3	v_4
v_5	v_6	v_7	v_8
v_9	v_{10}	v_{11}	v_{12}
v_{13}	v_{14}	v_{15}	v_{16}

Data



2x2 max pooling(

v_1	v_2	v_3	v_4
v_5	v_6	v_7	v_8
v_9	v_{10}	v_{11}	v_{12}
v_{13}	v_{14}	v_{15}	v_{16}

) =

m_1	m_2
m_3	m_4

$$m_1 = \max(v_1, v_2, v_5, v_6)$$
$$m_2 = \max(v_3, v_4, v_7, v_8)$$
$$m_3 = \max(v_9, v_{10}, v_{13}, v_{14})$$
$$m_4 = \max(v_{11}, v_{12}, v_{15}, v_{16})$$

`nn.MaxPool2d(2, 2)`



Feature map (220x220)

max pooling
(2x2) →



Feature map
(110x110)

max pooling
(2x2) →



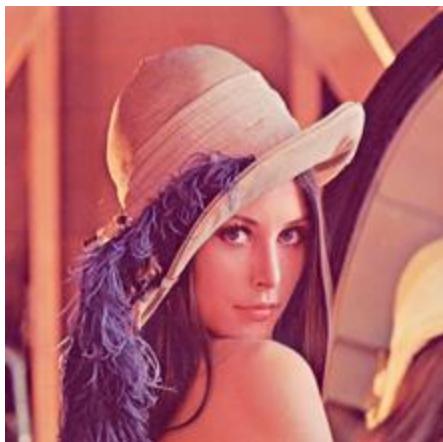
Feature map
(55x55)

Down-sample Feature Map

❖ Average pooling: Features are preserved

v_1	v_2	v_3	v_4
v_5	v_6	v_7	v_8
v_9	v_{10}	v_{11}	v_{12}
v_{13}	v_{14}	v_{15}	v_{16}

Data



average
pooling(

v_1	v_2	v_3	v_4
v_5	v_6	v_7	v_8
v_9	v_{10}	v_{11}	v_{12}
v_{13}	v_{14}	v_{15}	v_{16}

) =

m_1	m_2
m_3	m_4

$$m_1 = \text{mean}(v_1, v_2, v_5, v_6)$$

$$m_2 = \text{mean}(v_3, v_4, v_7, v_8)$$

$$m_3 = \text{mean}(v_9, v_{10}, v_{13}, v_{14})$$

$$m_4 = \text{mean}(v_{11}, v_{12}, v_{15}, v_{16})$$

`nn.AvgPool2d(2, 2)`



Feature map (220x220)

Average
Pooling (2x2) →



Feature map
(110x110)

Average
Pooling (2x2) →

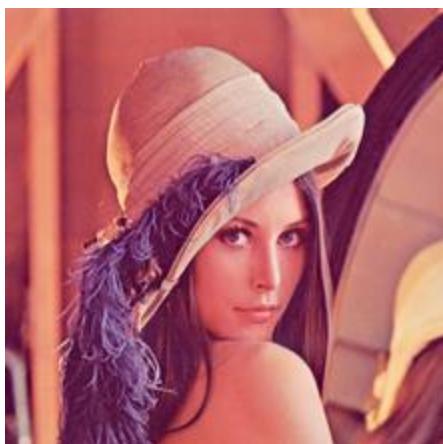


Feature map
(55x55)

Down-sample Feature Map

❖ Max pooling vs. Average pooling

`nn.MaxPool2d(2, 2)`



Feature map (220x220)

max pooling
(2x2)



max pooling
(2x2)



Feature map
(55x55)

`nn.AvgPool2d(2, 2)`



Feature map (220x220)

Average
Pooling (2x2)



Average
Pooling (2x2)



Feature map
(110x110)

Feature map
(55x55)

```

class CustomModel(nn.Module):
    def __init__(self):
        super(CustomModel, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, kernel_size=3)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=3)
        self.conv3 = nn.Conv2d(64, 128, kernel_size=3)
        self.conv4 = nn.Conv2d(128, 256, kernel_size=3)
        self.conv5 = nn.Conv2d(256, 512, kernel_size=3)
        self.relu = nn.ReLU()
        self.pool = nn.MaxPool2d(2, 2) # 2x2 Max pooling
        self.flatten = nn.Flatten()
        self.dense = nn.Linear(2*2*512, 10)

```

```

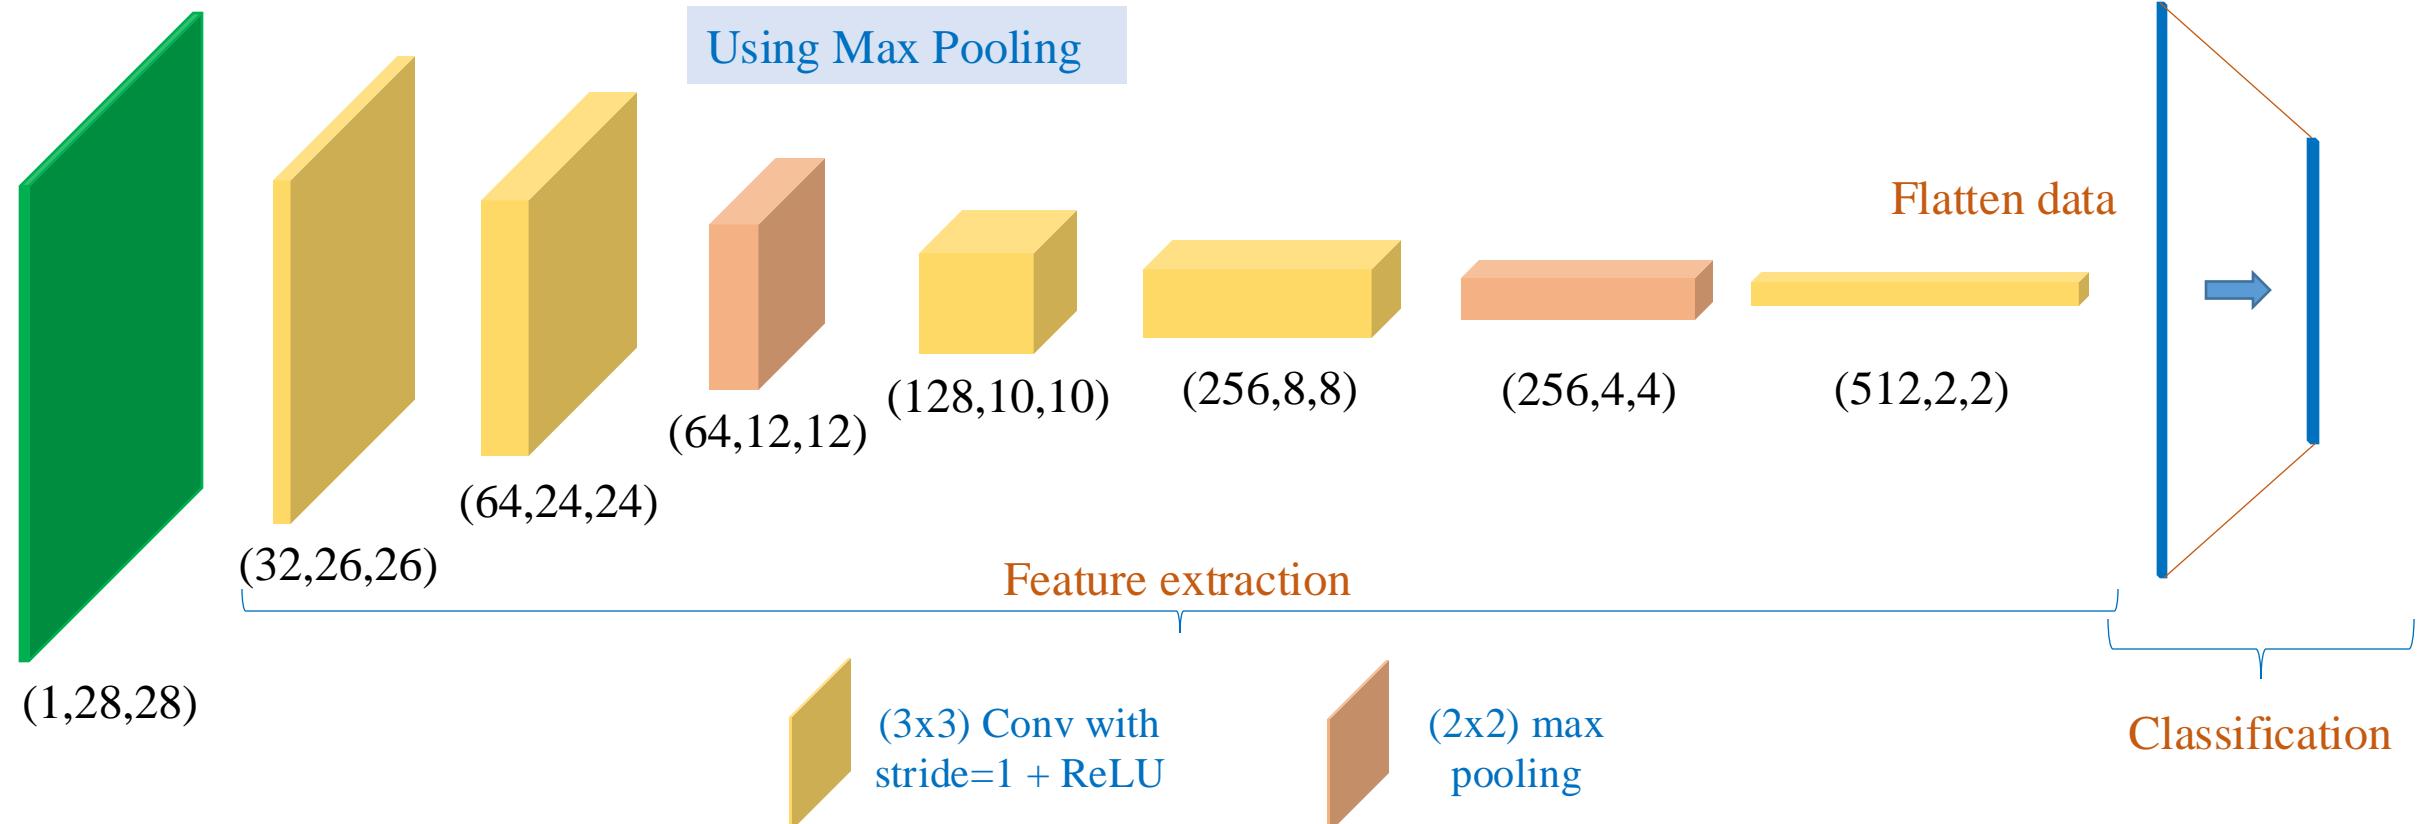
def forward(self, x):
    x = self.relu(self.conv1(x))
    x = self.relu(self.conv2(x))
    x = self.pool(x)

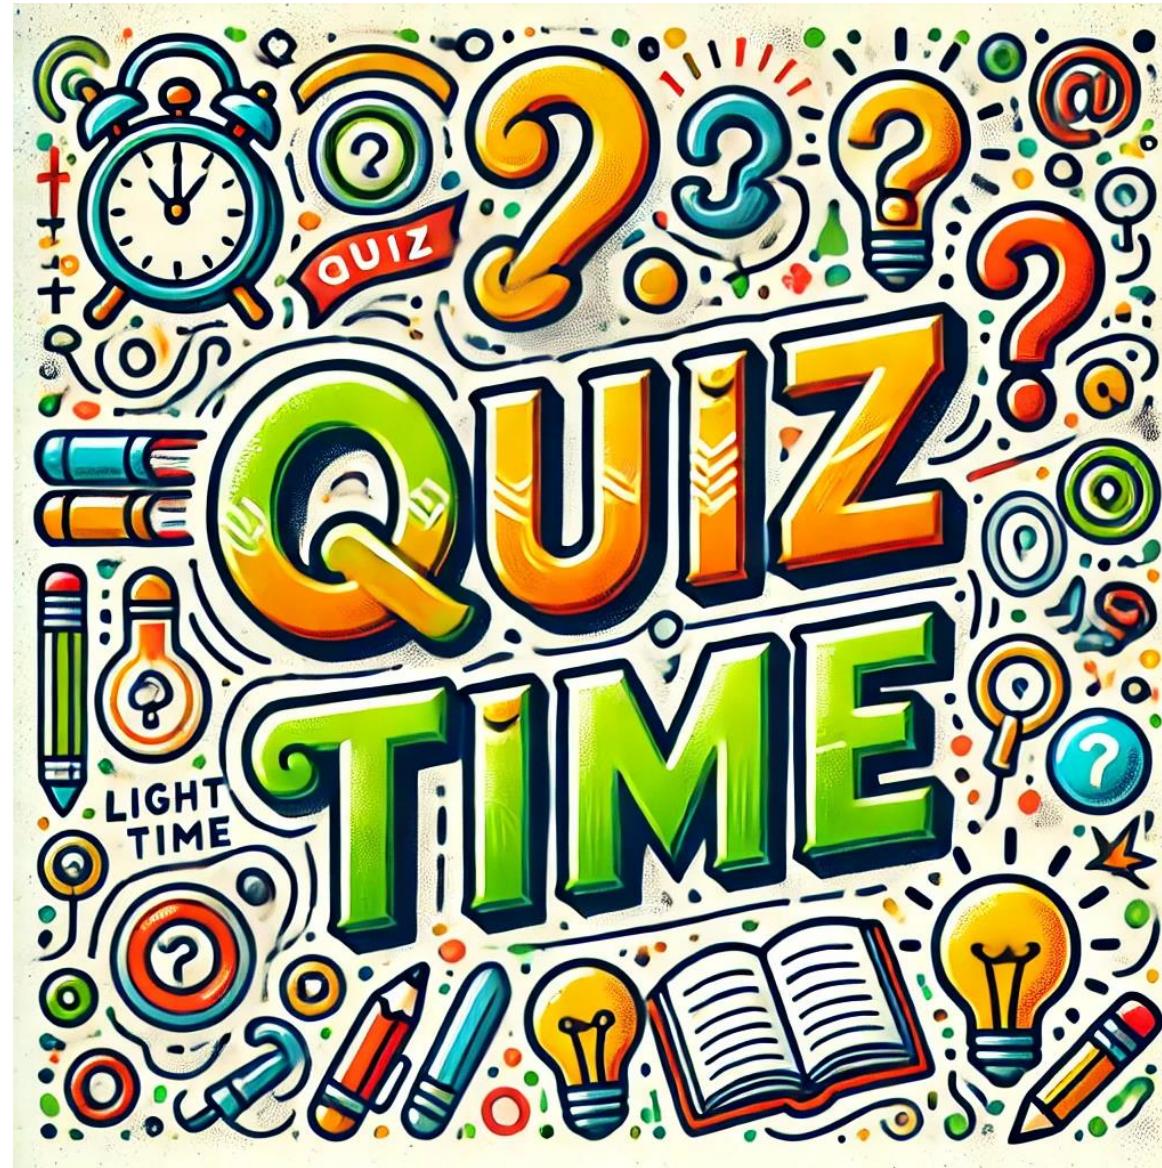
    x = self.relu(self.conv3(x))
    x = self.relu(self.conv4(x))
    x = self.pool(x)

    x = self.relu(self.conv5(x))
    x = self.flatten(x)
    x = self.dense(x)

    return x

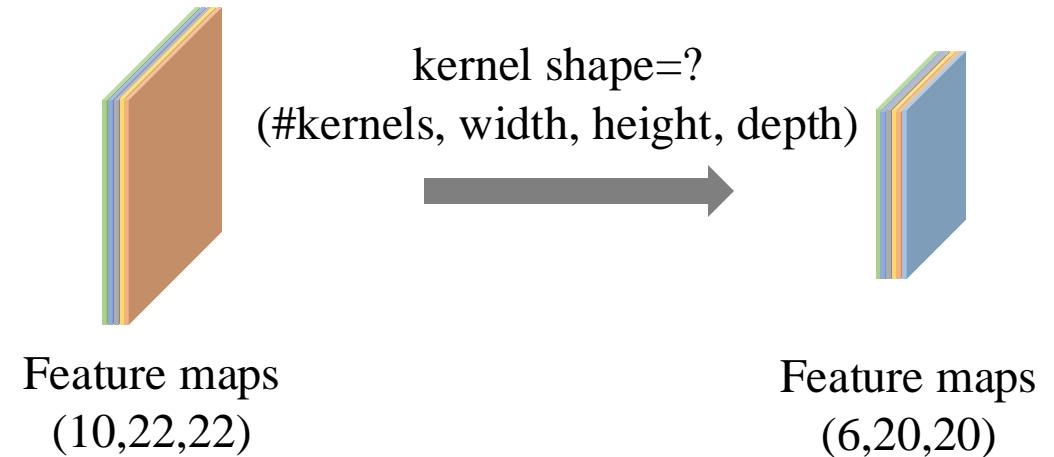
```





Question 1

❖ Giả sử chúng ta dùng kernel có width=3 và height=3, shape của kernel là gì (#kernels là số lượng kernel)?

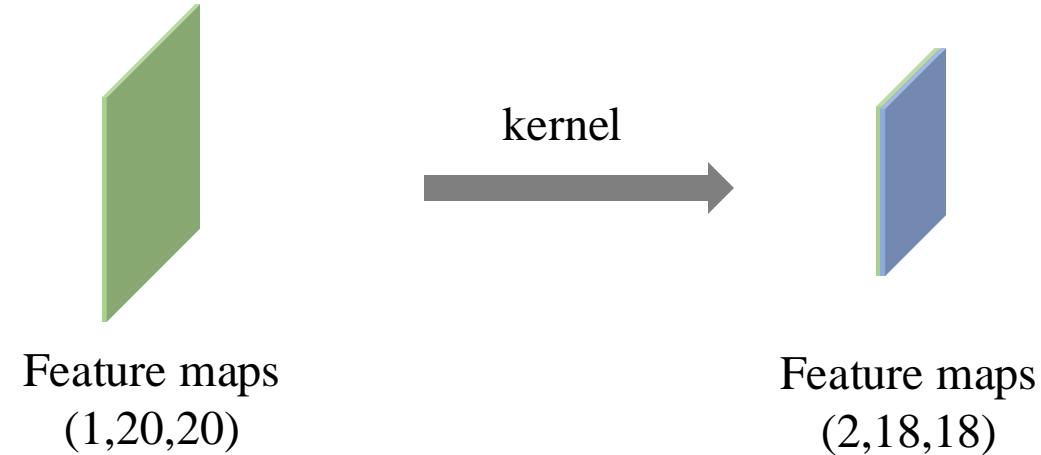


- a) (10, 3, 3, 10)
- c) (6, 3, 3, 10)

- b) (10, 3, 3, 6)
- d) Không xác định

Question 2

❖ Giả sử chúng ta dùng kernel có width=3 và height=3, tổng số lượng tham số của kernel là bao nhiêu?



a) Có 9 tham số

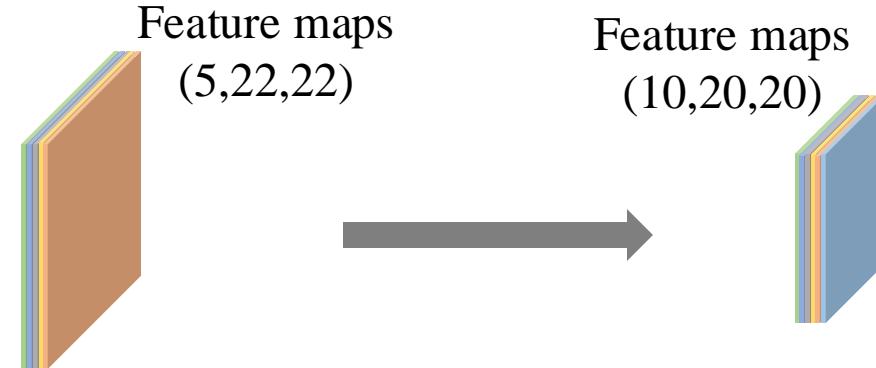
b) Có 10 tham số

c) Có 18 tham số

d) Có 20 tham số

Question 3

❖ Code Pytorch nào áp dụng đúng cho hình sau?



C1

`nn.Conv2d(5, 10, kernel_size=3)`

C2

`nn.Conv2d(10, 10, kernel_size=3)`

C3

`nn.Conv2d(5, 5, kernel_size=3)`

C4

`nn.Conv2d(10, 5, kernel_size=3)`

Question 4

❖ Input (cho model) có shape là $(1, 50, 50)$ thì output có shape là gì (bỏ qua phần batch size)?

```
class CustomModel(nn.Module):
    def __init__(self):
        super(CustomModel, self).__init__()
        self.conv = nn.Conv2d(1, 32, kernel_size=3)
        self.pool = nn.MaxPool2d(2, 2)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.conv(x)
        x = self.relu(x)
        x = self.pool(x)
        return x

model = CustomModel()
summary(model, (1, 50, 50))
```

a) $(1, 48, 48)$

b) $(1, 24, 24)$

c) $(32, 48, 48)$

d) $(32, 24, 24)$

Question 5

❖ Input (cho model) có shape là $(1, 50, 50)$ thì output có shape là gì (bỏ qua phần batch size)?

```
class CustomModel(nn.Module):
    def __init__(self):
        super(CustomModel, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, kernel_size=3)
        self.pool1 = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(32, 32, kernel_size=5)
        self.pool2 = nn.MaxPool2d(2, 2)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.relu(self.conv1(x))
        x = self.pool1(x)
        x = self.relu(self.conv2(x))
        x = self.pool2(x)
        return x

model = CustomModel()
summary(model, (1, 50, 50))
```

- a) $(32, 24, 24)$
- c) $(32, 8, 8)$

- b) $(32, 10, 10)$
- d) $(32, 6, 6)$

Question 6

❖ Input (cho model) có shape là $(1, 50, 50)$ thì output có shape là gì (bỏ qua phần batch size)?

```
class CustomModel(nn.Module):
    def __init__(self):
        super(CustomModel, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, kernel_size=3)
        self.conv2 = nn.Conv2d(1, 32, kernel_size=3)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.relu(self.conv1(x))
        x = self.relu(self.conv2(x))
        return x

model = CustomModel()
summary(model, (1, 50, 50))
```

a) shape=(32, 50, 50)

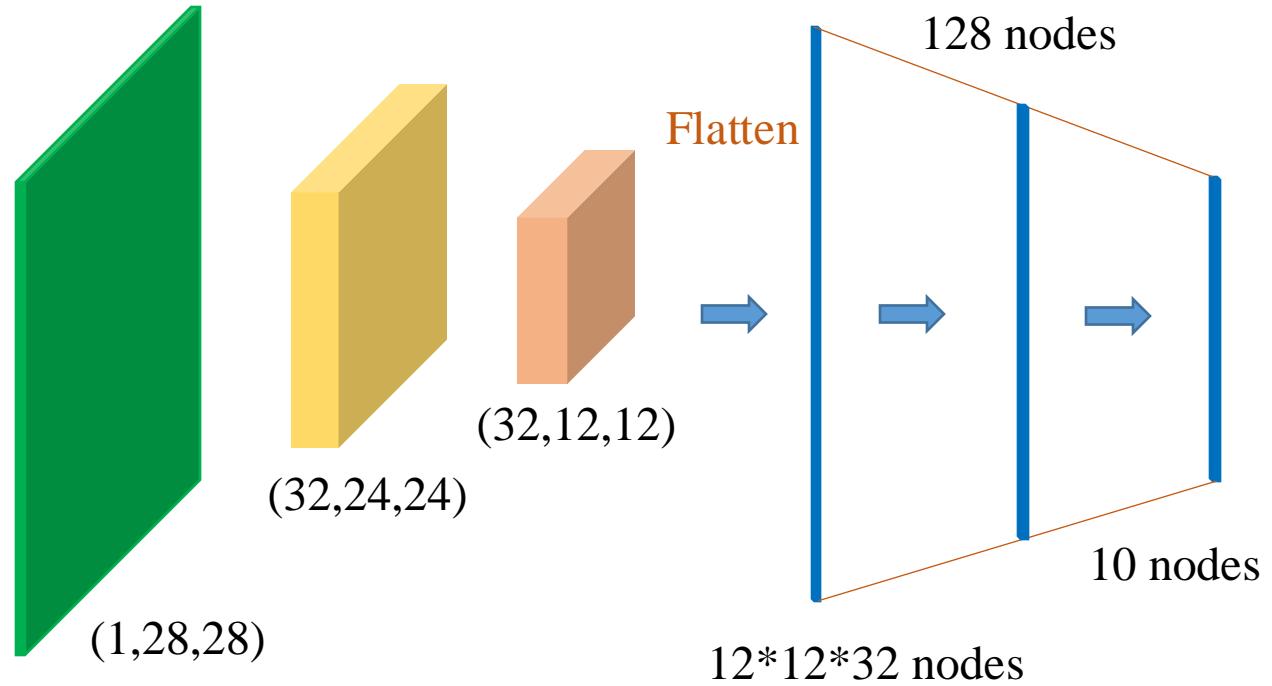
b) shape=(32, 48, 48)

c) shape=(32, 46, 46)

d) Có lỗi

Comparison of down-samplings

Model Design



(5x5) Conv with
stride=1 + ReLU

(2x2) down-
sampling

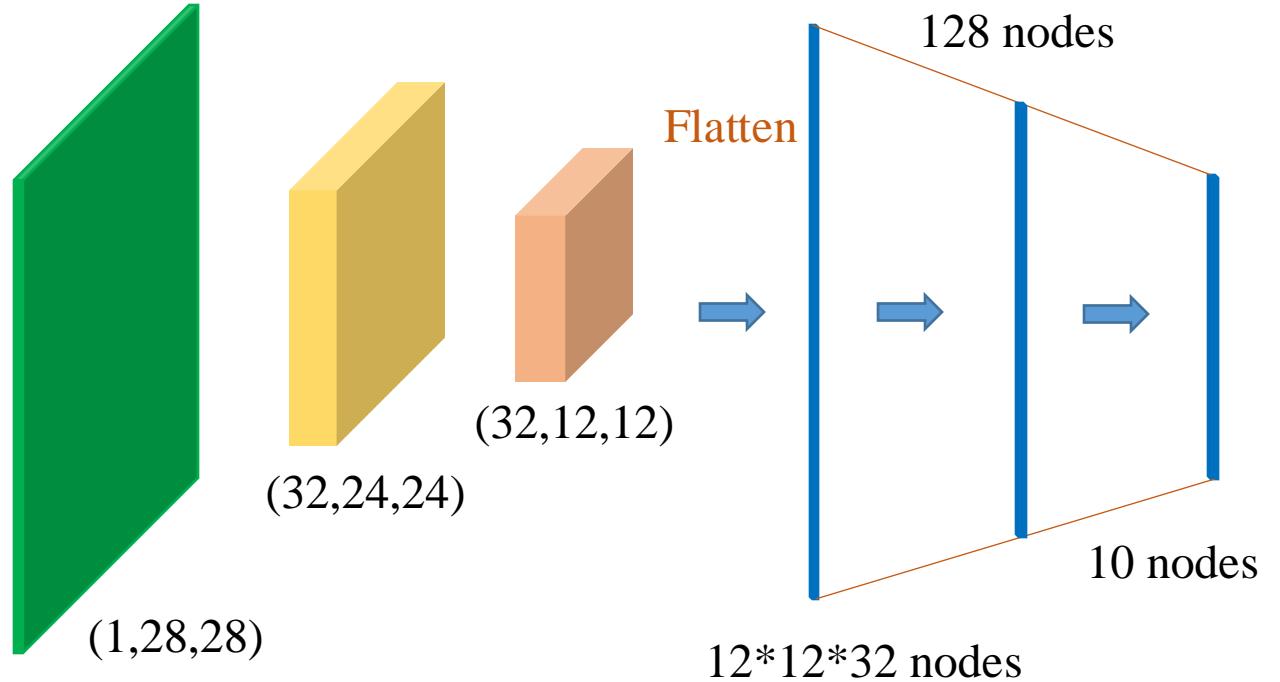
Implementation 1

```
class CustomModel(nn.Module):  
    def __init__(self):  
        super(CustomModel, self).__init__()  
        self.conv = nn.Conv2d(1, 32, kernel_size=5)  
        self.pool = nn.MaxPool2d(2, 2)  
        self.flatten = nn.Flatten()  
        self.dense1 = nn.Linear(12*12*32, 128)  
        self.dense2 = nn.Linear(128, 10)  
        self.relu = nn.ReLU()  
  
    def forward(self, x):  
        x = self.conv(x)  
        x = self.relu(x)  
        x = self.pool(x)  
        x = self.flatten(x)  
        x = self.relu(self.dense1(x))  
        x = self.dense2(x)  
        return x  
  
# create model  
model = CustomModel()  
model = model.to(device)  
  
# Loss and optimizer  
criterion = nn.CrossEntropyLoss()  
optimizer = optim.Adam(model.parameters(), lr=1e-5)
```

max pooling

Comparison of down-samplings

Model Design



(5x5) Conv with
stride=1 + ReLU

(2x2) down-
sampling

Implementation 2

```
class CustomModel(nn.Module):
    def __init__(self):
        super(CustomModel, self).__init__()
        self.conv = nn.Conv2d(1, 32, kernel_size=5)
        self.pool = nn.AvgPool2d(2, 2)
        self.flatten = nn.Flatten()
        self.dense1 = nn.Linear(12*12*32, 128)
        self.dense2 = nn.Linear(128, 10)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.conv(x)
        x = self.relu(x)
        x = self.pool(x)
        x = self.flatten(x)
        x = self.relu(self.dense1(x))
        x = self.dense2(x)
        return x
```

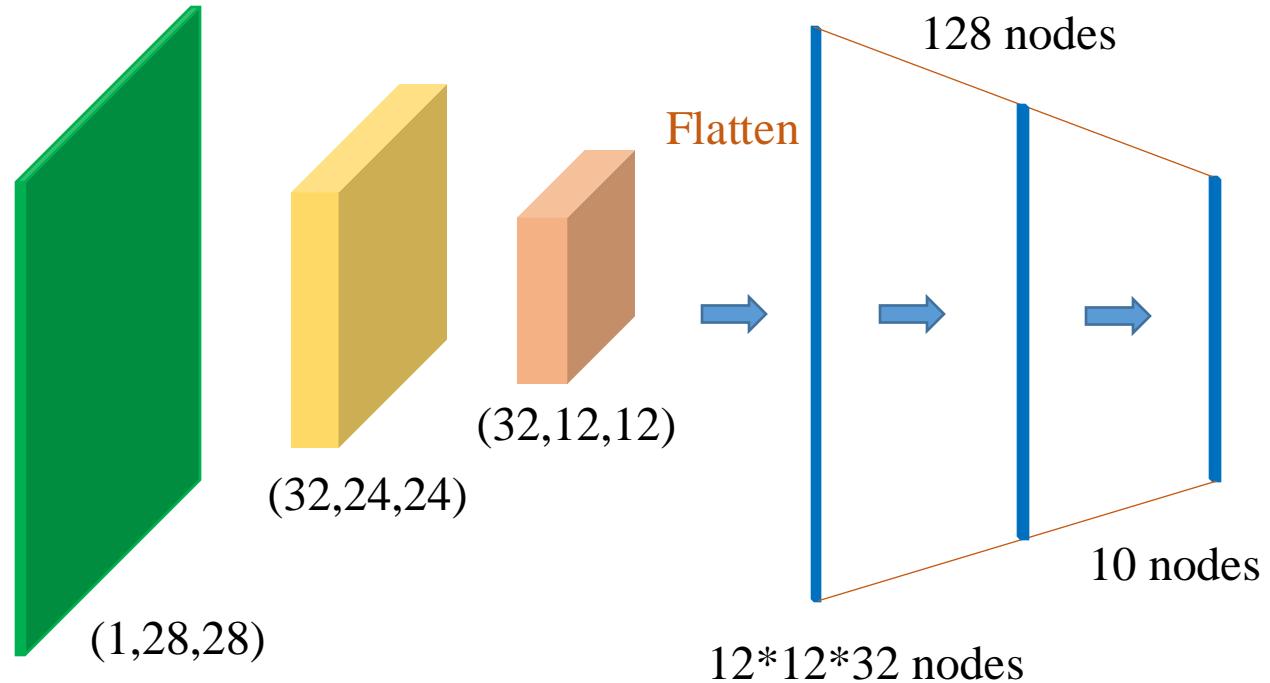
average pooling

```
# create model
model = CustomModel()
model = model.to(device)

# Loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=1e-5)
```

Comparison of down-samplings

Model Design



(5x5) Conv with
stride=1 + ReLU

(2x2) down-
sampling

```
class ResizeLayer(nn.Module):
    def __init__(self, scale_factor, mode='bilinear',
                 align_corners=False):
        super(ResizeLayer, self).__init__()
        self.scale_factor = scale_factor
        self.mode = mode
        self.align_corners = align_corners

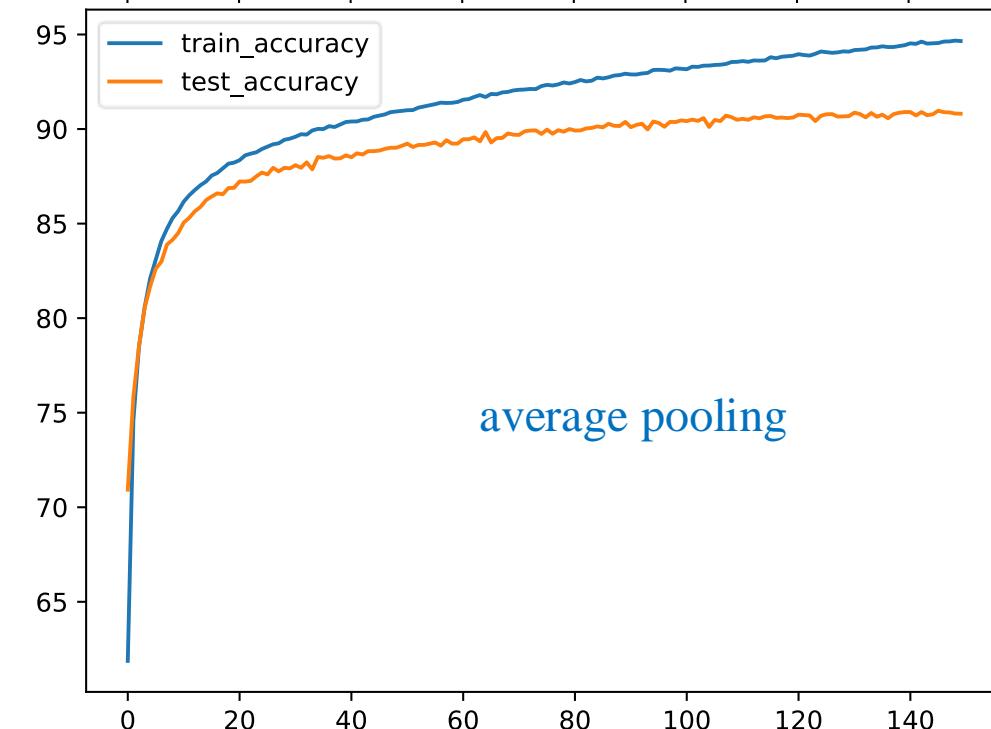
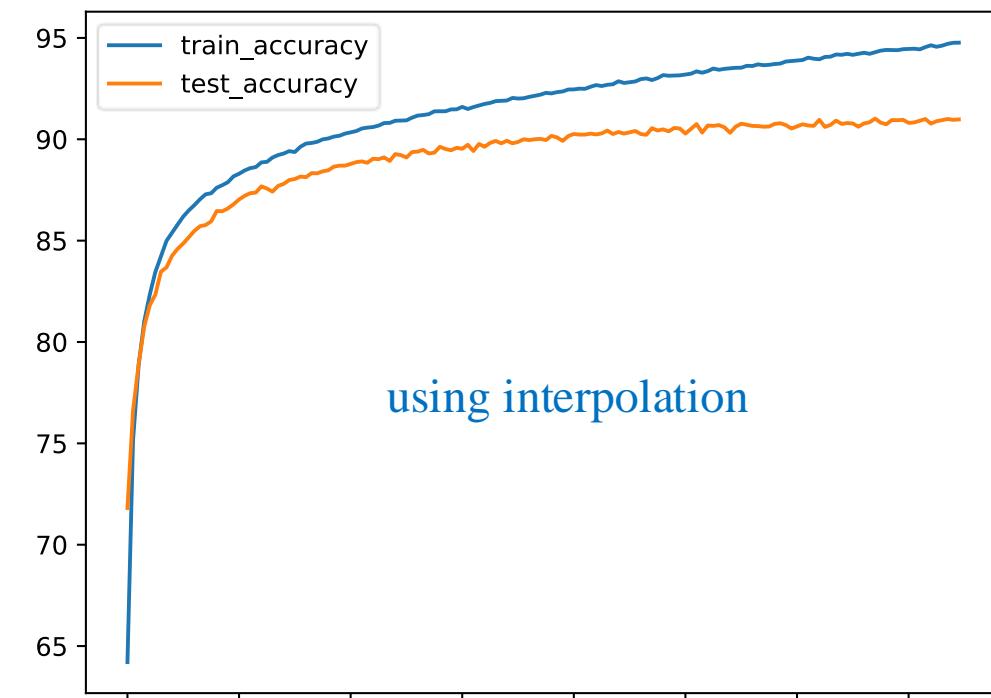
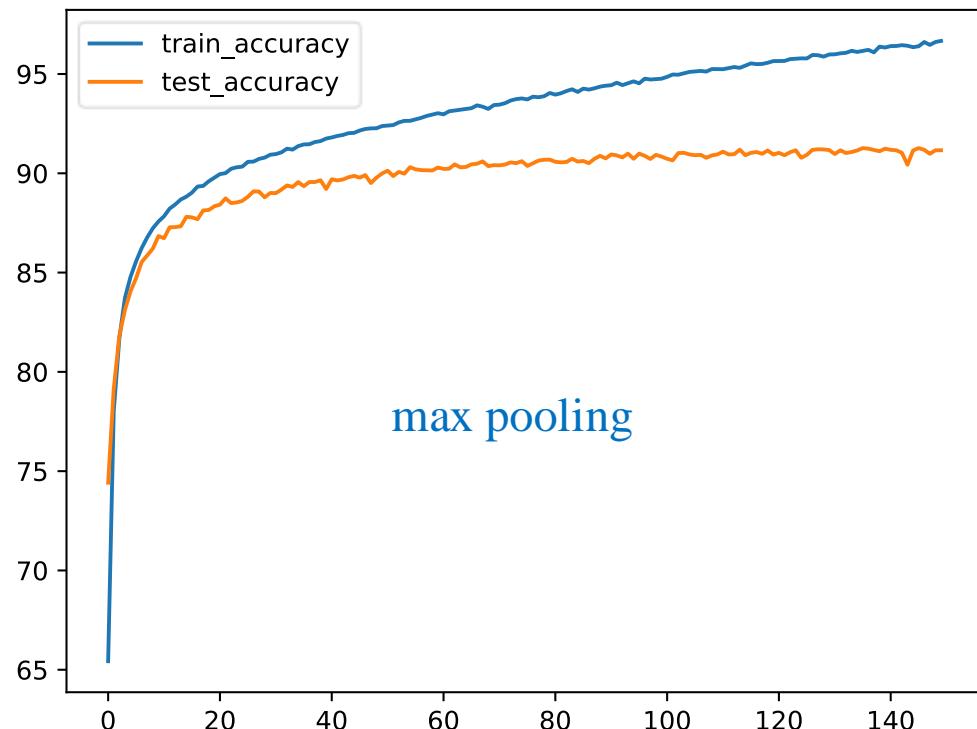
    def forward(self, x):
        return F.interpolate(x, scale_factor=self.scale_factor,
                            mode=self.mode,
                            align_corners=self.align_corners)
```

interpolation

```
class CustomModel(nn.Module):
    def __init__(self):
        super(CustomModel, self).__init__()
        self.conv = nn.Conv2d(1, 32, kernel_size=5)
        self.resize = ResizeLayer(0.5)
        self.flatten = nn.Flatten()
        self.dense1 = nn.Linear(12*12*32, 128)
        self.dense2 = nn.Linear(128, 10)
        self.relu = nn.ReLU()

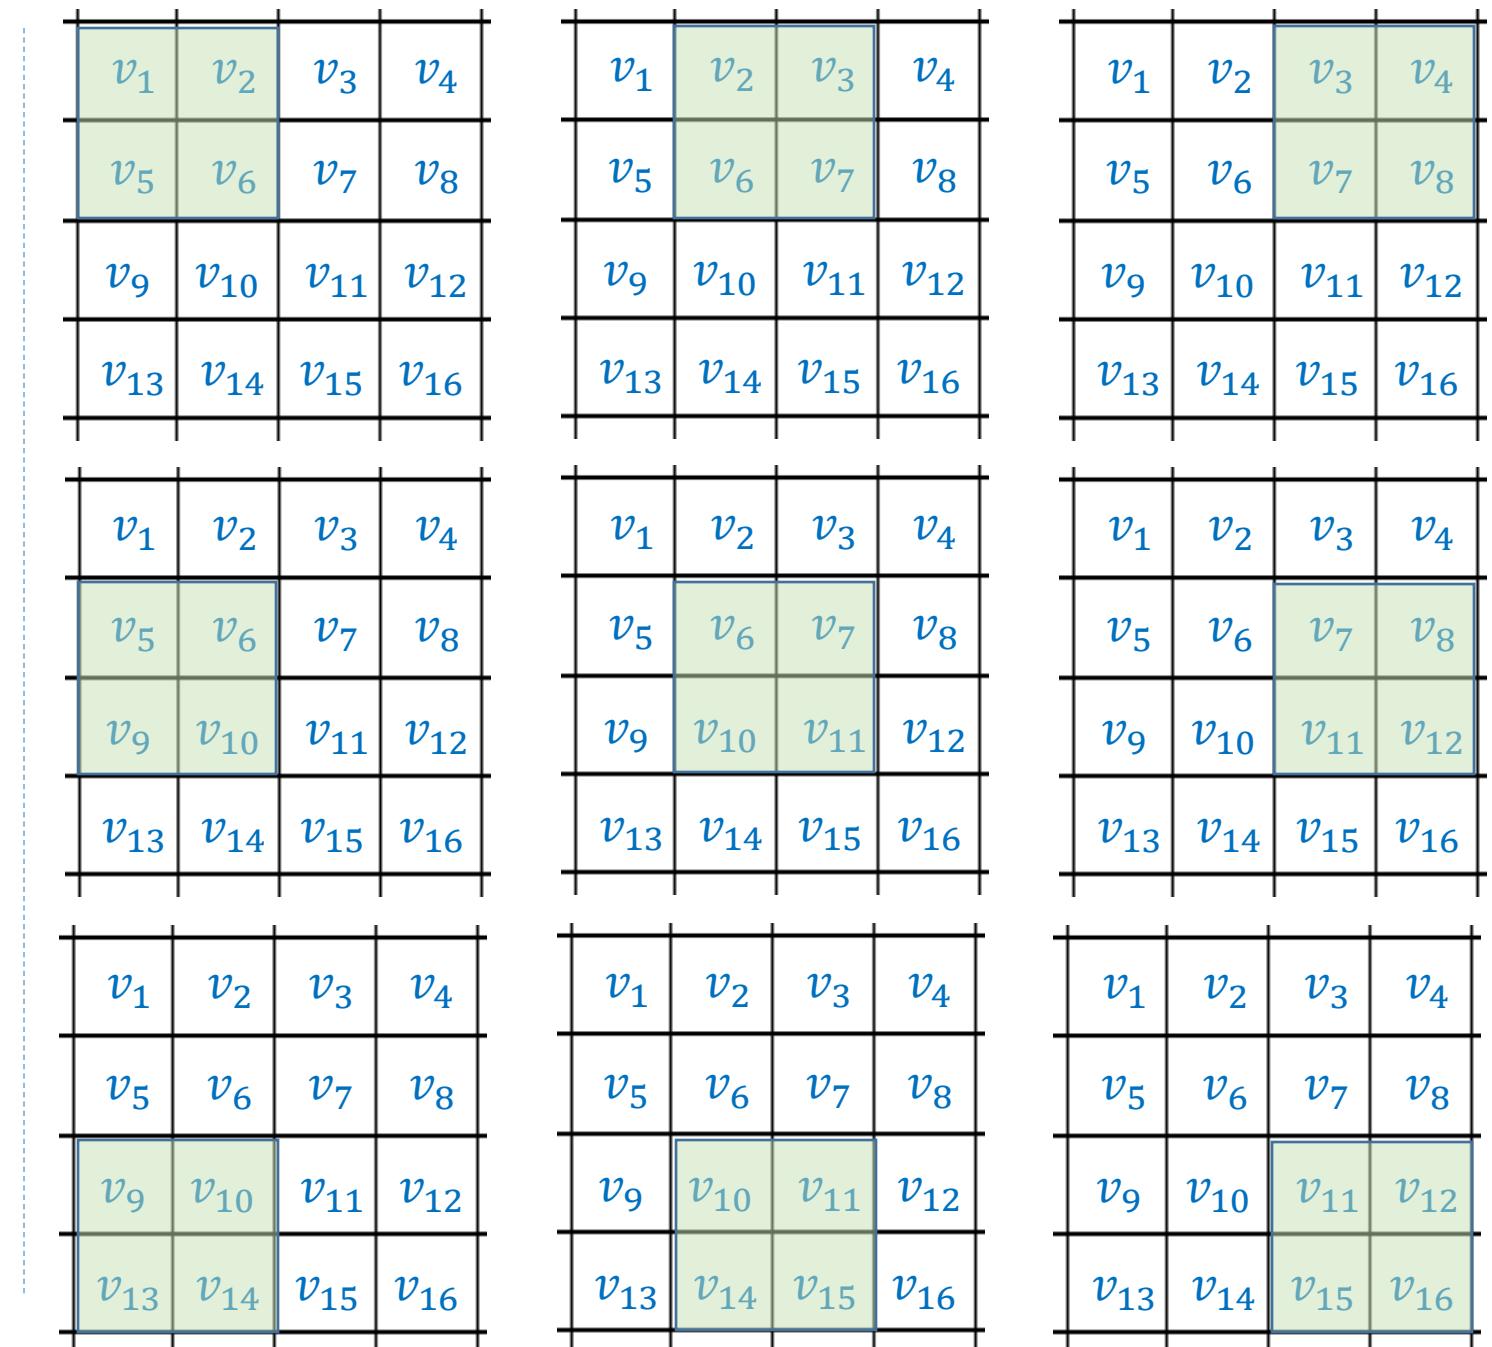
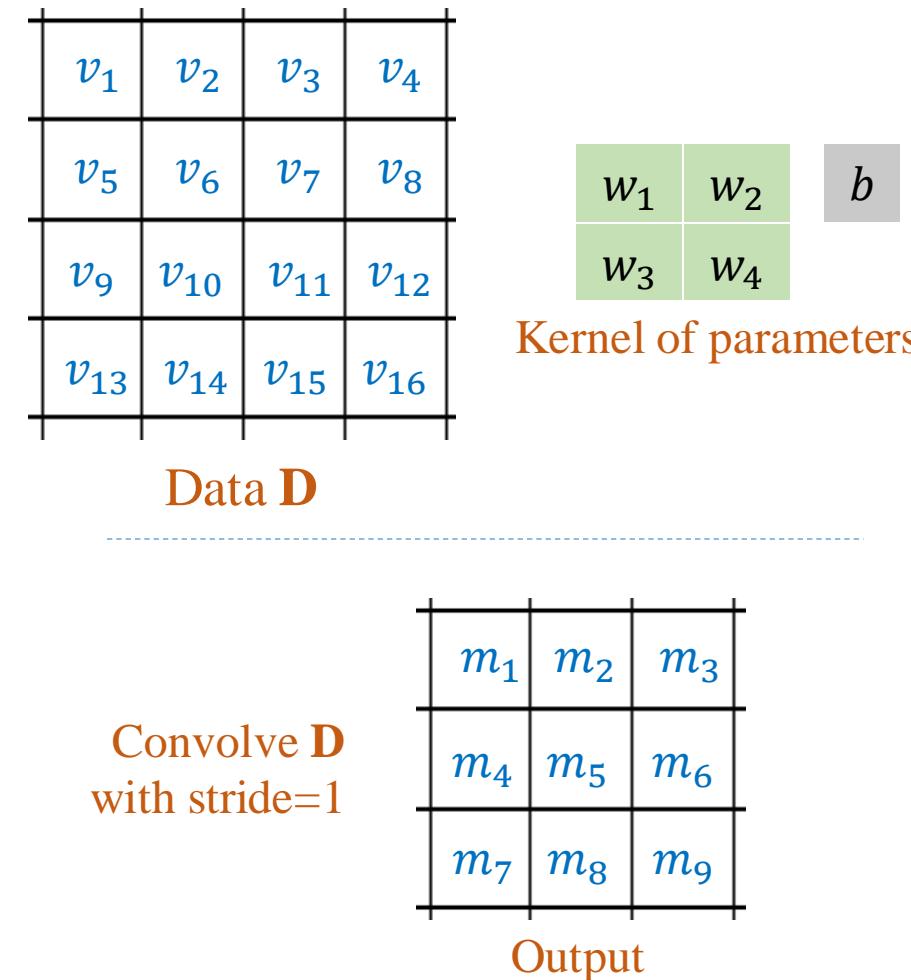
    def forward(self, x):
        x = self.relu(self.conv(x))
        x = self.resize(x)
        x = self.flatten(x)
        x = self.relu(self.dense1(x))
        x = self.dense2(x)
        return x
```

Comparison of down-samplings



Down-sample Feature Map

❖ Convolve with stride

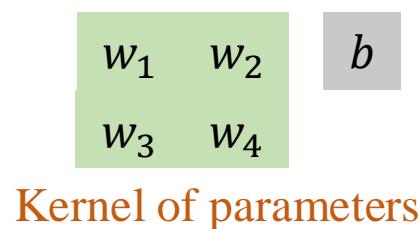


Down-sample Feature Map

❖ Convolve with stride

v_1	v_2	v_3	v_4
v_5	v_6	v_7	v_8
v_9	v_{10}	v_{11}	v_{12}
v_{13}	v_{14}	v_{15}	v_{16}

Data D



Convolve D
with stride=2

m_1	m_2
m_3	m_4

Output

v_1	v_2	v_3	v_4
v_5	v_6	v_7	v_8
v_9	v_{10}	v_{11}	v_{12}
v_{13}	v_{14}	v_{15}	v_{16}

v_1	v_2	v_3	v_4
v_5	v_6	v_7	v_8
v_9	v_{10}	v_{11}	v_{12}
v_{13}	v_{14}	v_{15}	v_{16}

v_1	v_2	v_3	v_4
v_5	v_6	v_7	v_8
v_9	v_{10}	v_{11}	v_{12}
v_{13}	v_{14}	v_{15}	v_{16}

v_1	v_2	v_3	v_4
v_5	v_6	v_7	v_8
v_9	v_{10}	v_{11}	v_{12}
v_{13}	v_{14}	v_{15}	v_{16}

```
nn.Conv2d(in_channels, out_channels, kernel_size, stride=1)
```

```
nn.Conv2d(in_channels, out_channels, kernel_size, stride=2)
```

```

class CustomModel(nn.Module):
    def __init__(self):
        super(CustomModel, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, kernel_size=3)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, stride=2)
        self.conv3 = nn.Conv2d(64, 128, kernel_size=3)
        self.conv4 = nn.Conv2d(128, 256, kernel_size=3, stride=2)
        self.conv5 = nn.Conv2d(256, 512, kernel_size=3)
        self.relu = nn.ReLU()
        self.flatten = nn.Flatten()
        self.dense = nn.Linear(2*2*512, 10)

```

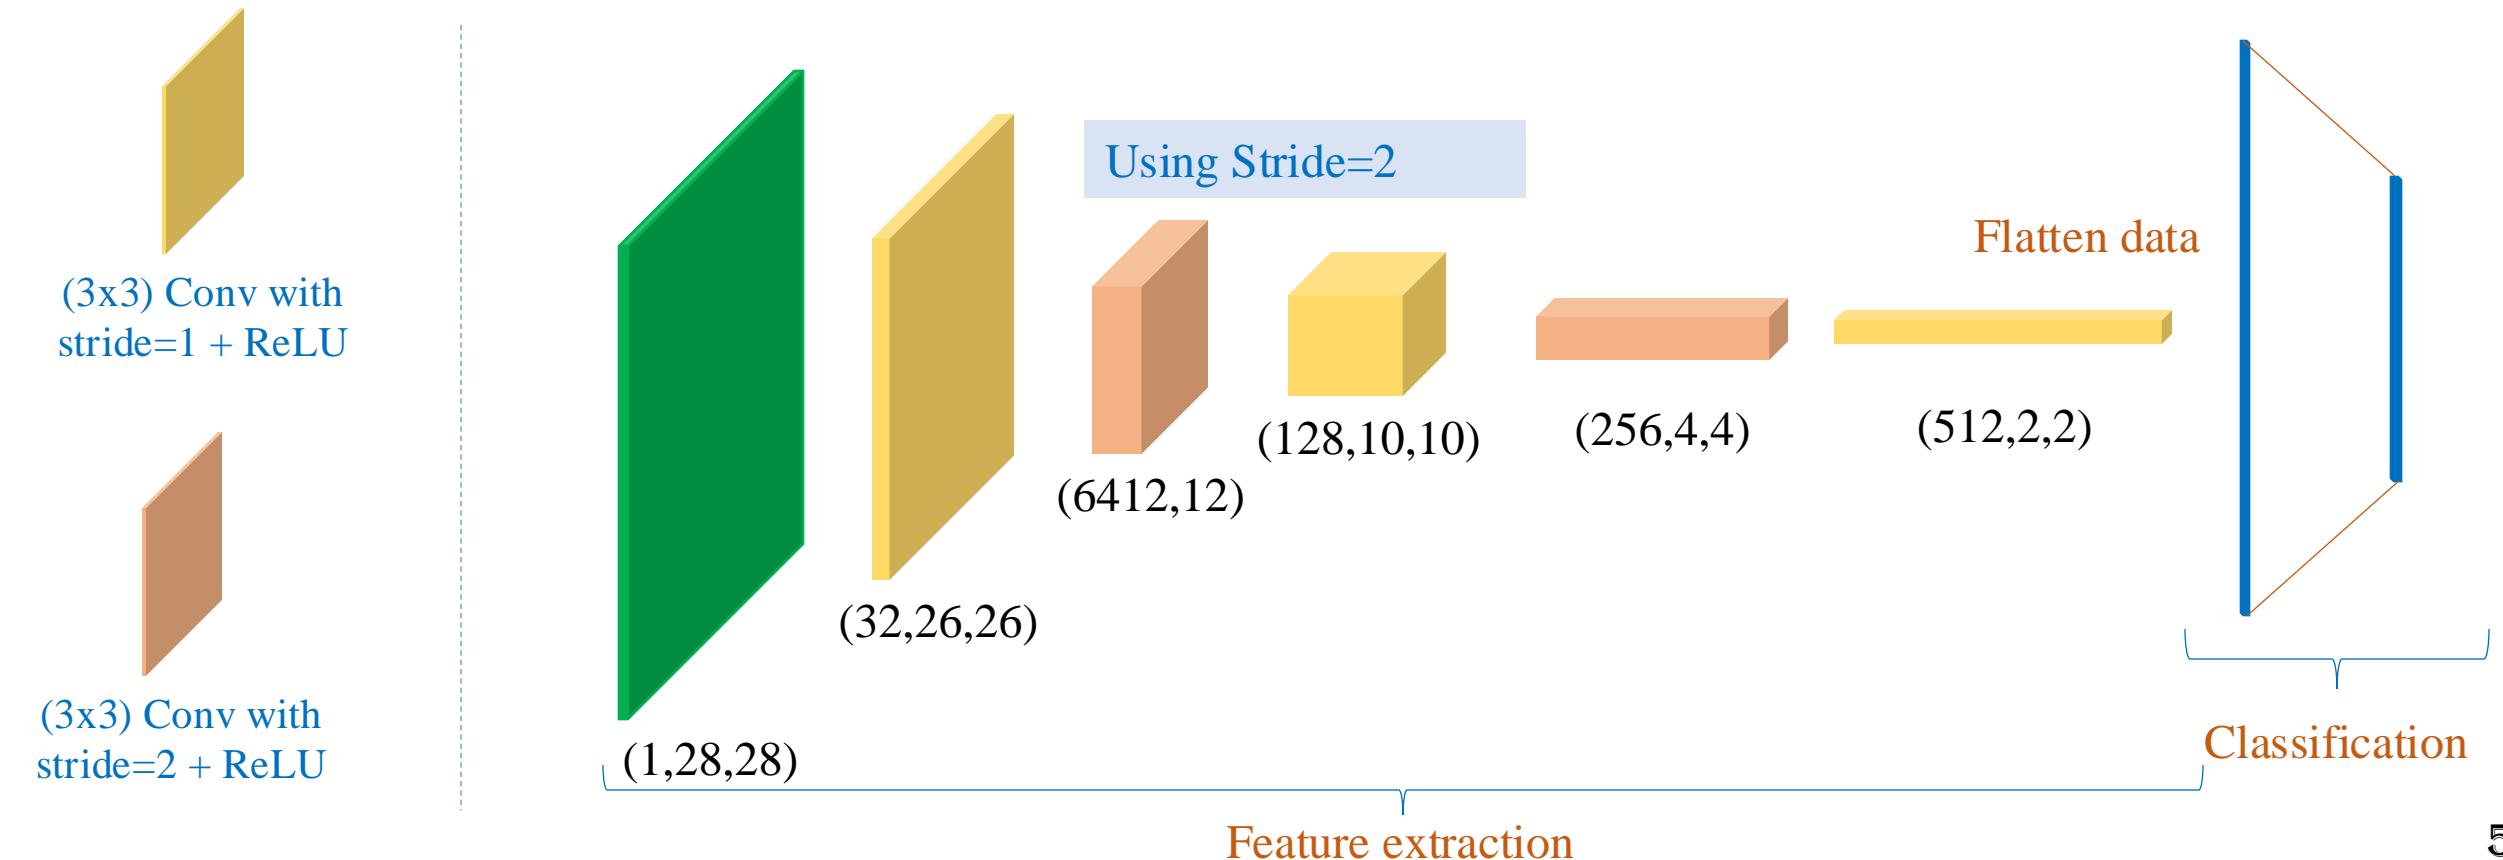
```

def forward(self, x):
    x = self.relu(self.conv1(x))
    x = self.relu(self.conv2(x))

    x = self.relu(self.conv3(x))
    x = self.relu(self.conv4(x))

    x = self.relu(self.conv5(x))
    x = self.flatten(x)
    x = self.dense(x)
    return x

```



Padding

$$S_o = \left\lfloor \frac{S_D - K + 2P}{S} \right\rfloor + 1$$

Keep resolution
of feature map

v_1	v_2	v_3	v_4
v_5	v_6	v_7	v_8
v_9	v_{10}	v_{11}	v_{12}
v_{13}	v_{14}	v_{15}	v_{16}

Data D
(4x4)

w_1	w_2	w_3
w_4	w_5	w_6
w_7	w_8	w_9

Kernel of parameters

1

Without using padding
or padding=0

Convolve
with stride=1 (D) =

m_1	m_2
m_4	m_5

Output
(2x2)

Padding = 1

v	v	v	v	v	v
v	v_1	v_2	v_3	v_4	v
v	v_5	v_6	v_7	v_8	v
v	v_9	v_{10}	v_{11}	v_{12}	v
v	v_{13}	v_{14}	v_{15}	v_{16}	v
v	v	v	v	v	v

Data D_p

Convolve
with stride=1 (D_p) =

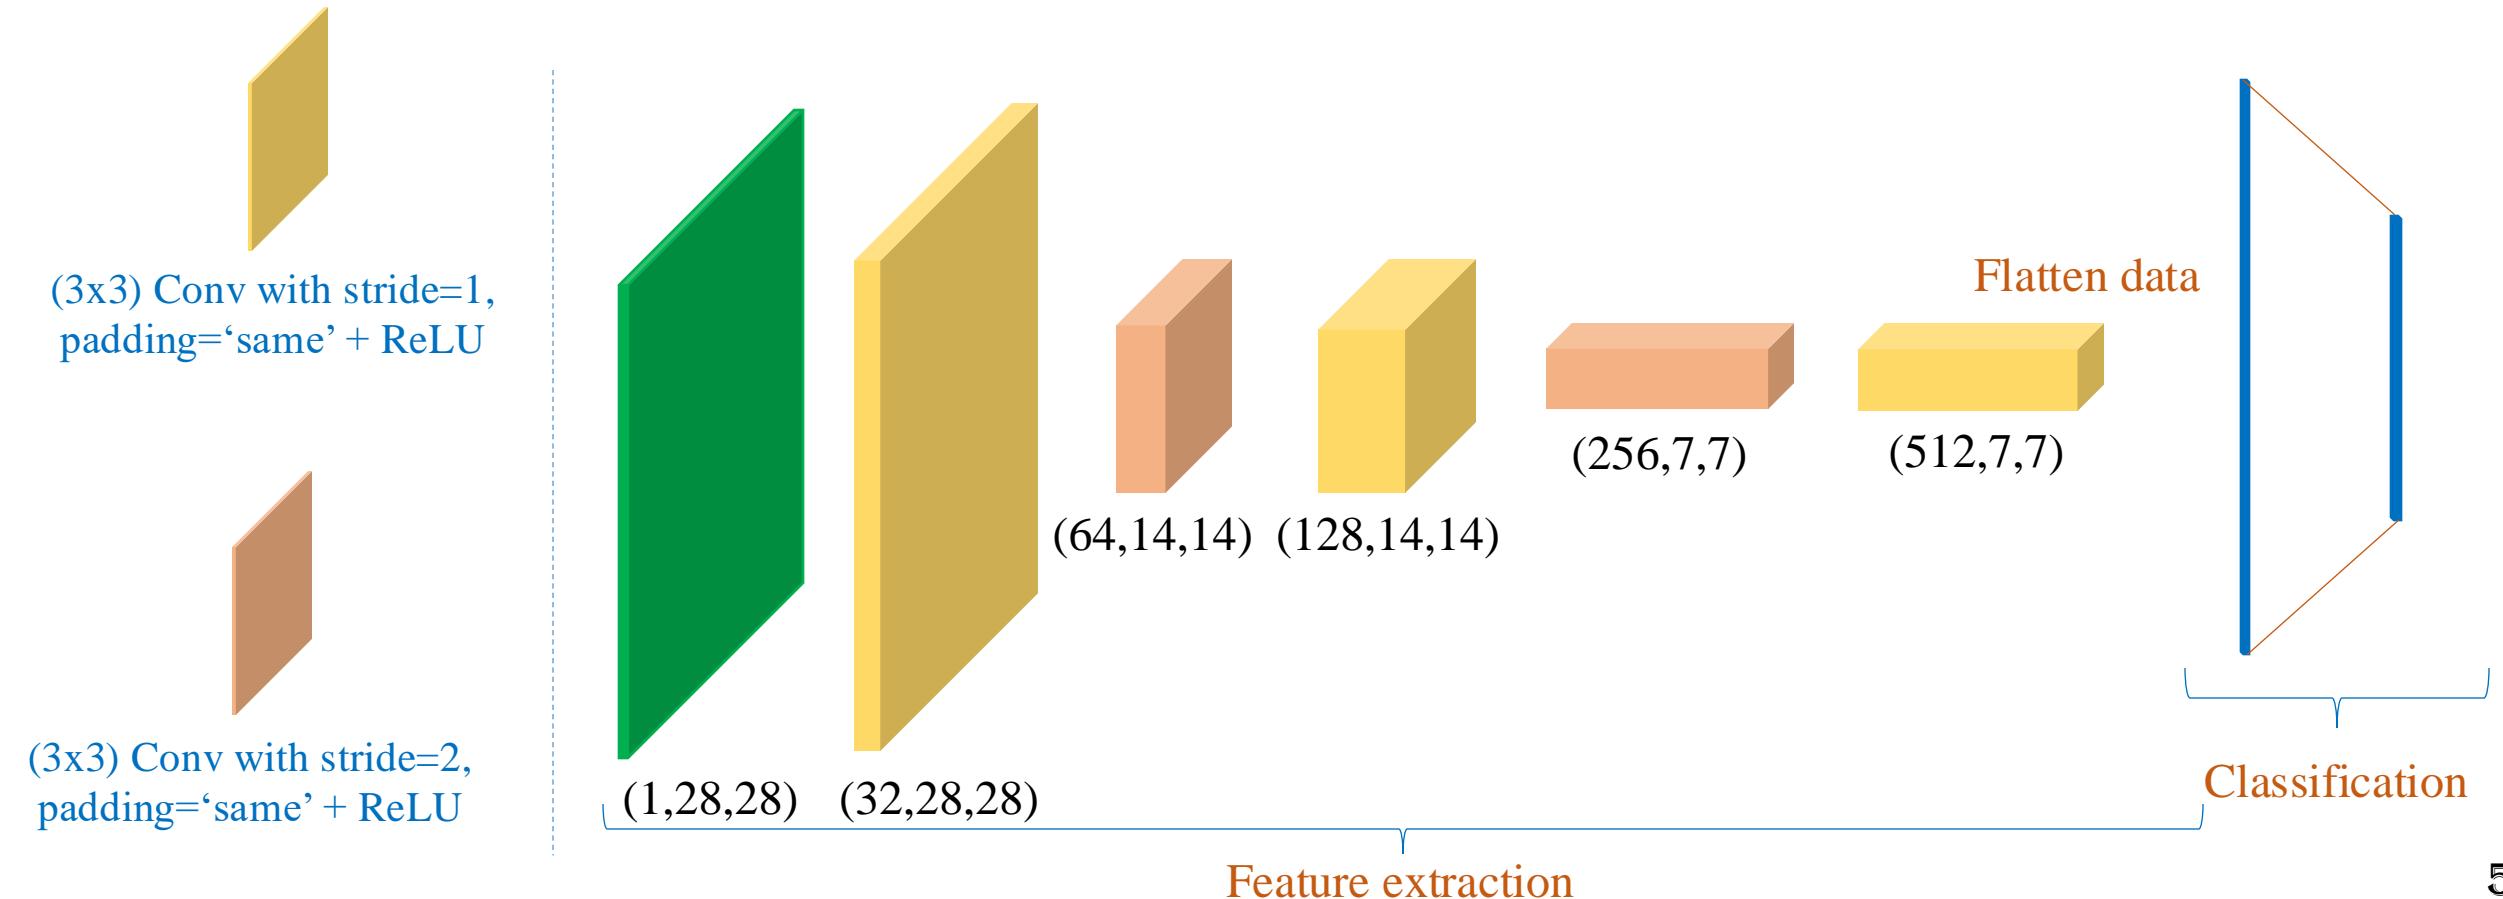
m_1	m_2	m_3	m_4
m_5	m_6	m_7	m_8
m_9	m_{10}	m_{11}	m_{12}
m_{13}	m_{14}	m_{15}	m_{16}

Output
(4x4)

Padding

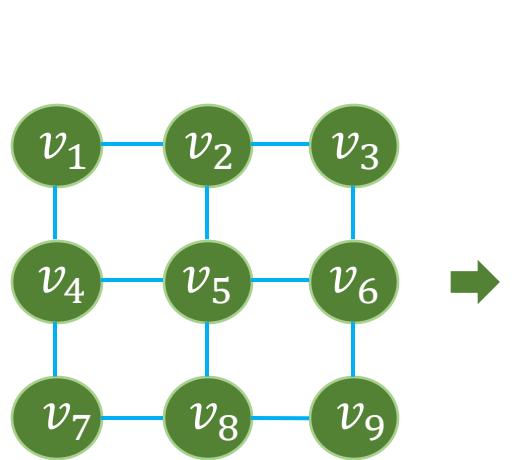
❖ Example

```
model = nn.Sequential(  
    nn.Conv2d(1, 32, kernel_size=3, padding='same', stride=1), nn.ReLU(),  
    nn.Conv2d(32, 64, kernel_size=3, padding=1, stride=2), nn.ReLU(),  
    nn.Conv2d(64, 128, kernel_size=3, padding=1, stride=1), nn.ReLU(),  
    nn.Conv2d(128, 256, kernel_size=3, padding=1, stride=2), nn.ReLU(),  
    nn.Conv2d(256, 512, kernel_size=3, padding=1, stride=1), nn.ReLU(),  
    nn.Flatten(), nn.Linear(7*7*512, 10)  
)
```



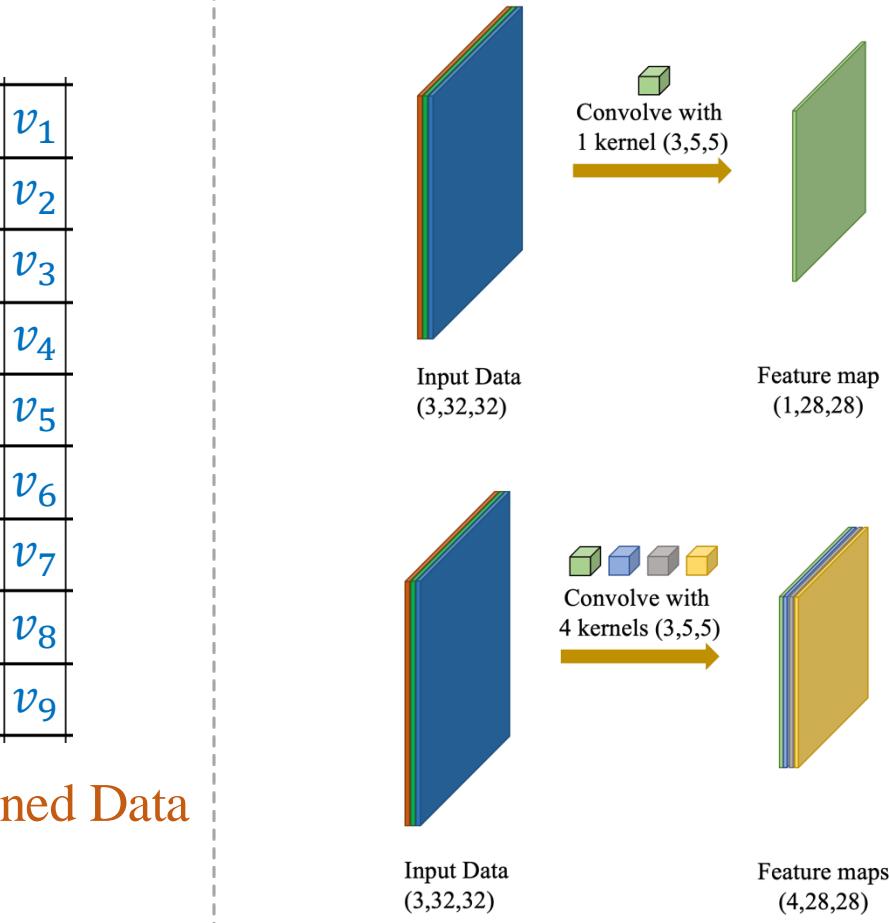
Summary

MLP Limitations

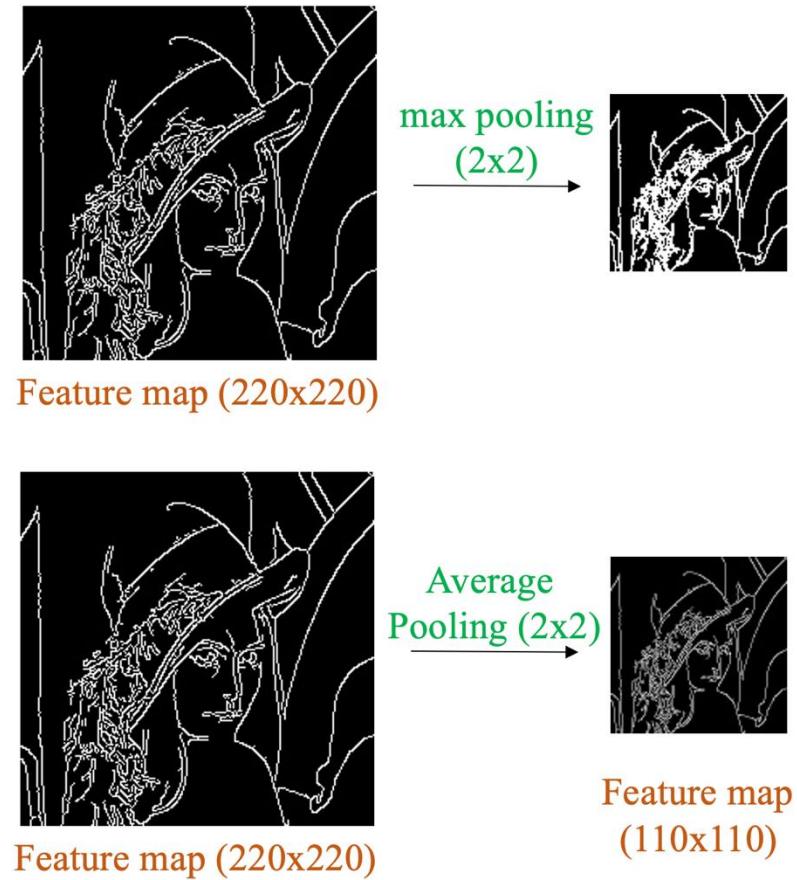


Flattened Data

Convolutional Layer



Standard CNNs



Further Reading

❖ Reading

<https://cs231n.github.io/convolutional-networks/>

<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

<https://stanford.edu/~shervine/teaching/cs-230/cheatsheet-convolutional-neural-networks>

